

# multiprocessing

## HPC Python

R. Todd Evans  
`rtevens@tacc.utexas.edu`

April 17, 2015

# What is Multiprocessing

- Process-based parallelism
- Not threading!
- Threads are light-weight execution units within a process that share memory
- Processes are more heavy-weight and do not share memory
- Processes communicate over some interprocess communication channel.

# Multiprocessing in Python - Why not threads?

## Python Threading Does Not Allow True Concurrency

- Python has a threading module and can use threads
- Python uses a global interpreter lock (GIL)
- In Python, only one thread executes Python code at a time
- GIL avoids concurrent access (race conditions) but no gain in performance for CPU-bound code with multiple threads

## Python Multiprocessing Does Allow True Concurrency

- multiprocessing uses subprocesses instead of threads to side-step GIL and allow concurrency in Python code!
- Parent process spawns child processes with own Python interpreter and own GIL
- Each child process inherits the data and program state from parent process

# Why Python Multiprocessing

## Parallelism!

- Python threads cannot work concurrently due to GIL
- Multiprocessing enables true concurrent (parallel) work across multiple cores
- As modern systems increase core counts, effective utilization of these cores is critical to performance
- If data is restricted to each process significant performance gains are possible
- Spawning processes and sharing data can have significant overhead
- Can scale well to a single node. (16 cores on Stampede & 20 cores on Maverick)
- Python is easy and allows for relatively fast prototyping and high productivity

# Multiprocessing Library

- Great implementation of multiprocessing on Python
- multiprocessing provides an interface similar to threading libraries such as OpenMP
- If you know threads, multiprocessing is easy.
- If you don't know threads, multiprocessing is still easy!
- You can communicate Python objects
- What you lose in performance (because Python can be slow), you gain in shorter development time

# Quick Example (Gets easier)

```
1 # multiprocessing_test.py
2 import random,os
3 import multiprocessing
4
5 def list_append(count, out_list):
6     """
7     Appends a
8     random number to the list 'count' number
9     of times. A CPU-heavy operation!
10    """
11    print os.getpid(), 'is working'
12    for i in range(count):
13        out_list.append(random.random())
14
15 if __name__ == "__main__":
16     size = 10000000    # Number of random numbers to add
17     procs = 2         # Number of processes to create
18
19     # Create a list of processes and define work for each process
20     process_list = []
21
22     for i in range(0, procs):
23         out_list = list()
24         process = multiprocessing.Process(target=list_append,
25                                           args=(size, out_list))
26         process_list.append(process)
27
28     # Start the processes (i.e. calculate the random number lists)
29     for p in process_list:
30         p.start()
31
32     # End all of the processes have finished
33     for p in process_list:
34         p.join()
35
36     print "List processing complete."
```

# Multiprocessing Workflow

## Basic Workflow

1. Create Process objects (aka child processes or **worker**) and assign target functions w/ arguments to work on
2. Spawn processes & processes do work
3. Wait for processes to terminate

## Basic Syntax

1. `p = Process(target=func, args=(arg0,arg1,...))`
2. `p.start()`
3. `p.join()`

# Functionality

## Basically same as threading

- Exchange objects/data between processes
  - Queues: multiple producers and consumers of work
  - Pipes: sends data between two processes (`send()`, `recv()`)
- Synchronization - locks etc.
- Sharing data between Processes
  - Shared memory: `Value` (scalar-like) `Array` (vector-like)
  - Server process (Manager): creates and manages shared objects/data
- Pool of Workers: sets up a pool of processes and gives them tasks

Manager + Pool = The safe, easy, and typical approach to multiprocessing. We will focus on this approach.



# Process Pools

## `multiprocessing.pool.Pool(processes)`: Basic methods

- `apply(func[, args, [kwargs]])`: call func using one worker from Pool. Blocks until complete (not concurrent)
- `apply_async(func[, args[, kwargs]])`: apply that does not block (allows concurrency)
- `map(func, iterable)`: iterable is any object than can be iterated over such as a list or dict. the iterable is divided among workers as chunks. When worker finishes it grabs new chunk. Blocks until complete.
- `map_async(func, iterable)`: map that does not block

# apply(): simple example

## apply() vs apply\_async(): blocking vs non-blocking

```
1 # apply_test.py
2 import time
3 from multiprocessing import Pool
4
5 def f():
6     start = time.time()
7     time.sleep(2)
8     end = time.time()
9     return end-start
10
11 p = Pool(processes=1)
12
13 # apply function
14
15 result = p.apply(f) # blocking
16 print "apply is blocking"
17 print 'total time',result
18
19 # apply_async function
20 result = p.apply_async(f) # non-blocking
21 print "apply_async is non-blocking"
22 while not result.ready():
23     time.sleep(0.5)
24     print 'working on whatever else I want...'
25 print 'total time',result.get() # but get() is blocking
```

# map(): simple example

```
1 # map_test.py
2 import time
3 from multiprocessing import Pool
4
5 def f(x):
6     return x**3
7
8
9 y = range(int(1e7))
10
11 p = Pool(processes=4)
12
13 # map function
14 start = time.time()
15 results = p.map(f,y) # blocking
16 end = time.time()
17 print "map blocks"
18 print "time",end-start
19
20 # map_async
21 start = time.time()
22 results = p.map_async(f,y) # non-blocking
23 end = time.time()
24 print "map_async is non-blocking"
25 output = results.get() # but get() is blocking
26 print "time",end-start
```

# map(): Monte Carlo Integration

```
1 # multiproc_mc.py
2 # integrate f(x) from a to b
3
4 from multiprocessing import Pool
5 from numpy import random
6 import time,math
7
8 def f(x):
9     r,s=x
10    var = random.random()*(s-r)+r # [0,1] -> [a,b]
11    return var**2 # x^2
12
13 a = 0 # lower integration bound
14 b = 1 # upper integration bound
15 N = 10000000 # number of samples
16 irange = N*[a,b]
17
18 # Scalar
19 random.seed(0)
20 start = time.time()
21 samples = map(f,irange) # serial map
22 end = time.time()
23 I = (b-a)*sum(samples)/N # compute integral
24 print 'scalar result',I,end-start
25
26 # Multiprocessing
27 p = Pool(processes=4) # 4 process Pool
28 random.seed(0)
29 start = time.time()
30 samples = p.map(f,irange) # parallel map
31 end = time.time()
32 I = (b-a)*sum(samples)/N # compute integral
33 print 'parallel result',I,end-start
```

# map(): Prime Factorization

Try with different number of processes N=1,2,3,...

python multiproc\_prime.py N

```
1 # multiproc_prime.py
2 from multiprocessing import Pool
3 import time, sys
4
5 def is_prime(n):
6     if n < 2: return False
7     if n < 4: return True
8     maxfactors = int(n ** 0.5) + 1
9     for i in range(2,maxfactors):
10         is_factor = n % i == 0
11         if is_factor:
12             return False
13     return True
14
15 numbers_to_test = range(int(1e5))
16 # Serial
17 start = time.time()
18 results = map(is_prime,numbers_to_test)
19 end = time.time()
20 print 'Time for serial prime factorization',end-start
21
22 # Parallel
23 processes = int(sys.argv[1])
24 p = Pool(processes)
25 # map determines
26 start = time.time()
27 results = p.map(is_prime,numbers_to_test)
28 end = time.time()
29 print 'Time for parallel prime factorization',end-start
```

# Inter-process Data Sharing via Manager()

- `map()` map doesn't share data between processes
- Managing shared data can be tricky and error prone.

## Let the Multiprocessing Manager do it!

```
manager = Manager()
```

```
d = manager.dict()
```

```
l = manager.list()
```

`d` and `l` are visible to every process and manager keeps them sync'd

# Shared data are seen by all processes

```
1 # multiprocessing_manager.py
2 from multiprocessing import Manager, Pool
3 import os
4
5 def f(l,d):
6     l.append('worker')
7     d[str(os.getpid())] = 'worker'
8 manager = Manager()
9 pool    = Pool(2)
10
11 # private_l and private_d only visible to local process
12 private_l = list()
13 private_d = dict()
14
15 # shared_l and shared_d visible to every process
16 shared_l = manager.list()
17 shared_d = manager.dict()
18
19 # manager process can see this change
20 private_l.append('manager')
21 private_d[str(os.getpid())] = 'manager'
22
23 # manager process can see this change
24 shared_l.append('manager')
25 shared_d[str(os.getpid())] = 'manager'
26
27 # changes child processes makes are lost
28 pool.apply(f,args=(private_l,private_d))
29 pool.apply(f,args=(private_l,private_d))
30 print "try to add to private data",private_l,private_d
31
32 # changes child processes makes are kept
33 pool.apply(f,args=(shared_l,shared_d))
34 pool.apply(f,args=(shared_l,shared_d))
35 print "try to add to shared data",shared_l,shared_d
```

# Shared data are seen by all processes

## Save those primes!

```
1 # multiproc_prime_manager.py
2 from multiprocessing import Pool, Manager
3 from functools import partial
4 import time, sys
5
6 def shared_is_prime(n,p):
7     if n < 2: return False
8     if n < 4: return True
9     maxfactors = int(n ** 0.5) + 1
10    for i in range(2,maxfactors):
11        is_factor = n % i == 0
12        if is_factor:
13            return False
14    p.append(n)
15    return True
16
17 numbers_to_test = range(int(1e5))
18
19 # Parallel
20 processes = int(sys.argv[1])
21 p = Pool(processes)
22 manager = Manager()
23 primes = manager.list()
24 # map determines
25 start = time.time()
26 partial_shared_is_prime = partial(shared_is_prime, p = primes) # trick to pass more than 1 arg
27 results = p.map(partial_shared_is_prime,numbers_to_test)
28 end = time.time()
29 print primes
30 print 'Time for parallel prime factorization',end-start
```



# License

©The University of Texas at Austin, 2015

This work is licensed under the Creative Commons Attribution Non-Commercial 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/>

When attributing this work, please use the following text: "HPC Python", Texas Advanced Computing Center, 2015. Available under a Creative Commons Attribution Non-Commercial 3.0 Unported License.

