

Applications of Machine Learning to Analysis and Computation for Particle Accelerators

Jonathan Edelen and Auralee Edelen

Andreas Adelmann, Daniel Brommelsiek, Dean Edstrom, Aliaksei Halavanau, Nicole Neveu, Philippe Piot, Sasha Romanov, Jinhao Ruan, and Sasha Valishev.

Software Engineering Assembly - Improving Scientific Software

9 April 2019 – NCAR



Outline

- Overview of accelerators, their applications, and why we're interested in ML
- Current efforts to apply ML to accelerators
 - *Anomaly detection and data cleaning*
 - *Virtual diagnostics and on-line modeling*
 - *Surrogate modeling and optimization*
- Summary

Particle accelerators and their applications

Electron Accelerators

E-Beam sources

(Electron sources for a range of applications)

- Cancer treatment
- Polymer linking
- Wastewater treatment

Free electron lasers

(Highly tunable coherent photon sources ranging from IR to Hard X-Ray)

- Basic research
- Directed energy
- Industrial research

Synchrotron light sources

(Tunable photon sources usually in the x-ray)

- Basic research (NSLS-II, APS)
- Materials science, biology, etc.

Hadron accelerators

Proton accelerators

- High flux neutron sources (SNS)
- Accelerator driven subcritical reactors
- Basic research (fixed target and colliders)
- Cancer treatment

Ion accelerators

- Basic research (nuclear physics)
- Cancer treatment
- Isotope production

Particle accelerators and their applications

Electron Accelerators

E-Beam sources

(Electron sources for a range of applications)

- Cancer treatment
- Polymer linking
- Wastewater treatment

Free electron lasers

(Highly tunable coherent photon sources ranging from IR to Hard X-Ray)

- Basic research
- Directed energy
- Industrial research

Synchrotron light sources

(Tunable photon sources usually in the x-ray)

- Basic research (NSLS-II, APS)
- Materials science, biology, etc.

Hadron accelerators

Proton accelerators

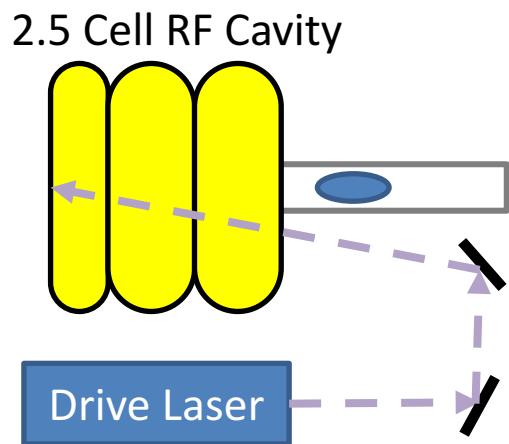
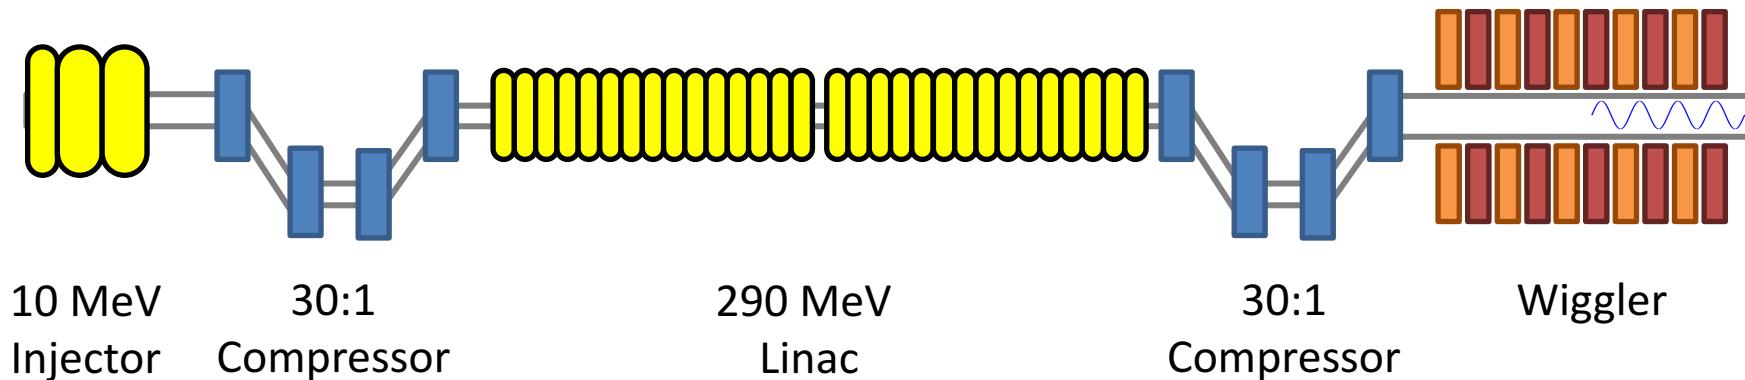
- High flux neutron sources (SNS)
- Accelerator driven subcritical reactors
- Basic research (fixed target and colliders)
- Cancer treatment

Ion accelerators

- Basic research (nuclear physics)
- Cancer treatment
- Isotope production

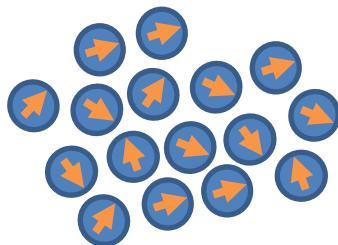
See DOE Report “Accelerators for America’s Future”
(<https://science.energy.gov/~media/hep/pdf/accelerator-rd-stewardship/Report.pdf>)

Electron sources for context

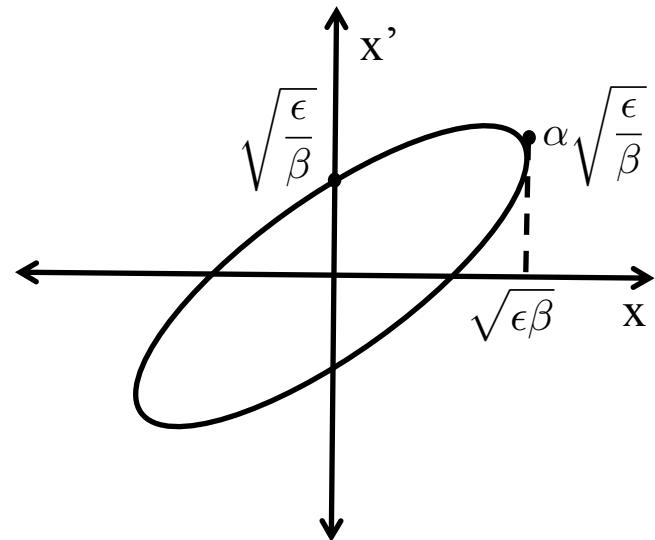


Emittance: a measure of beam quality

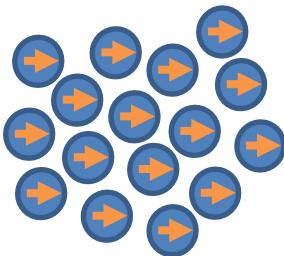
High degree of disorder – high emittance



- Beam ellipse in $x-x'$ phase space with important parameters noted:
- Beta – beam envelope function
- Alpha – change in the beam envelope function
- Epsilon – beam emittance



Low degree of disorder – low emittance



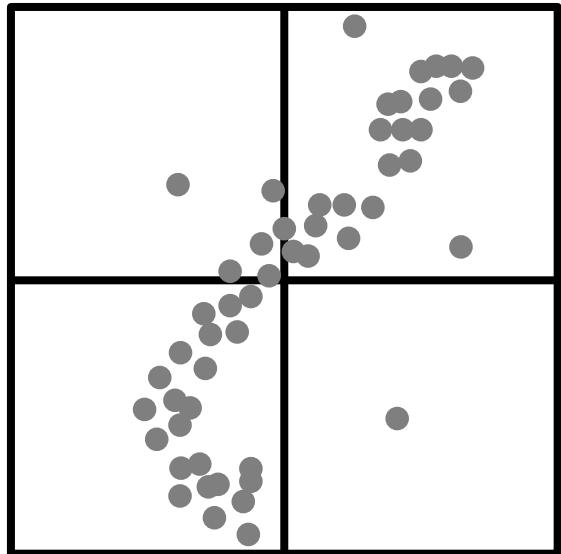
Why are we interested in machine learning?

- Accelerator facilities rely heavily on human operators for tuning/control
- Modeling and control of these machines is challenging
 - *Nonlinear systems with large parameter spaces*
 - *Variety of diagnostics (e.g. beam images), but these are limited and number, and some are not continuously available for use*
 - *Time-varying/ non-stationary behavior*
- Strong incentives for improving control (and understanding system)
 - *High user demand → want to switch between custom user requests quickly*
 - *High cost for unintended down-time → personnel time, user time, scientific output*
 - *Achieve challenging beam setups for new science goals*

Fast, accurate accelerator models → aid operation, model-based control, prototyping and experiment planning, design of new setups

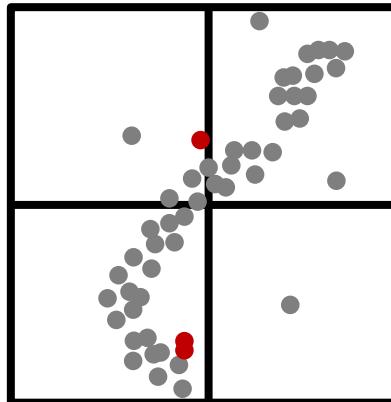
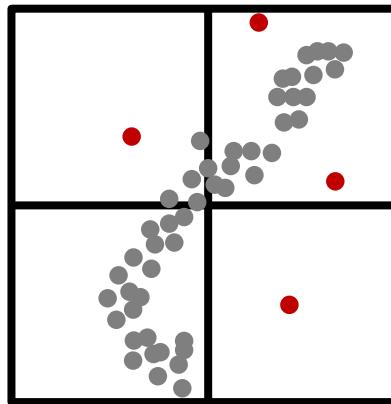
Application: Data Cleaning

What do we mean by data cleaning?

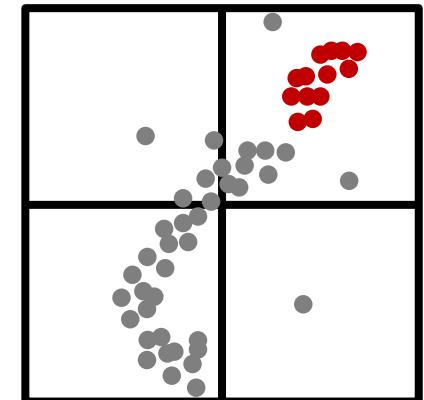


Raw data
(measurement or
simulation)

Obvious Outliers
(bad data or simulation run)



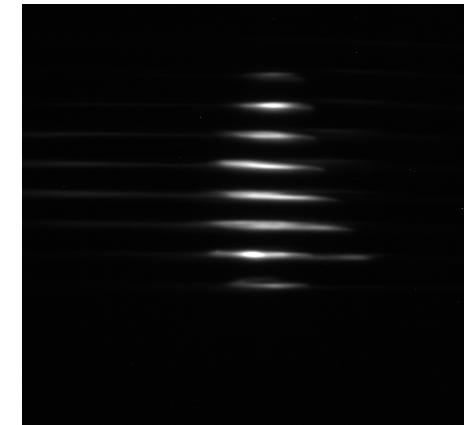
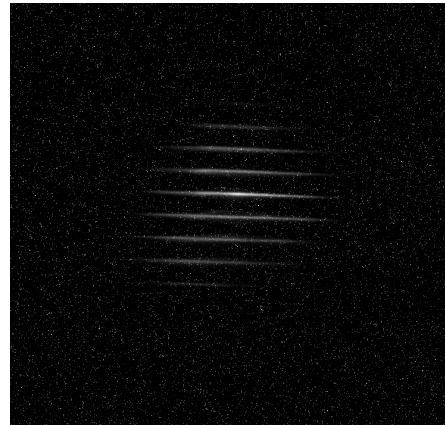
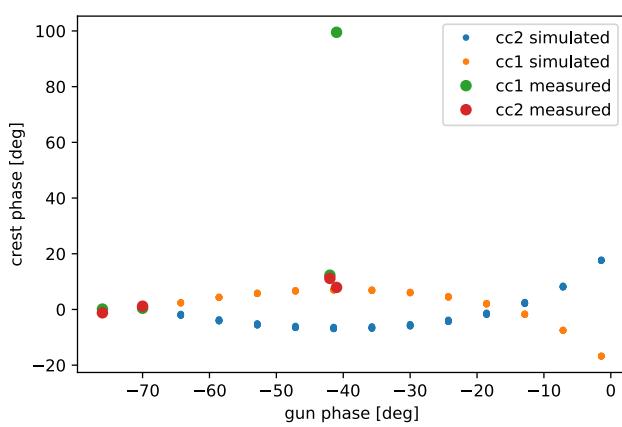
Subtle Outliers
(possible anomalies)



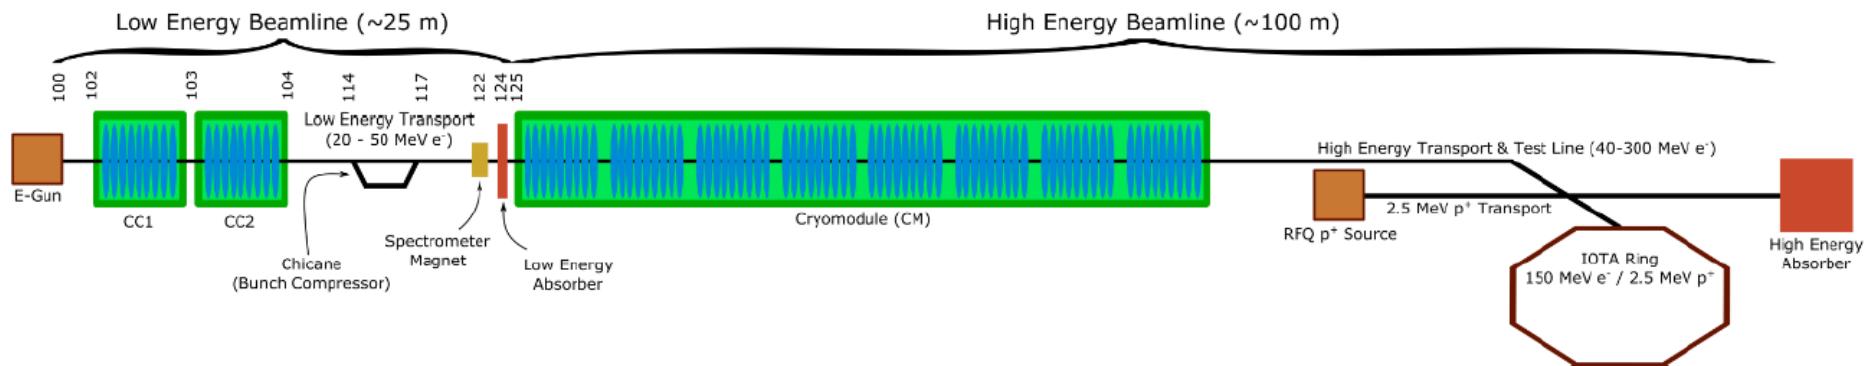
Discontinuities
(change in calibration or
machine state)

Case study: Simulations of the FAST Accelerator

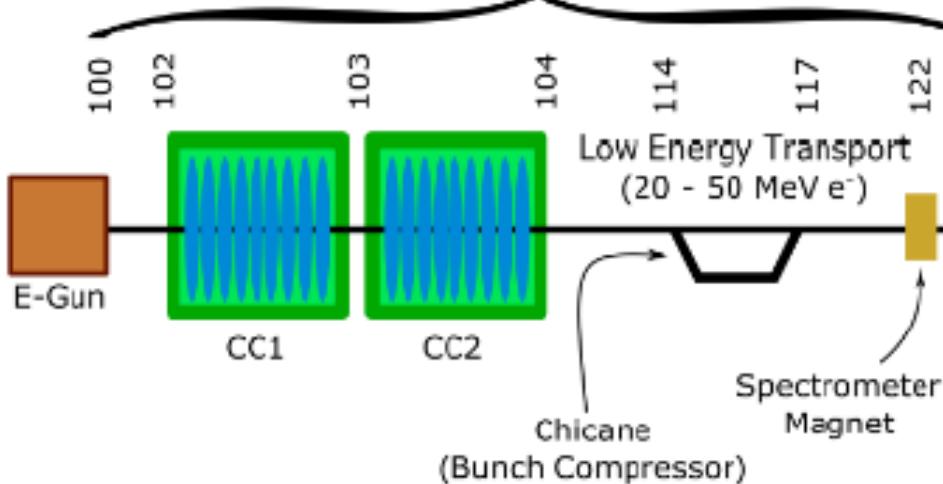
- Explore the use of unsupervised learning for automatic data cleaning using case studies:
 - *Start simple*
 - Batch simulation scans
 - *Increase complexity:*
 - Classification of machine drift
 - Multi-slit emittance measurements (images)



The Fermilab Science and Technology Facility



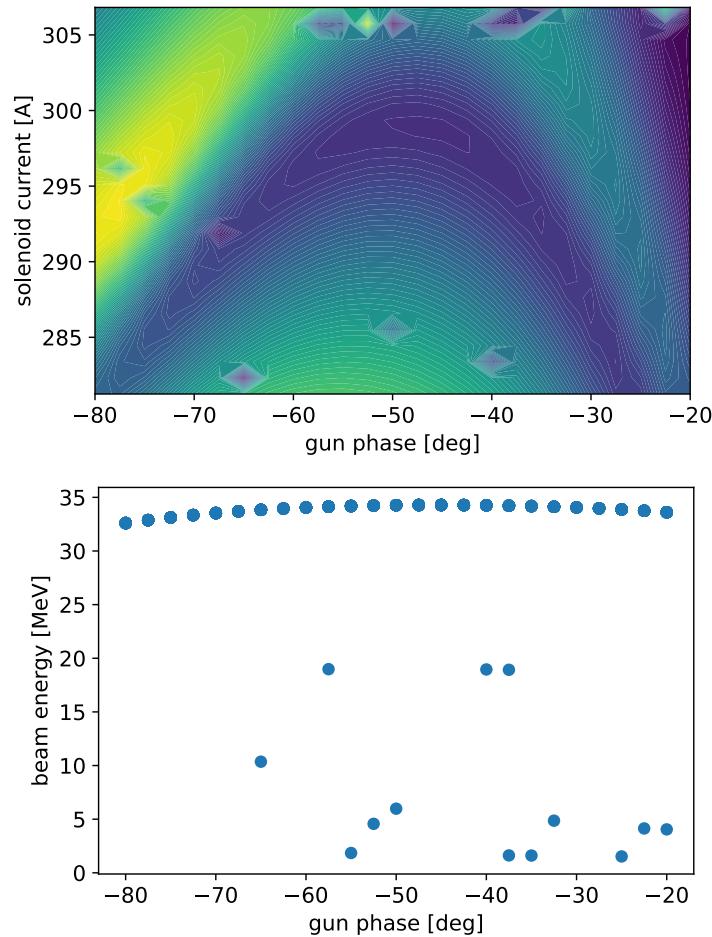
Low Energy Beamline (~25 m)



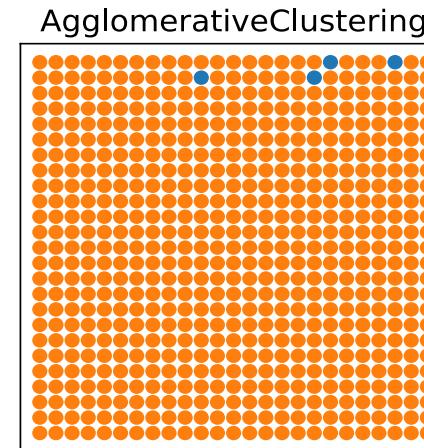
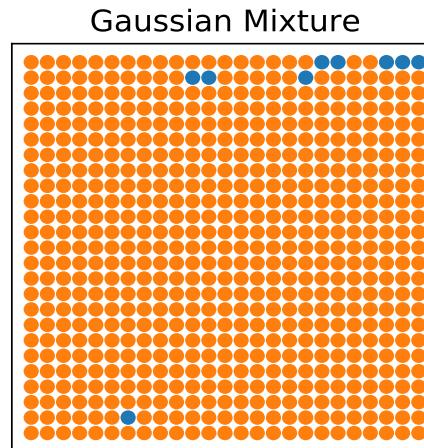
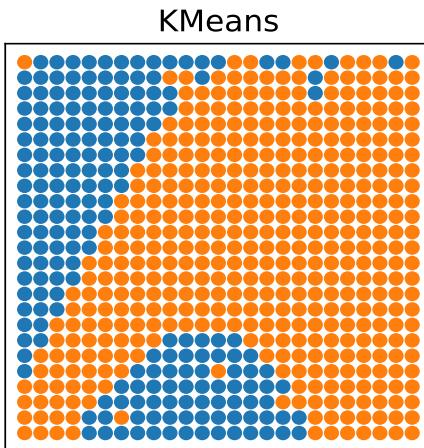
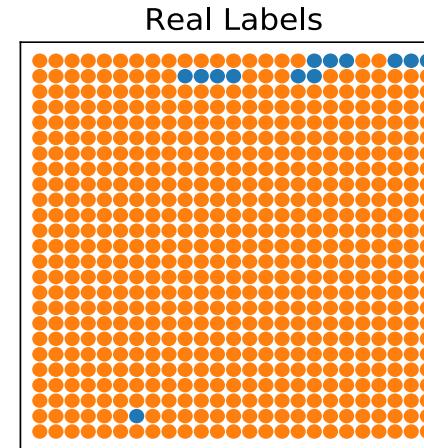
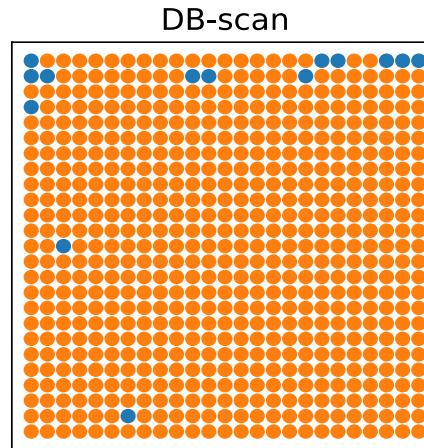
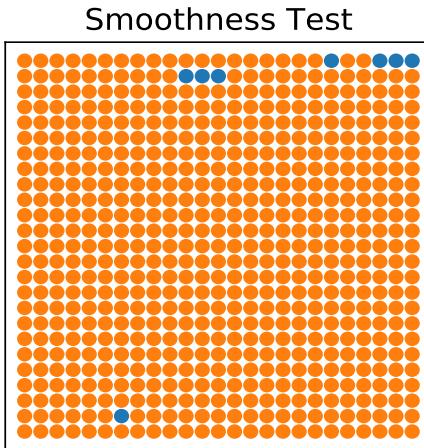
- Simulations start at the cathode and run up to x117
- Photo-cathode RF Electron Gun accelerates the beam to 4.5 MeV
- Superconducting RF cavities accelerate the beam to ~40 MeV

2-D Simulation scan data

- 2-D scan of gun-phase and solenoid strength
 - Run on high performance computing, (Linux cluster with 100 cores)
 - Some simulations terminated unexpectedly
 - Remove unwanted data from dataset
- Energy is the cleanest indicator of good vs. bad
 - Use this to label dataset but exclude from clustering analysis



Identification of bad simulation runs



Orange: Identified good run

Blue: Identified bad run

Identification of bad simulation runs

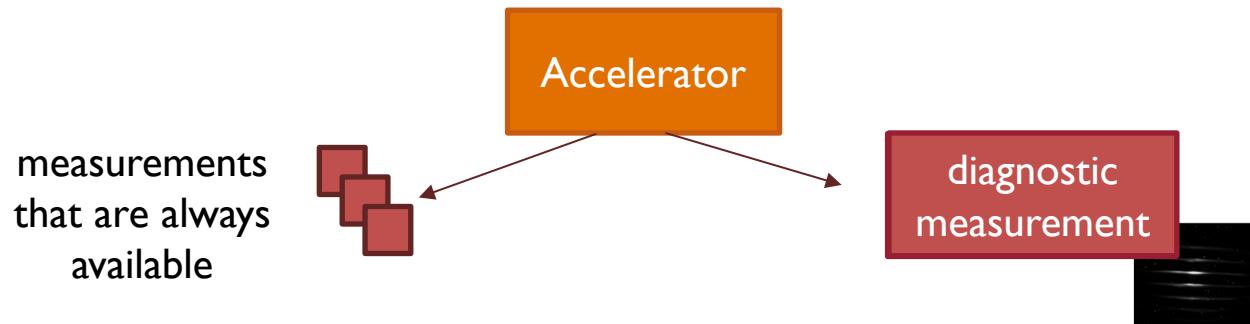
	K-means	DB-Scan	Gaussian Mix	Agglo	Smoothness
Percentage Correct	67.2%	98.6%	99.4%	98.6%	99.2%
Correctly Identified Bad Runs	3/13	9/13	9/13	4/13	8/13
False Positive	10/13	4/13	4/13	9/13	5/13
False Negative	195/612	5/612	0/612	0/612	0/612

False Positive: Predicted to be good but are actually bad. False Negative: Predicted to be bad but are actually good

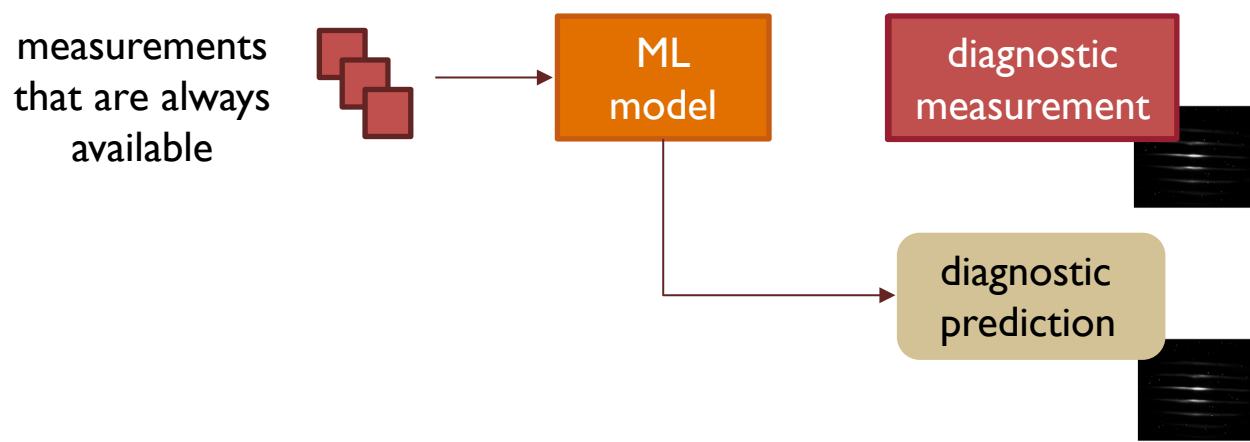
- DB-Scan/Gaussian Mixture/Smoothness have similar performance
- K-means and Agglomerative are both poor performers
- Gaussian mixture is a very good option for this dataset as specification of hyper-parameters is easiest and zero false positives

Application: Virtual Diagnostics and On-line Modeling

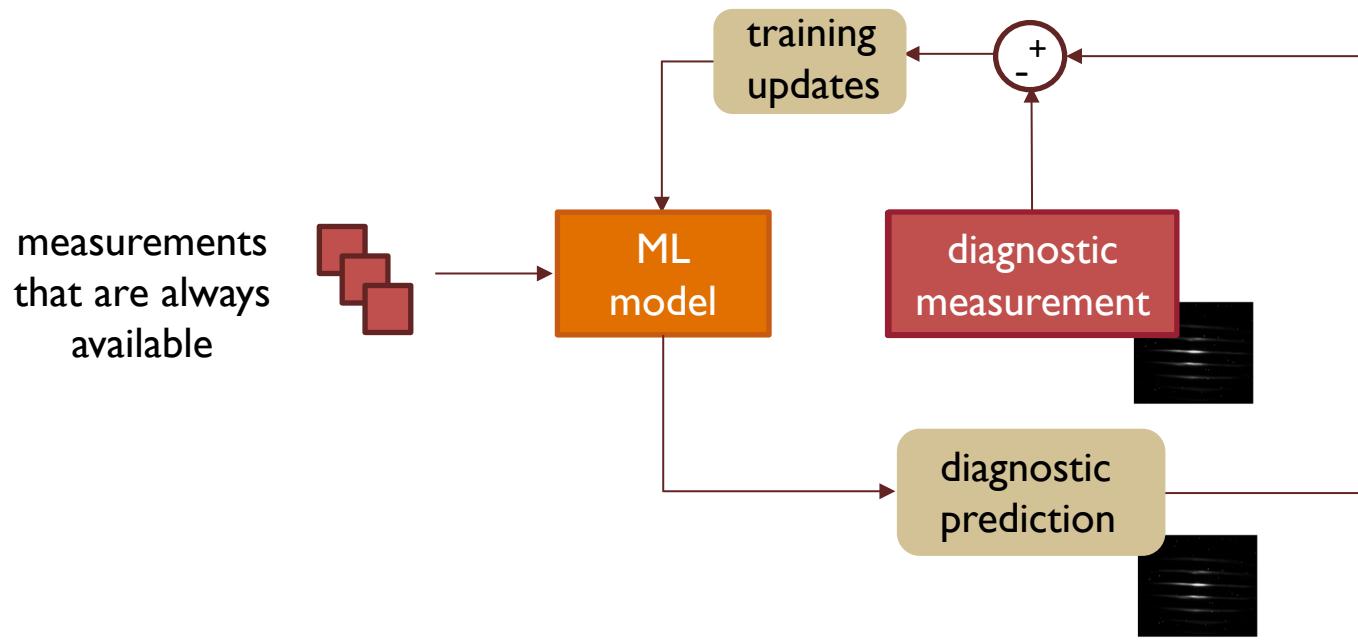
What is a virtual diagnostic?



What is a virtual diagnostic?



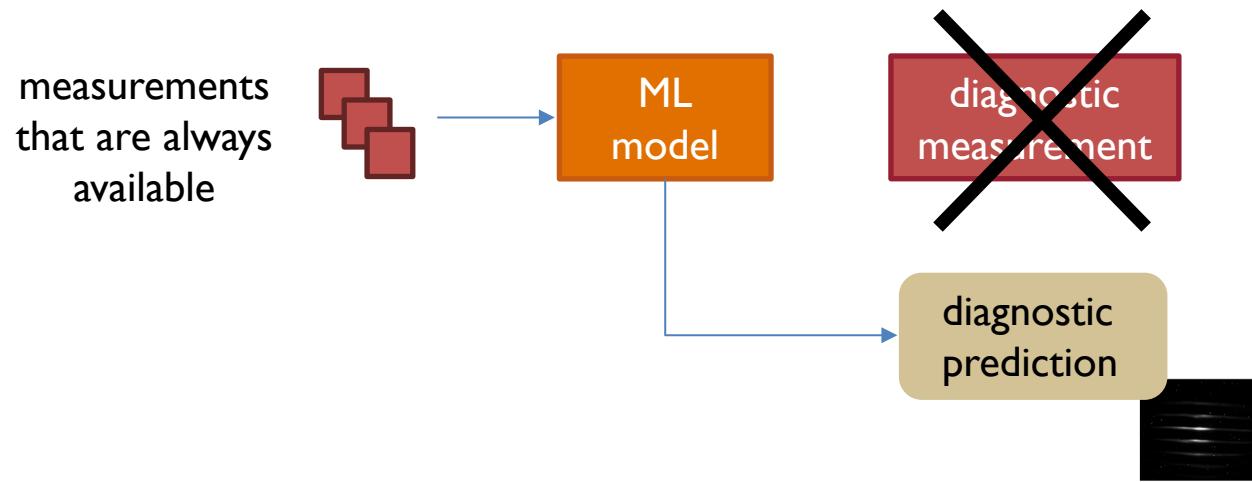
What is a virtual diagnostic?



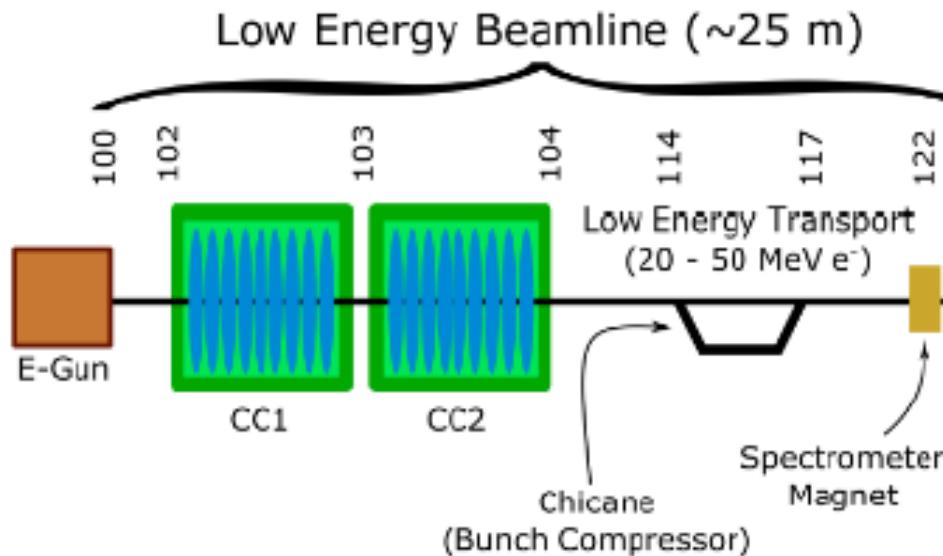
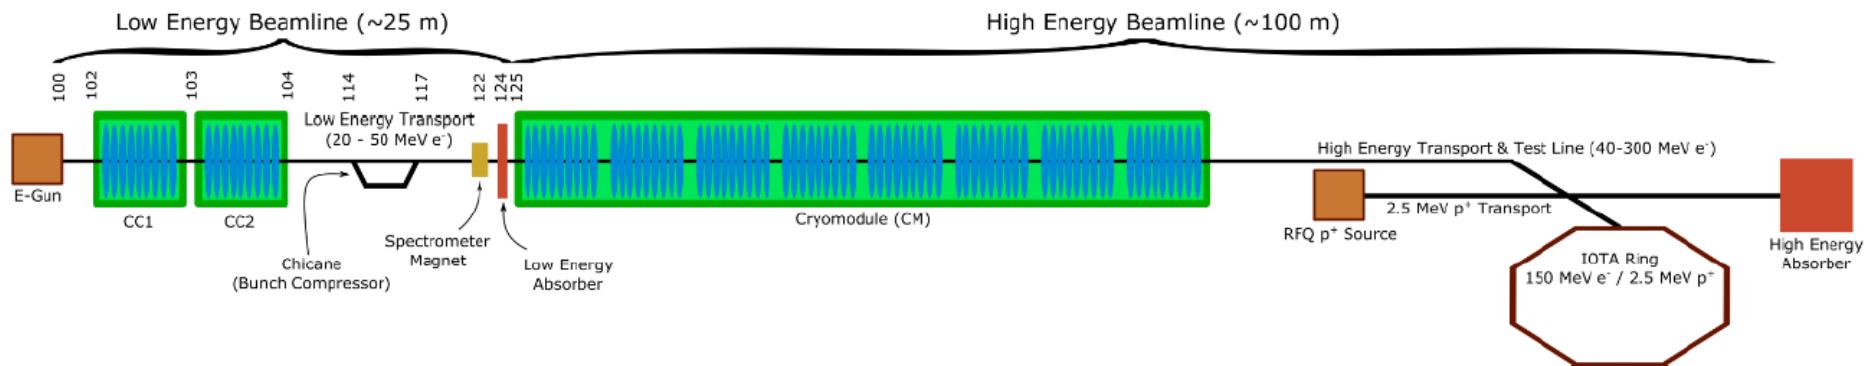
What is a virtual diagnostic?

Real diagnostic no longer available:

- destructive, cannot use during normal operations
- slower update rate than desired
- moved to another location (e.g. cost constraints)

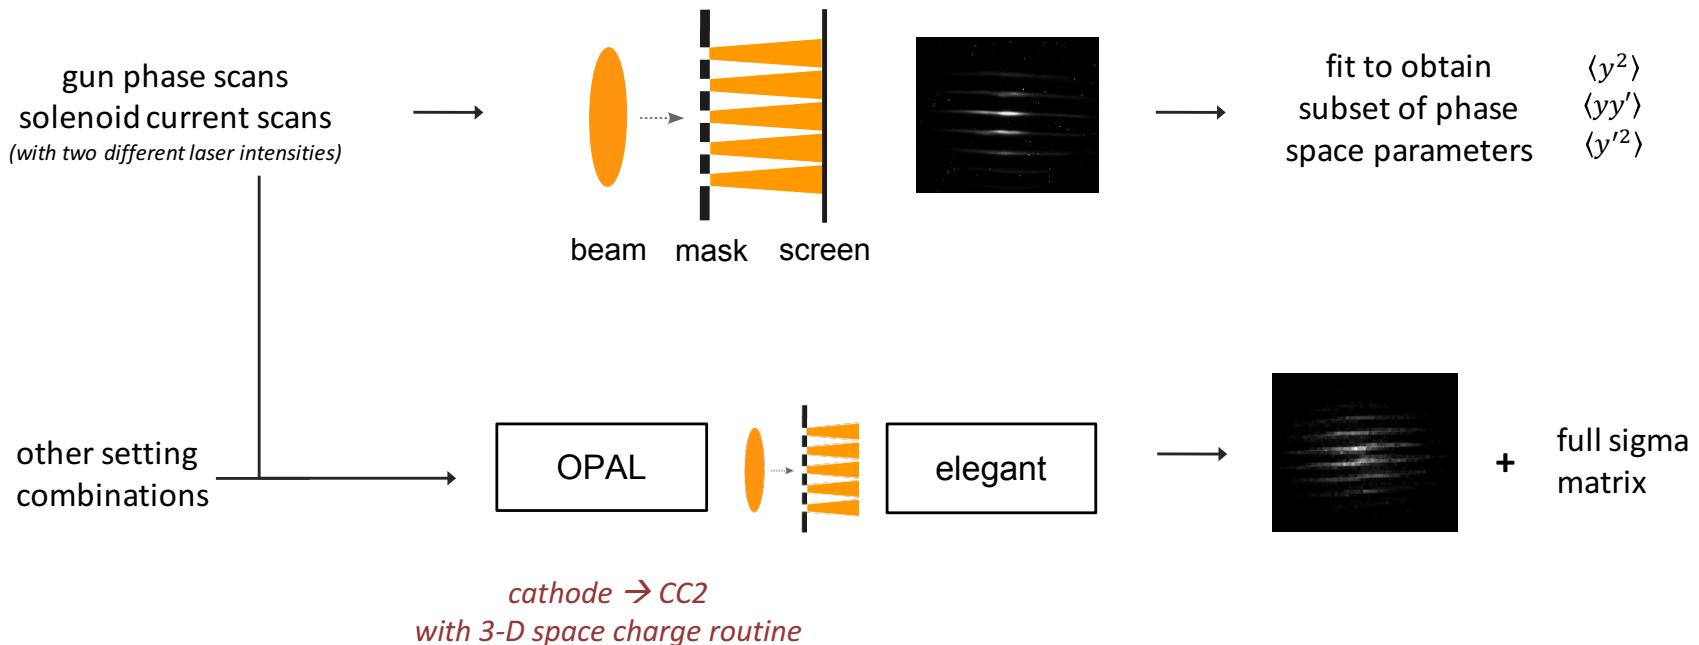


The Fermilab Science and Technology Facility



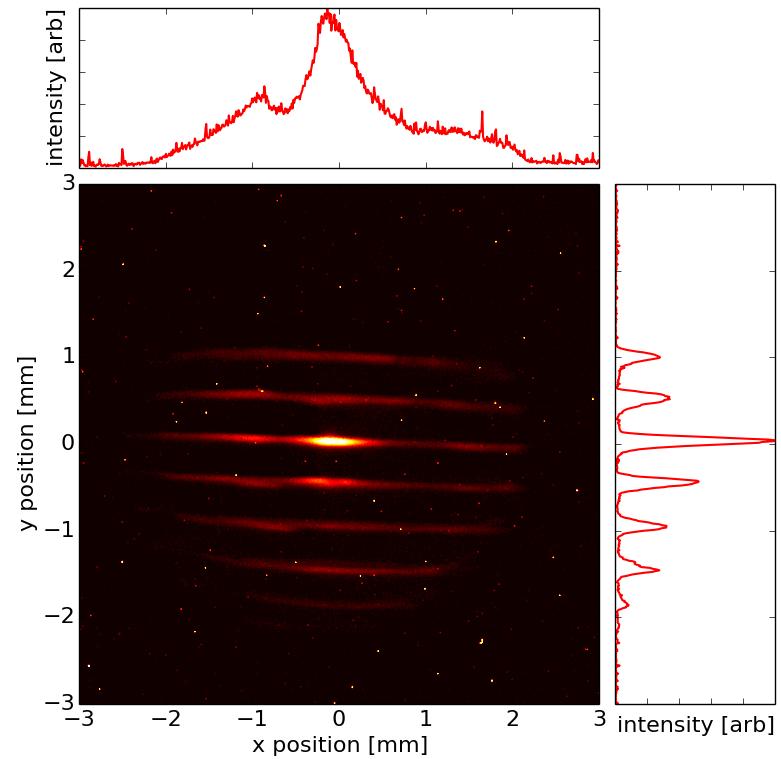
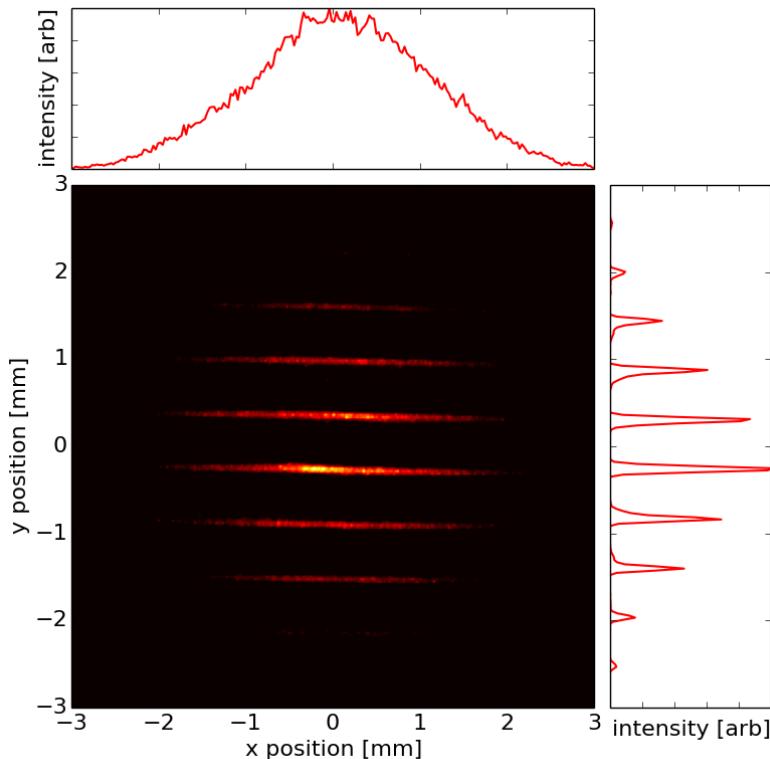
- 1.3 GHz photocathode RF gun
 - PITZ style gun with solenoid and bucking coil
 - Beam accelerated to ~4 MeV
- 1.3 GHz 9-cell Tesla type cavities
 - Beam accelerated to ~35 MeV

Developing our simulation model



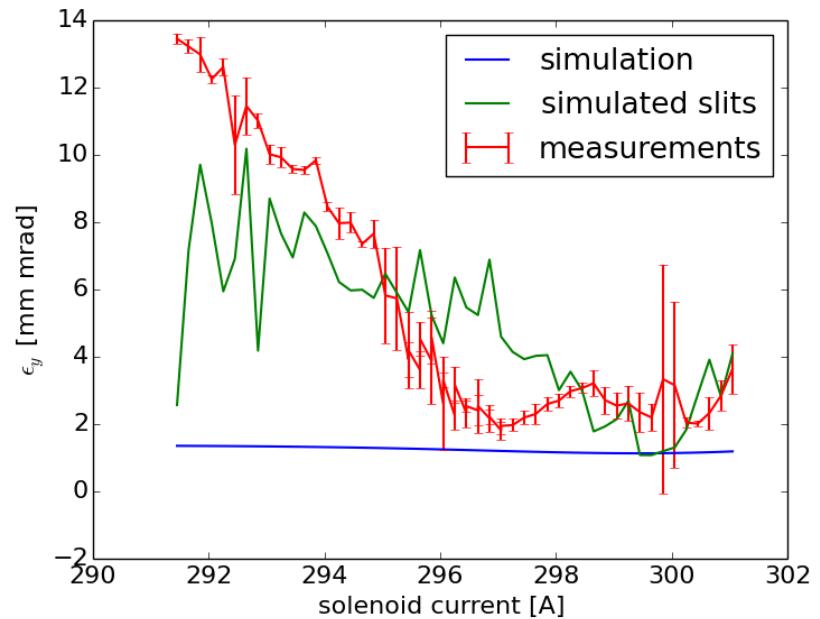
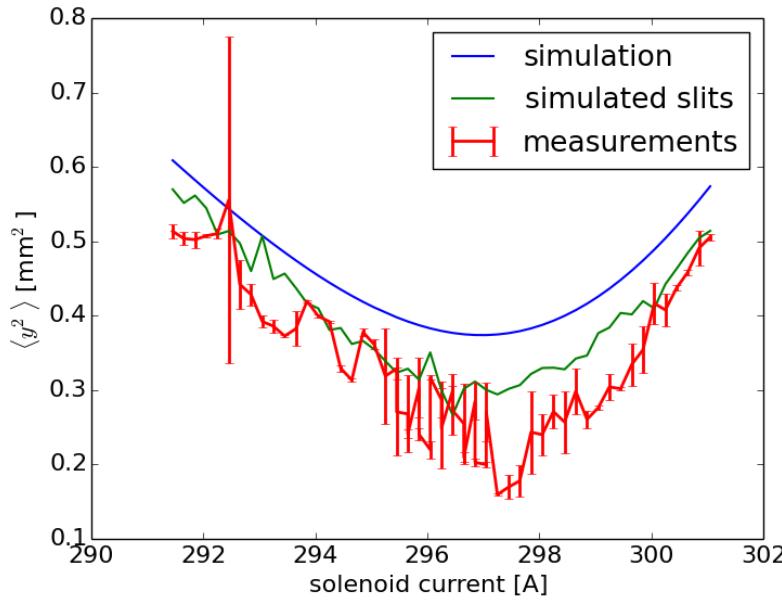
Comparison of simulations and measurements

- Simulating the multi-slit diagnostic
 - Export beam distribution at X107, apply mask, propagate to X111
 - Generate simulated images from 2-d histograms
 - Process images in the same manor as is done on the machine
 - Compare simulated images with measured images and compare processed results

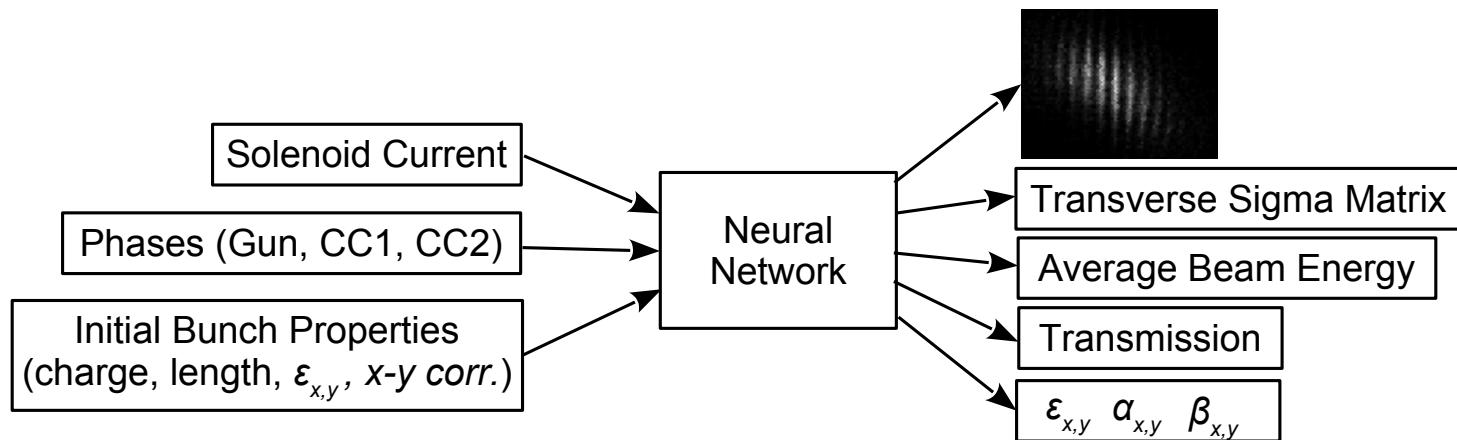


Comparison of simulations and measurements

- Simulating the multi-slit diagnostic
 - Export beam distribution at X107, apply mask, propagate to X111
 - Generate simulated images from 2-d histograms
 - Process images in the same manor as is done on the machine
 - Compare simulated images with measured images and compare processed results

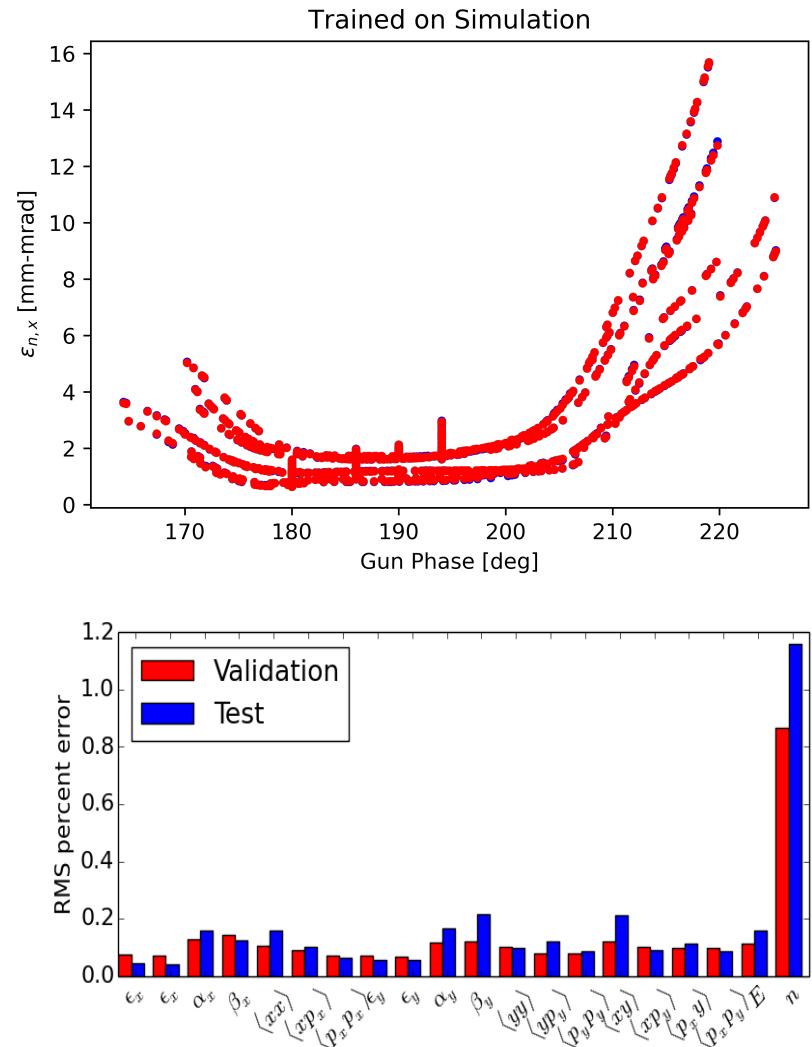


Neural Network Modeling



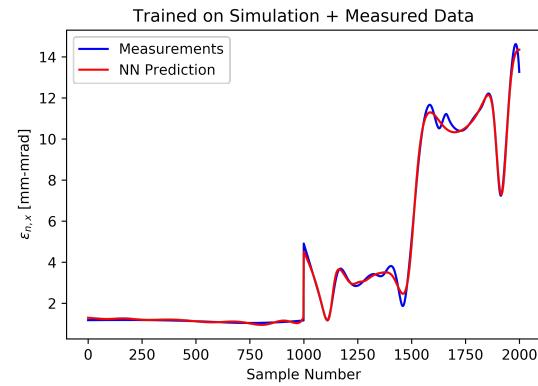
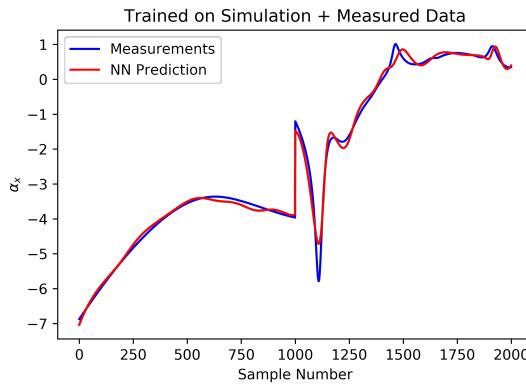
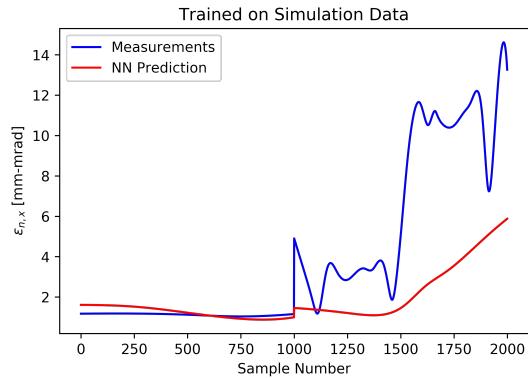
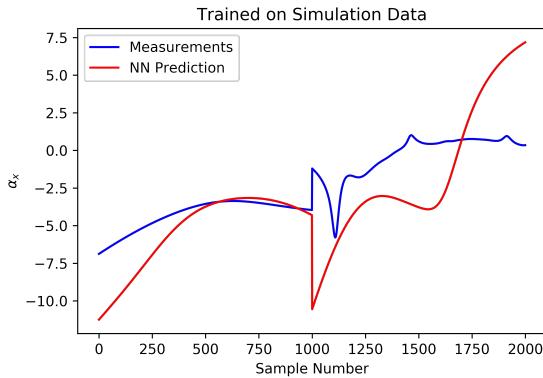
NN Architecture and performance

- Data separated into Training, Validation, and Test sets
 - *Training set: used directly in training*
 - *Validation set: not used in weight updates, but used during training to assess model fit (e.g. avoid overfitting)*
 - *Test set: data not used at all during training, sometimes outside range of training data*
- Noise added to the data before training
- Performance across validation and test set
 - *Top: prediction and simulation as a function of gun phase*
 - *Bottom: rms percent error between neural network and simulations*
- All output parameters perform well except transmission
 - *All transmission is 100% in our range of simulations so this is dominated by noise added during training*



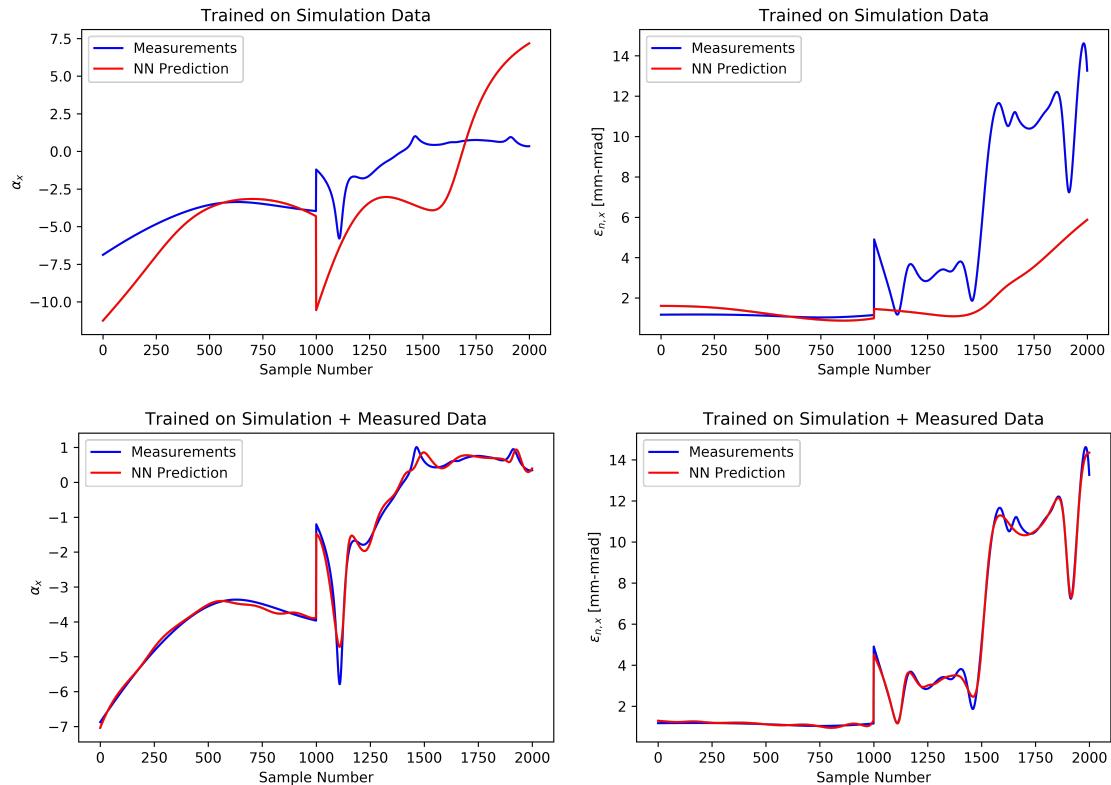
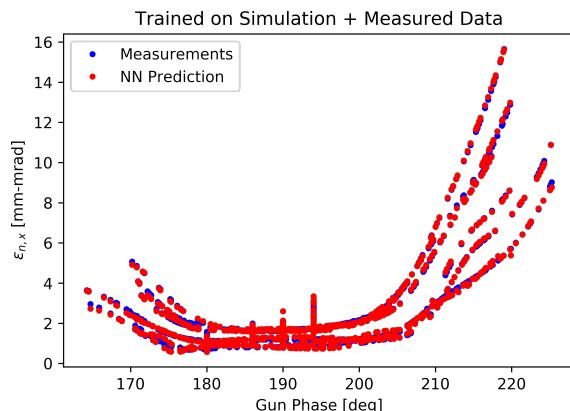
Updating the network with measurements

- Compare network trained on simulations to measurements: top
 - Note this model was trained on rms parameters from simulations, not the simulated multi-slit measurement
- Updating with measurements: bottom
 - Bottom Left: Alpha as a function of sample number for updated dataset
 - Bottom Right: Normalized emittance as a function of sample number for updated dataset



Updating the network with measurements

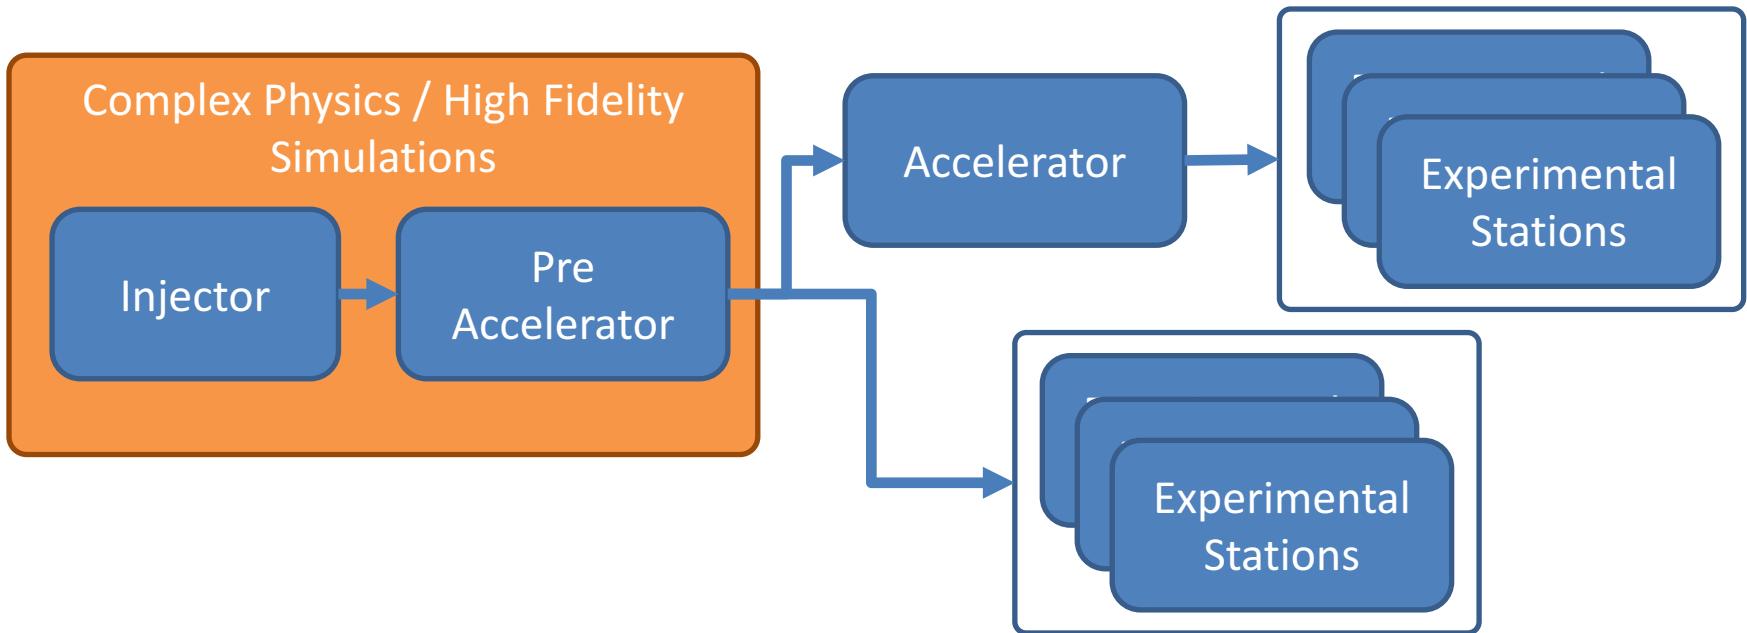
- Compare network trained on simulations to measurements: top
 - Note this model was trained on rms parameters from simulations, not the simulated multi-slit measurement
- Updating with measurements
 - Bottom Left: Alpha as a function of sample number for updated dataset
 - Bottom Right: Normalized emittance as a function of sample number for updated dataset



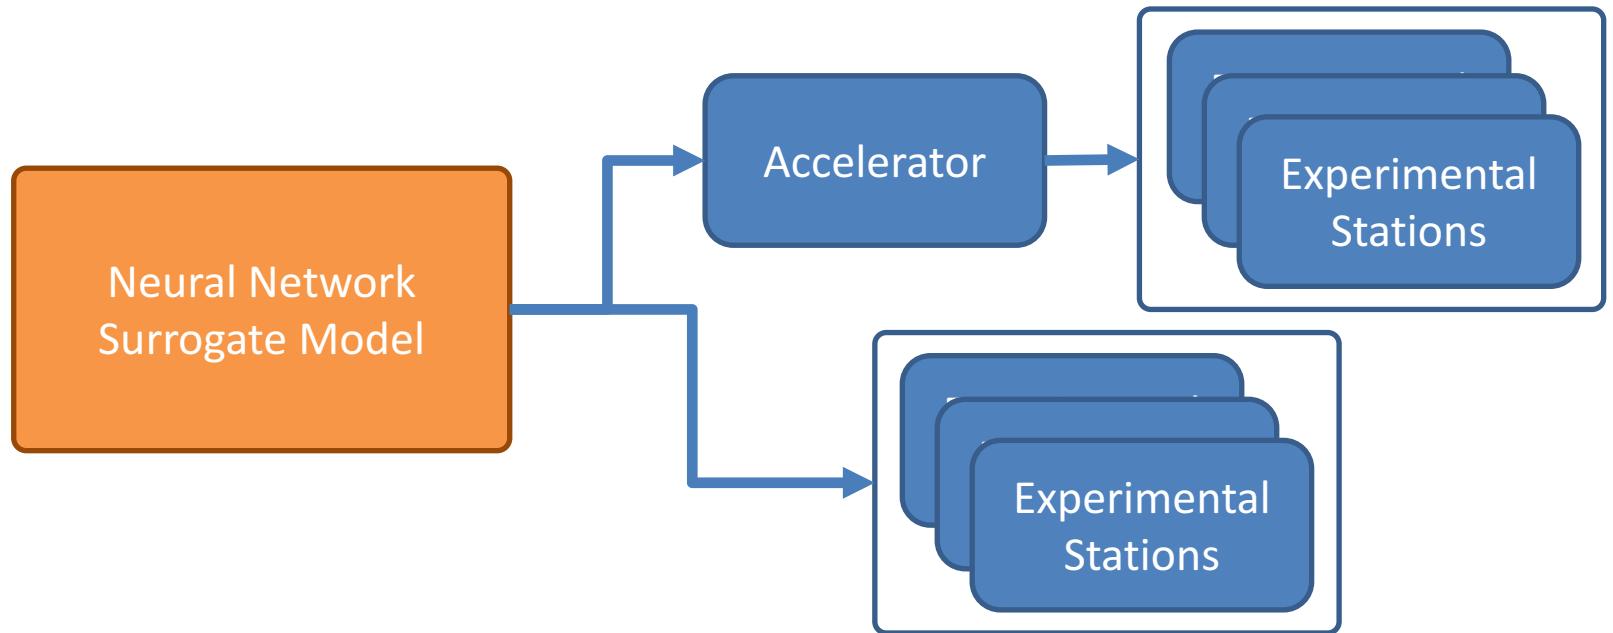
- Network retains the information from the simulations
 - Left: comparison of network prediction for phase scan data from before and after updating with measurements

Application: Surrogate Modeling and Optimization

Surrogate Modeling

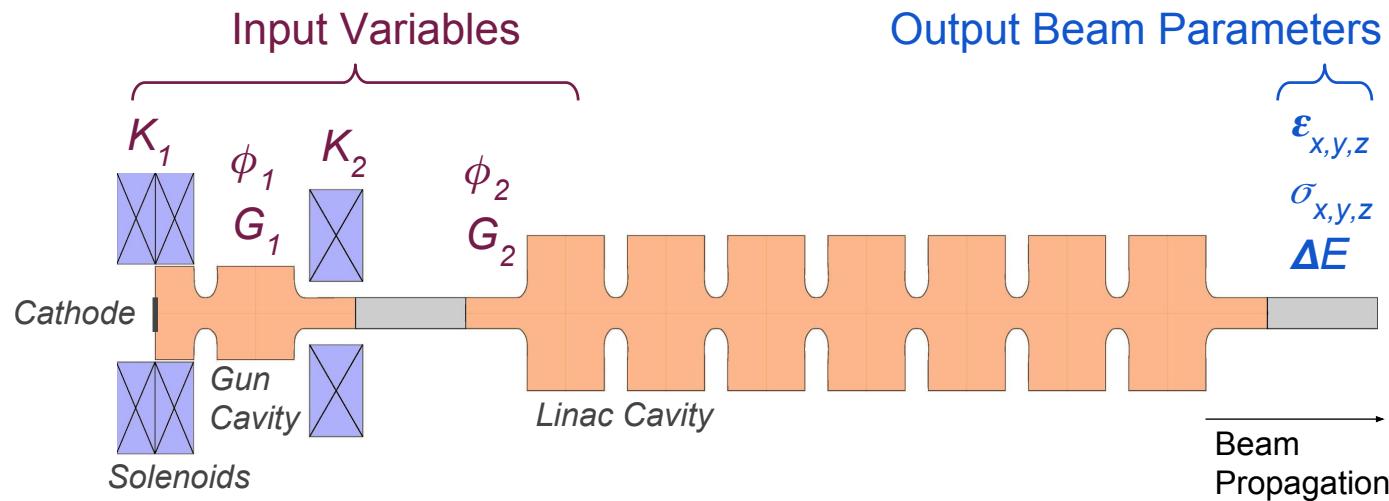
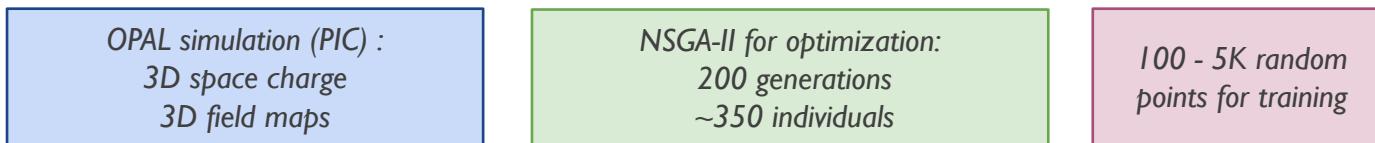


Surrogate Modeling



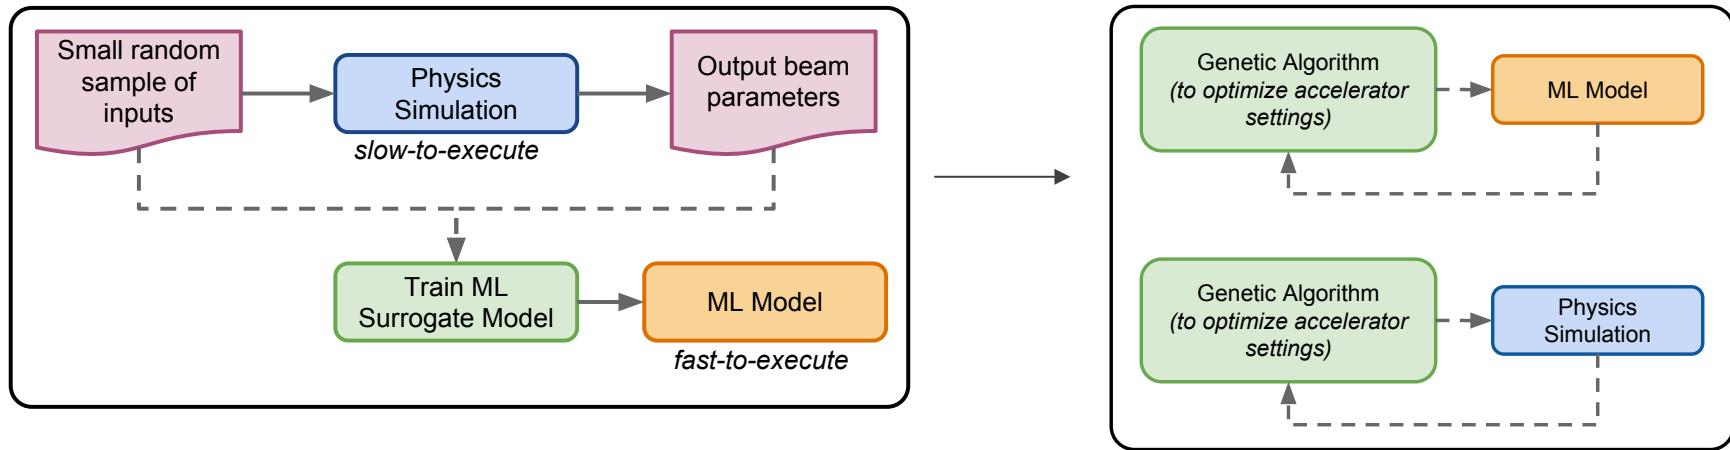
Injector Surrogate Modeling and Optimization

Test Case: Argonne Wakefield Accelerator Injector



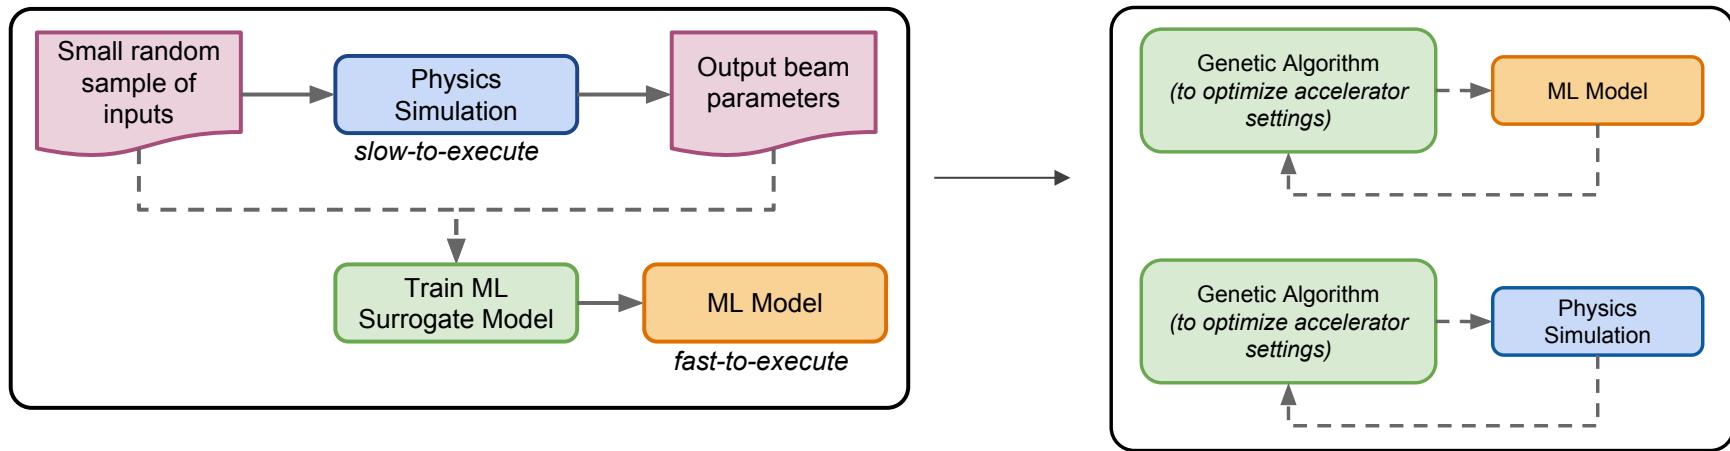
Injector Surrogate Modeling and Optimization

Generate ML Model using Sparse Random Sample

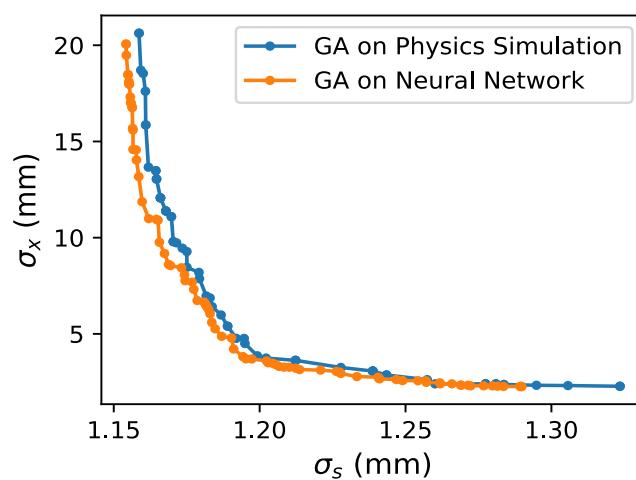


Injector Surrogate Modeling and Optimization

Generate ML Model using Sparse Random Sample



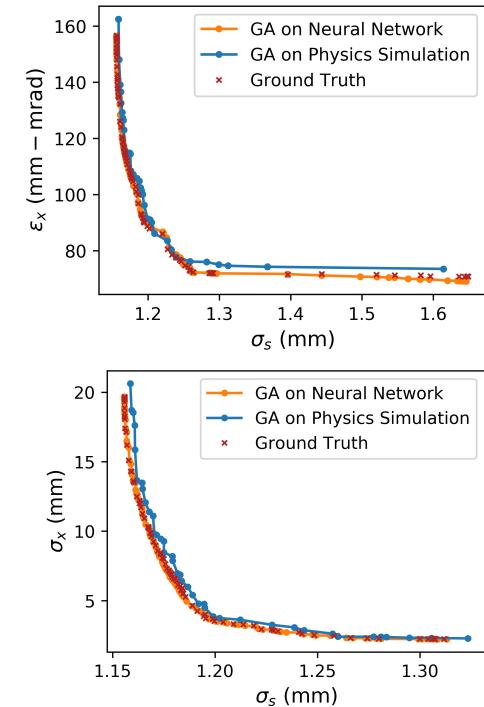
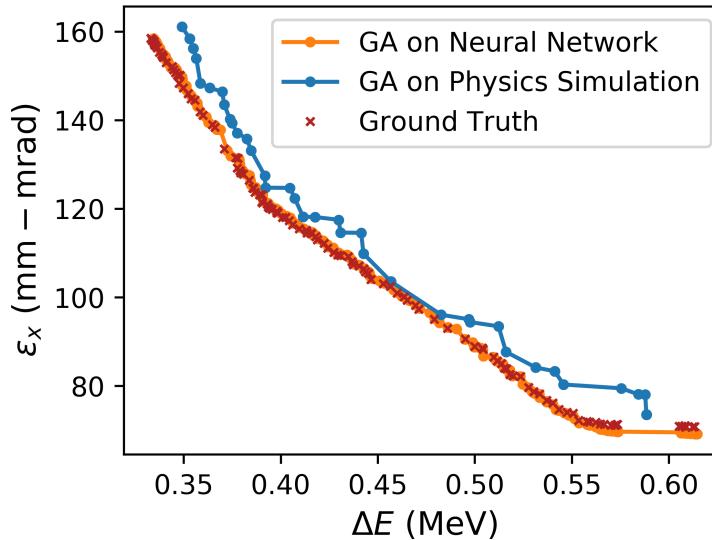
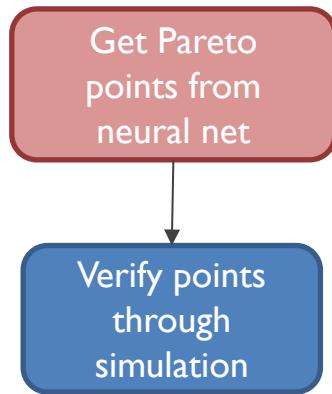
Compare Resulting Pareto Fronts



Verifying pareto front from the neural network

In some cases, optimization over simulation takes too long to converge

→ validate Pareto front from neural network more directly



Neural network gave a **better approximation of the Pareto front** than the naïve GA on the simulation

Required **13x fewer simulations** and had **10^6 times faster execution** in the optimization

→ In addition to creating a surrogate model, this is a way to speed up optimization for new designs

Accelerator workshops on ML



2nd ICFA Workshop on Machine Learning for Charged Particle Accelerators

26 February 2018 to 1 March 2018
Villigen PSI
Europe/Zurich timezone

<https://indico.psi.ch/event/6698/>

The ICPA website header features a dark background with a large, semi-transparent network graph. At the top right, there are navigation links: HOME, REGISTRATION, PROGRAMME, COMMITTEES, and CONTACT. The title "ICPA" is in yellow. Below the title, the text "INTELLIGENT CONTROLS FOR PARTICLE ACCELERATORS" is in white. The main event information is displayed in large white text: "30-31 JAN 2018" and "DARESBURY LABORATORY". A yellow button labeled "TELL ME MORE" is located to the right of the event details.

<https://www.cockcroft.ac.uk/events/ICPA/>

A screenshot of the SLAC workshop page. The title "Machine Learning Applications for Particle Accelerators" is in white on a blue header. Below it, text indicates the workshop is chaired by Christopher Mayes (SLAC National Accelerator Laboratory), Daniel Bowring (Fermilab), Daniel Ratner (SLAC), and Auralee Edelen (Fermilab, CSU). The event runs from Tuesday, February 27, 2018 at 07:30 to Friday, March 2, 2018 at 18:00 (US/Pacific) at SLAC Building 53 (Trinity). The address is 2575 Sand Hill Rd, Menlo Park, CA 94025. A "Description" section states the goal is to build a world-wide community of researchers interested in applying machine learning techniques to particle accelerators, mentioning four sequential topics: tuning, optimization, control; prognostics, alarm handling, anomaly-breakout detection; data analysis; and simulations, modeling. A note below says the workshop was broadcast via Zoom.

<https://indico.fnal.gov/event/16327/other-view?view=standard>

Publications

J. P. Edelen A. L. Edelen, and D. Edstrom "Methods for Data Cleaning" 2nd ICFA Mini Workshhop on Machine Learning for Particle Accelerators (Feb 2019)

J. P. Edelen, A. L. Edelen, and D. Edstrom "Neural Network Based Virtual Diagnostics at FAST" Contributed Oral at the Fermilab Workshop on Megawatt Rings (May 2018)

A.L. Edelen et al. "Results and Discussion of Recent Applications of Neural Network-Based Approaches to the Modeling and Control of Particle Accelerators" Proc. IPAC 2018 (THYGBE2)

A.L. Edelen et al " Neural Network Virtual Diagnostic and Tuning for the FAST Low Energy Beamline" IPAC 2018 (SUSPL054)

J.P. Edelen, A.L. Edelen & D. Edstrom, "Neural network modeling and virtual diagnostics at FAST," presented at ICFA Beam Dynamics Mini-Workshop: Machine Learning Applications for Particle Accelerators (SLAC, 2018).

A. Edelen et al. "Opportunities in Machine Learning for Particle Accelerators". In: (2018) arXiv: 1811.03172 [physics.acc-ph].

A. L. Edelen et al. "Using Neural Network Control Policies For Rapid Switching Between Beam Parameters in a Free Electron Laser". In: Proceedings of the 2017 Deep Learning for Physical Sciences workshop at the Neural Information Processing Systems Conference. 2017.

A. L. Edelen et al. "Neural Networks for Modeling and Control of Particle Accelerators" In: IEEE Transactions on Nuclear Science 63.2 (Apr. 2016), pp. 878–897. issn: 0018-9499.

A.L. Edelen, S.G. Biedron, S.V. Milton & J.P. Edelen, , "First steps towards incorporating image based diagnostics into particle accelerator control systems using convolutional neural networks," Proc. North American Part. Accel. Conf., TUPOA51 (2016)