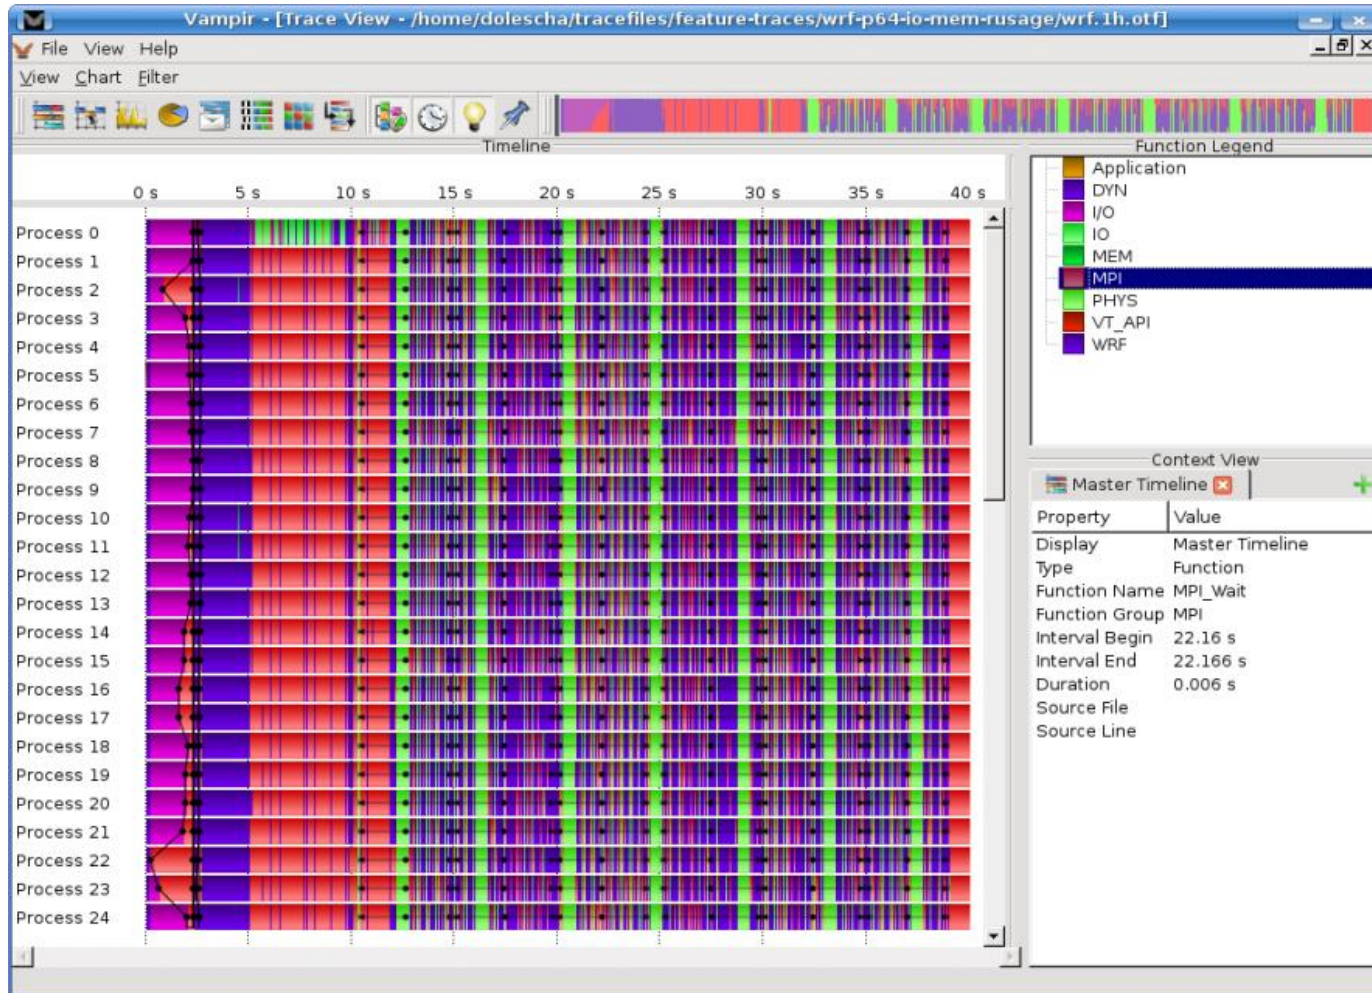


The Scalasca Performance Analysis Toolset

Markus Geimer
Jülich Supercomputing Centre
m.geimer@fz-juelich.de
April 3, 2013

“A picture is worth 1000 words...”

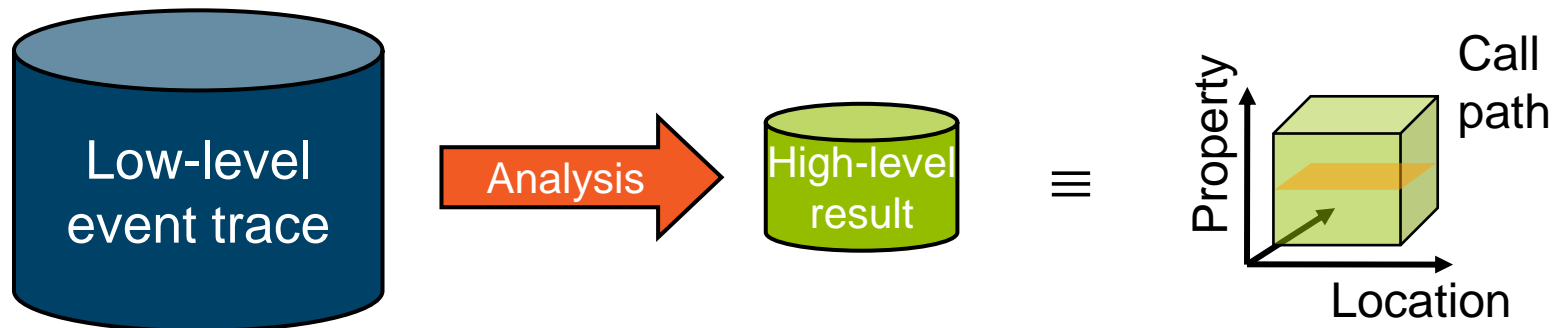




Automatic trace analysis

- Idea

- Automatic search for patterns of inefficient behavior
- Classification of behavior & quantification of significance



- Advantages

- Guaranteed to cover the entire event trace
- Quicker than manual/visual trace analysis
- Helps to identify hot-spots for in-depth manual analysis
 - Complements the functionality of other tools



Higher degrees of parallelism

Number of Cores share for TOP 500 November 2012

NCore	Count	Share	Σ Rmax	Share	Σ NCore
1025-2048	1	0.2%	122 TF	0.1%	1,280
2049-4096	2	0.4%	155 TF	0.1%	7,104
4097-8192	81	16.2%	8,579 TF	5.3%	551,624
8193-16384	206	41.2%	24,543 TF	15.1%	2,617,986
> 16384	210	42.0%	128,574 TF	79.4%	11,707,806
Total	500	100%	161,973 TF	100%	14,885,800

Average system size: **29,772 cores**

Median system size: **15,360 cores**

Higher degrees of parallelism (II)

- Also new demands on **scalability** of software tools
 - Familiar tools cease to work in a satisfactory manner for large processor/core counts
- Optimization of applications more difficult
 - Increasing machine complexity
 - Every doubling of scale reveals a new bottleneck
- Need for scalable performance tools
 - Efficient to meet performance expectations
 - Effective to use so that programmer productivity is maximized



The Scalasca project

- Project started in 2006
 - Follow-up to pioneering KOJAK project (started 1998)
 - Automatic pattern-based trace analysis
 - Initial funding by Helmholtz Initiative & Networking Fund
 - Many follow-up projects
- Objective:
 - Development of a **scalable** performance analysis toolset
 - Specifically targeting **large-scale** parallel applications
 - such as those running on IBM Blue Gene or Cray XT with 10,000s or 100,000s of processes
- Now joint project of
 - Jülich Supercomputing Centre
 - German Research School for Simulation Sciences



The Scalasca toolset

Scalasca 1.4.3

- Custom instrumentation & measurement system
- Scalasca trace analysis components based on custom trace format EPILOG
- Analysis report explorer & algebra utilities CUBE v3
- New BSD license

Scalasca 2.0 β

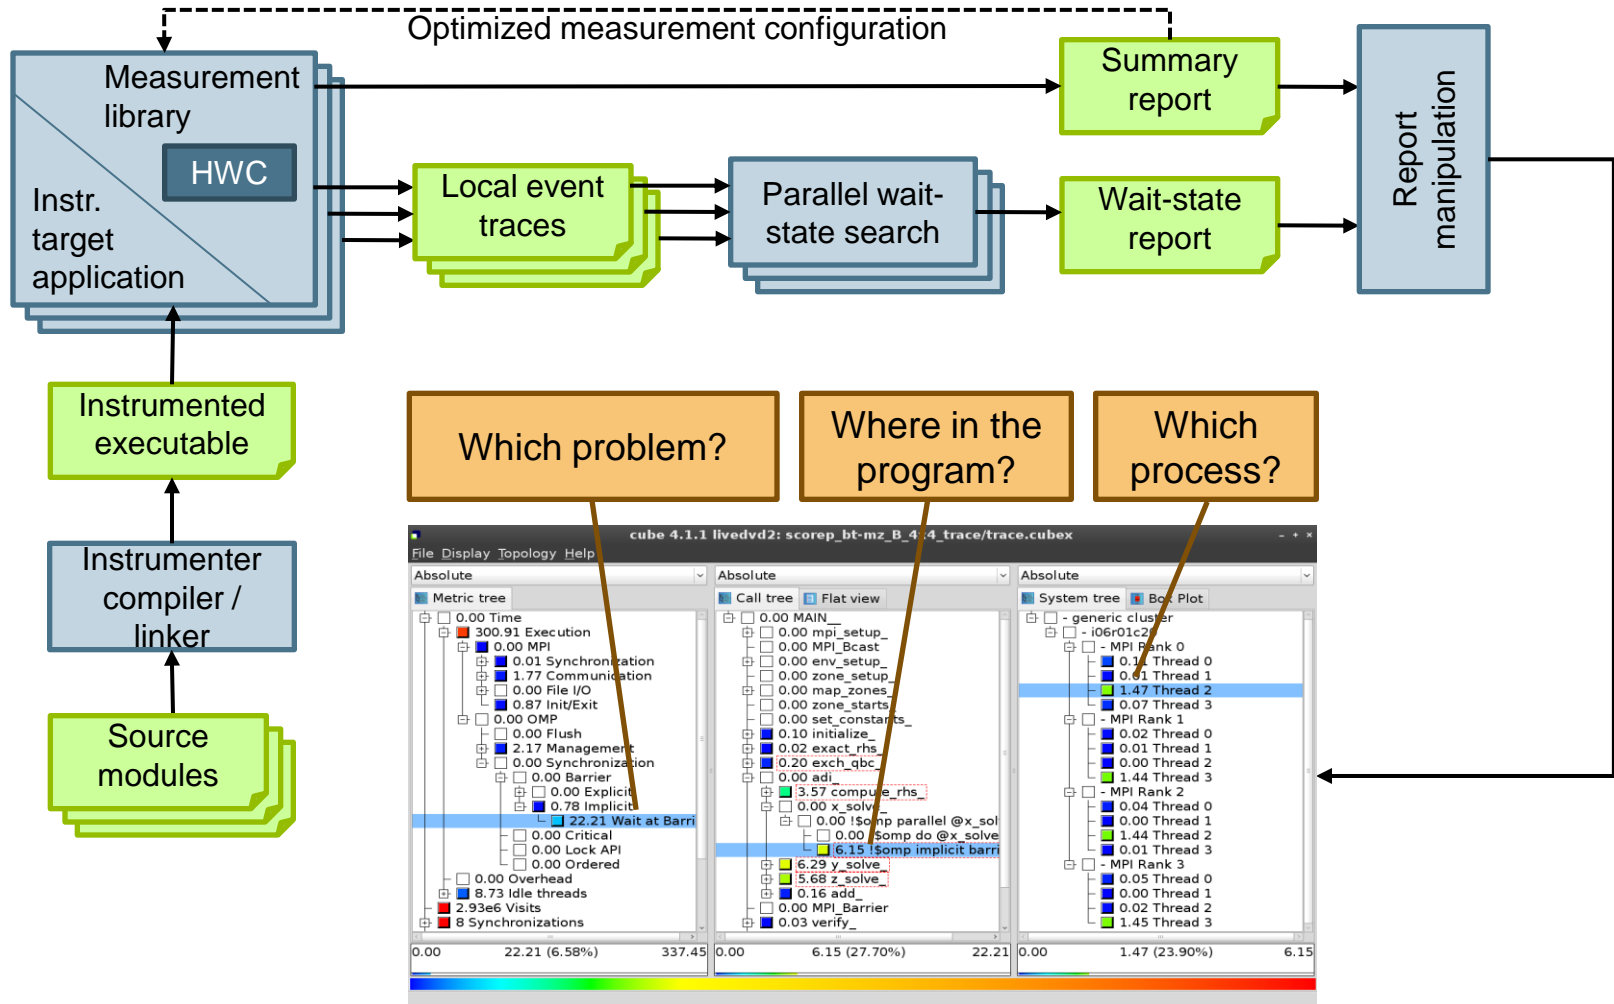
- Community instrumentation & measurement system **Score-P**
- Scalasca trace analysis components based on community trace format **OTF2**
- Analysis report explorer & algebra utilities **CUBE v4**
- New BSD license

<http://www.scalasca.org>


Scalasca features

- Open source
- Portable
 - Blue Gene/Q, Blue Gene/P, IBM SP & blade clusters, SGI Altix, Solaris & Linux clusters
 - Scalasca 1.4.3 only: Cray XT, NEC SX, K Computer, Fujitsu FX10
- Supports parallel programming paradigms & languages
 - MPI, OpenMP & hybrid MPI+OpenMP
 - Fortran, C, C++
- Scalable trace analysis
 - Automatic wait-state search
 - Parallel replay exploits memory & processors to deliver scalability


Scalasca workflow




Example: MPI patterns




(a) Late Sender



(b) Late Receiver




(c) Late Sender / Wrong Order



(d) Wait at N x N


ENTER
 LEAVE
 SEND
 RECV
 COLLECTIVE

Example: Late Sender



Sender:

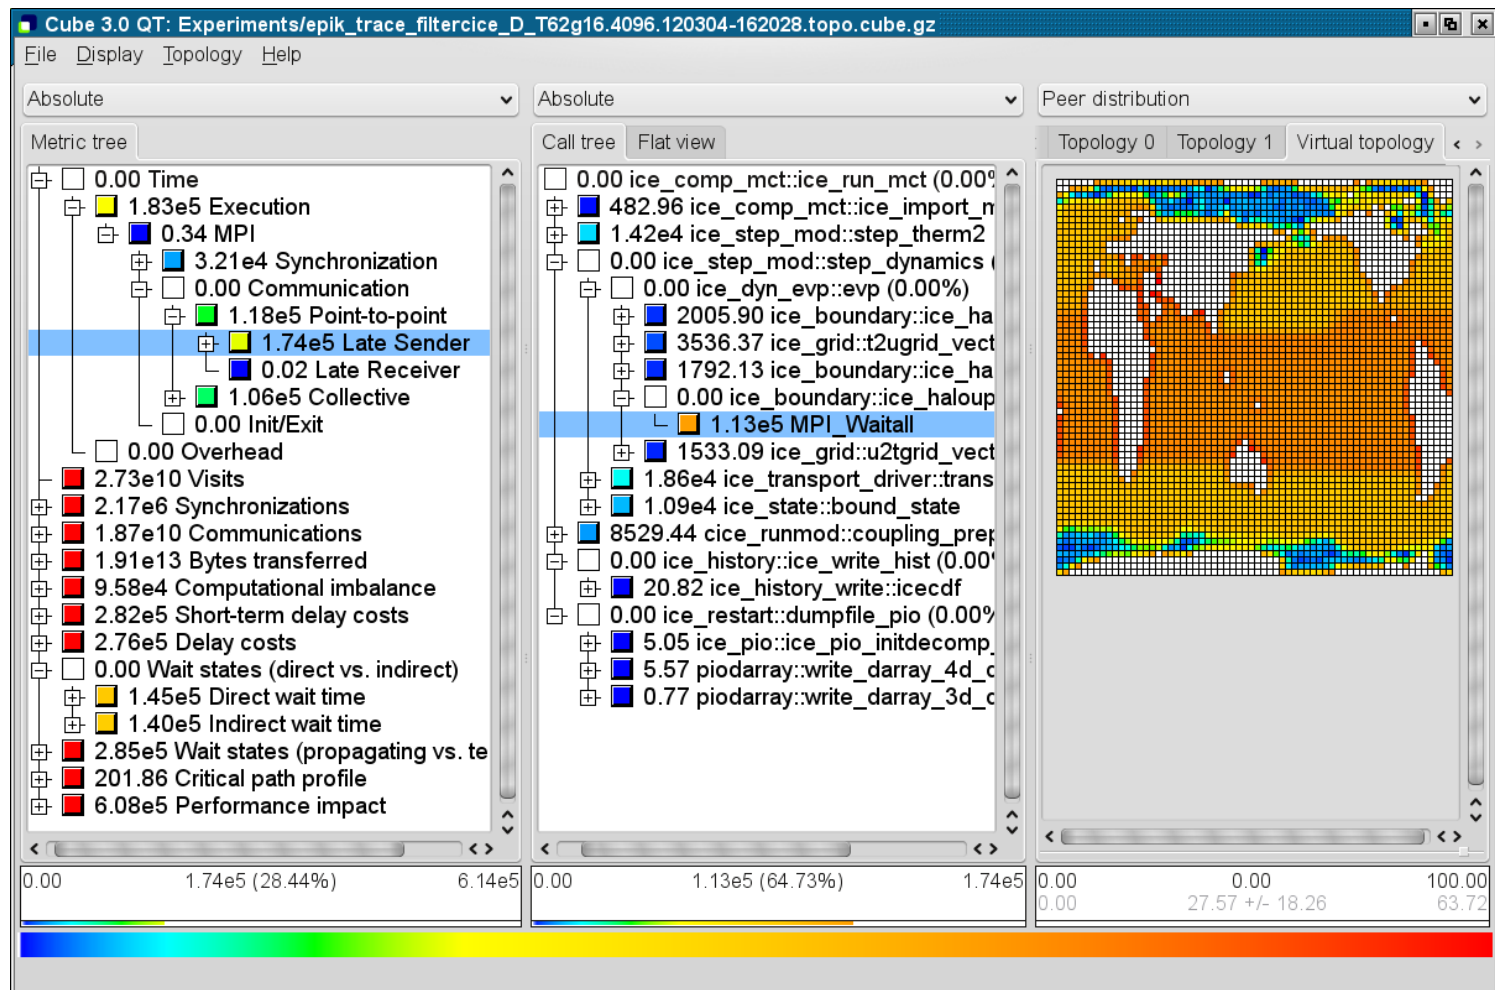
Triggered by send event
 Determine enter event
 Send both events to receiver



Receiver:

Triggered by receive event
 Determine enter event
 Receive remote events
 Detect *Late Sender* situation
 Calculate & store waiting time

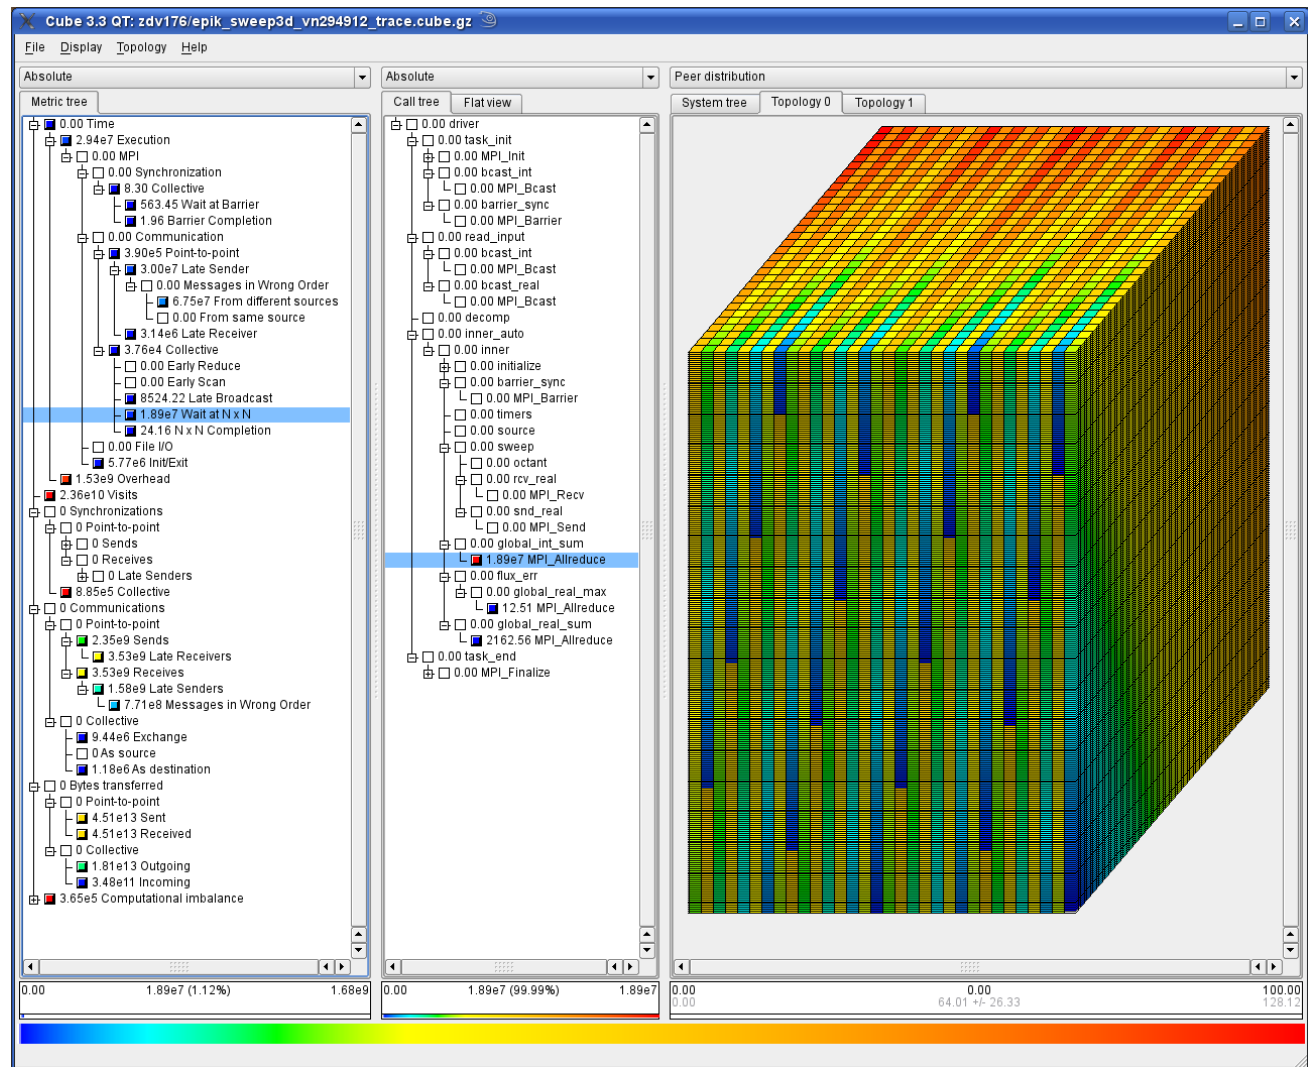
CESM Sea Ice Module – Late Sender Analysis



Scalasca trace analysis sweep3D@294,912 BG/P

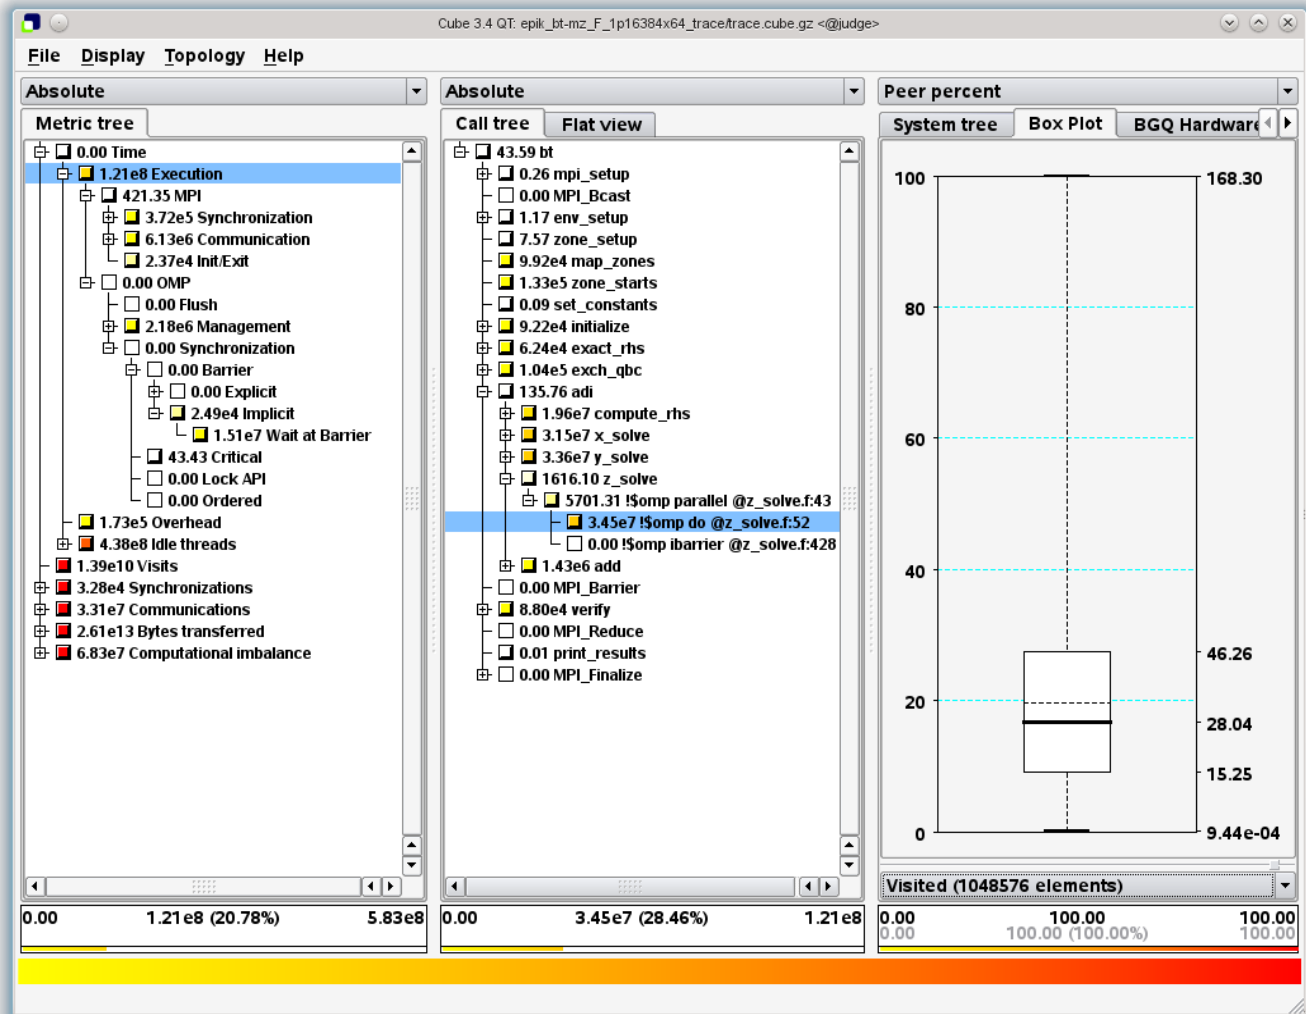
- 10 min sweep3D runtime
- 11 sec analysis
- 4 min trace data write/read (576 files)
- 7.6 TB buffered trace data
- 510 billion events

B. J. N. Wylie, M. Geimer,
B. Mohr, D. Böhme,
Z. Szabenyi, F. Wolf:
Large-scale performance
analysis of Sweep3D with
the Scalasca toolset.
Parallel Processing Letters,
20(4):397-414, 2010.



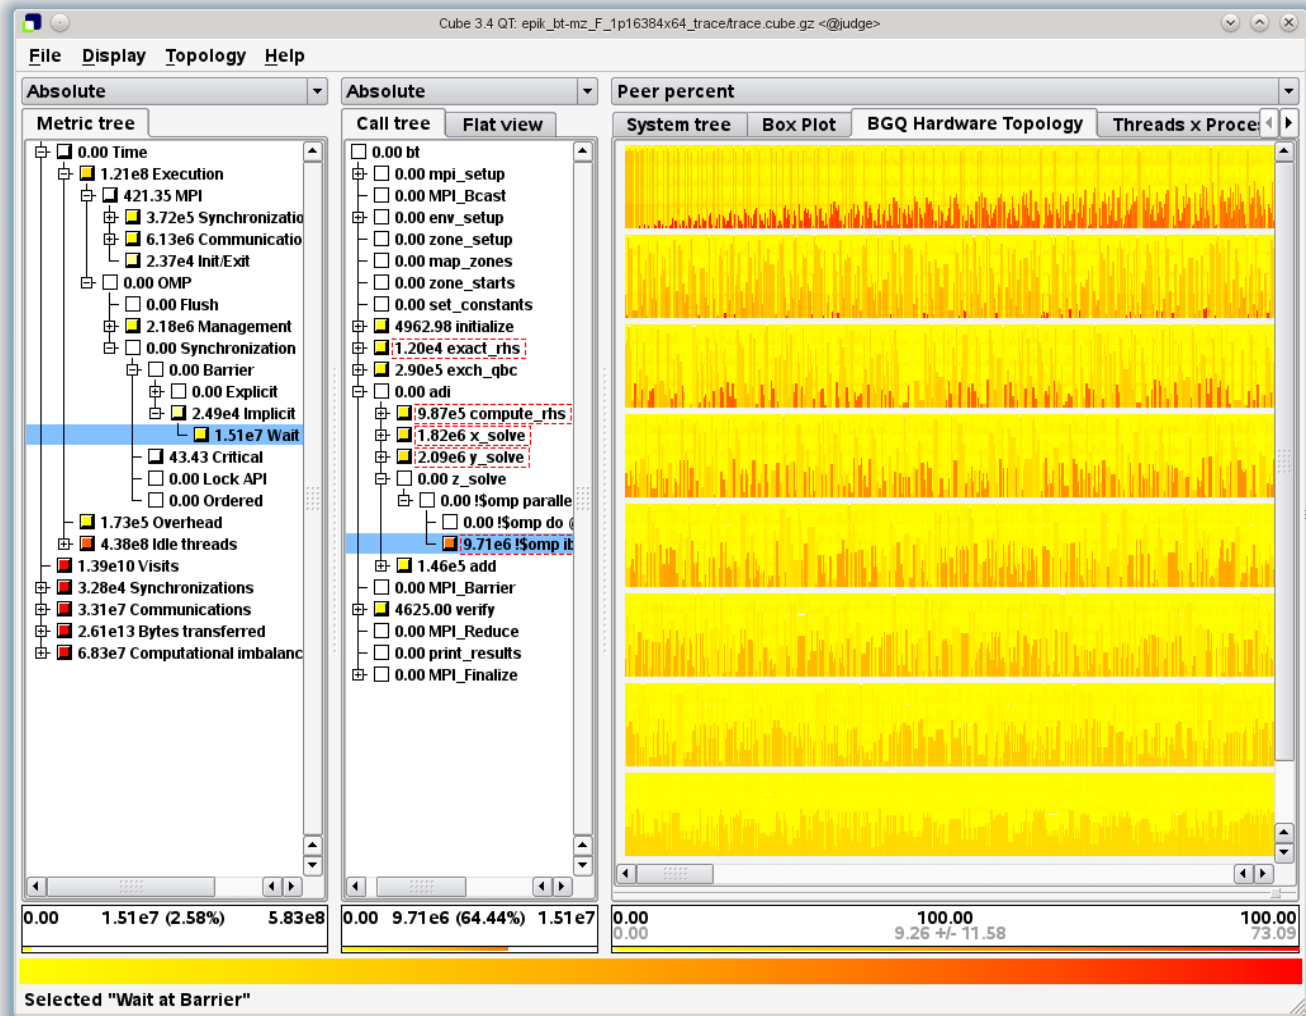
Scalasca trace analysis bt-mz@1,048,704 BG/Q

Execution
imbalance
“z_solve”



Scalasca trace analysis bt-mz@1,048,704 BG/Q

Wait at
implicit
barrier
“z_solve”



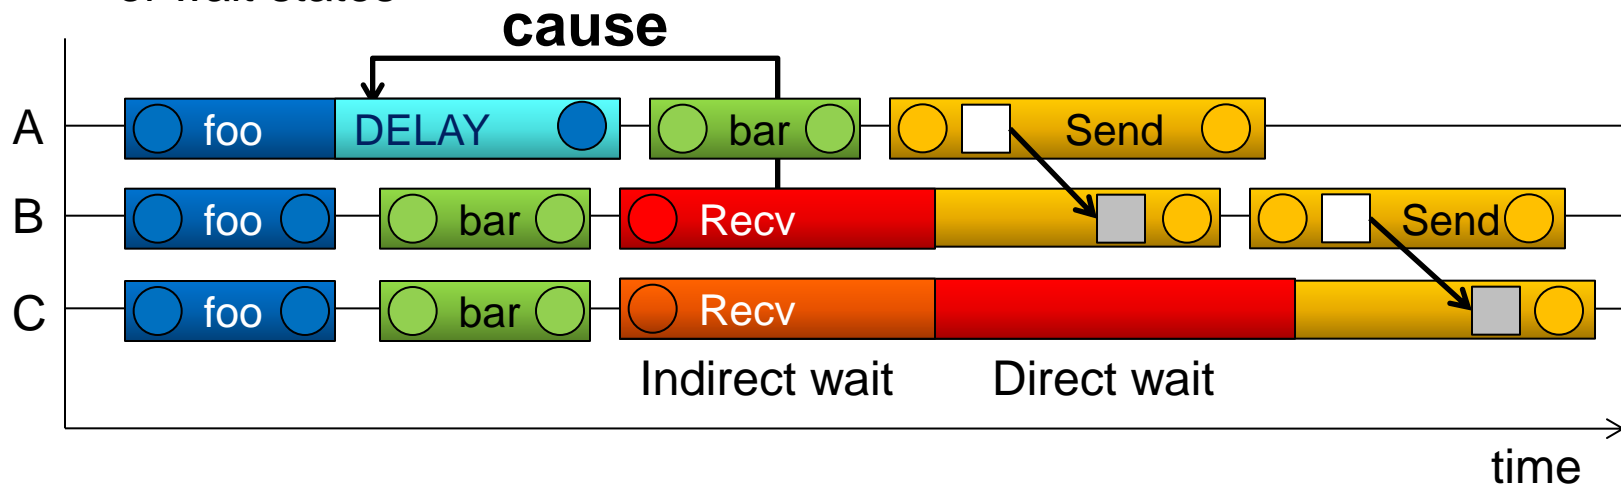
Research: Root Cause Analysis

Root-cause analysis

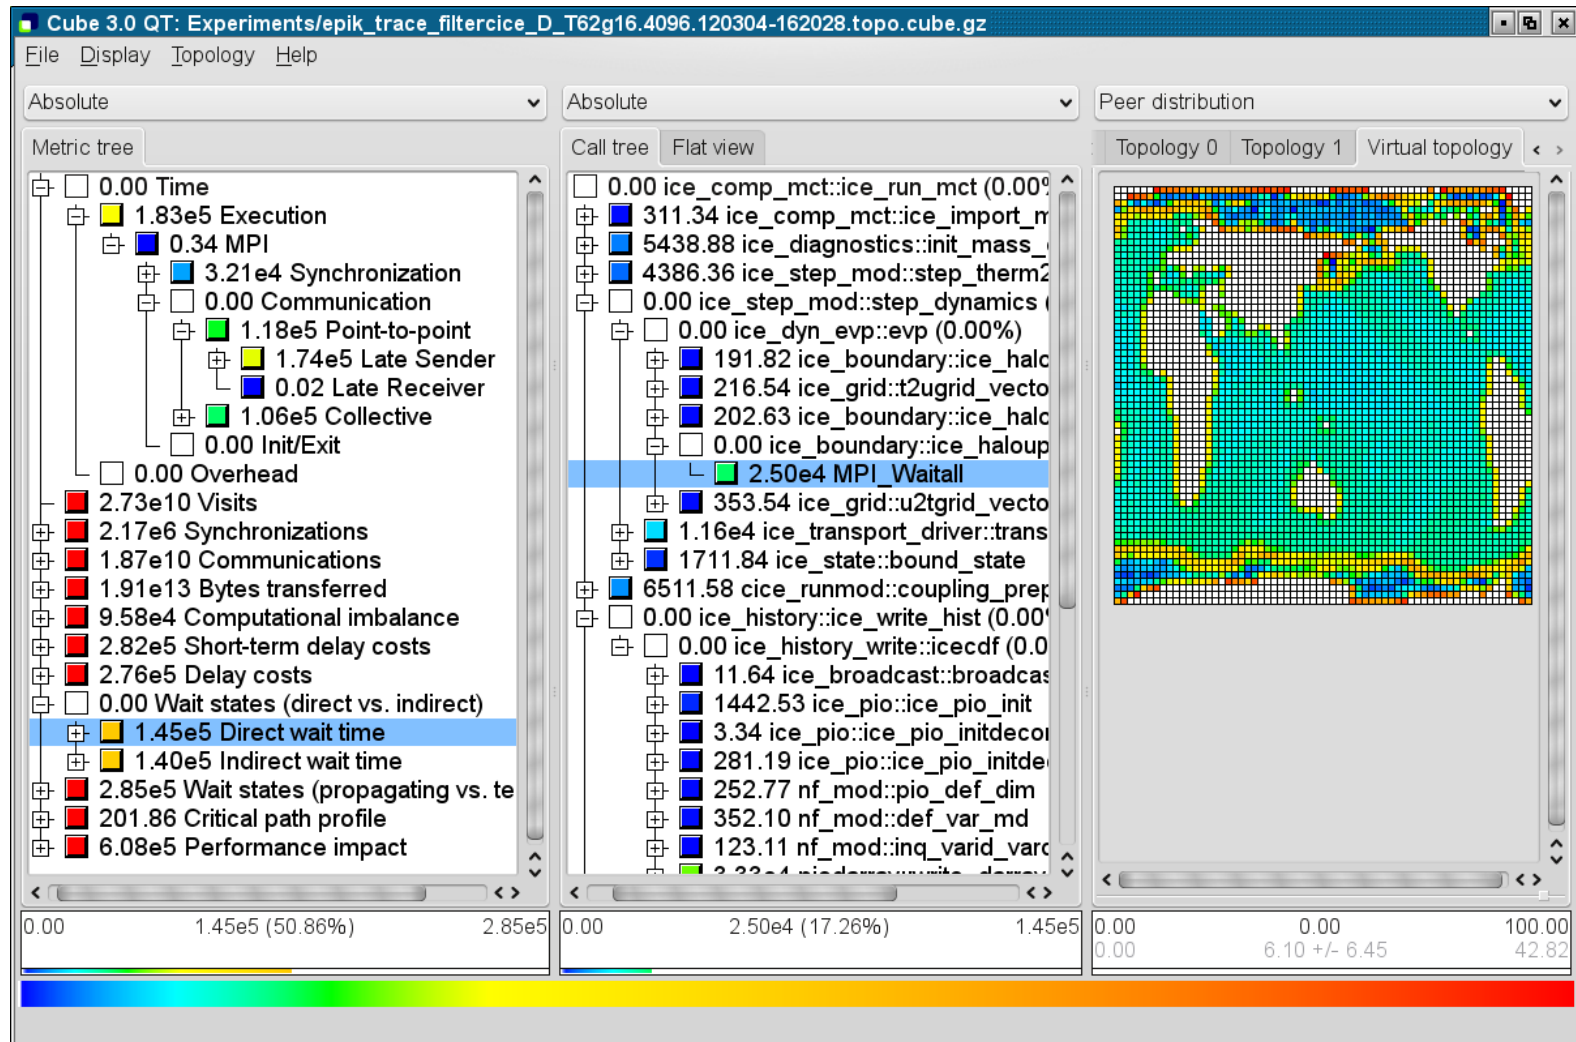
- Wait states typically caused by load or communication imbalances earlier in the program
- Waiting time can also propagate (e.g., indirect waiting time)
- Goal: Enhance performance analysis to find the root cause of wait states

Approach

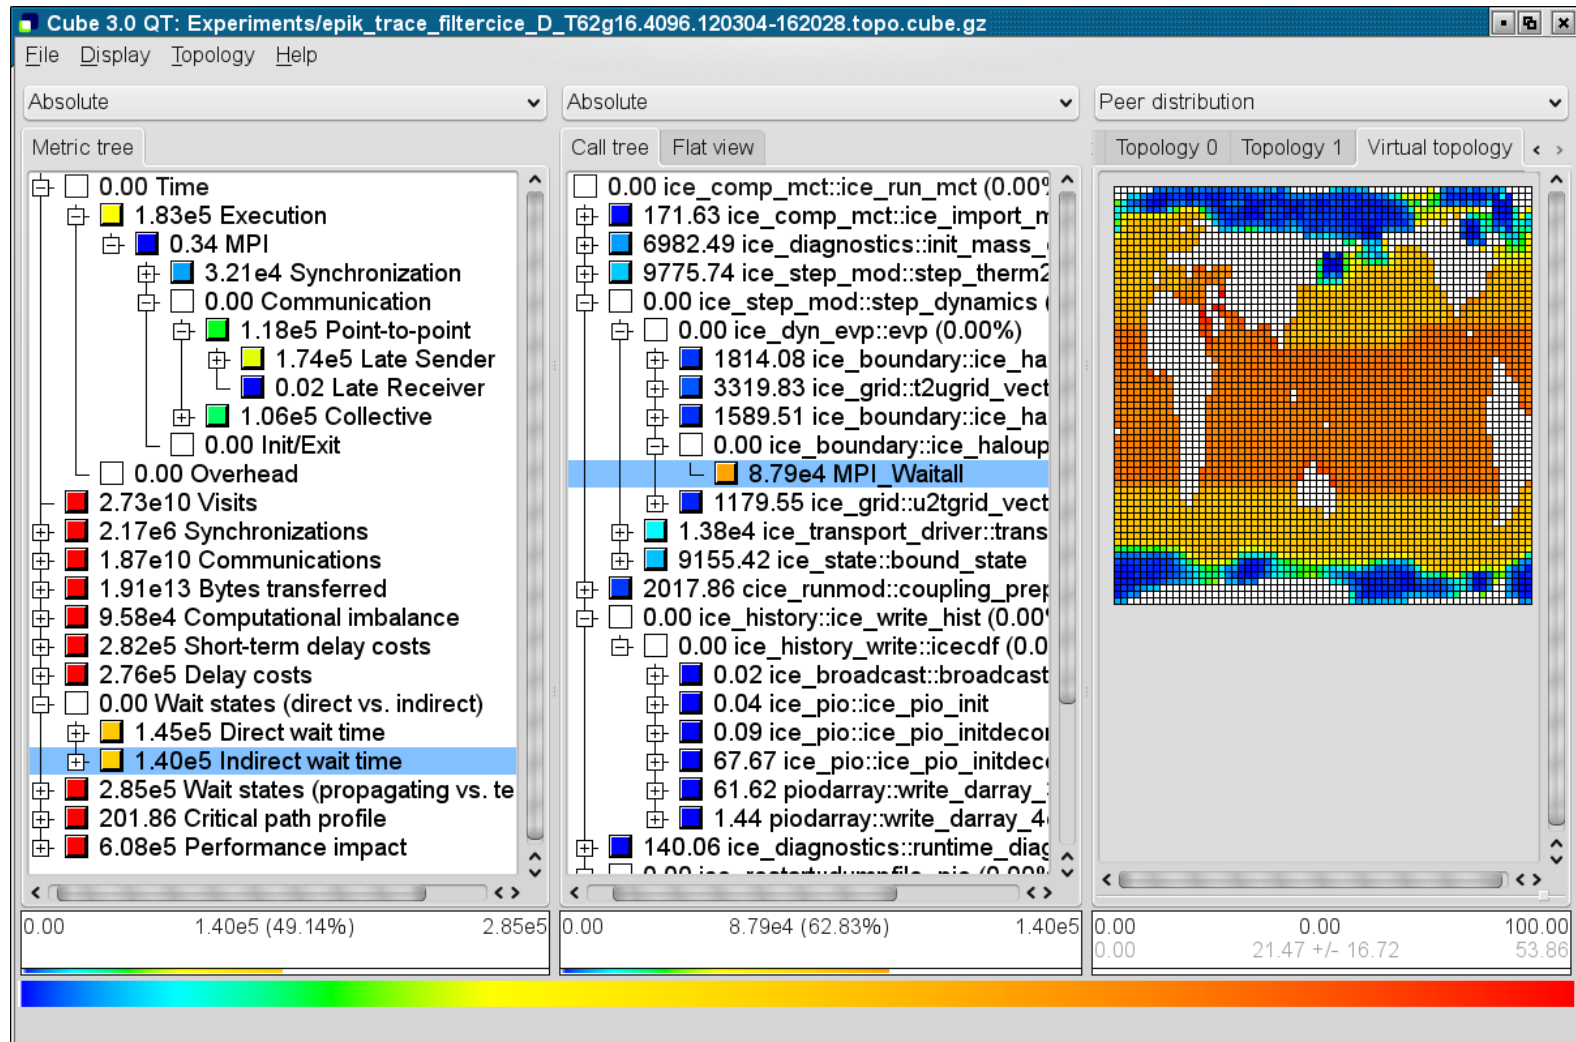
- Distinguish between direct and indirect waiting time
- Identify call path/process combinations delaying other processes and causing first order waiting time
- Identify original **delay**



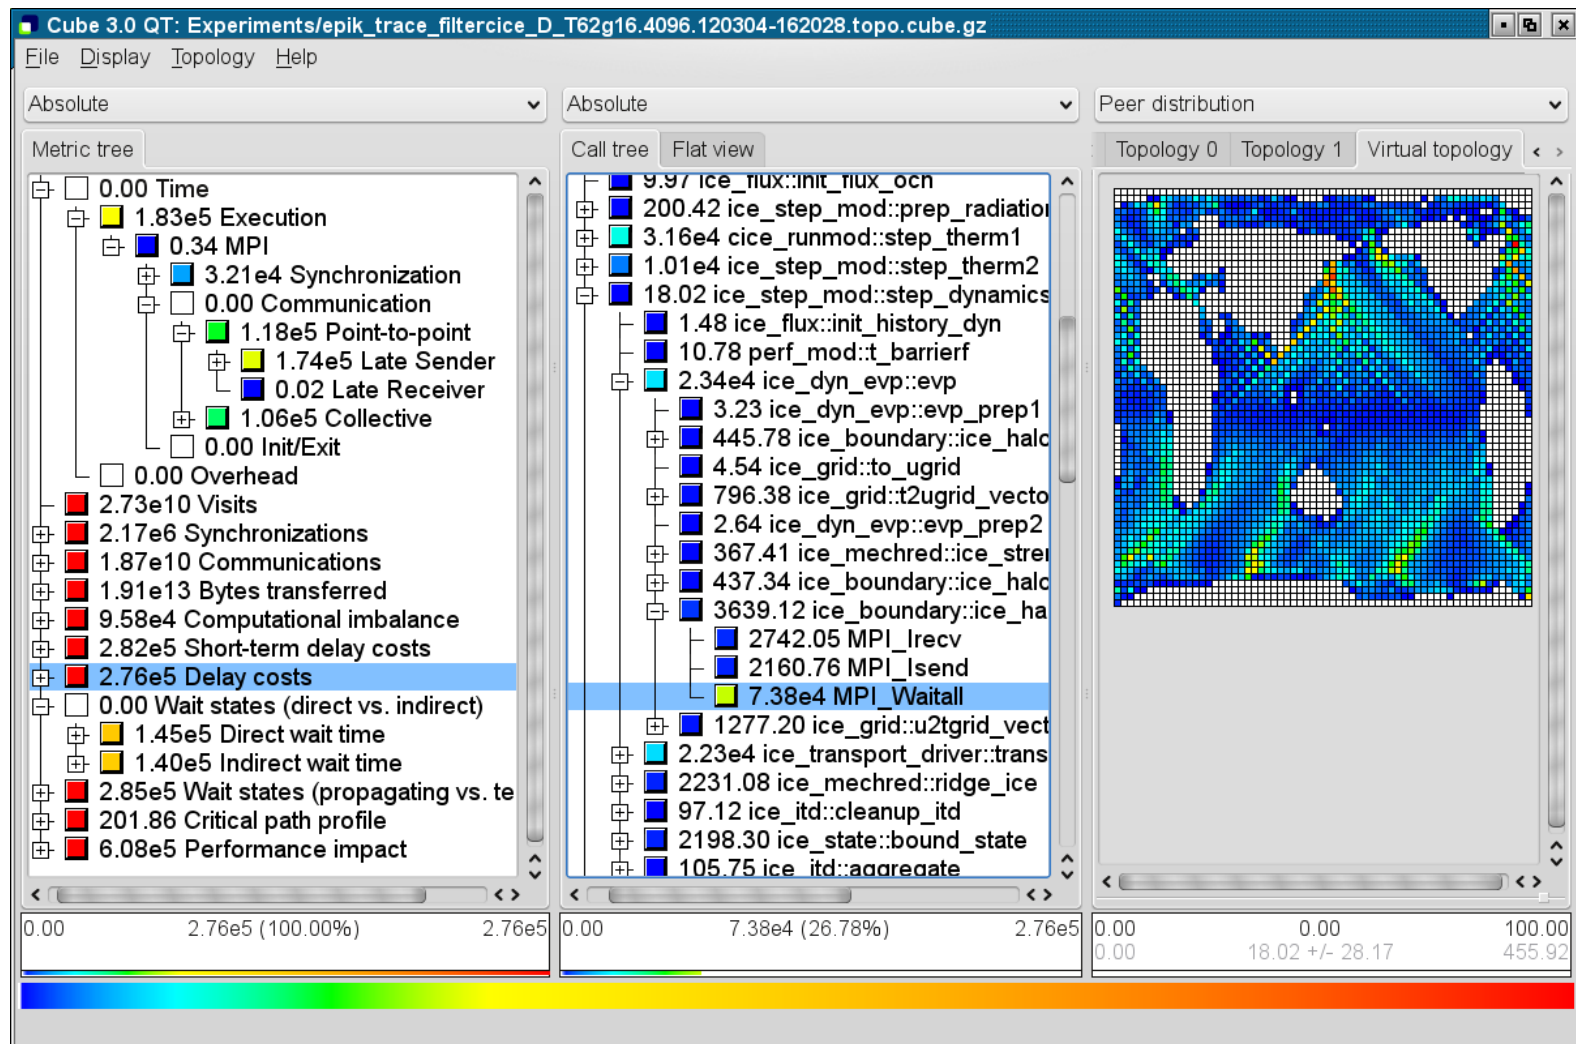
CESM Sea Ice Module – Direct Wait Time



CESM Sea Ice Module – Indirect Wait Time



CESM Sea Ice Module – Delay Costs



Acknowledgements

Scalasca team (JSC)



(GRS)



Sponsors



Thank you!

<http://www.scalasca.org>