
KGEN:

Fortran Kernel Generator

Youngsung Kim, John Dennis, Raghu R. Kumar, and Amogh Simha

National Center for Atmospheric Research (NCAR)

Contents

- Introduction
- A kernel generation example
- Kernel generation from large-scale app.
- MG2 kernel(CESM) demo.
- Use-cases
- Development status and plans

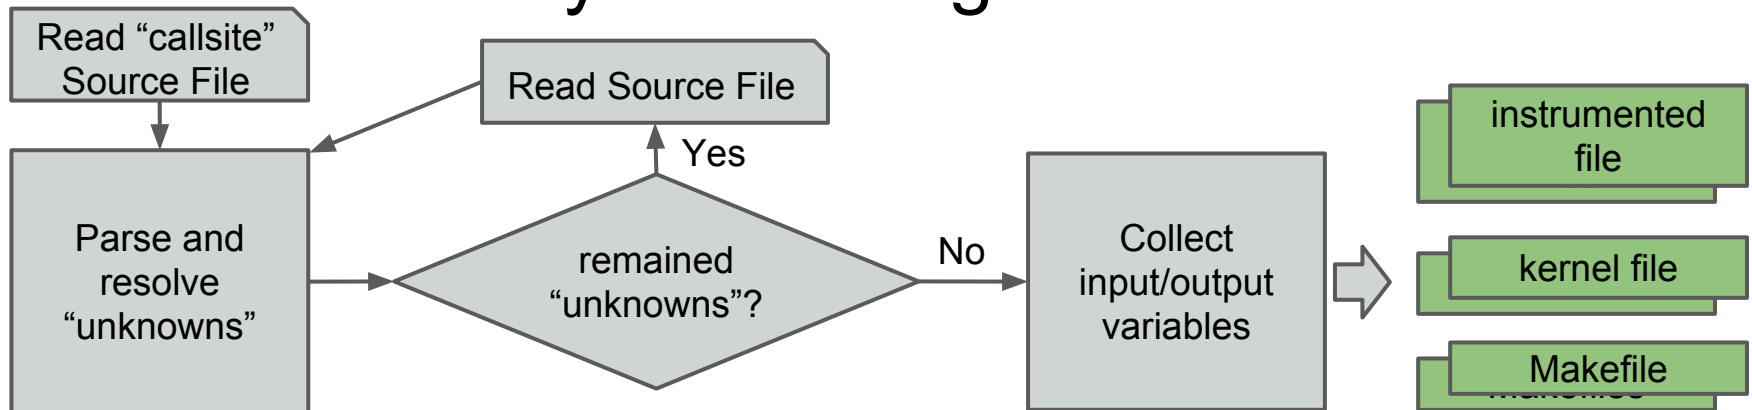
KGEM: Key Features

- KGEM extracts a Fortran subprogram as a stand-alone software out of a large software application such as CESM*
- In addition, it generates instrumented files that save input & output data for the generated kernel
- Correctness check and timing measurement are included in the generated kernel

KGEN: Implementation Overview

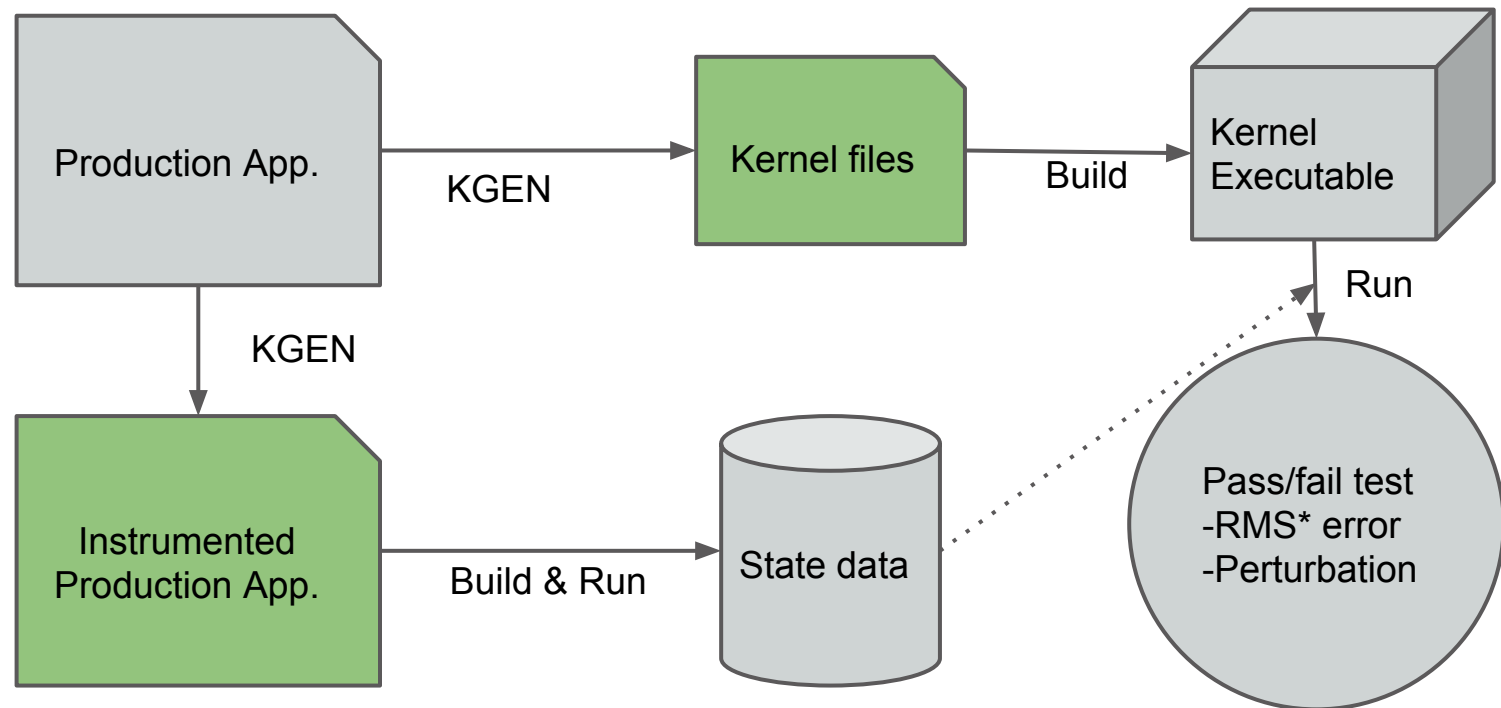
- KGEN is:
 - written in Python 2.6.6 as an extension of F2PY* Fortran parsers. No need for other external modules
 - capable of resolving “unknowns” by searching Abstract Syntax Tree generated by F2PY* parsers

- KGEN Activity Block diagram



KGEN: Workflow

- Two parts: kernel generation and state data generation



A kernel generation example

./program.F90

```
PROGRAM demo
  USE update_mod, &
    only : update
  INTEGER t
  DO t=1,10
    CALL update
  END DO
END PROGRAM
```

./update_mod.F90

```
MODULE update_mod
  USE calc_mod, only : calc
  PUBLIC update
  CONTAINS
    SUBROUTINE update()
      INTEGER :: i, j
      INTEGER :: output(4,4)
      DO i=1,4
        DO j=1,4

          CALL calc(i, j, output)

        END DO
      END DO
    END SUBROUTINE
END MODULE
```

./calc_mod.F90

```
MODULE calc_mod
  PUBLIC calc
  CONTAINS
    SUBROUTINE calc(i, j, output)
      INTEGER, INTENT(IN) :: i, j
      INTEGER, INTENT(OUT), &
        dimension(:,:) :: output
      output(i,j) = i + j
    END SUBROUTINE
END MODULE
```

A kernel generation example - cont.

./program.F90

```
PROGRAM demo
  USE update_mod, &
    only : update
  INTEGER t
  DO t=1,10
    CALL update
  END DO
END PROGRAM
```

./update_mod.F90

```
MODULE update_mod
  USE calc_mod, only : calc
  PUBLIC update
CONTAINS
  SUBROUTINE update()
    INTEGER :: i, j
    INTEGER :: output(4,4)
    DO i=1,4
      DO j=1,4
        !$kgen callsite calc
        CALL calc(i, j, output)
      END DO
    END DO
  END SUBROUTINE
END MODULE
```

Call-site

./calc_mod.F90

```
MODULE calc_mod
  PUBLIC calc
CONTAINS
  SUBROUTINE calc(i, j, output)
    INTEGER, INTENT(IN) :: i, j
    INTEGER, INTENT(OUT), &
      dimension(:,:) :: output
    output(i,j) = i + j
  END SUBROUTINE
END MODULE
```

Kernel

A kernel generation example - cont.

./program.F90

```
PROGRAM demo
  USE update_mod, &
    only : update
  INTEGER t
  DO t=1,10
    CALL update
  END DO
END PROGRAM
```

./update_mod.F90

```
MODULE update_mod
  USE calc_mod, only : calc
  PUBLIC update
CONTAINS
  SUBROUTINE update()
    INTEGER :: i, j
    INTEGER :: output(4,4)
    DO i=1,4
      DO j=1,4
        !$kgen callsite calc
        CALL calc(i, j, output)
      END DO
    END DO
  END SUBROUTINE
END MODULE
```

Call-site

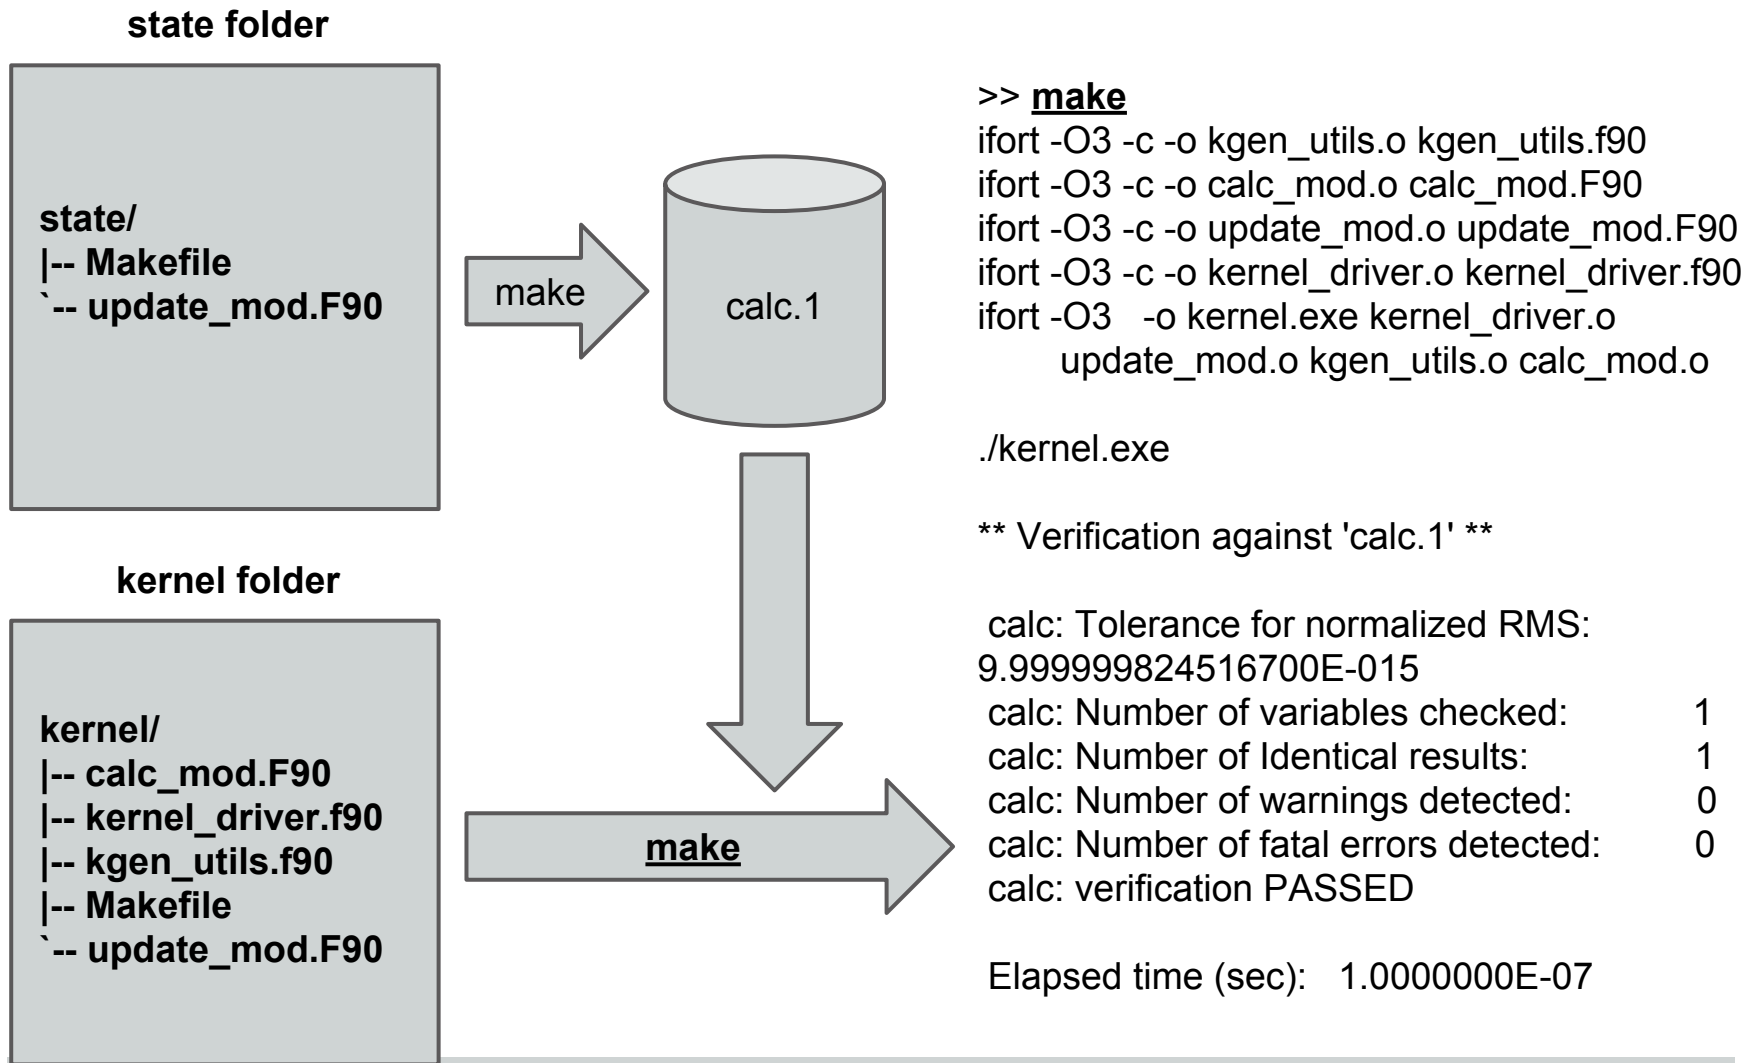
./calc_mod.F90

```
MODULE calc_mod
  PUBLIC calc
CONTAINS
  SUBROUTINE calc(i, j, output)
    INTEGER, INTENT(IN) :: i, j
    INTEGER, INTENT(OUT), &
      dimension(:,:) :: output
    output(i,j) = i + j
  END SUBROUTINE
END MODULE
```

Kernel

```
>> export KGEN=/glade/u/tdd/asap/contrib/kgen/std/src/kgen.py
>> python ${KGEN} ./update_mod.F90
```


A kernel generation example - cont.



KGEN for Large-scale Applications

- Preprocessing

- Pre-processing using “fpp” or “cpp”
 - macros are specified using “-D” or “-i <path>” flag
 - include paths are specified using “-I” or “-i <path>”

```
>> python ${KGEN} -D NC=4,NP=4,PLEV=30 -i include.ini \  
    ${SRC_DIR}/micro_mg_cam.F90
```

KGEN for Large-scale Applications

- Preprocessing

- Pre-processing using “fpp” or “cpp”
 - macros are specified using “-D” or “-i <path>” flag
 - include paths are specified using “-I” or “-i <path>”

Macro definitions

ini file for macros
and include paths

```
>> python ${KGEN} -D NC=4,NP=4,PLEV=30 -i include.ini \  
    ${SRC_DIR}/micro_mg_cam.F90
```

include.ini

```
[include]  
/ncar/opt/intel/12.1.0.233/impi/4.0.3.008/intel64/include =  
  
[~/cam5/components/cam/src/physics/cam/micro_mg_cam.F90]  
PSUBCOLS = 1  
HAVE_F2003_PTR_BND_REMAP = 1  
...  
[~/cam5/components/cam/src/physics/cam/micro_mg_cam.F90]  
...
```

KGEN for Large-scale Applications

- Excluding unnecessary searching

- Excluding unnecessary “searching”
 - Intrinsic subroutines are not searched as default
 - User-specified variables or subprograms can be ignored

```
>> python ${KGEN} -e exclude.ini ${SRC_DIR}/micro_mg_cam.F90
```

KGEN for Large-scale Applications

- Excluding unnecessary searching

- Excluding unnecessary “searching”
 - Intrinsic subroutines are not searched as default
 - User-specified variables or subprograms can be ignored

ini file for exclusion

```
>> python ${KGEN} -e exclude.ini ${SRC_DIR}/micro_mg_cam.F90
```

exclude.ini

```
[common]
endrun = comment
t_initf = comment
t_setLogUnit = comment
t_getLogUnit =comment
```

KGEN for Large-scale Applications

- Data generation from MPI app.

- State data generation from MPI application
 - User can specify MPI ranks and Nth invocation to save input data to and output data from the kernel

```
>> python ${KGEN} \  
    --ordinal-numbers 1,10,20 \  
    --mpi ranks=0,100,300 \  
    ${SRC_DIR}/micro_mg_cam.F90
```

KGEN for Large-scale Applications

- Data generation from MPI app.

- State data generation from MPI application
 - User can specify MPI ranks and Nth invocation to save input data to and output data from the kernel

>> **python \${KGEN} **

Nth invocation to kernel for data generation

**--ordinal-numbers 1,10,20 **

**--mpi ranks=0,100,300 **

MPI ranks for data generation

\${SRC_DIR}/micro_mg_cam.F90

MG2 Kernel(CESM) Demo.

MG2 (Morrison-Gettleman) Kernel

- MG2 is “a new two-moment stratiform cloud microphysics scheme in a general circulation model.”¹
- 10% of the cost of CAM² on Yellowstone
- Being optimized by multiple experts from various compiler vendors and NCAR using the MG2 kernel generated from KGEN.

1: Hugh Morrison and Andrew Gettelman, 2008: A New Two-Moment Bulk Stratiform Cloud Microphysics Scheme in the Community Atmosphere Model, Version 3 (CAM3). Part I: Description and Numerical Tests. *J. Climate*, **21**, 3642–3659.

doi: <http://dx.doi.org/10.1175/2008JCLI2105.1>

2: Community Atmospheric Model

An example: Creating a kernel of MG2 microphysics from CESM

SUBROUTINE micro_mg_tend (..... 114 dummy arguments)

- Around 2300 lines of Fortran code in this subroutine only.
- Other source files provide sub-programs and specifications such as type declarations and parameters required to compile this subroutine
- Need data to drive this kernel and to verify the correctness of its result

END SUBROUTINE

KGEN command line for MG2 kernel

```
#!/bin/bash
# kgen_run.sh

CASE_DIR := ${CAM5_HOME}/cime/scripts/FC5-cam5-mg2-SNB
SOURCE_MODS := ${CASE_DIR}/SourceMods/src.cam
SRC := ${SRC_DIR}/micro_mg_cam.F90

python ${KGEN} \
  -i include.ini \
  --outdir ${OUTPUT_DIR} \
  --ordinal-numbers 10,50,100 \
  --mpi ranks=0:100:300,comm=mpicom,use="spmd_utils:mpicom" \
  --kernel-compile FC=ifort,FC_FLAGS='-xHost -O2' \
  ${SRC}:micro_mg_cam.micro_mg_cam_tend.micro_mg_tend2_0
```

User-provided macros definitions and included paths

```
; include.ini
[include]
/ncar/opt/intel/12.1.0.233/impi/4.0.3.008/intel64/include =

[~/cam5_3_74/components/cam/src/physics/cam/micro_mg_cam.F90]
PSUBCOLS = 1
HAVE_F2003_PTR_BND_REMAP = 1
HAVE_SLASHPROC = 1
HAVE_NANOTIME = 1
...
[~/cam5_3_74/components/cam/src/physics/cam/micro_mg_cam.F90]
...
```

NOTE: KGEN provides a python script that generates include paths and macro definitions from CESM log files.

Screen output

```
>> kgen_run.sh
```

Pre-processing is done

Reading ~/cam5/components/cam/src/physics/cam/micro_mg_cam.F90

Call-site location is found

Reading ~/cam5/components/cam/src/physics/cam/micro_mg2_0.F90

Reading ~/cam5/components/cam/src/physics/cam/micro_mg_utils.F90

Reading ~/cam5/components/cam/src/physics/cam/wv_sat_methods.F90

Reading ~/cam5/cime/share/csm_share/shr/shr_spfn_mod.F90

Reading ~/cam5/cime/share/csm_share/shr/shr_kind_mod.F90

Kernel information is collected

Instrumented files are generated

Kernel files are generated

Makefiles are generated

Post-processing is done

Completed.

Generated files in output folders

kernel

- |-- kernel_driver.f90
- |-- kgen_utils.f90
- |-- Makefile
- |-- micro_mg2_0.F90
- |-- micro_mg_cam.F90
- |-- micro_mg_utils.F90
- |-- shr_kind_mod.F90
- |-- shr_spfn_mod.F90
- `-- wv_sat_methods.F90

state

- |-- Makefile
- |-- micro_mg2_0.F90
- |-- micro_mg_cam.F90
- |-- micro_mg_utils.F90
- `-- wv_sat_methods.F90

Potential use-cases

- Optimization and Porting
 - Original use case
 - Enable fast cycle of workflow
- Debugging
 - Potentially no need for queue
 - Focusing on the target code
 - Quick cycle-time
- Unit-test
 - Verification of new-updates in production code with data generated from the production code
- Benchmark test
 - allows organization to generate their own collection of test cases
- And more
 - New compiler verification
 - Automated optimization
 - Training

Development status and plans

- Development status

- Funded by Intel Parallel Computing Center focused on Weather and Climate Simulation (IPCC-WACS)
- Has been applied to mostly CAM of CESM with Intel compiler
- Available to “early-adopters” TODAY
 - svn co <https://proxy.subversion.ucar.edu/pubasap/kgen/trunk>
 - Available in “/glade/u/tdd/asap/contrib/kgen/std” on Yellowstone
 - Welcome your opinion and comments on using KGEN.

If you have one, please send it to “kgen@ucar.edu”

- Plans

- Improving supports for Fortran specification
- Expanding tests to components of CESM and other applications
- Expanding supports for various compilers
- Supporting for “early-adopters” of KGEN

Thank you!

Funded by Intel Parallel Computing Center focused on
Weather and Climate Simulation (IPCC-WACS)

Youngsung Kim, John Dennis, Raghu R. Kumar, and Amogh Simha
(kgen@ucar.edu)