



NCAR

Bursts and Cascades: Scaling Up Scientific Data Analysis

Seth McGinnis

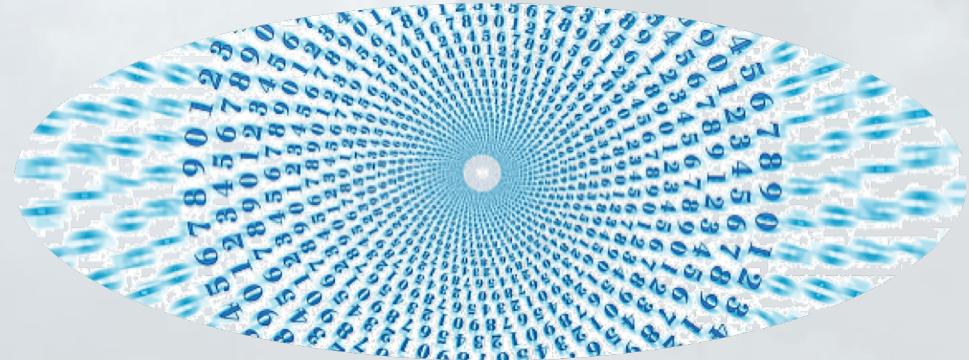
2018-04-03

SEA Conference 2018

National Center for Atmospheric Research

What This Talk Is About

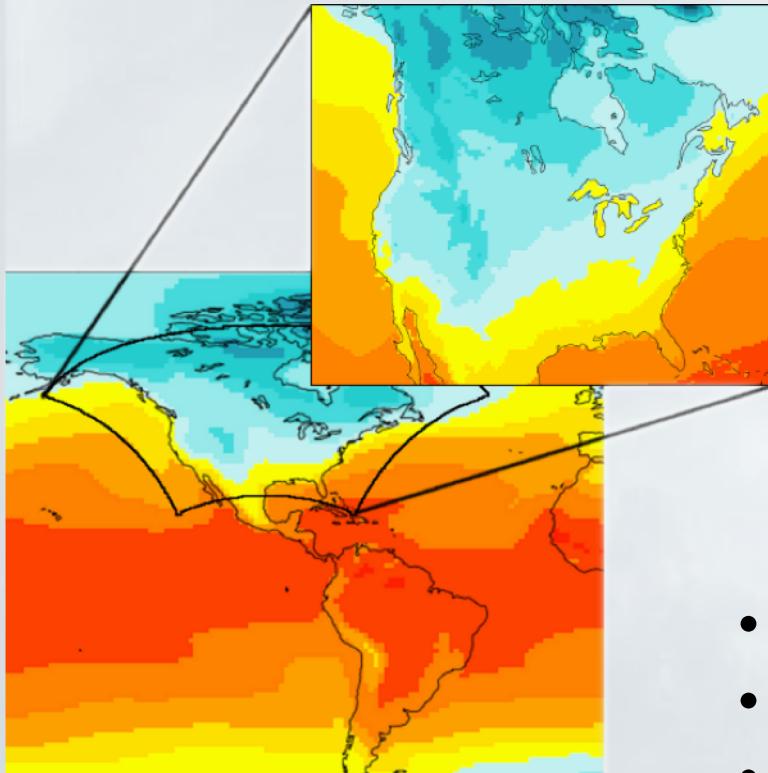
- Developing an analysis / derived data product:
 - Start with a test case you understand
 - Develop code
 - Apply it to the entire dataset
- Big Data happens when data volume (etc) is so large that normal approaches don't work
- Two successful approaches to scaling up:
 - Parallel burst
 - Hierarchical cascade



Parallel Burst: Bias-Correcting NA-CORDEX



Regional Climate Modeling in NA-CORDEX



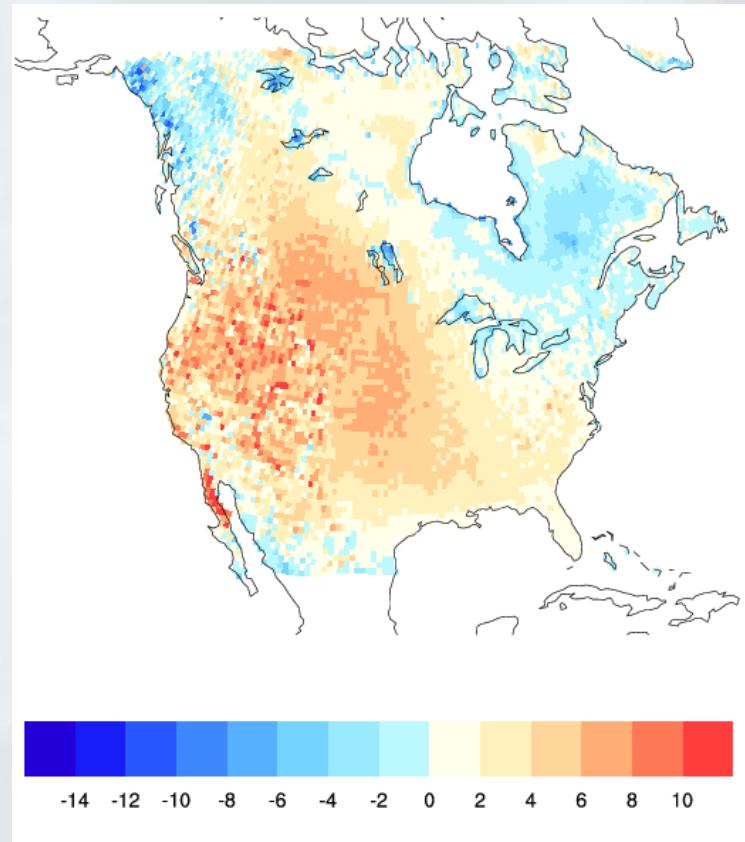
Nest high-resolution regional models (RCMs) inside coarser global models (GCMs) over North America

- Better mountains & coasts
- Better land surface detail
- Better convective processes
- Closer to impacts scale

Bias Correction

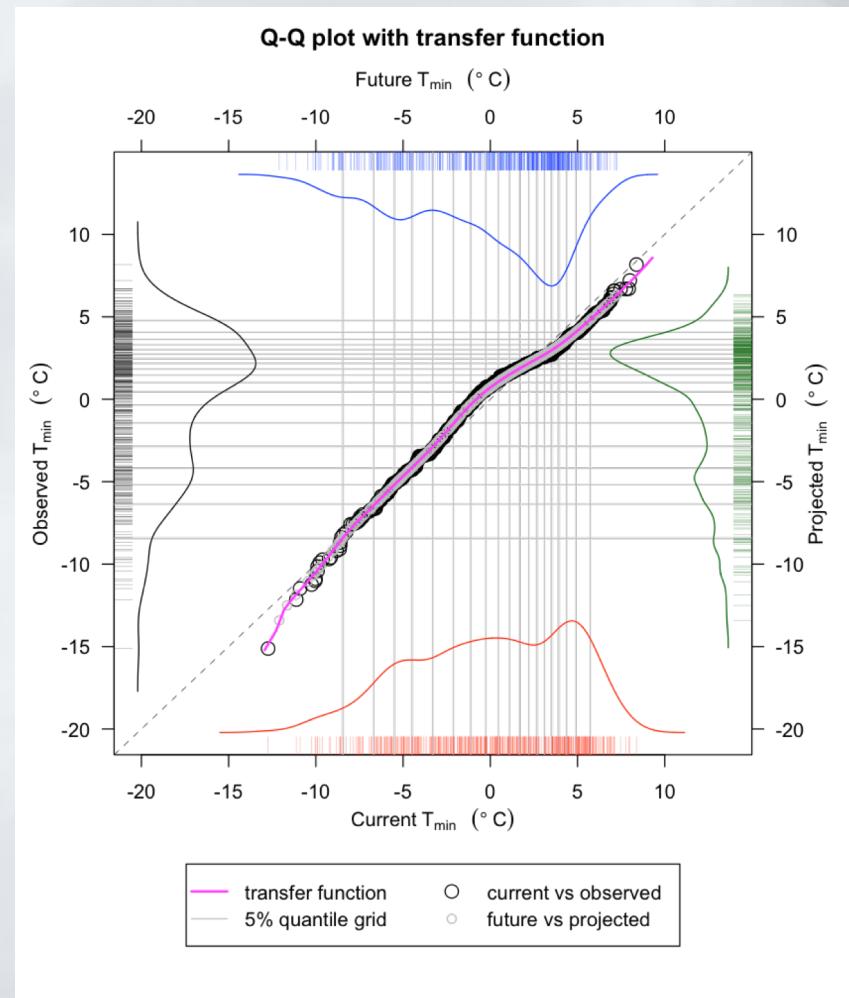
- RCMs & GCMs have systematic errors
 - Finite resolution
 - Parameterizations
 - Imperfect knowledge
- Correcting bias makes data more useable

Summer RCM Temperature Bias (°C)



BC Method: Distribution Mapping

- Adjust individual values so model PDF matches PDF of observations
- R script
- 20-25 seconds runtime
- 335-350 MB memory
- Reads 3 files (3.86 MB)
- Writes 2 files (3.1 MB)



Bias-Correcting (Part of)* the Dataset

*Full dataset is 5-10x bigger

3 variables
3 GCMs
2 RCMs
~~600 x 258~~ gridcells
232 x 154 (CONUS only)

x _____

643,104 timeseries

- 643k x 22 seconds
 - = 14 million CPU-sec
 - = 3,930 CPU-hours
 - = 23+ weeks wallclock
-
- 2.36 TB in, 1.9 TB out

Conclusion: better run it in parallel!

Running in Parallel

GOOD

- Ridiculously parallel problem structure

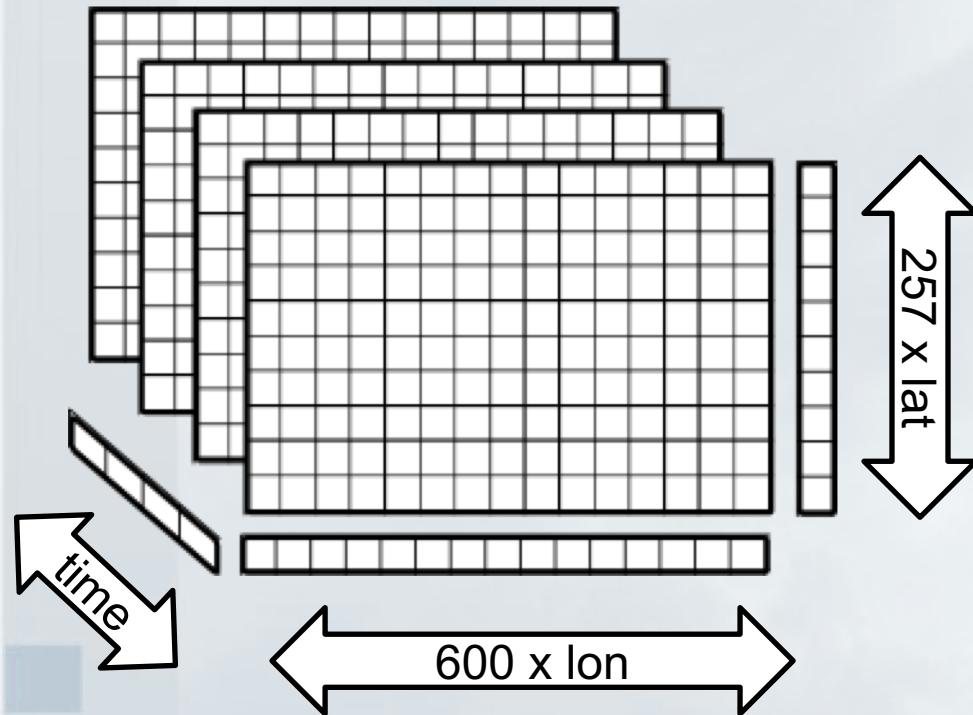
BAD

- No good tools for parallel I/O in R (esp. NetCDF)
- Can't rewrite code in another language
 - No time; code uses R packages extensively

UGLY

- Rmpi package: version mismatch problems
- Supervisor process too big to fit on normal node
- Debugging parallel code is really hard

Solution (pt. 1): Disassemble NetCDF



- Data cube → data rods
- Components:
 - Global attributes
 - Ancillary (lat, lon, etc.)
 - Time coordinate
 - Main data variable
- Slice data by lon
- Slice each slice into rods (timeseries) by lat

Solution (pt. 2): MPMD Parallelism

- Use scheduler to distribute tasks from batch command file:
 - one task per line
 - one task per core
- Wrap commands in shell one-liner to capture output:

```
tcsh -c "Rscript --vanilla script.R arg1 arg2..."
```
- 1 job = 600 cores looping over 258 rods / slice

(This works so well I wrote a generic script for running batch files in parallel.)

Result: x2300 speedup with minimal code modifications

Lessons Learned

- Use generator scripts & templates to create batch files
 - Parallelize disassembly / reassembly steps, too
 - wallclock limit < 15 minutes + #cores < 1000
 - = nigh-instant scheduling (backfilling empty holes)
- *Some problems don't show up until you try your code on the entire dataset*
- *Big error fixes can mean rerunning EVERYTHING*
- *I used up 20x extra CPU-hours getting it working correctly*

Moving to the Cloud?

- This problem is well-shaped for cloud
 - Burst demand
 - Highly parallel
 - Small task footprint
- AWS recommendation:
Lambda service
 - Serverless computation triggered by file arrival
 - Est. cost: O(\$100s)

Complications

- In the cloud:
 - file = immutable blob
 - ? File dis- / re-assembly
 - ? NetCDF → timeseries DB
 - ? transpose problem
- Data transfer in & out
 - 10s of TB is a lot
 - AWS “snowball” @\$200?
- Only in Python

Hierarchical Cascade: Model Evaluation in FACETS

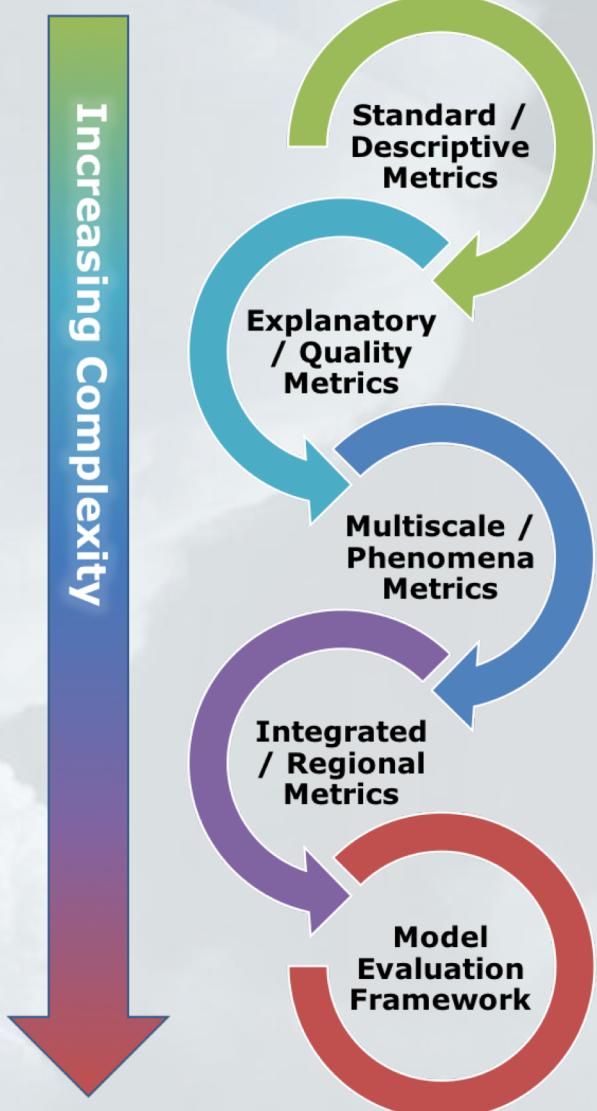


U.S. Department of Energy
FACETS
Framework for Assessing Climate's Energy-Water-Land
Nexus using Targeted Simulations

Hierarchical Evaluation Framework: A Cascade of Metrics

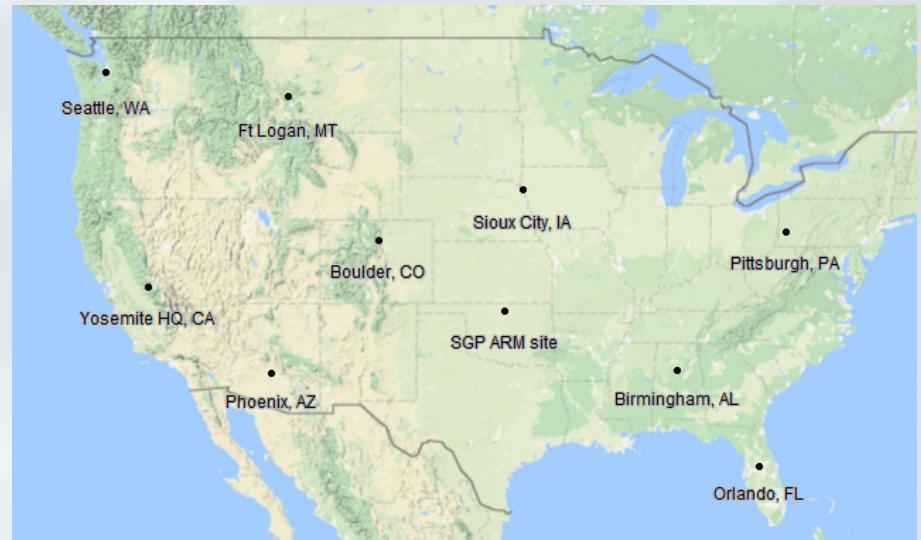
1. Statistical Analysis
2. Process Analysis
3. Process Interaction Analysis
4. Credibility Analysis
5. Differential Credibility

*Simpler analyses are distilled
into quantitative metrics used
to target more complex analysis*



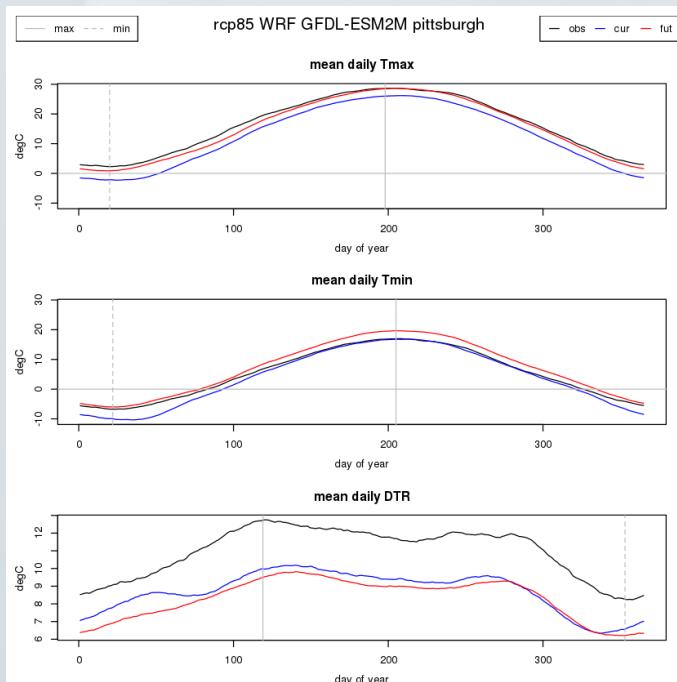
Evaluating NA-CORDEX RCMs

- 2 RCMs:
 - RegCM4 & WRF
- 3 CMIP5 GCMs
 - HadGEM2-ES
 - MPI-ESM-LR
 - GFDL-ESM2M
- Downscaling to 25 km
- Daily precip, Tmin, Tmax



10 test locations
with different climates
(1st stage of targeting)

Statistical Analysis → Numeric Metrics



period	GCM	RCM	tmaxcor	tmincor	dtrcor	tmaxmad	tminmad	dtrmad
hist	GFDL-ESM2M	RegCM4	0.995	0.998	0.515	7.86	5.6	1.99
hist	HadGEM2-ES	RegCM4	0.998	0.997	0.446	5.82	3.13	2.1
hist	MPI-ESM-LR	RegCM4	0.993	0.999	0.626	4.14	3.13	1.6
hist	GFDL-ESM2M	WRF	0.999	0.998	0.93	5.24	1.25	3.57
hist	HadGEM2-ES	WRF	0.998	0.996	0.724	2.68	3.02	3.06
hist	MPI-ESM-LR	WRF	0.999	0.996	0.572	6.77	3.33	3.32
rcp85	GFDL-ESM2M	RegCM4	0.994	0.999	0.743	5.13	2.29	1.37
rcp85	HadGEM2-ES	RegCM4	0.999	0.995	0.815	2.9	3.24	1.79
rcp85	MPI-ESM-LR	RegCM4	0.996	0.999	0.844	2.24	1.26	1.35
rcp85	GFDL-ESM2M	WRF	0.999	0.999	0.987	1.21	2.71	3.97
rcp85	HadGEM2-ES	WRF	0.998	0.997	0.919	2.99	6.41	3.76
rcp85	MPI-ESM-LR	WRF	0.999	0.996	0.785	3.03	5.04	2.86

“Low-Bar” Statistical Metrics

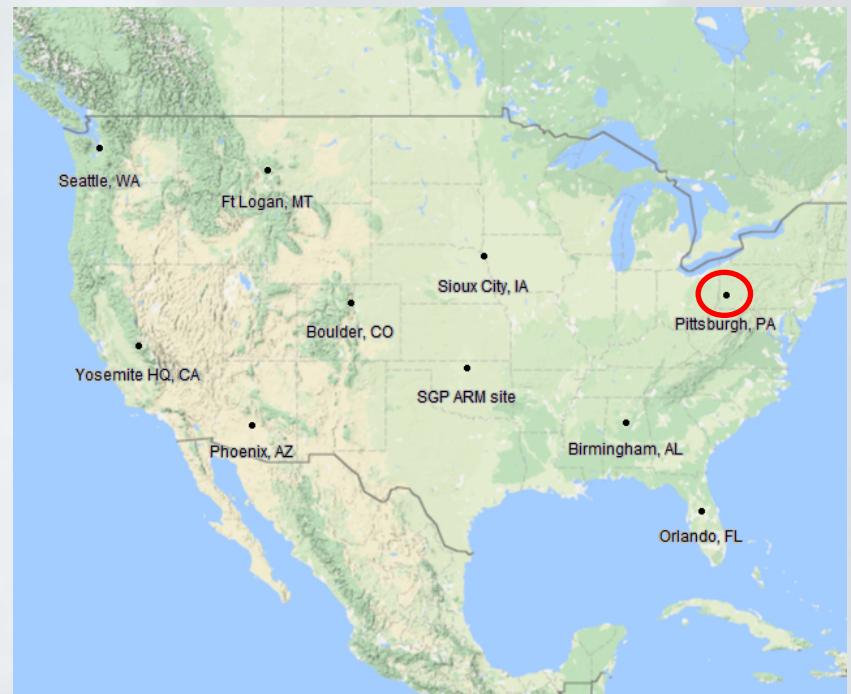
- Annual cycle (T, P)
- Monthly PDFs (T, P)
- Extreme value analysis of precipitation
- Seasonal precip intensity spectra
- Joint distribution of temp & precip
- Most informative: annual cycle of precipitation
 - Frequency
 - Intensity
 - Total
- Because everything affects precipitation

Pittsburgh, PA

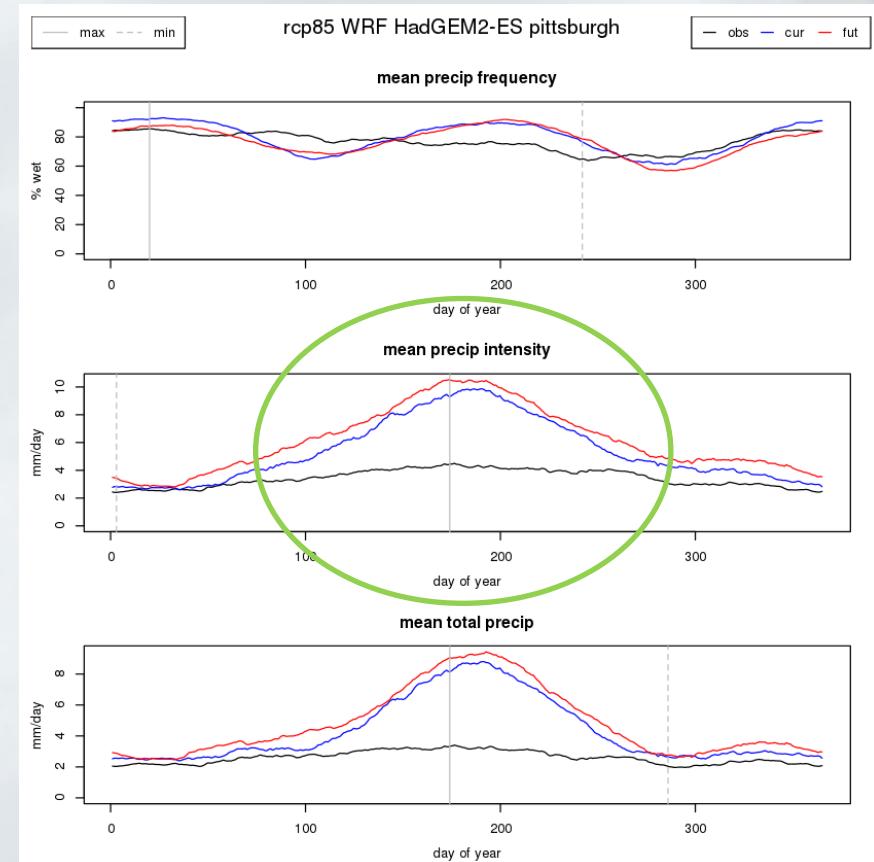
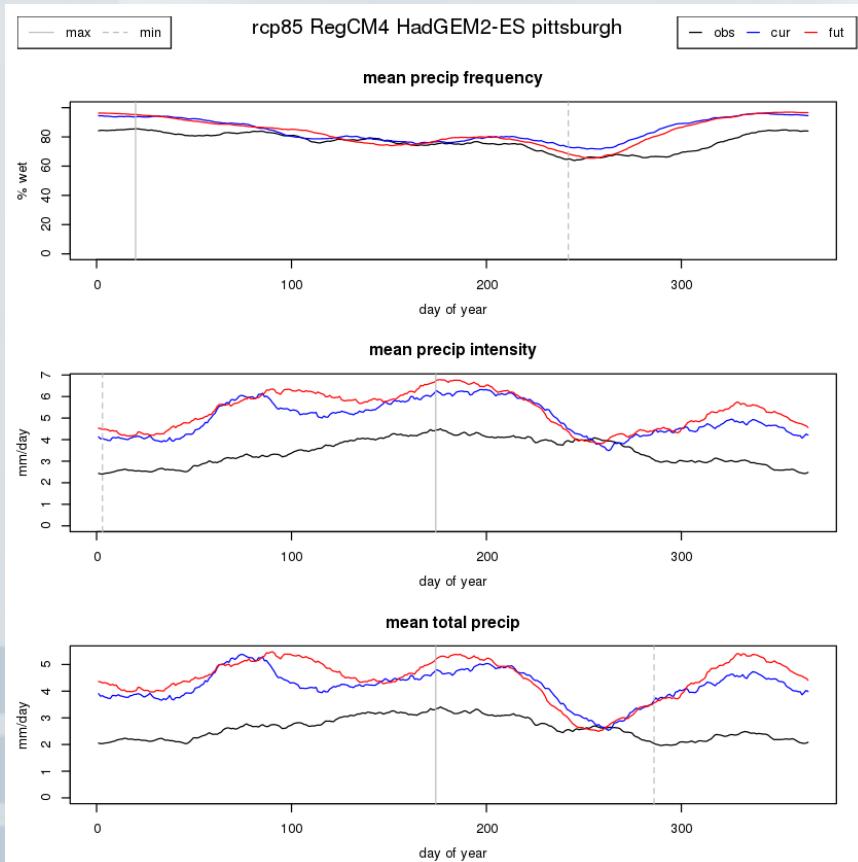
- Simple precipitation cycle

In general:

- Summer: mostly convective
- Winter: heavily forced
- Simulations with the same RCM look very similar



Pittsburgh: summer WRF bias ⇒ investigate convective parameterizations

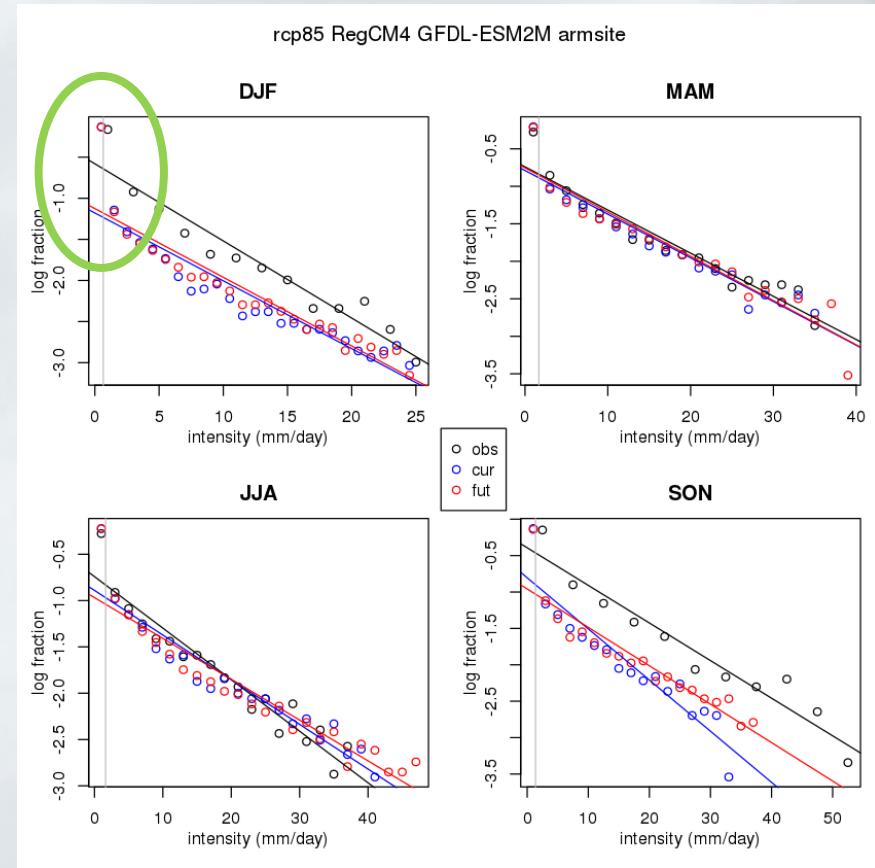
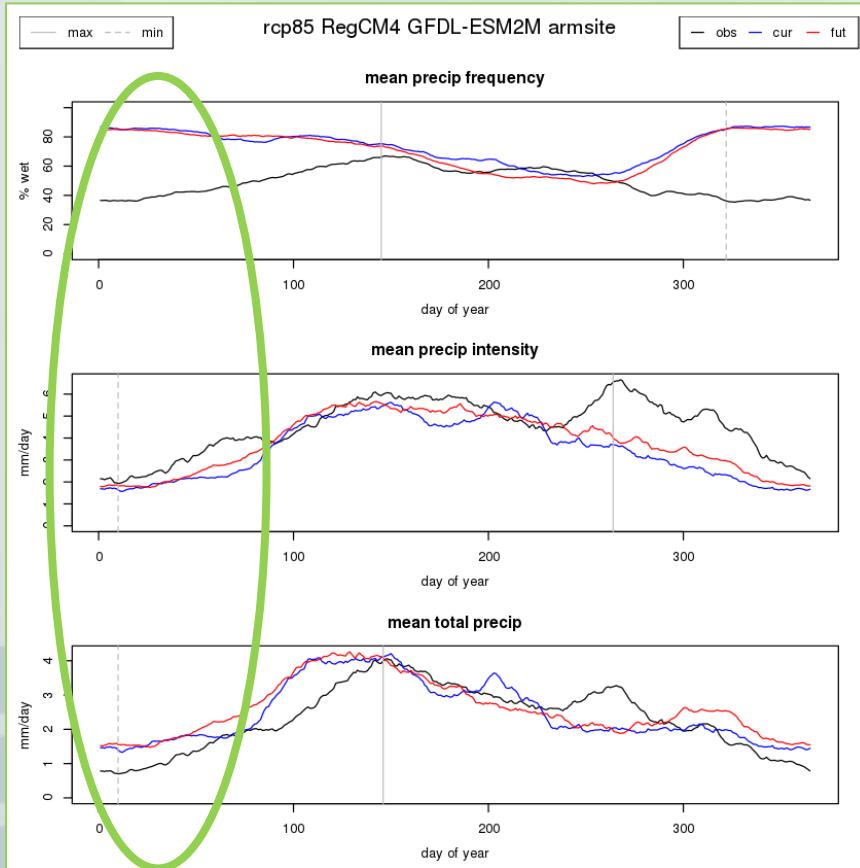


SGP ARM site, OK

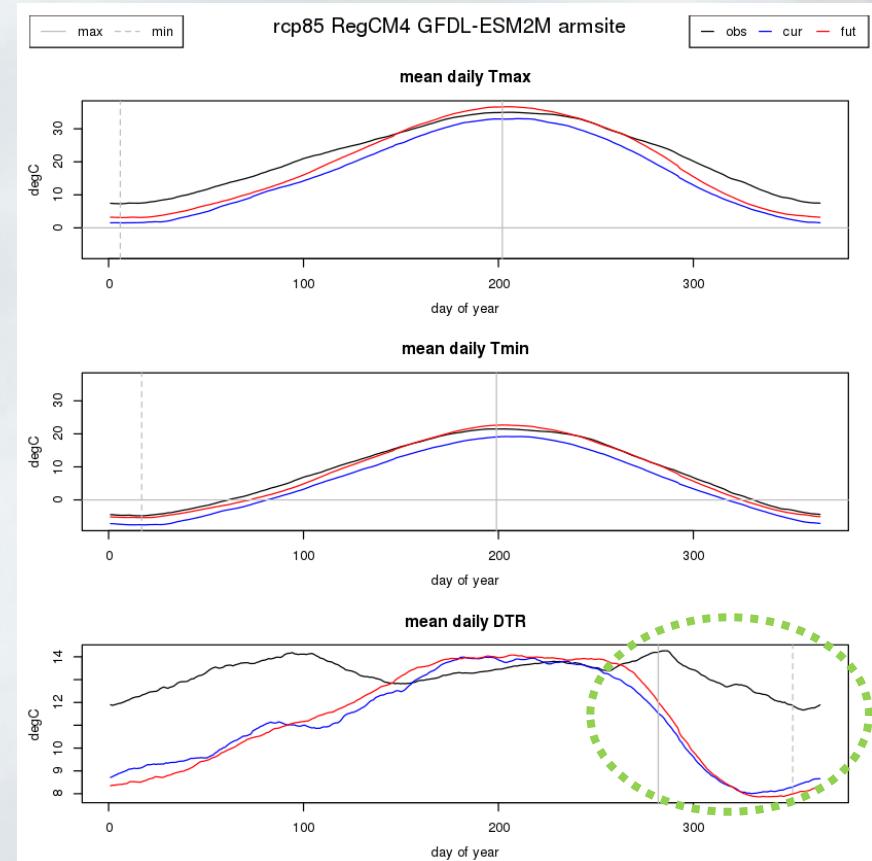
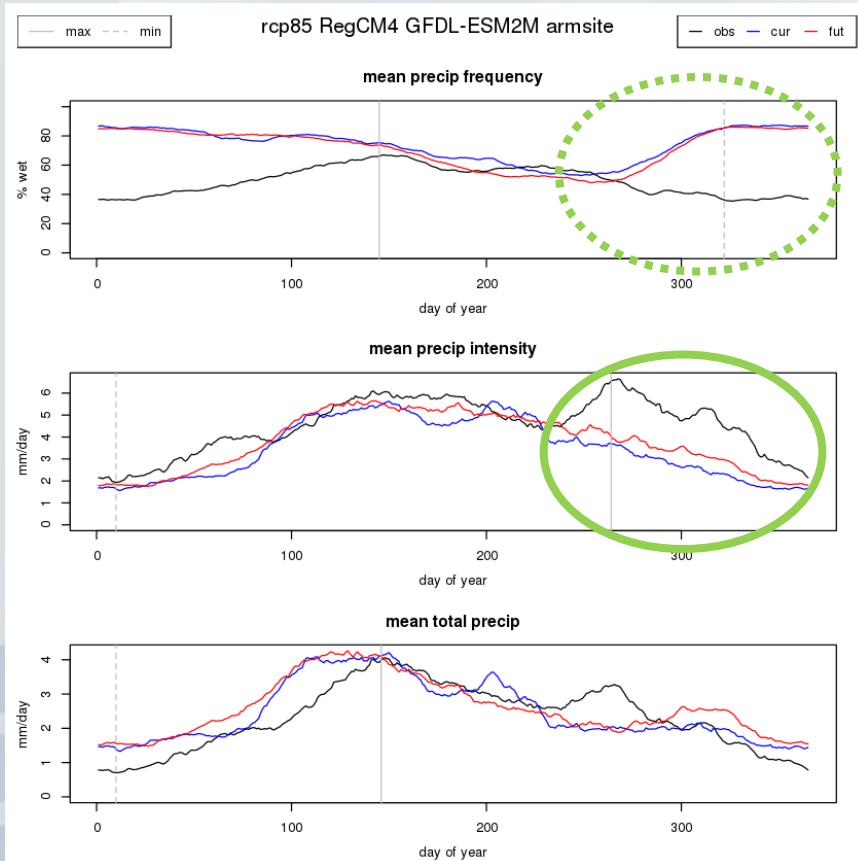
- DoE Climate Research Facility
- All WRF simulations look reasonable
- All RegCM4 simulations have similar problems



ARM site (1): Too much cold season drizzle in RegCM4 ↔ no ice microphysics



ARM site (2): Missing cold-front convective lines ⇒ drizzle erodes instability in RegCM4



More complex analysis



More focused targets

- Basic climatology helps you pick test locations
- Differences in numeric metrics tell you which statistical analyses to examine
- Applying meteorological knowledge to the statistical analyses identifies candidates for process analysis
- Process analysis leads to region-specific phenomena-based metrics for use in interaction analysis
- Etc.

Acknowledgements

This work supported in part by:



DoD ESTCP RC-201666



DoD SERDP grant #2516



DoE award DE-SC0016438



NSF award CMMI 1635638/1635686

The National Center for Atmospheric Research
is sponsored by the National Science Foundation

Questions?

