

FROM RESEARCH TO INDUSTRY



TOWARDS AN EFFICIENT COMMUNICATION OVERLAP THROUGH HARDWARE OFFLOADING

SEA conference 2018, UCAR
Julien JAEGER,
CEA, DAM, DIF, F-91297 Arpajon, France

4 APRIL 2018

www.cea.fr

- **Team overview**

- Runtime system and software stack for HPC
- Team as of March 2018 (CEA and CEA/Intel/UVSQ ECR Lab)
 - 2 research scientists, 3 PhD students, 3 engineers
 - Contact: patrick.carribault@cea.fr or julien.jaeger@cea.fr
 - Partner: Paratools (1,5 MY)
- Available software
 - MPC framework
 - MALP (performance analysis tool)
 - JCHRONOSS (job scheduler for test suite on production machines)
 - Wi4MPI (MPI abstraction)
- Website for team work: <http://hpcframework.com>

- **MPC framework**


- Unified parallel runtime for clusters of NUMA machines
 - Idea: one process per node, compute units exploited by user-level threads
- Integration with other HPC components
 - Parallel memory allocator, compilers, debuggers, topology tool...
- Homepage: <http://mpc.hpcframework.com>

MPC

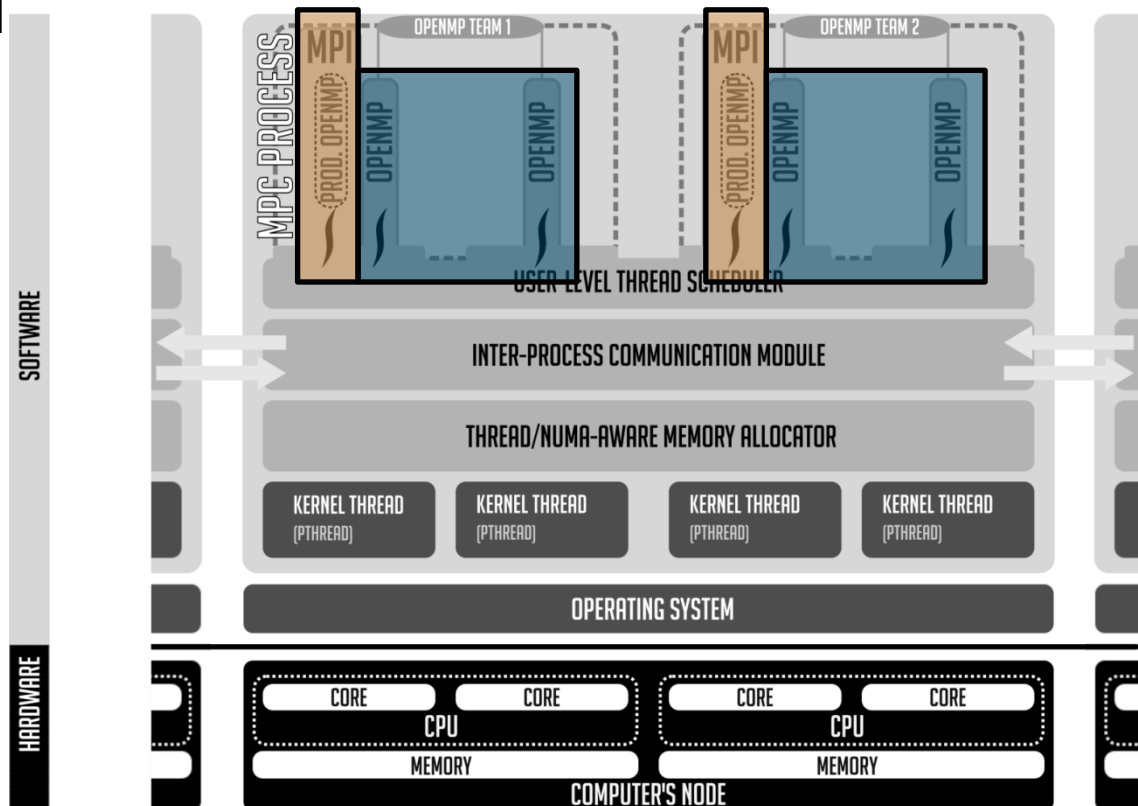


multi-**P**rocessor **C**omputing

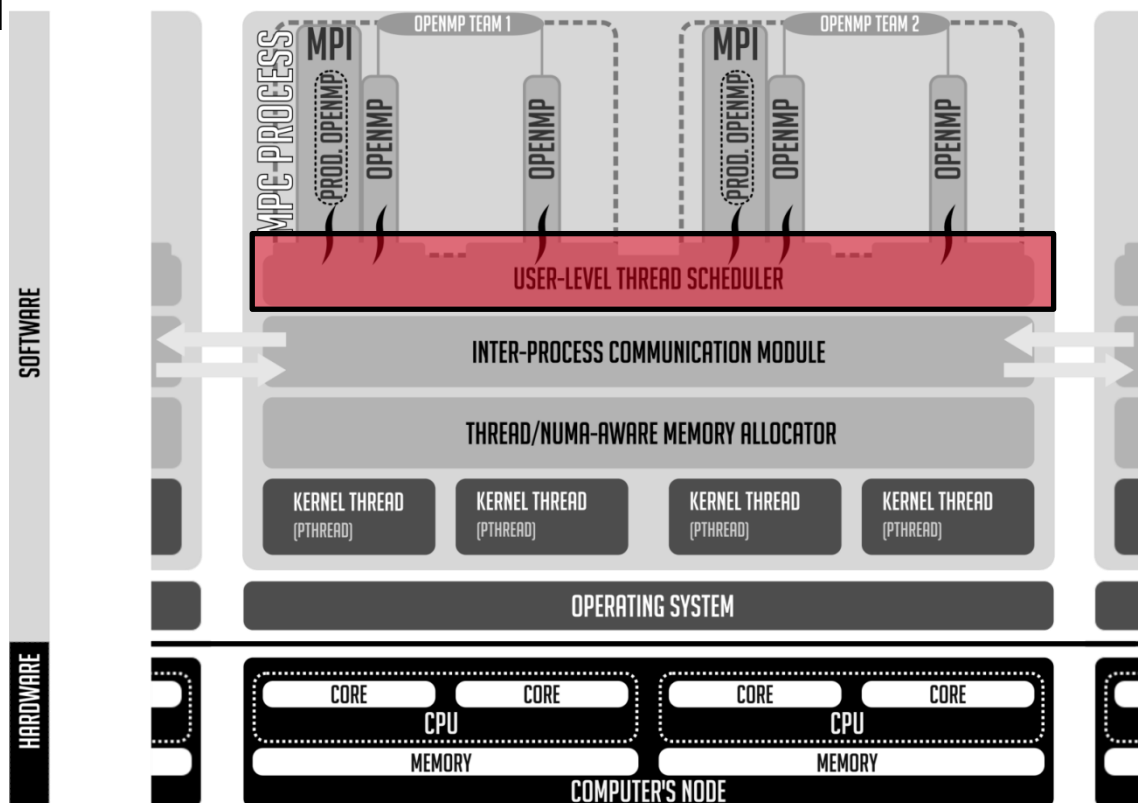
- MPC is...



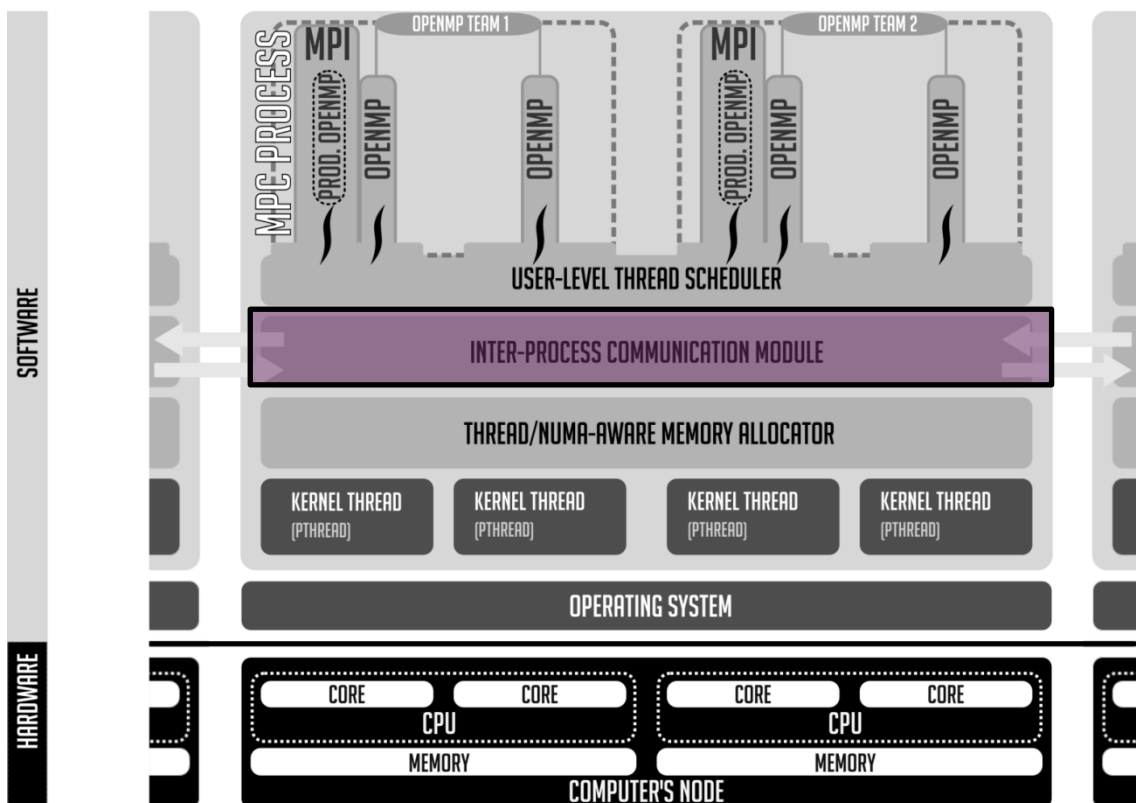
- MPC is...
- ... MPI, OpenMP, pthread implementations...



- MPC is...
- ... MPI, OpenMP, pthread implementations...
- ... based on the same **thread scheduler** allowing interoperability, ...




- MPC is...
- ... **MPI**, **OpenMP**, pthread implementations...
- ... based on the same **thread scheduler** allowing interoperability, ...
- ...sharing the **communication module** for collaborative polling...



- MPC is...

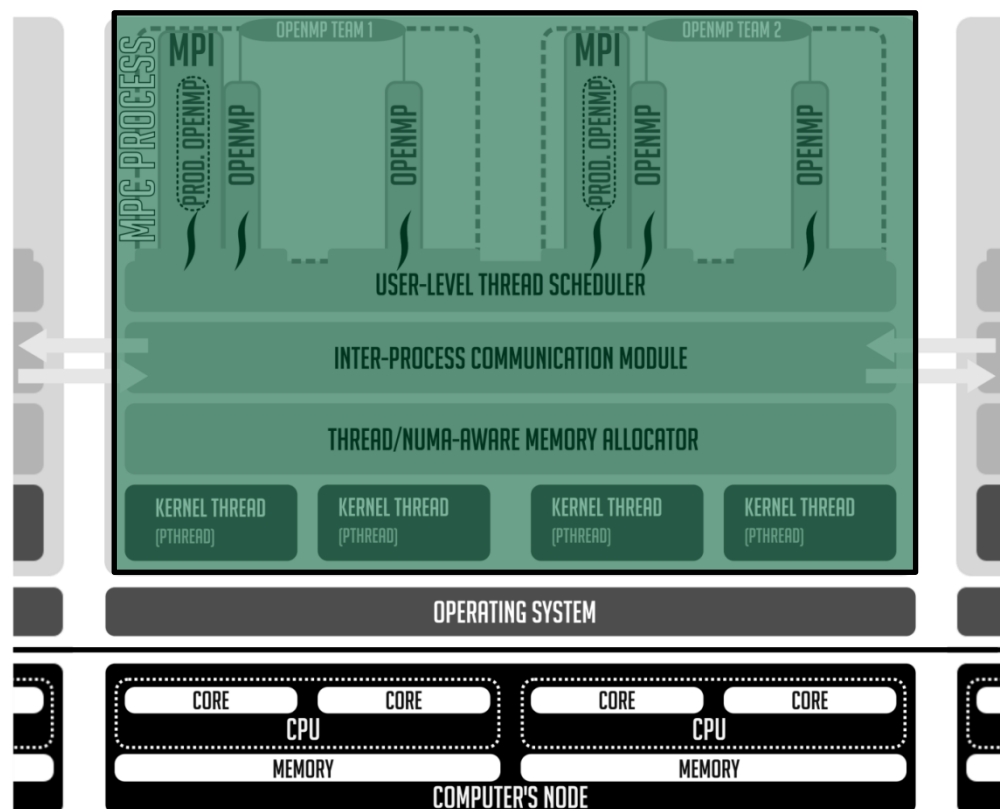
- ... MPI, OpenMP, pthread implementations...
- ... based on the same **thread scheduler** allowing interoperability, ...
- ...sharing the **communication module** for collaborative polling...
- ...and a **memory allocator** avoiding NUMA effects...



- MPC is...
- ... **MPI**, **OpenMP**, pthread implementations...
- ... based on the same **thread scheduler** allowing interoperability, ...
- ...sharing the **communication module** for collaborative polling...
- ...and a **memory allocator** avoiding NUMA effects...
- ...for **efficient hybrid MPI+X programming**

SOFTWARE

HARDWARE




MPC – UNIFIED USER-LEVEL THREAD SCHEDULER



- Why user level threads ?
- Easier development
 - compare to kernel development
- Optimizations:
 - Keep only useful HPC functionalities
 - Less or no system calls (no system calls if no signal support)
- Portable:
 - OS independent
- Drawbacks:
 - Hard to debug: need specific debugger support (see tools section)
 - Architecture dependencies:
 - Optimized ASM context/switch, spinlocks, ...
 - Topology detection and binding (easier thanks to HWLOC)

MPC User-Level Thread Scheduler

- MPC User-Level Thread Scheduler



- **MPC User-Level thread scheduler features**
- MxN thread scheduler
 - M user-level threads $>$ N kernel threads
- Optimizations
 - ASM context switches
 - Link with MPC memory allocator to ensure data locality
- Topological binding:
 - Static a priori distribution (MPI scatter/OpenMP compact)
 - On demand migration (link with memory allocator)
- MPI optimized scheduler
 - Internal dedicated task engine for message progression
- Optimized busy waiting policy
 - Use the thread scheduler to dynamically adapt busy waiting policy according to node workload
 - Busy waiting delegation to thread scheduler (e.g. “smart” mwait/futex)
- Modular approach in order to evaluate new scheduling policies
- No preemption

- Useful for use 3rd part tools/runtimes
- **Features**
 - Standard PThread API
 - Signal support
 - Futexes
 - No preemption
- **Examples**
 - Intel TBB (included in the MPC Framework package)
 - First port of Non Blocking Collectives (progression thread)
- **Debugging**
 - Provides a patched GDB
 - Support from Allinea DDT debugger

MPC – MPI IMPLEMENTATION



MPC – MPI: Purpose and Supported Features

- **Goals**

- Smooth integration with multithreaded model
- Low memory footprint
- Deal with unbalanced workload

- **Supported Features**

- Fully MPI 3.1 compliant
- Handle up to MPI_THREAD_MULTIPLE level (max level)
- MPI I/O
- Non-blocking collectives
- Neighborhood collectives


- **Inter-node communications**

- TCP, InfiniBand & Portals4

- **Tested up to 80,000 cores with various HPC codes**


MPC – MPI Process-Based Execution Model

- Application with 4 MPI processes in 4 processes



MPC – MPI Thread-Based Execution Model

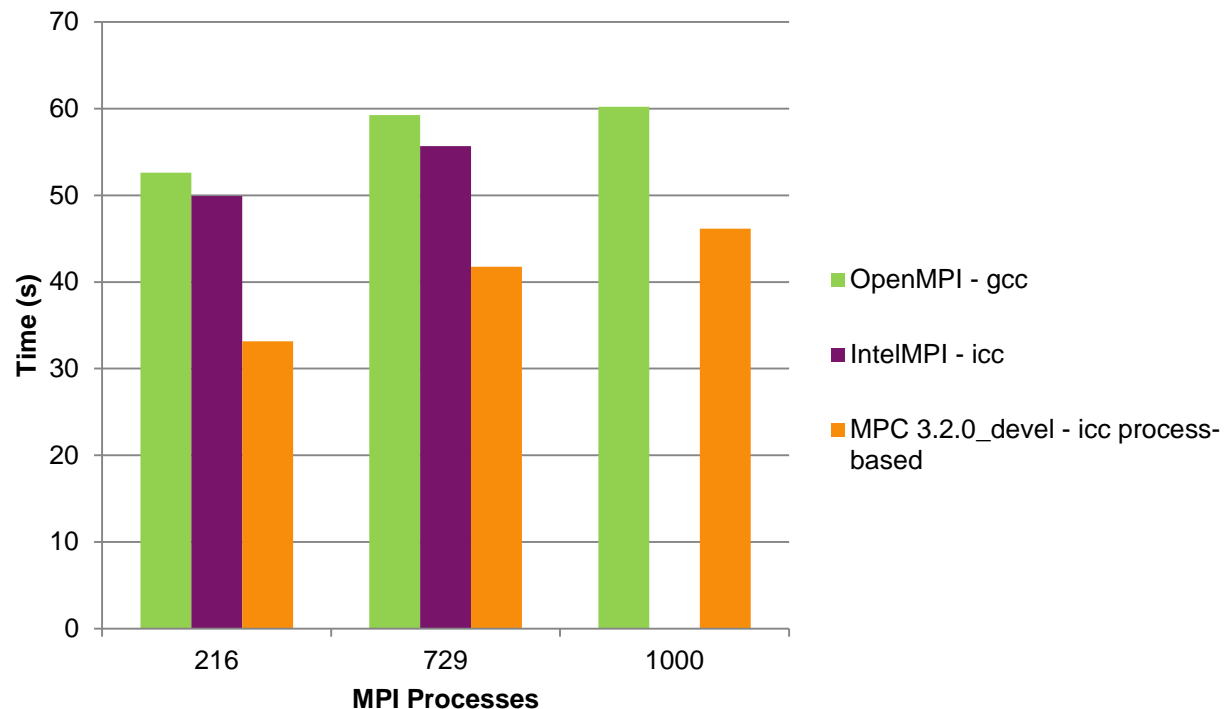
- Application with 4 MPI processes in 1 process (with 4 threads)



MPI Experiment

- Architecture: Intel KNL
- Application: LULESH MPI
- Evaluation of multiple MPI implementations (OpenMPI, IntelMPI, MPC with either GNU or Intel compiler)

Lulesh 30x30x30




MPC – OPENMP IMPLEMENTATION



- **Impact of runtime stacking on OpenMP Layer**
- Strong performance in fine-grain and coarse-grain approaches
 - Fine-grain
 - Optimization of launching/stopping a parallel region
 - Optimization of performing loop scheduling
 - Coarse-grain
 - Optimization of synchronization constructs (barrier, single, nowait...)
- Thread placement according to available cores (job manager + MPI runtime)
 - Oversubscribing → need to avoid busy waiting
- **OpenMP runtime design and implementation**
- Goal: design of OpenMP runtime fully integrated into MPI runtime dealing with Granularity and Placement
- Implementation in MPC unified with optimized MPI layer

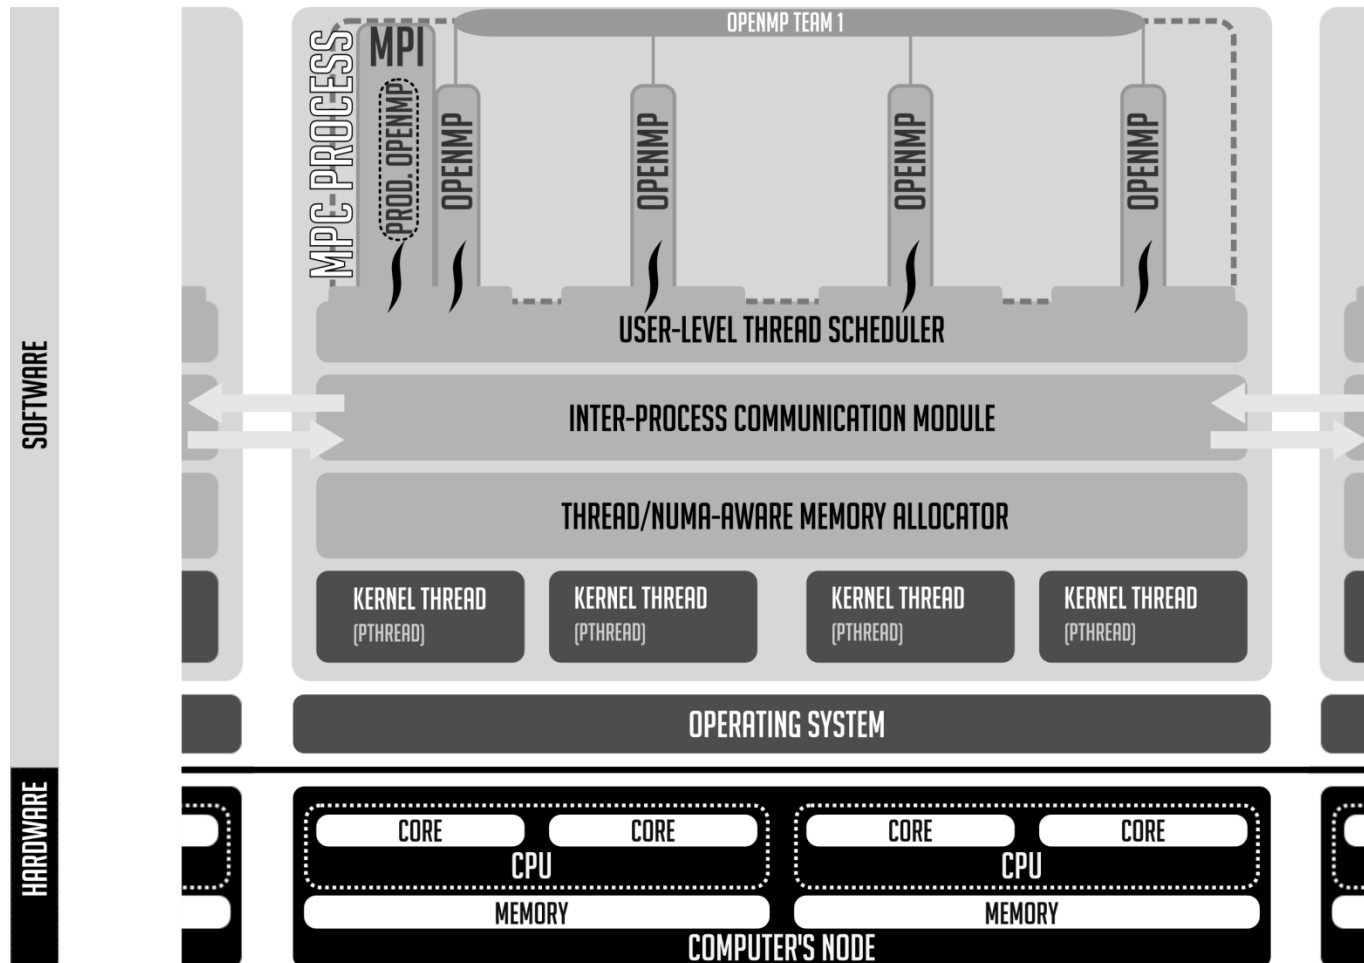
MPC Execution Model: Full MPI

- 4 MPI tasks in one process




MPC Hybrid Execution Model: MPI+OpenMP (1)

- 1 MPI process + 4 OpenMP threads in one MPC process



MPC Hybrid Execution Model: MPI+OpenMP (2)

- 2 MPI processes + 4 OpenMP threads in one MPC process




User-Level Thread Scheduler: OpenMP

- **OpenMP threads are User-Level thread**
- Smart binding thanks to information sharing with MPI runtime through thread scheduler
- Automatic MPI process and OpenMP thread placement on the node
 - Topology inheritance
- **Benefits**
- Avoid busy waiting in multi-programming model context
 - Useful for fast thread wakeup/sleep (entering/leaving parallel regions, ...)
- Data locality (link scheduler \Leftrightarrow memory allocator)
- **Drawbacks**
- Unable to use standard TLS for #pragma threadprivate
 - Solution with extended-TLS

OpenMP Experiment

- Architecture: dual-socket 16-core Haswell
- Application: LULESH OpenMP compiled with Intel 16
- Evaluation of Intel OpenMP runtime vs. MPC OpenMP runtime (same compiler)

LULESH 50x50x50



MPC – EXTENDED TLS AND AUTOMATIC PRIVATIZATION



Process virtualization

- Use threads (e.g. MPI processes) instead of OS processes
 - Convert standard OS process to MPC thread
 - Encapsulate OS process within threads
- Difficulties
 - How to handle global variables
 - Unable to use standard TLS for `#pragma threadprivate`
 - Deal with non-thread safe libraries
 - User libraries: HDF5, ...
 - System libraries: getopt, ...
 - Compiler libraries: libgfortran,
- ➔ Automatic privatization thanks to compiler support
- ➔ Provide patched version if not compatible yet with automatic privatization
- ➔ Provide non-thread safe system libraries

- **Solution: Automatic privatization**
 - Automatically convert any MPI code for thread-based MPI compliance
 - Duplicate each global variable
- **Design & Implementation**
 - Completely transparent to the user
 - When parsing or creating a new global variable: flag it as MPI thread-local
 - `#pragma threadprivate` : flag it as OpenMP thread-local
 - Generate runtime calls to access such variables (extension of TLS mechanism)
 - Linker optimization for reduce overhead of global variable access
- **Compiler support**
 - New option to GCC C/C++/Fortran compiler (`-fmpc-privatize`)
 - Patched GCC provided with MPC (4.8.0, on going on GCC 4.9.x and 5.x)
 - ICC support automatic privatization with same flag (`-fmpc-privatize`)
 - ICC 15.0.2 and later
 - On-going work for PGI compiler support

- Official support of MPC since Intel Compiler v15.02

■ Extracted from the icc/icpc/ifort man page

```
> man icc
...
Feature: Privatization of static data for the MPC unified parallel runtime Requirement:
Appropriate elements of the MultiProcessor Computing (MPC) framework For more information,
see http://mpc.sourceforge.net/
...
-fmpc-privatize (L*X only) / -fno-mpc-privatize (L*X only)
Enables or disables privatization of all static data for the MultiProcessor Computing
environment (MPC) unified parallel runtime.
Architecture Restriction: Only available on Intel(R) 64 architecture
Arguments: None
Default: -fno-mpc-privatize
The privatization of all static data for the MPC unified parallel runtime is disabled.
Description:
This option enables or disables privatization of all static data for the MultiProcessor
Computing environment (MPC) unified parallel runtime.
Option -fmpc-privatize causes calls to extended thread-local-storage (TLS) resolution, run-
time routines that are not supported on standard Linux* OS distributions.
This option requires installation of another product. For more information, see Feature
Requirements.
```

Conclusion & Future Work

—

Overview of MPC

- **MPC**

- Unified user-level thread scheduler
 - With an inter-process communication module and a thread/numa-aware allocator

- **Programming models**

- Provide widely spread standards: MPI 3.1, OpenMP 3.1+, Pthread, TBB
- Available at <http://mpc.hpcframework.com> (version 3.2 available)

- **Runtime optimization**

- Provide unified runtime for MPI + X applications
- New mechanism to mix thread-based programming models: Extended TLS

- **Support**

- Architecture: x86, x86_64, MIC, arm (in progress)
- Network: TCP, Infiniband and Portals4 with multi-rail
- Resource manager: Slurm & Hydra

- **Tools**


- Use HWLOC to detect topology
- Debugger support (Allinea DDT), Profiling
- Compiler support (Intel, GCC, PGI)

SUPPORTING
ASYNCHRONOUS
PROGRESSION
WITH
HARDWARE
OFFLOADING

HARDWARE MESSAGE MATCHING


- As seen on these production MPI (using default configuration) overlap is not to be taken for granted
 - Even when using « asynchronous » calls.
 - It is important to progress these calls.
- Reasons for this are multiple (non exhaustively):
 - Optimization for latency (noise from progress)
 - No standard support for thread-multiple (multiple entities progressing the same requests)
 - Need for remote information when progressing messages (aka matching)

MPI Matching



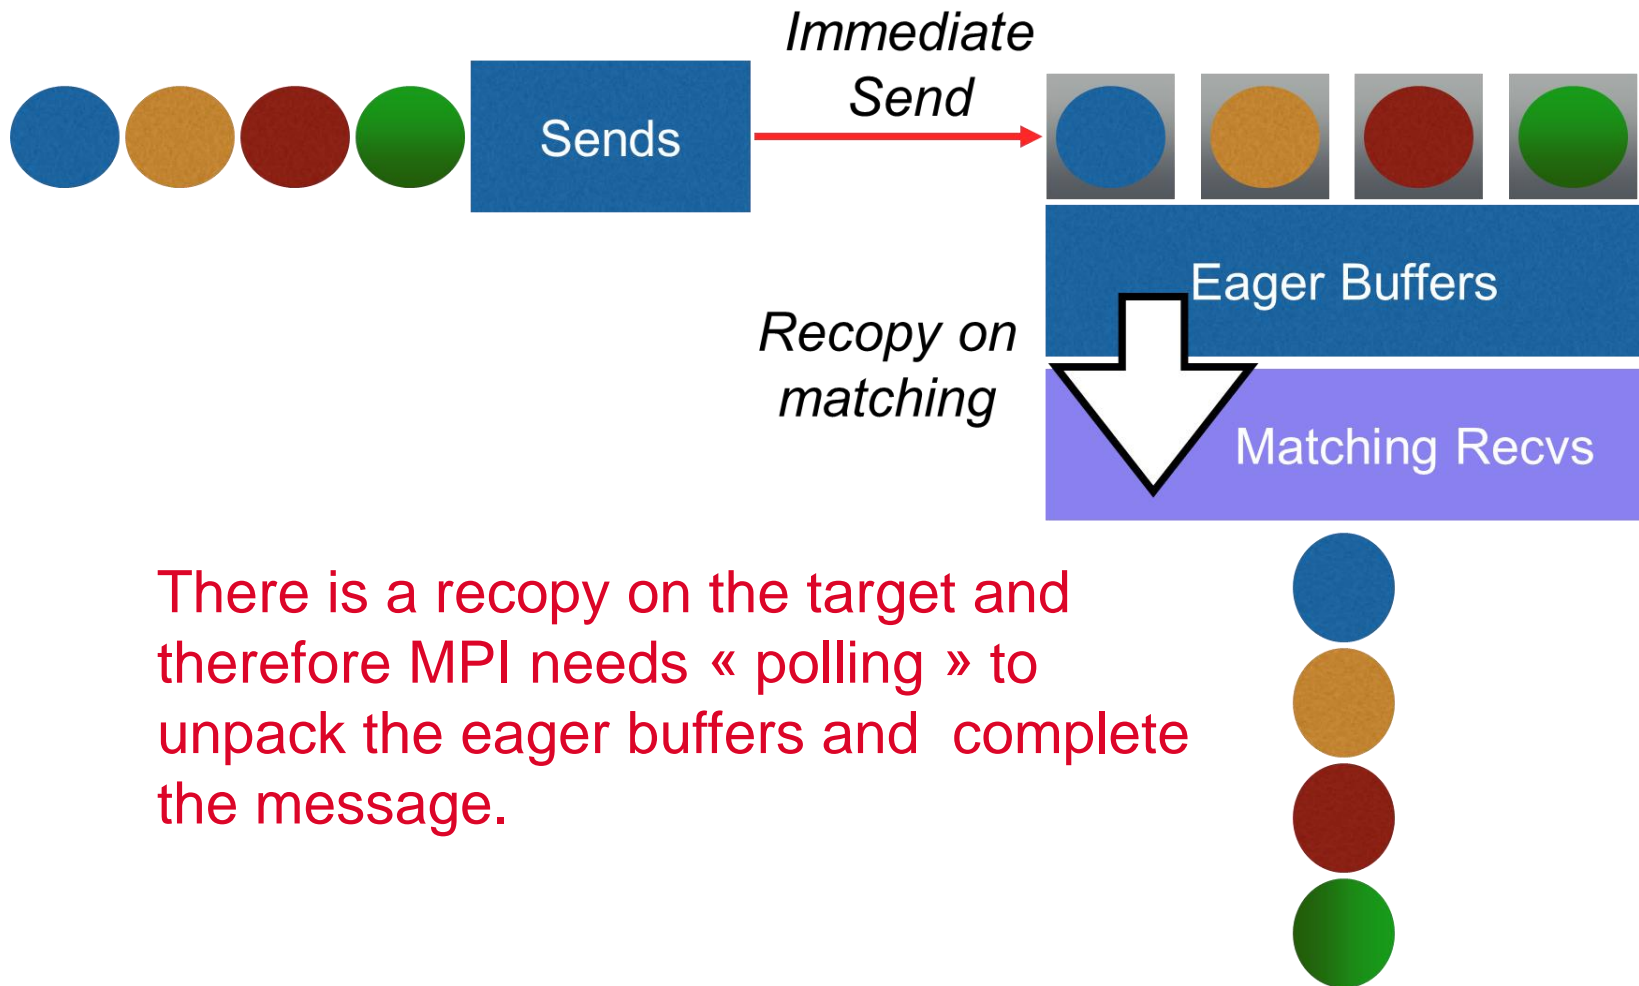
- Before copying any data, MPI has to make sure that the source and target buffers are correctly identified

- In the continuity of their corresponding Send and Recv.



- It makes « single copy » (aka « zero-copy » or « zero-recopy ») messages more difficult to implement.
- In practice, messages involve some recopy overhead:
- Latency optimized eager (recopied on target)
- Rendez-vous protocol requires a previous message to initiate the synchronization.

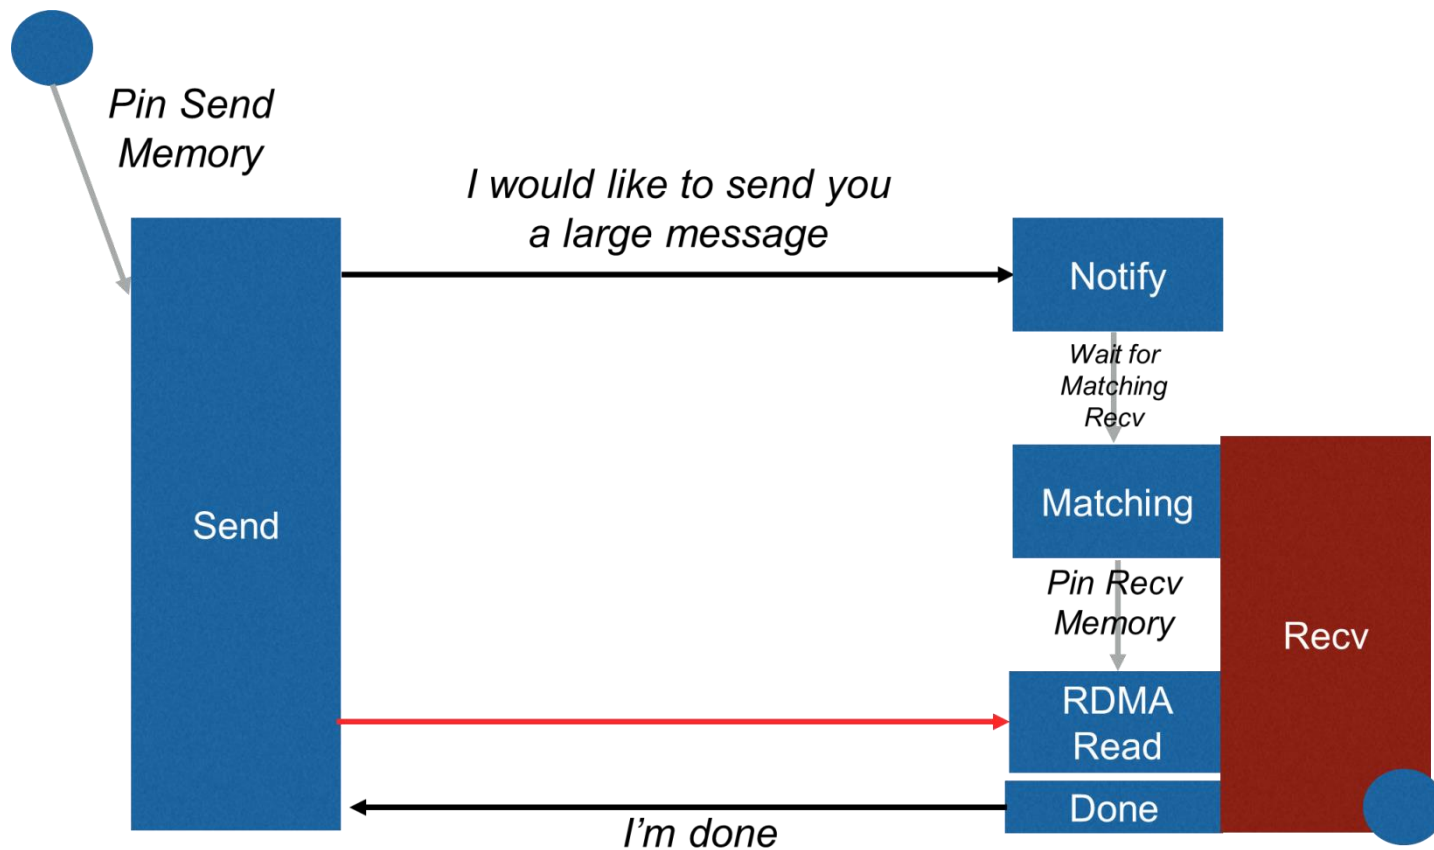
Eager Messages



- There is a recopy on the target and therefore MPI needs « polling » to unpack the eager buffers and complete the message.

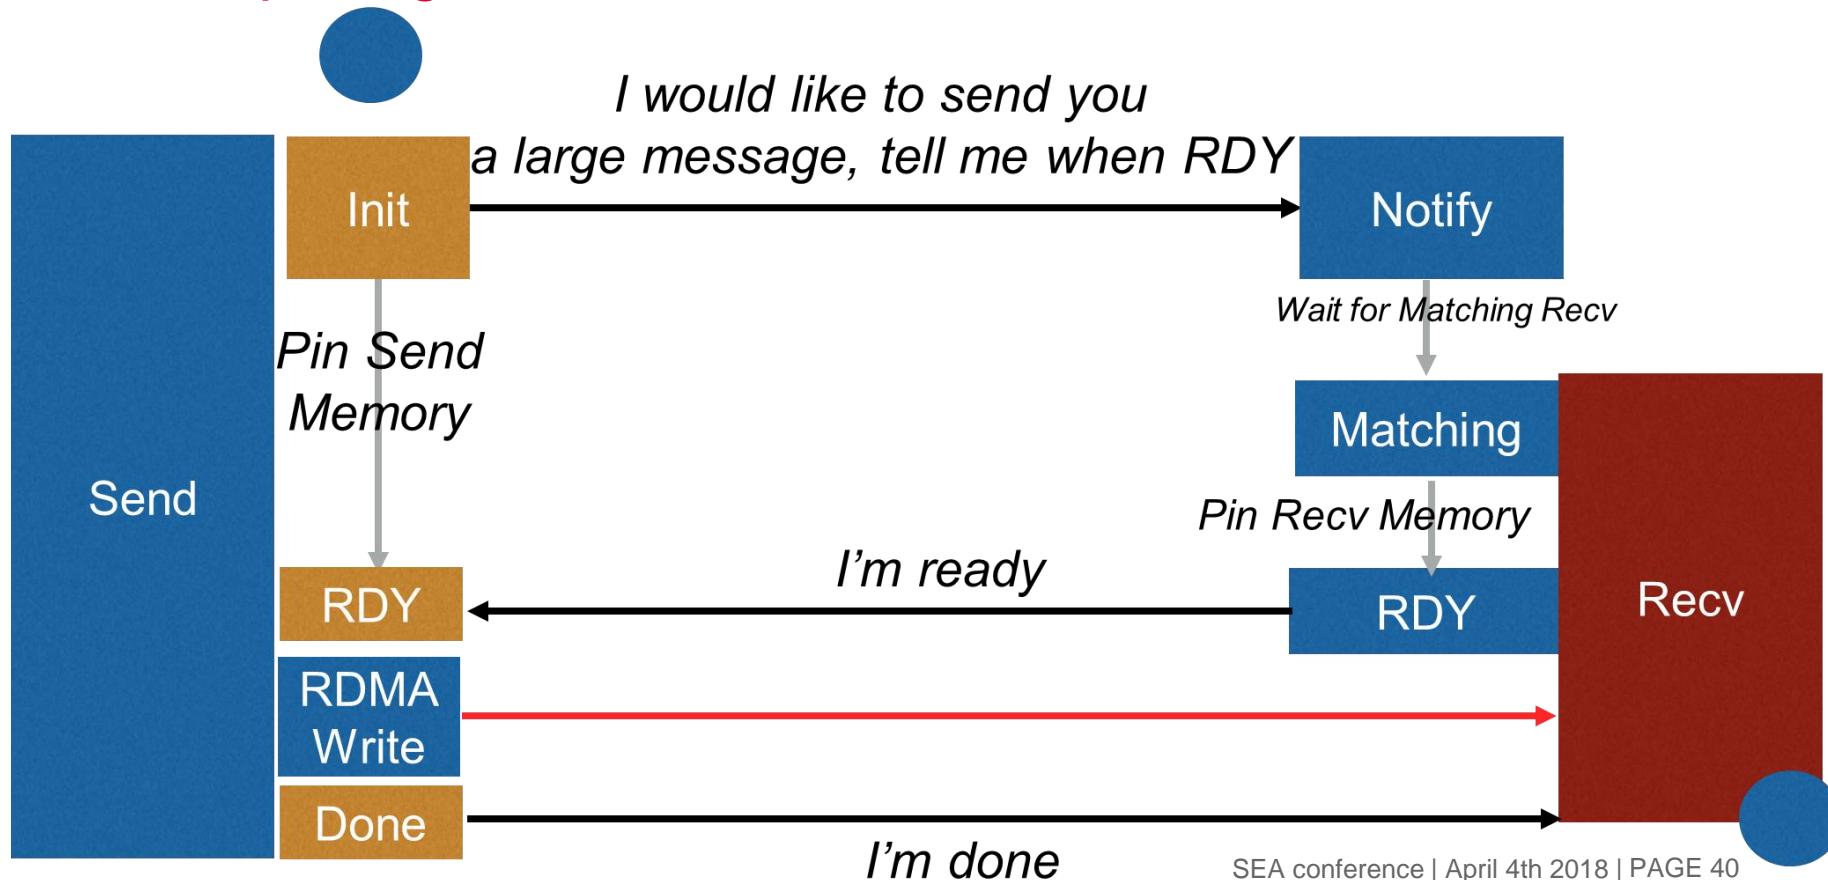
Rendez-Vous Read Example

- Payload movement is « zero-copy » but pinning and meta-data exchange to setup the rendez-vous needs polling.



Rendez-Vous Read Example

- Payload movement is « zero-copy » but pinning and meta-data exchange to setup the rendez-vous needs polling.




The progression issue

- How to progress messages?

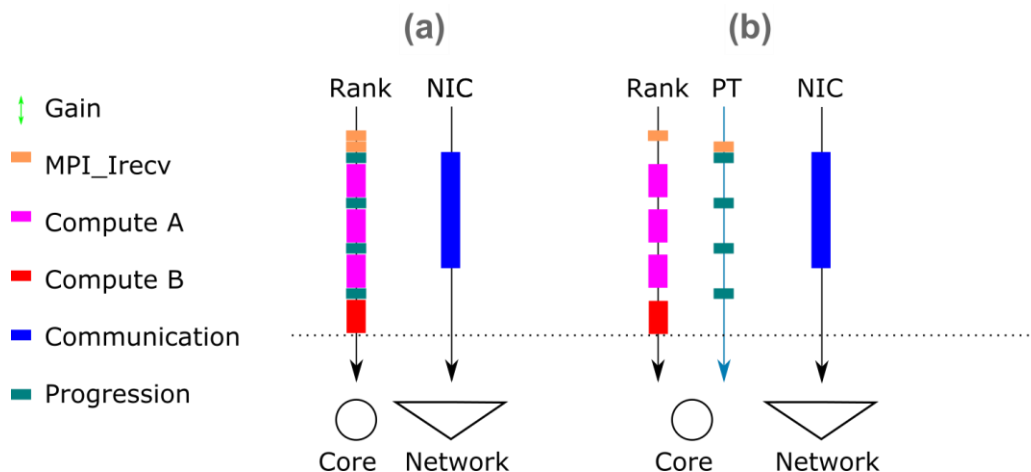
The progression issue

- How to progress messages?
- a) on the same thread
 - need to call functions to progress the message



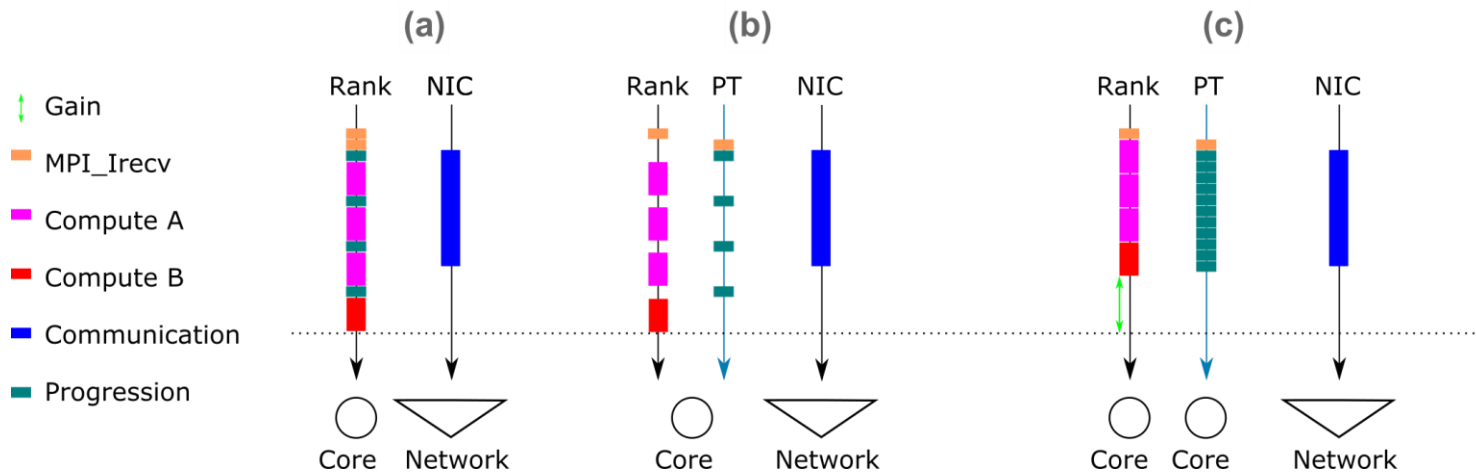
The progression issue

- **How to progress messages?**
 - a) on the same thread
 - need to call functions to progress the message
 - b) on another thread
 - only progression on thread but competition for resources



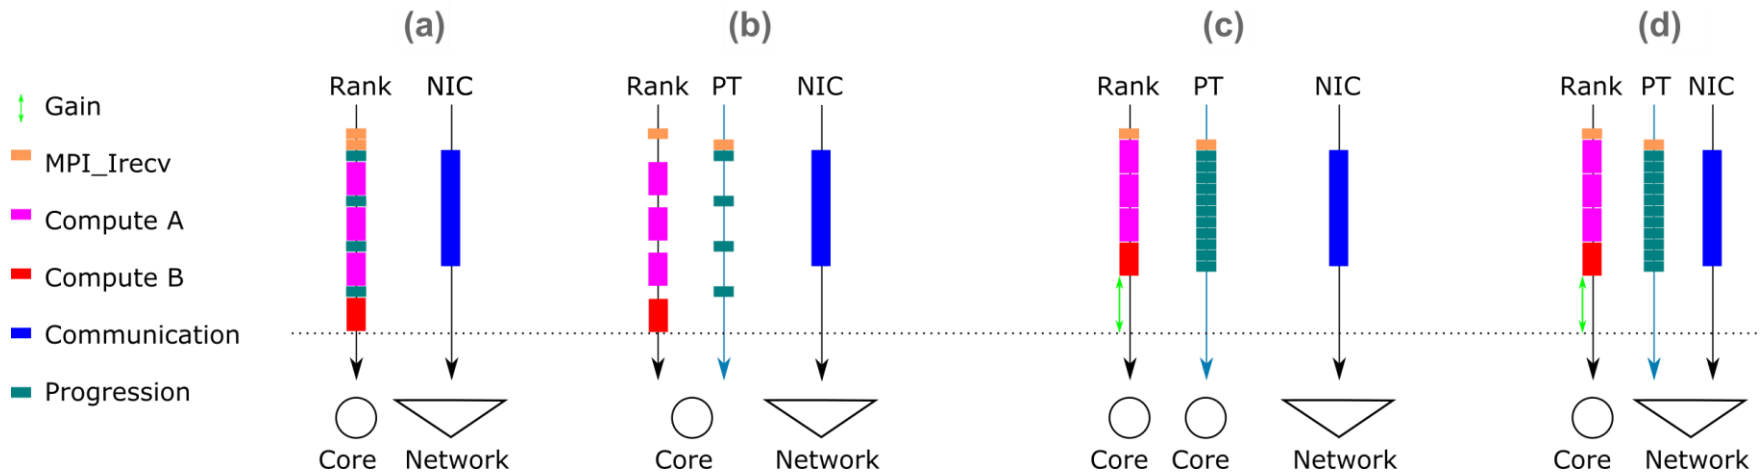
The progression issue

- **How to progress messages?**
 - a) on the same thread
 - need to call functions to progress the message
 - b) on another thread
 - only progression on thread but competition for resources
 - c) on another core
 - start and finish communications as soon as possible



The progression issue

- **How to progress messages?**
 - a) on the same thread
 - need to call functions to progress the message
 - b) on another thread
 - only progression on thread but competition for resources
 - c) on another core
 - start and finish communications as soon as possible
 - d) on the network card
 - if the functionality is available



MPI Matching and Overlap

- Matching prevents modern HPC network cards to fully express their potential
 - Requires information only available on the target process
- A software mechanism has to be involved to associate two MPI buffers
 - « meta-data » have to be exchanged prior to moving the actual message data.
 - This is (one of) the reasons why MPI needs to poll its message queues to enable efficient progress.
- A possible solution to this is to implement MPI matching in the hardware
 - As we are now going to present in the context of the Portals 4 message layer.

- **Portals 4**
 - A high-performance messaging interface
 - Developed by Sandia National Laboratories since 2017.
 - Provides enhanced semantics when compared to low-level Verbs (HW Matching) and features enabling the implementation of PGAS languages (such as Put notification).



The Bull Exascale Interconnect (BXI)

- The Bull Exascale Interconnect is the first implementation of the Portals 4 communication model.



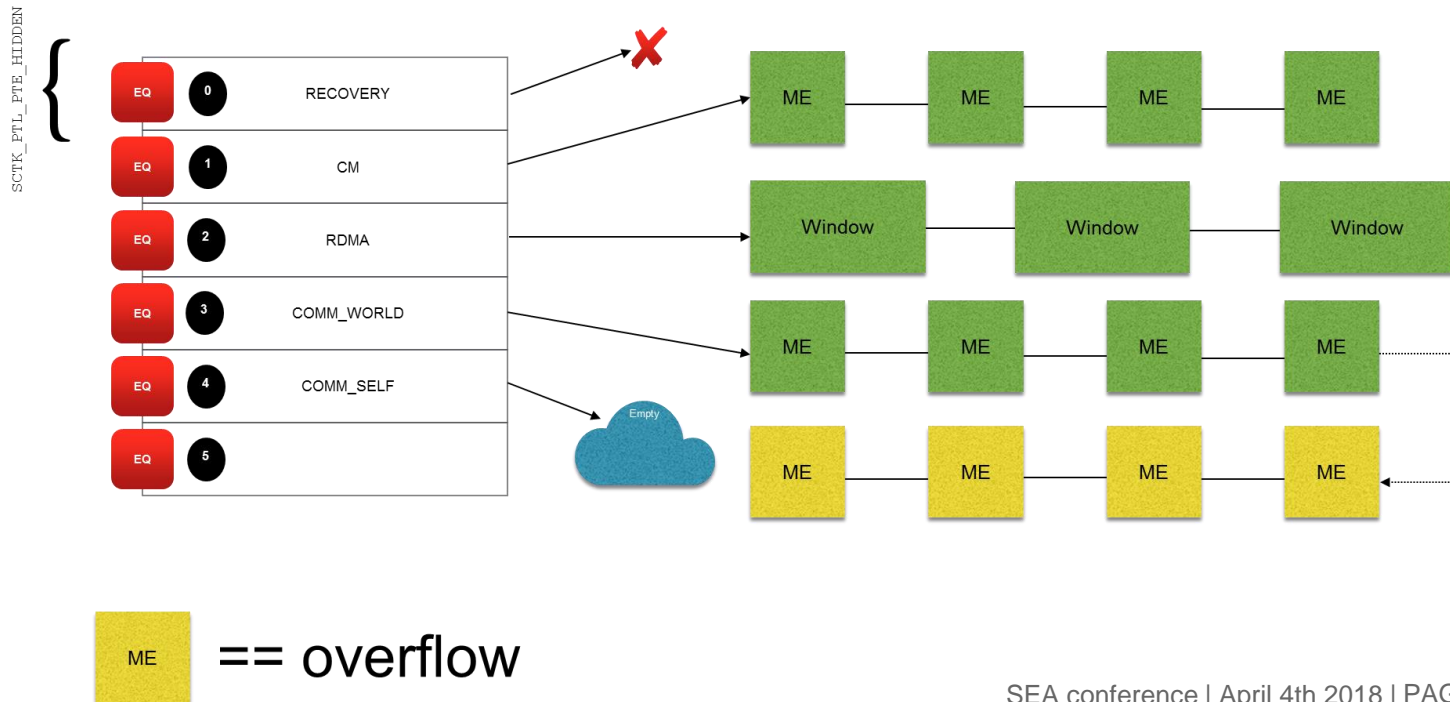
The BXI card




A BXI Switch

Portals 4 in the MPC Runtime

- The Portals 4 communication model with the help of the BXI has been integrated to the MPC runtime.
- To do so MPC relied on a mapping of communicators (round-robin) to the Portal 4 table
- ME=Matching List Entry




- The 64 bits matchbit is handled as follows in MPC:
 - No need for communicator as it is mapped to the table entry (currently limits communicator number to table size)
 - 32 bits for the tag
 - 16 bits for the source rank
 - 16 bits for the sequence number



- It allows a full offload of the matching in the HCA
 - Matching a Send and a Recv without requiring CPU arbitration.

Hardware Matching

- A key advantage of Portals 4's design is that it allows the matching to take place in the hardware.
- It is therefore possible to route a message end-to-end in an asynchronous manner.
- MD = Memory Descriptor



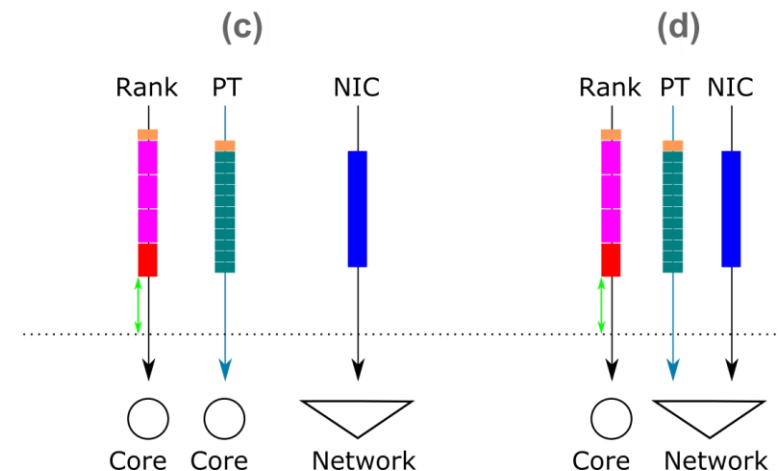

- Target message address is resolved within the hardware, allowing the put event to **directly** target the user-space MPI buffer.
- All steps involved in progressing the message are done by the HCA.

TOWARDS EFFICIENT HARDWARE COLLECTIVE OFFLOAD

Offloading collectives

- In addition to message matching, it will also be possible to offload collective communications on the network card
 - All point-to-point communications involved in the collective
 - Also handling the unfolding of the collective algorithm for all MPI processes
- Problems that may occur
 - Collective algorithm for all MPI processes are handle by one core
 - Not only communications but also compute power is now required and may take time on the core
 - May infer artificial serialization/ordering between messages

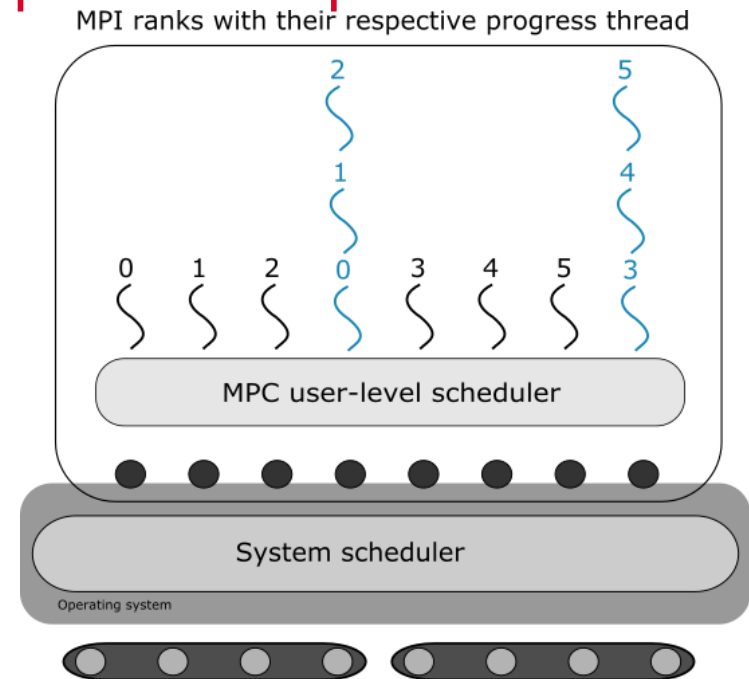
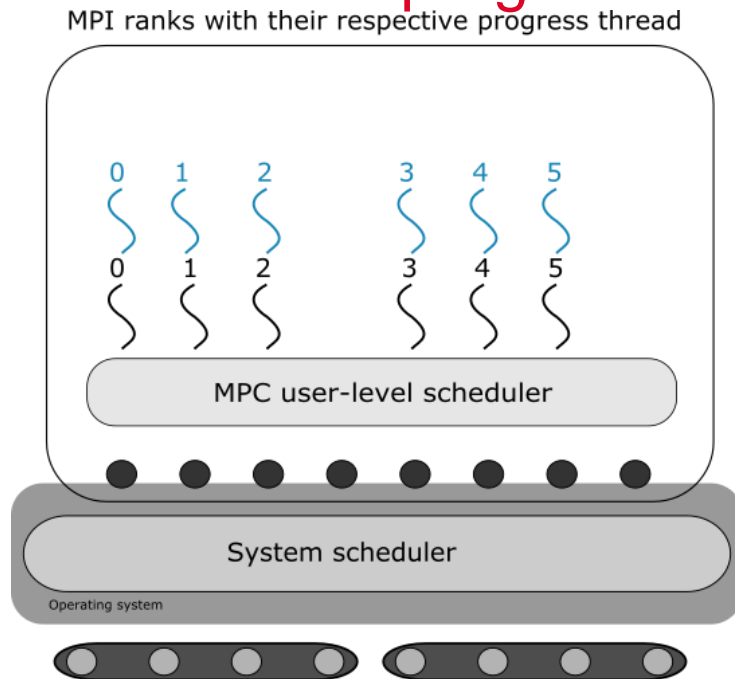
- To simulate this behavior, we offloaded collective handling to reserved core
- Seems similar to collective messages being grouped on network card processor



- What happens when messages are folded on a small number of reserved core

Dedicated resources for progress threads

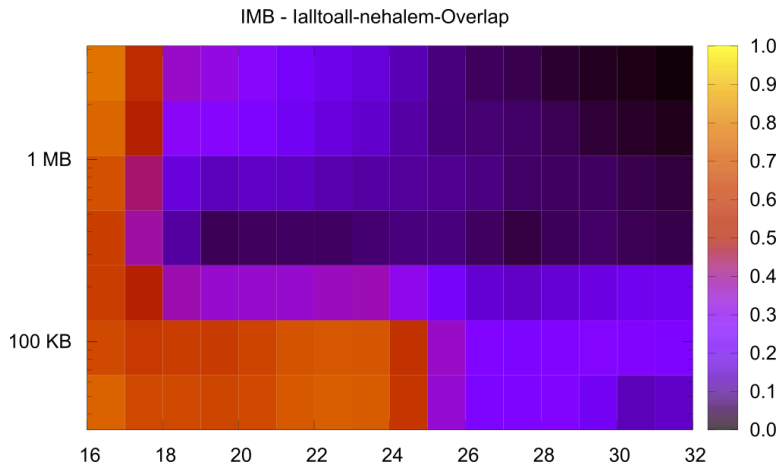
- MPC thread-based implementation allowed us to try different progress thread placement policies



- Left: threads are bound to the same core (can be another hyperthread)
- Right: threads are bound to the closest available core

Dedicated resources for progress threads

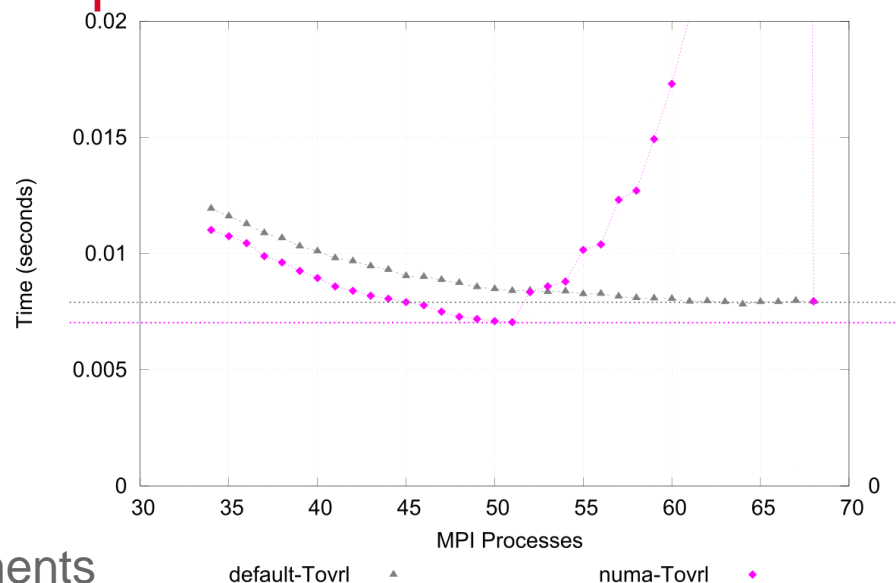

- Having dedicated resources increase overlap ratio, and may also provide better performance



- Left: overlap with different placements
 - 16: best - all progress threads have their own cores
 - 24: good compromise – 3 progress threads per available core
 - >24: some progress threads are bound to the core of their MPI processes

Dedicated resources for progress threads


- Having dedicated resources increase overlap ratio, and may also provide better performance



- Left: overlap with different placements
 - 16: best - all progress threads have their own cores
 - 24: good compromise – 3 progress threads per available core
 - >24: some progress threads are bound to the core of their MPI processes
- Right: alltoall performance on KNL
 - Grey: regular alltoall – Purple: ialltoall with thread placement according to left
 - 3 PTs per core provides overall best performances




Dedicated resources for progress threads

- Huge impact when all communications are folded on few cores
- Especially for intra-nodes comm
- Similar to what will happen when collectives will be offloaded to NIC
- May still not be a viable solution for full MPI or MPI+MPI applications




A solution: split-tree to reduce PTs cumul overhead

- Do the “heavy” parts of the tree on the MPI processes
- Try to reduce the overhead of having numerous PTs on few cores



Collective offload: hardware + software support

- To realize collective offload on NIC, hardware support is needed...
- ...but might be not enough to observe performance gain
- Necessity to adapt algorithm to the workload and number of cores available to handle offloaded collectives



CONCLUSION

- In this presentation we have shown various points limiting asynchronous progression...
- ...along with solutions by offloading services to a device (e.g., the BXI and Portals 4 through the MPC framework) to improve overlap
 - Hardware matching limiting the necessity of polling on CPUs
 - Collectives offloading to limit CPU time necessary to unfold the algorithm
- These evolutions in network cards which are bound to provide matching in a more generalized manner should open a lot of opportunities to HPC hardware.
 - But, as always, may also bring new issues to handle properly to obtain best performances.



Commissariat à l'énergie atomique et aux énergies alternatives

T. +33 (0)1 69 26 40 00

Etablissement public à caractère industriel et commercial | RCS Paris B 775 685 019