

Collective Mind

a collaborative curation tool for program optimization

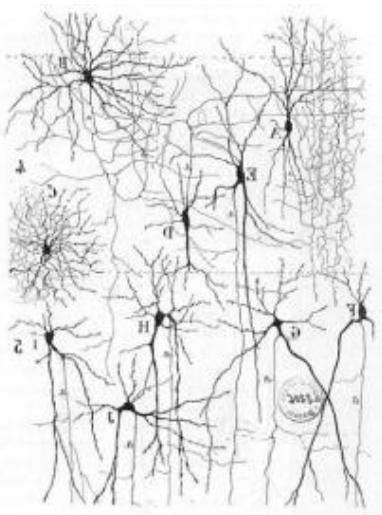
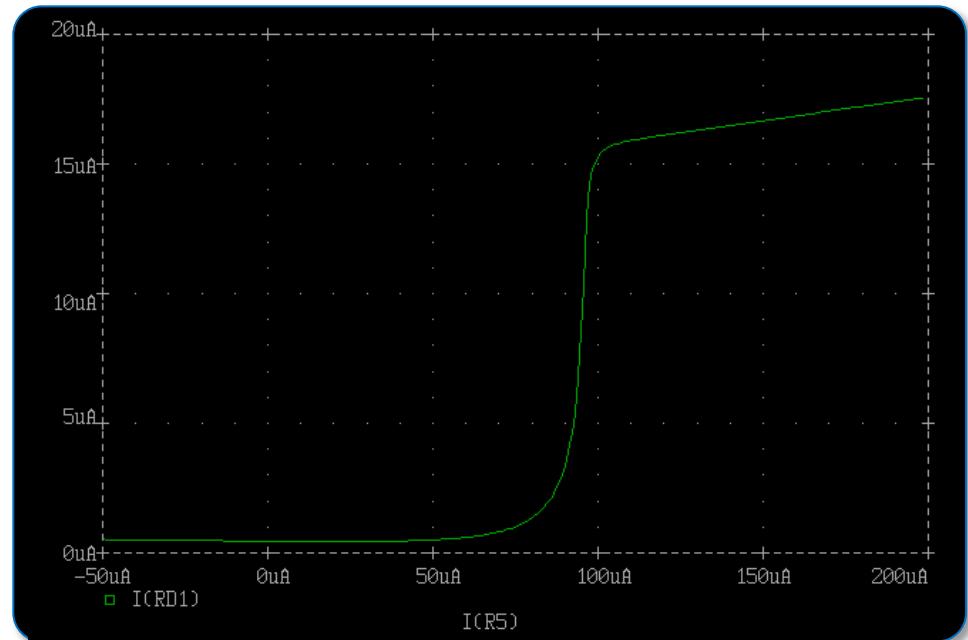
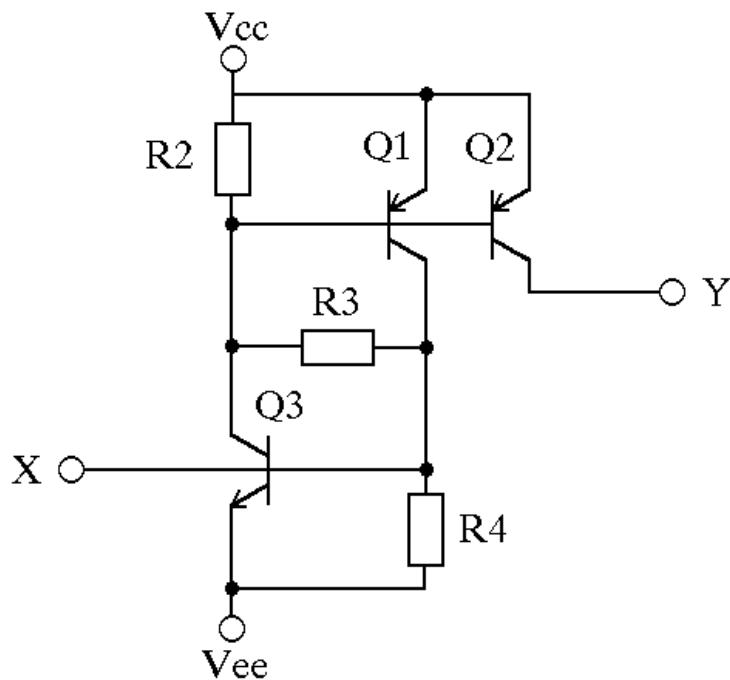
*Towards collaborative, systematic and reproducible
design and optimization of computer systems*

Grigori Fursin
INRIA, France

SEA 2014, Boulder, USA
April 2014

- Interdisciplinary background
- Back to basics: major problems in computer engineering
- Collaborative and Reproducible Research and Development:
 - *cTuning.org framework and repository*
 - *Exposed problems during practical usage*
 - *Collective Mind framework and repository*
 - *Collaborative research usage scenarios*
 - *Collaborative curation for program optimization*
 - *Changing R&D methodology and publication model*
- Conclusions and future work

Back to 1993: physics, electronics, machine learning

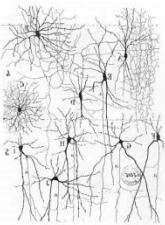


Semiconductor neuron -
base of neural accelerators
and possible future neuro computers
*Modeling and understanding
brain functions*

***My problem
with modeling:***

- Slow
- Unreliable
- Costly

Back to basics: task pipeline (workflow)

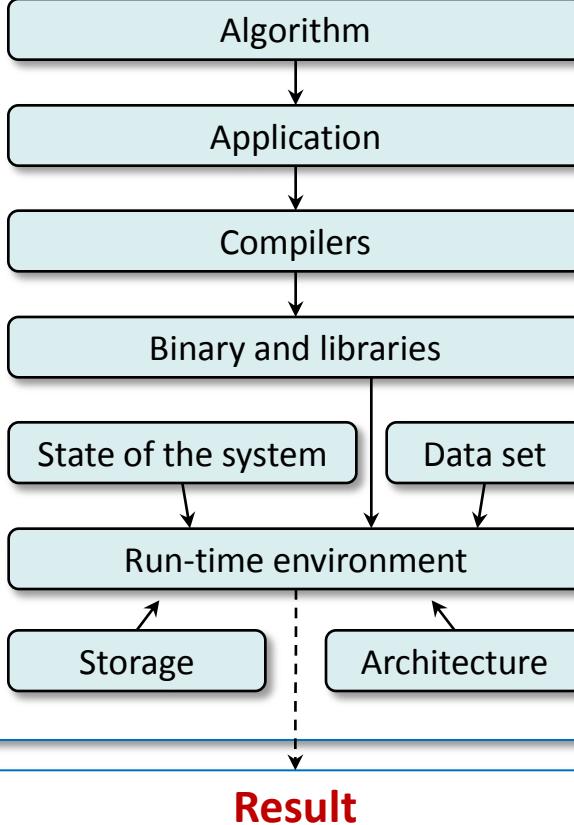


End
User
task



I needed help from computer scientists
to find the best solutions for my tasks!

Available solutions



Service/application providers

(cloud computing, supercomputers, mobile systems)

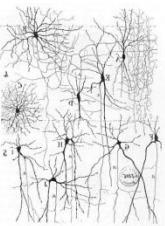
User requirements:

*minimize all costs (characteristics)
(execution time, power consumption, price,
size, faults, etc)*

*guarantee real-time constraints
(bandwidth, QoS, etc)*

Hardware and software designers
(processors, memory, interconnects,
languages, compilers, run-time systems)

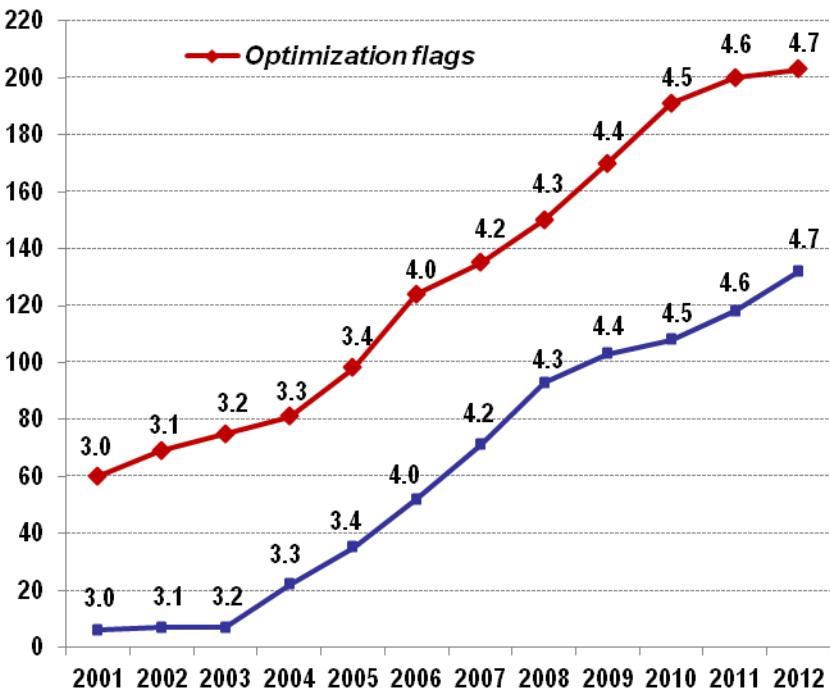
Back to basics: ever rising complexity of computer systems



End
User
task



GCC optimizations



Result

Deliver universal and optimal optimization heuristic is often impossible!

- 1) Too many design and optimization choices at all levels
- 2) Always multi-objective optimization: performance vs compilation time vs code size vs system size vs power consumption vs reliability vs return on investment
- 3) Complex relationship and interactions between ALL software and hardware components
- 4) Users and developers often have to resort to empirical auto-tuning

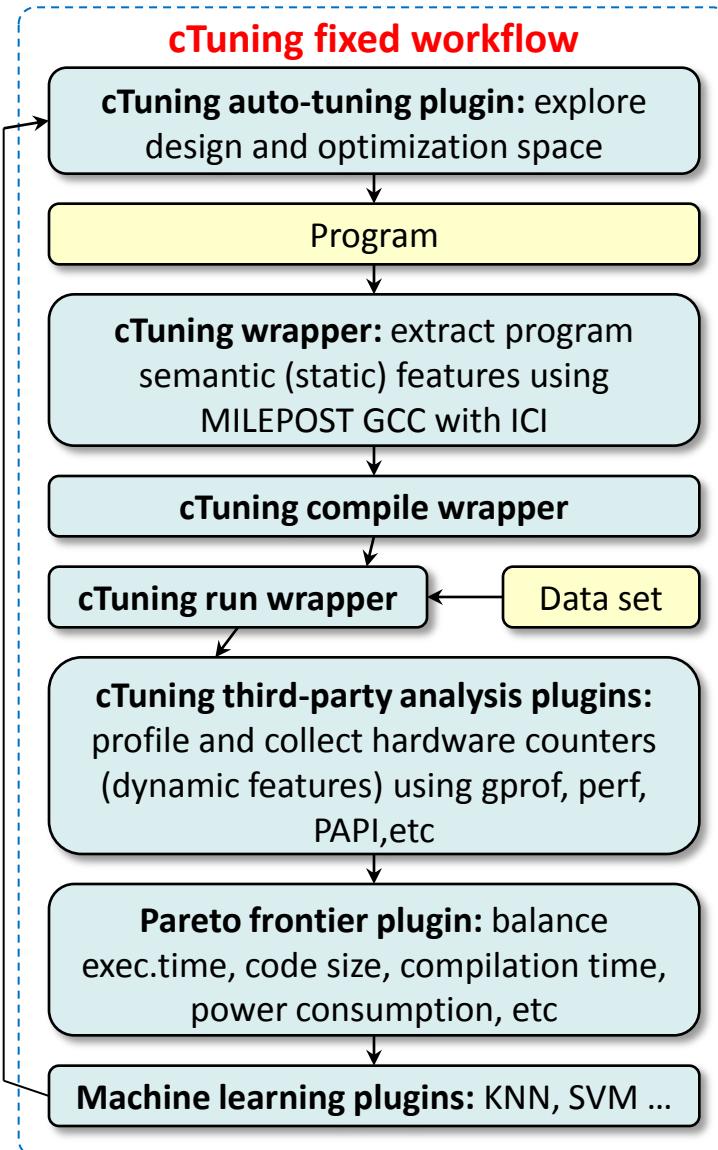
Empirical auto-tuning is too time consuming, ad-hoc and tedious to be a mainstream!

Many end users and service providers run numerous applications on different architectures with different data sets, run-time systems, compilers and optimizations!



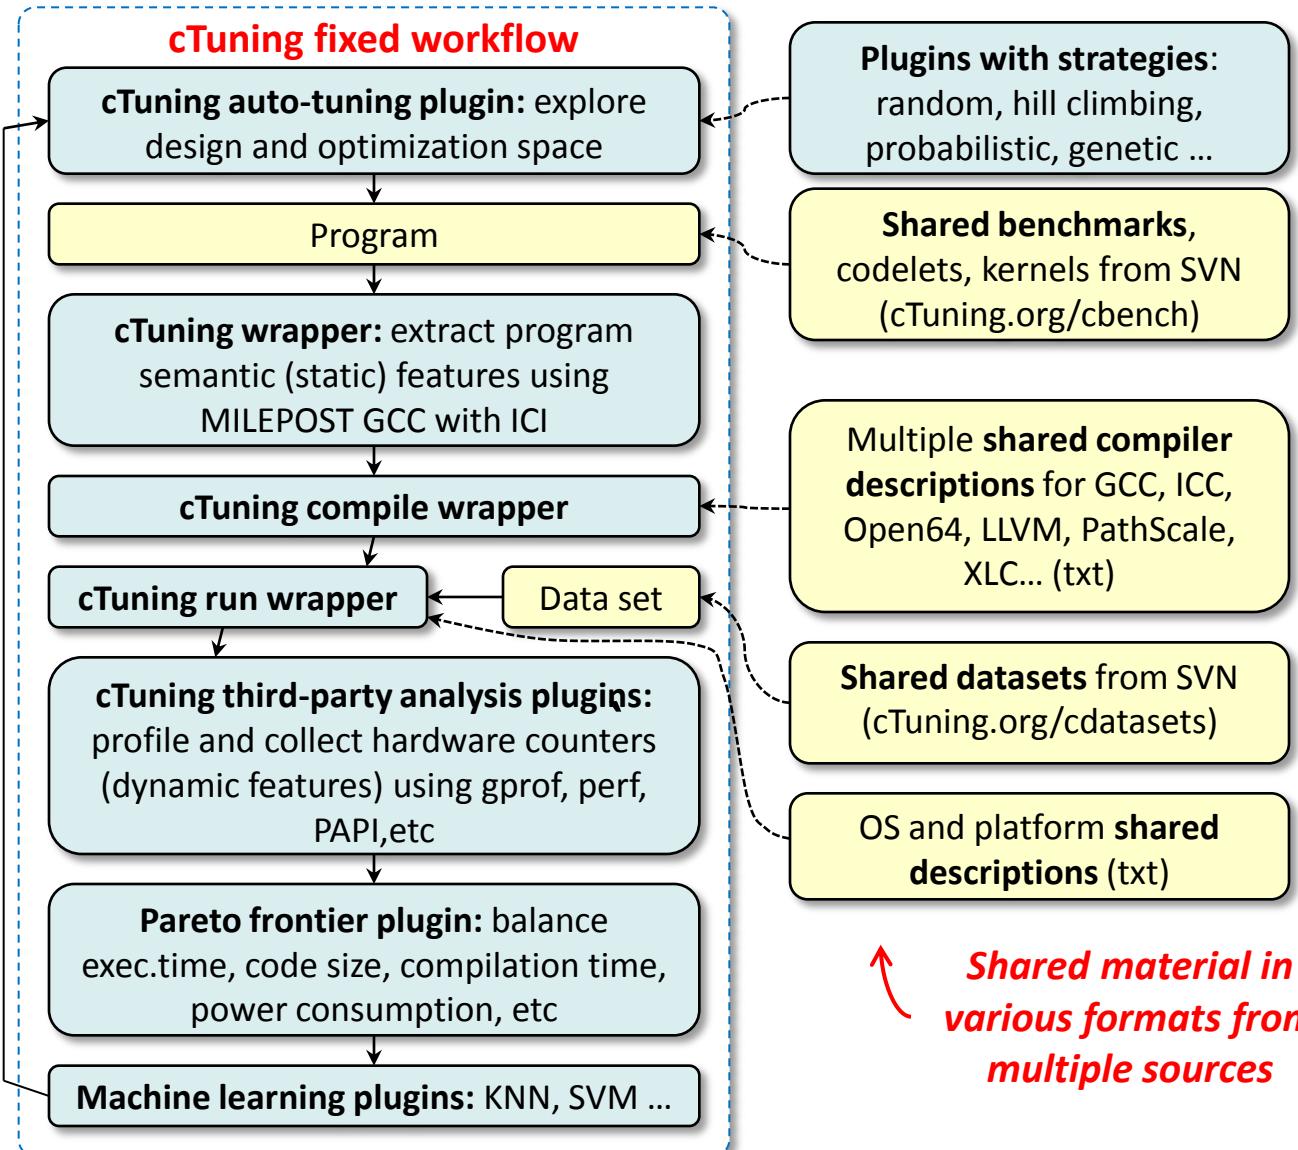
- Can we leverage their experience and computational resources?
- Can we build common research and experimentation infrastructure connected to a public repository of knowledge?
- Can we extrapolate collected knowledge using machine learning to predict optimal program optimizations, hardware designs and run-time adaptation scenarios?

Proposed and developed in the MILEPOST/cTuning project (2006-2009) to automate and crowdsource training of a machine learning based self-tuning compiler for reconfigurable processors (SW/HW co-design).



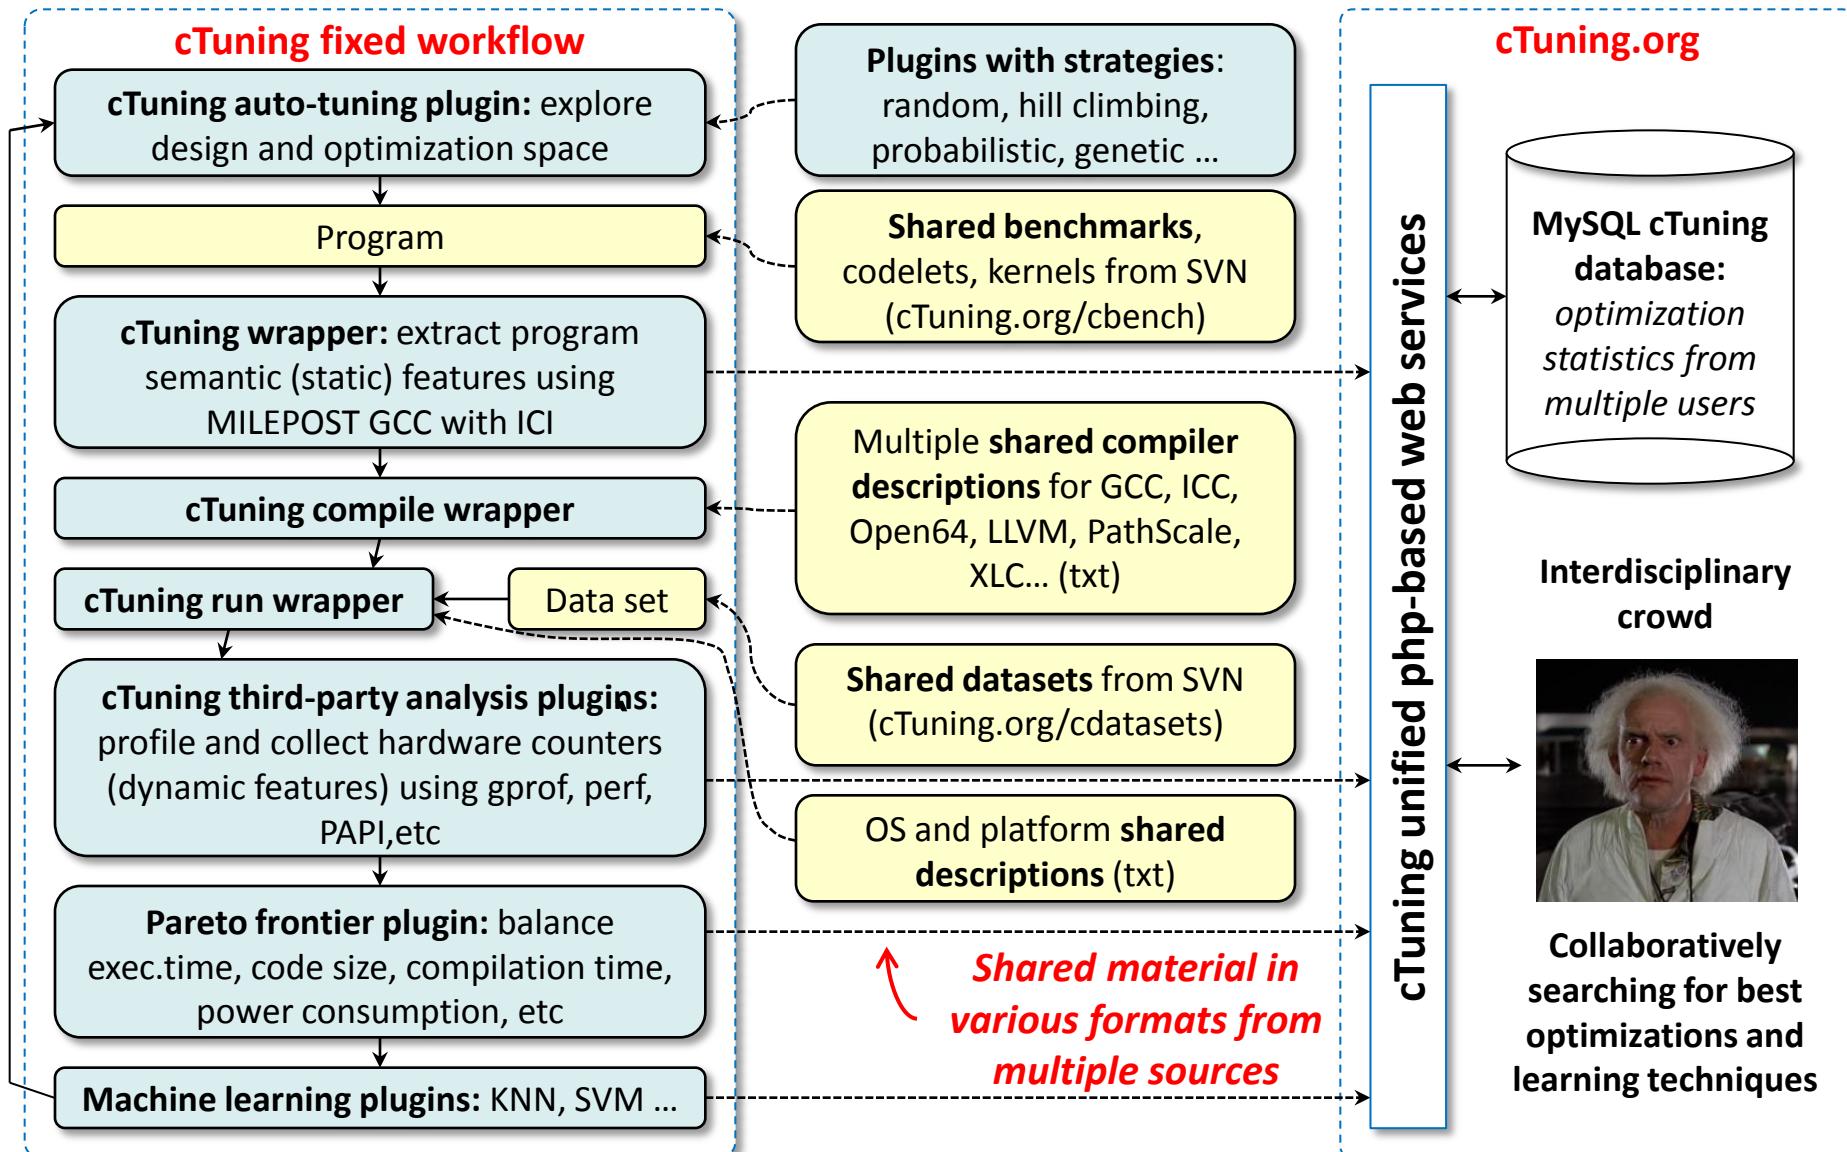
Common open-source cTuning CC framework (<http://cTuning.org/ctuning-cc>)

cTuning.org: plugin-based auto-tuning and learning framework



Common open-source cTuning CC framework (<http://cTuning.org/ctuning-cc>)

cTuning.org: plugin-based auto-tuning and learning framework



Common open-source cTuning CC framework (<http://cTuning.org/ctuning-cc>)

cTuning fixed workflow

cTuning auto-tuning plugin: explore design and optimization space

Plugins with strategies:
random, hill climbing,
probabilistic, genetic ...

cTuning.org

In 2009, we managed to automatically tune customer benchmarks and compiler heuristics for a range of real platforms from IBM and ARC (Synopsis)



Everything is solved?

Machine learning plugins: KNN, SVM ...

multiple sources

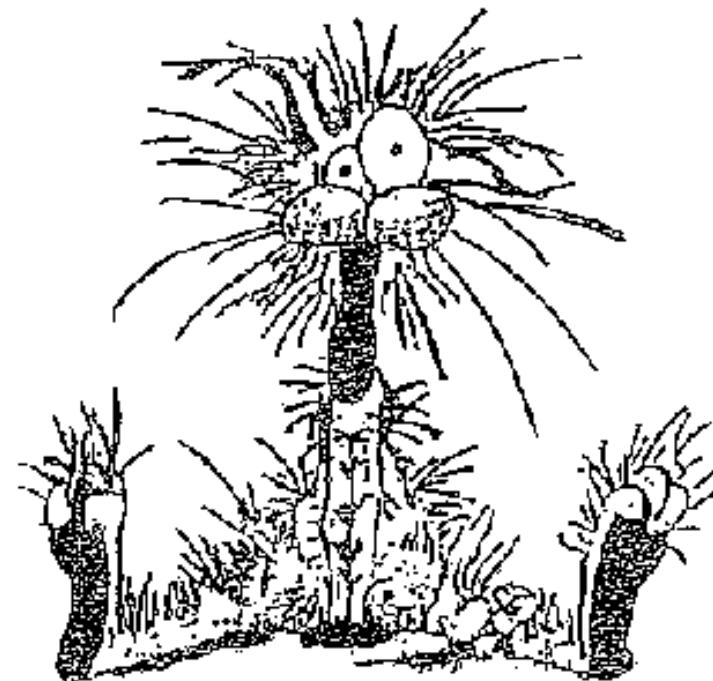
learning techniques

Common open-source cTuning CC framework (<http://cTuning.org/ctuning-cc>)

cTuning practical usage exposed a few problems

Technological chaos

LLVM 3.4	GCC 4.8.x	GCC 4.1.x	AVX	ATLAS
frequency	genetic algorithms	CUDA 5.x	LLVM 2.8	hardware counters
ARM v8	ICC 11.1	function level	reordering	perf
OpenMP	MPI	OpenCL	gprof	
reliability	CUDA 4.x	prof	Intel SandyBridge	
ARM v6	GCC 4.2.x	per phase		
	LLVM 3.0	reconfiguration		
GCC 4.5.x	MVS 2013	HMPP	memory size	
execution time		polyhedral	XLC	
SSE4		transformations	KNN	
MKL	GCC 4.3.x	predictive	ICC 12.0	Scalasca
bandwidth	LLVM 2.6	scheduling	ICC 11.0	
LLVM 2.9	ICC 12.1	SimpleScalar	Codelet	PAPI
process		Testarossa	Amplifier	Jikes
oprofile		HDD size	loop-level	Open64
GCC 4.7.x	LLVM 2.7	SVM	ISA	
threads		LTO	VTune	
Phoenix	program	algorithm-	threads	
cache size	level	level	IPA	TAU
			TBB	
				algorithm precision



**Many thanks to international academic
and industrial users for their feedback**

<http://cTuning.org/lab/people>

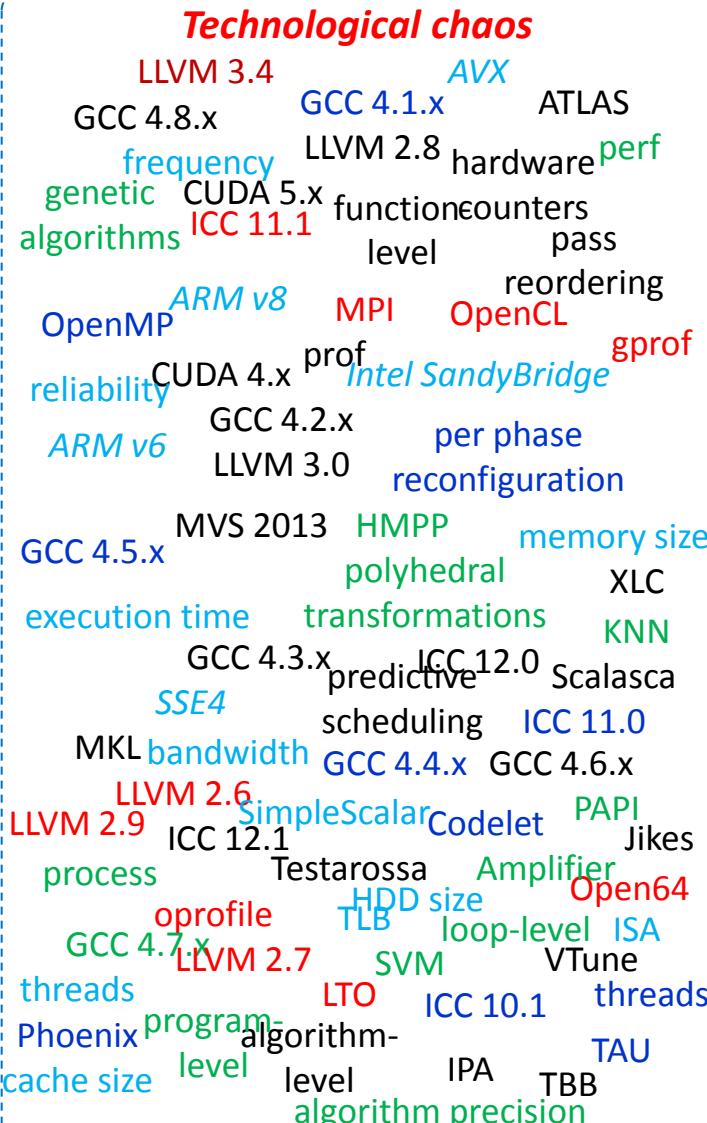
**Thanks to Davide Del Vento
and his interns from UCAR**

cTuning practical usage exposed a few problems

Technological chaos		
LLVM 3.4	GCC 4.8.x	AVX
frequency	GCC 4.1.x	ATLAS
genetic algorithms	LLVM 2.8	hardware perf
CUDA 5.x	function counters	
ICC 11.1	level	pass
ARM v8	reordering	
OpenMP	MPI	OpenCL
reliability	CUDA 4.x	gprof
ARM v6	prof	Intel SandyBridge
GCC 4.2.x		
LLVM 3.0	per phase	
	reconfiguration	
MVS 2013	HMPP	memory size
GCC 4.5.x	polyhedral	XLC
execution time	transformations	KNN
SSE4	GCC 4.3.x	Scalasca
MKL	bandwidth	ICC 12.0
LLVM 2.6	ICC 12.1	predictive
LLVM 2.9	SimpleScalar	Codelet
process	Testarossa	PAPI
oprofile	HDD size	Jikes
GCC 4.7.x	LLVM 2.7	loop-level
threads	LTO	ISA
Phoenix	SVM	VTune
program	IPA	threads
cache size	algorithm level	TAU
level	TBB	
	algorithm	
	precision	

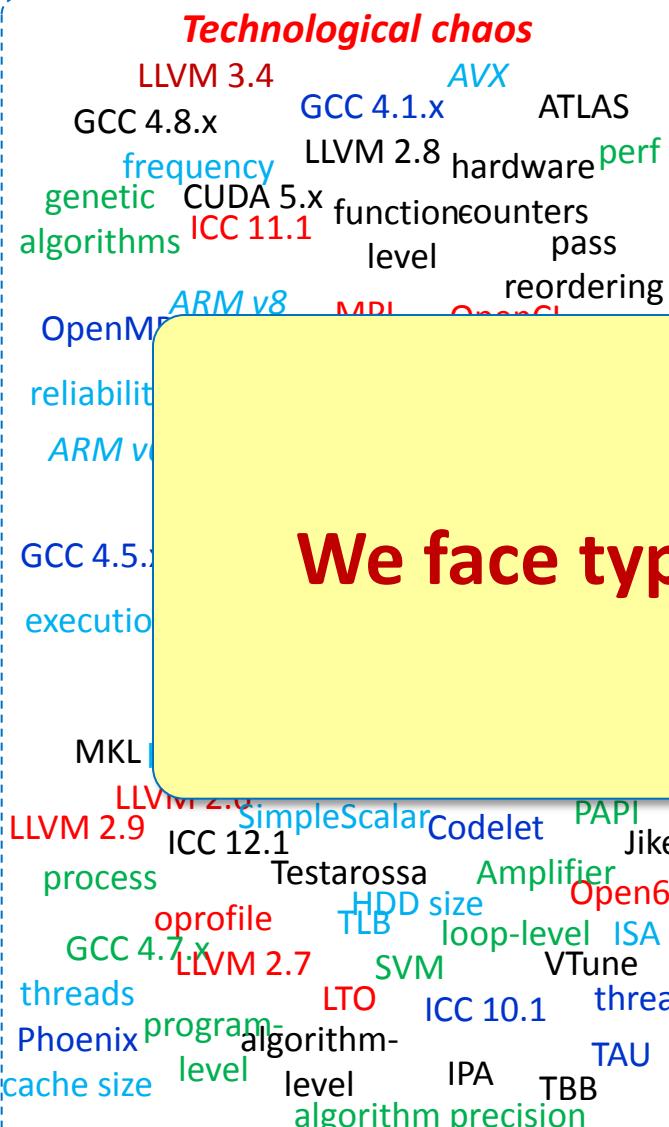
- Ever changing environment
(SW and HW choices, features, tools and libraries versions, interfaces, data formats ...)
 - Difficult to reproduce experimental results from different users
 - Difficult to keep cTuning up-to-date - by the end of development cycle, new software and hardware is often available
- Complex deployment and maintenance
 - Complex, manual configuration of the framework, repository, plugins, hardware and all dependencies
 - Many used languages (C, C++, PHP, Java) and data exchange formats (txt, csv, xls, ...)
 - Different versions of benchmarks, data sets, plugins, models, compilers, third-party tools, etc are scattered around many directories and disks

cTuning practical usage exposed a few problems



- Difficult or impossible to reproduce and reuse techniques from existing publications
 - *No culture of sharing code and data in computer engineering unlike other sciences*
 - *Impossible to publish negative results or just validation of existing works*
 - *Academia focuses more and more on quantity (number of publications and citations) rather than quality and reproducibility of results*
- No simple and scalable mechanism to preserve, share and reuse all heterogeneous experimental material
 - *Centralized, MySQL-based repository got quickly saturated when collecting and processing data from many users*
 - *MySQL repository is difficult to adapt to continuous changes in the system*
 - *Difficult to preserve different versions of benchmarks, data sets, plugins and tools, and make them co-exist using SVN repositories*

cTuning practical usage exposed a few problems



- Difficult or impossible to reproduce and reuse techniques from existing publications
 - *No culture of sharing code and data in computer engineering unlike other sciences*
 - *Impossible to publish negative results or just validation of existing works*

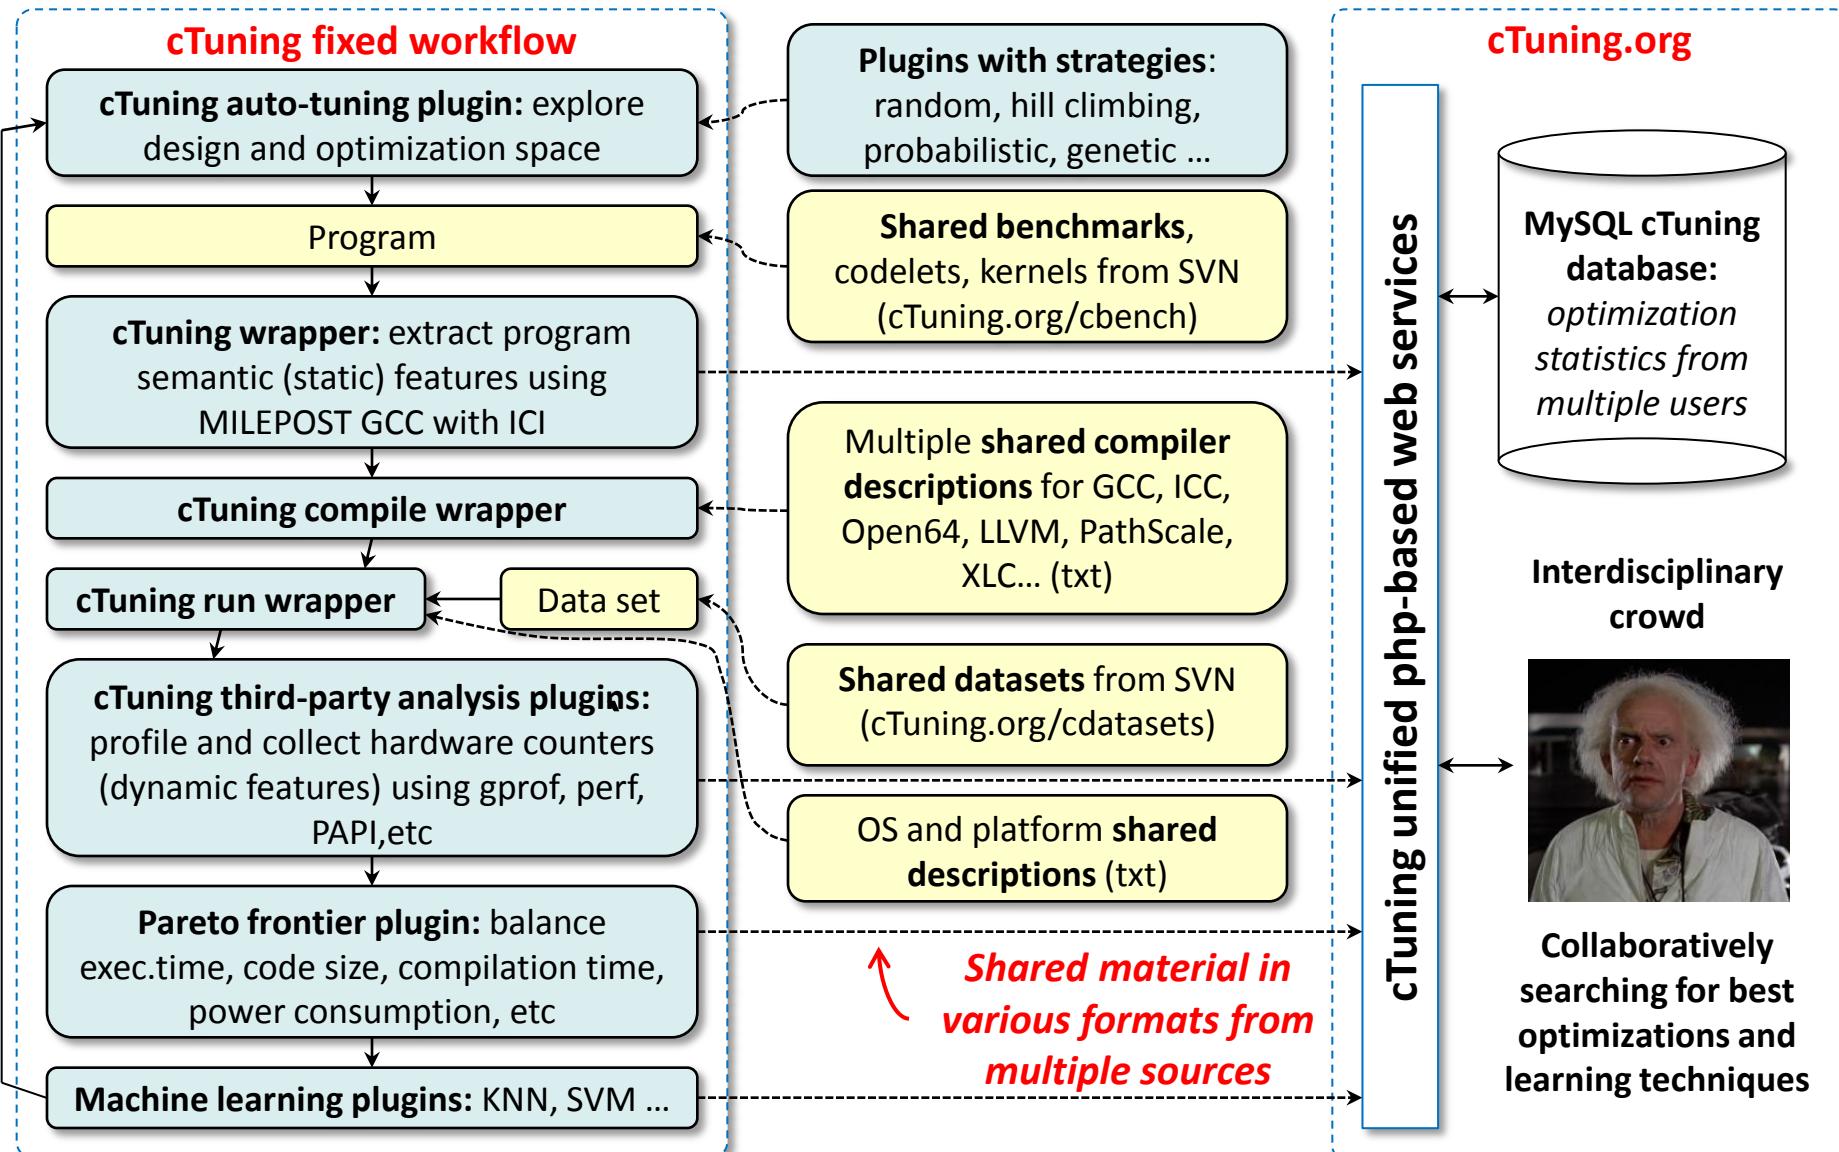
quantity
) rather
ults
serve,
mental

got

quickly saturated when collecting and processing data from many users

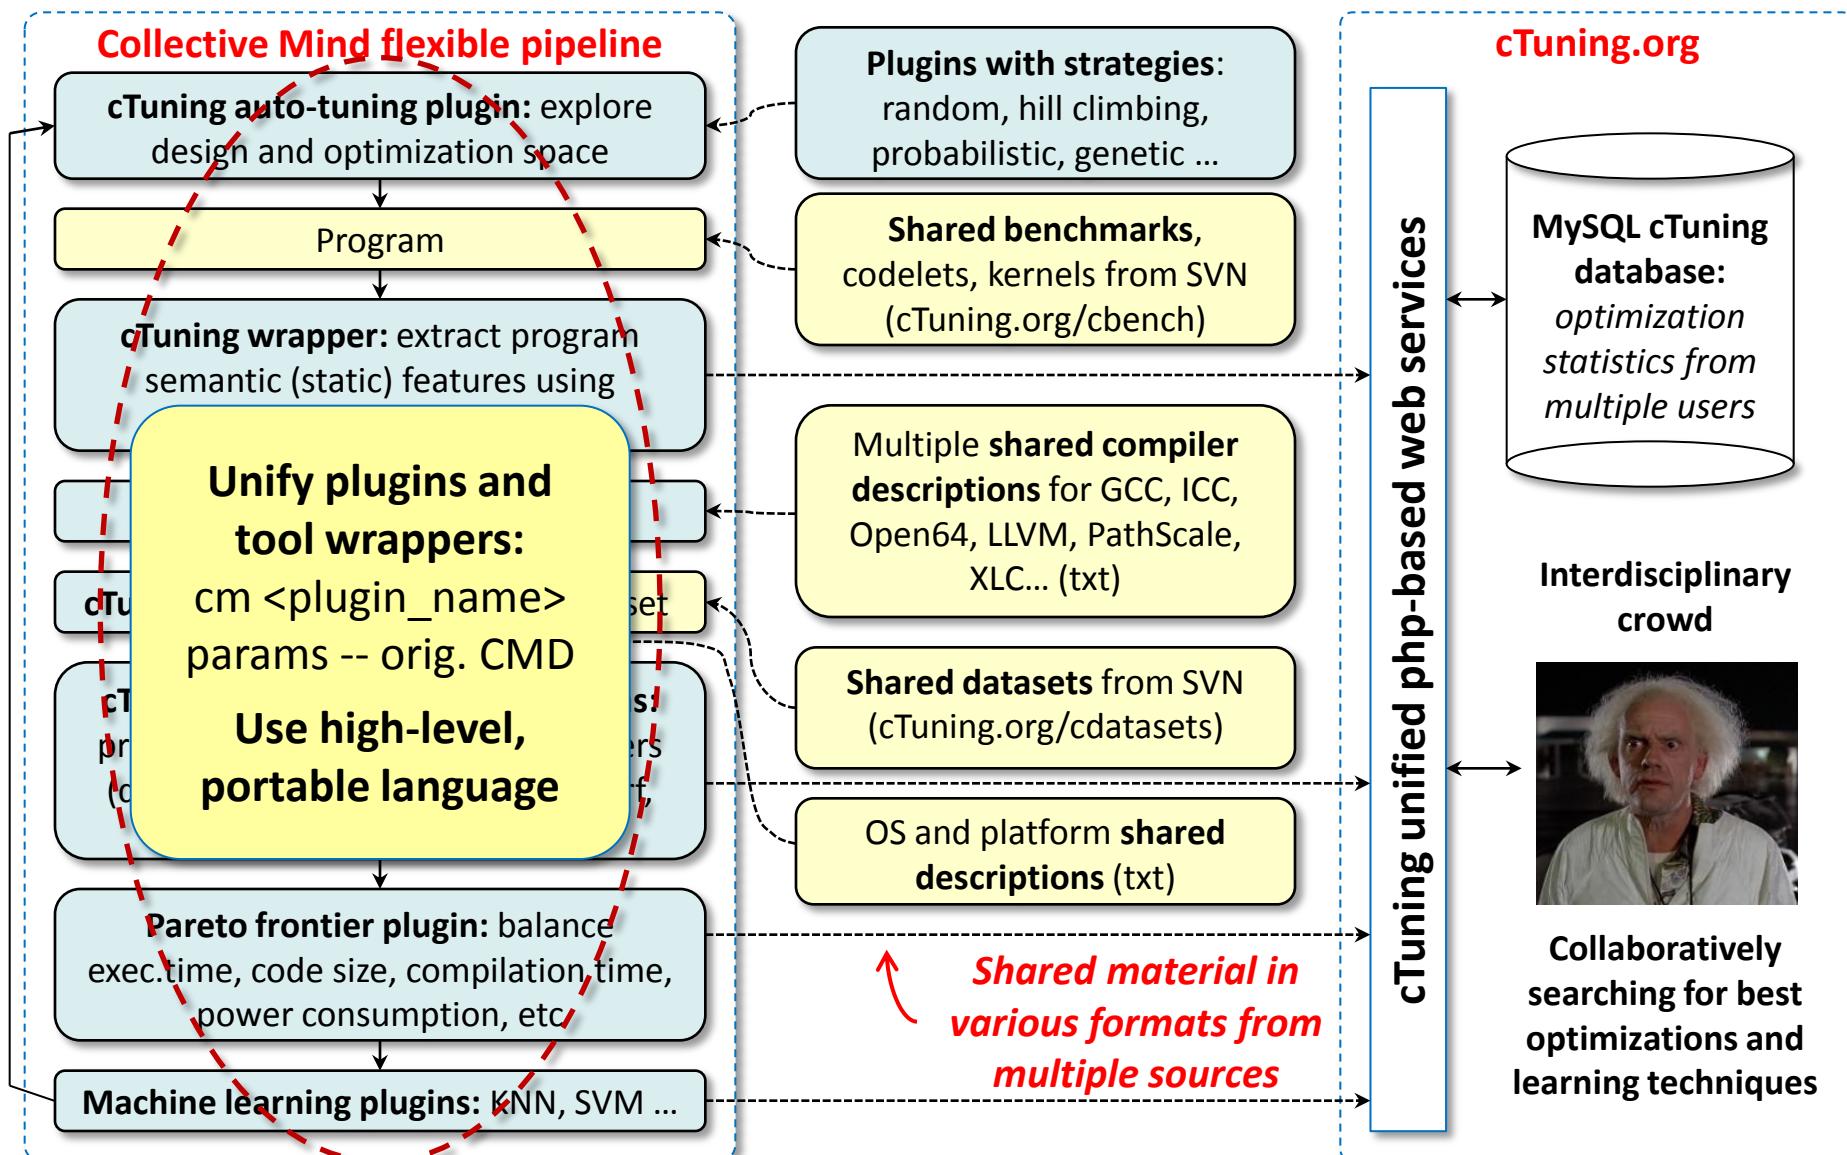
- MySQL repository is difficult to adapt to continuous changes in the system
- Difficult to preserve different versions of benchmarks, data sets, plugins and tools, and make them co-exist using SVN repositories

Collective Mind: attempt to fix cTuning problems



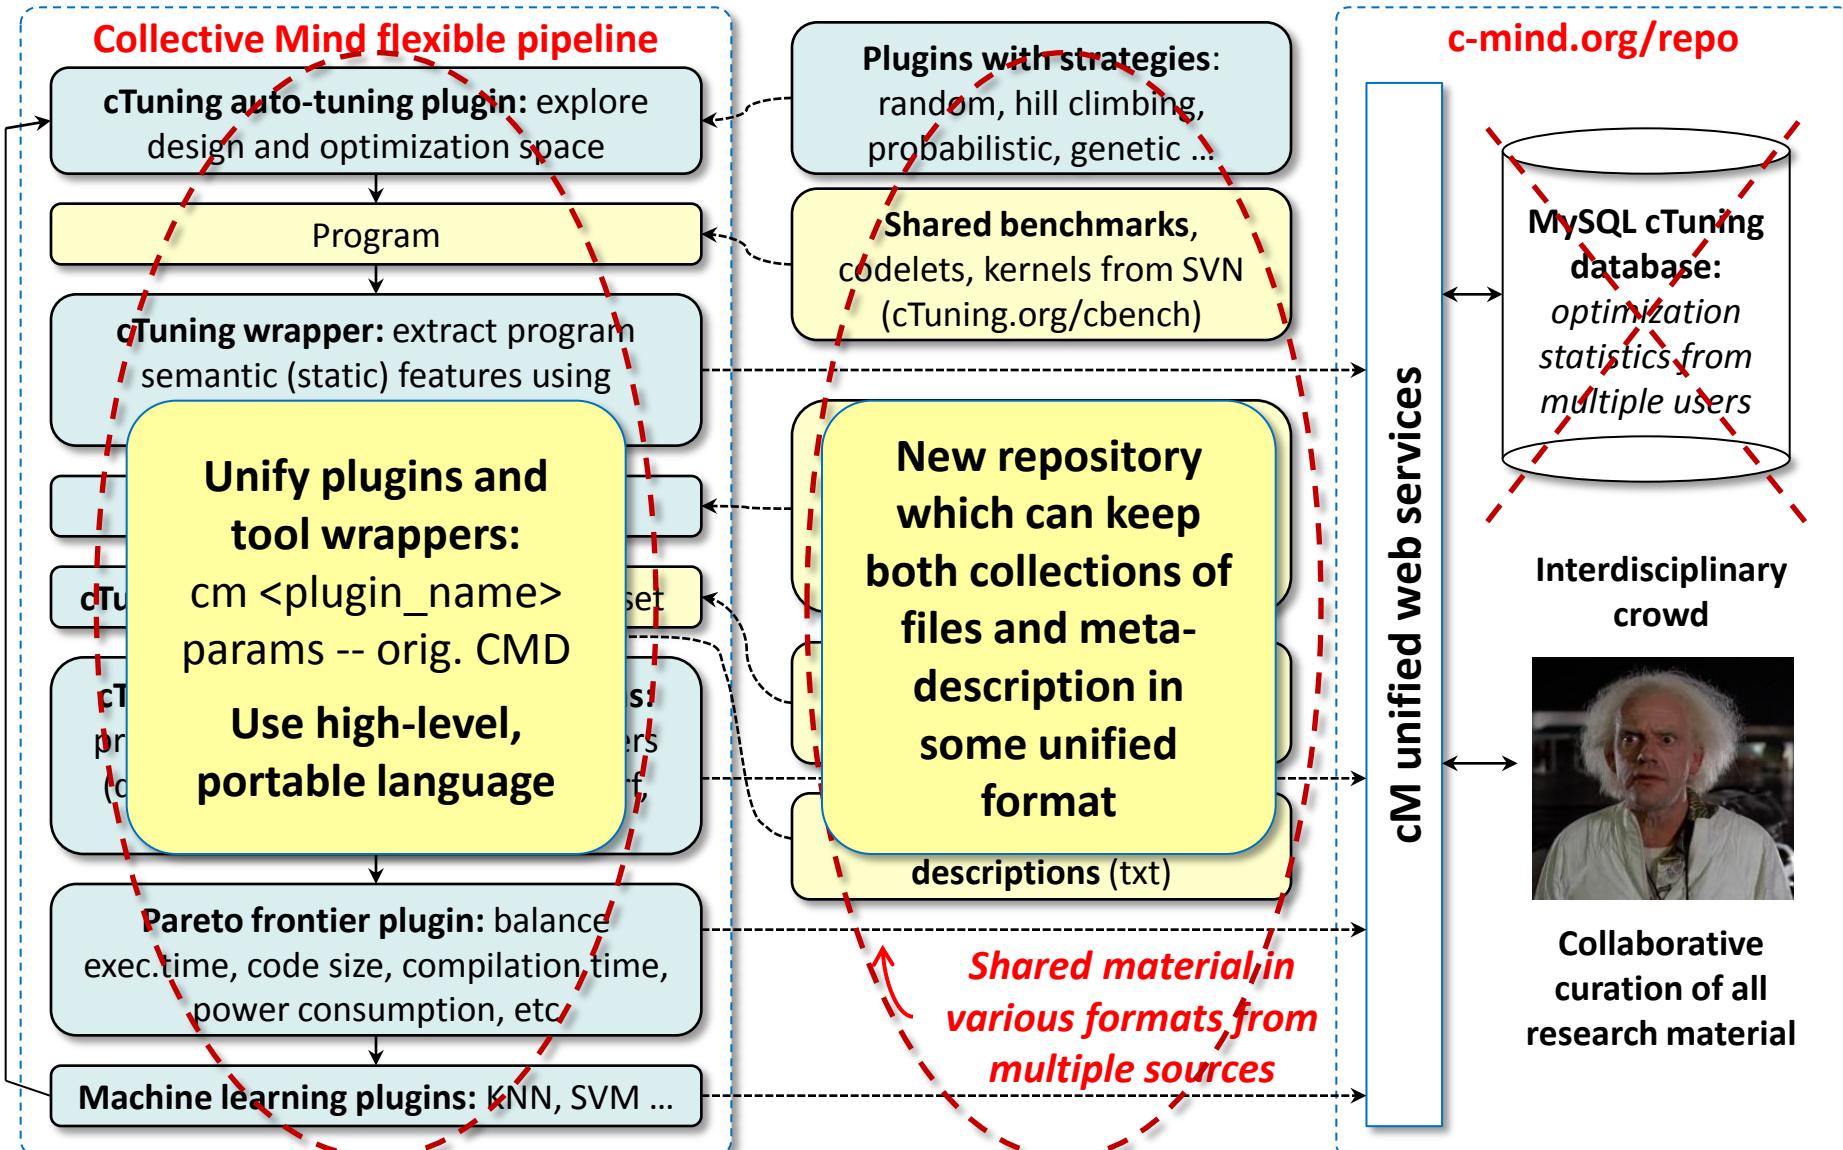
Common open-source Collective Mind framework (<http://c-mind.org>)

Collective Mind: attempt to fix cTuning problems



Common open-source Collective Mind framework (<http://c-mind.org>)

Collective Mind: attempt to fix cTuning problems



Collective Mind: which technology to use?

JSON: extensible and human-readable format

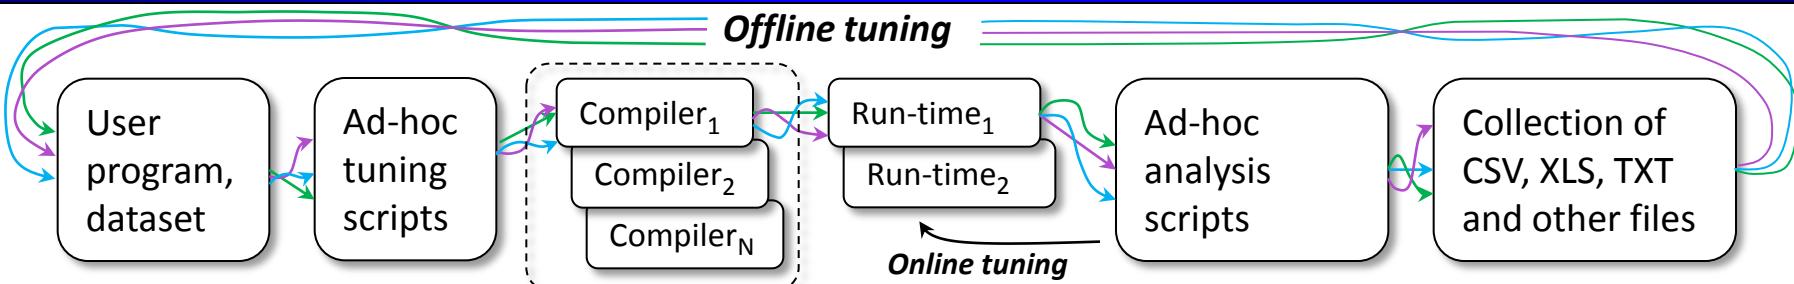
```
file data.json = {  
    "characteristics": {  
        "execution times": ["10.3", "10.1", "13.3"],  
        "code size": "131938", ...},  
    "choices": {  
        "os": "linux", "os version": "2.6.32-5-amd64",  
        "compiler": "gcc", "compiler version": "4.6.3",  
        "compiler_flags": "-O3 -fno-if-conversion",  
        "platform": {"processor": "intel xeon e5520",  
                    "l2": "8192", ...}, ...},  
    "features": {  
        "semantic features": {"number_of_bb": "24", ...},  
        "hardware counters": {"cpi": "1.4" ...}, ... }  
    "state": {  
        "frequency": "2.27", ...}}}
```

*Can be directly indexed using web-based,
distributed [ElasticSearch](#) (Apache2 license)*

Python: high-level, portable, popular language with JIT, JSON support and many statistical and machine learning libraries

```
import json  
  
f=file('data.json')  
array=json.loads(f.read())  
f.close()  
  
et=array['characteristics']['execution_times']  
  
for tm in et:  
    print 'Execution time:', tm
```

Convert ad-hoc experiments to cM python modules

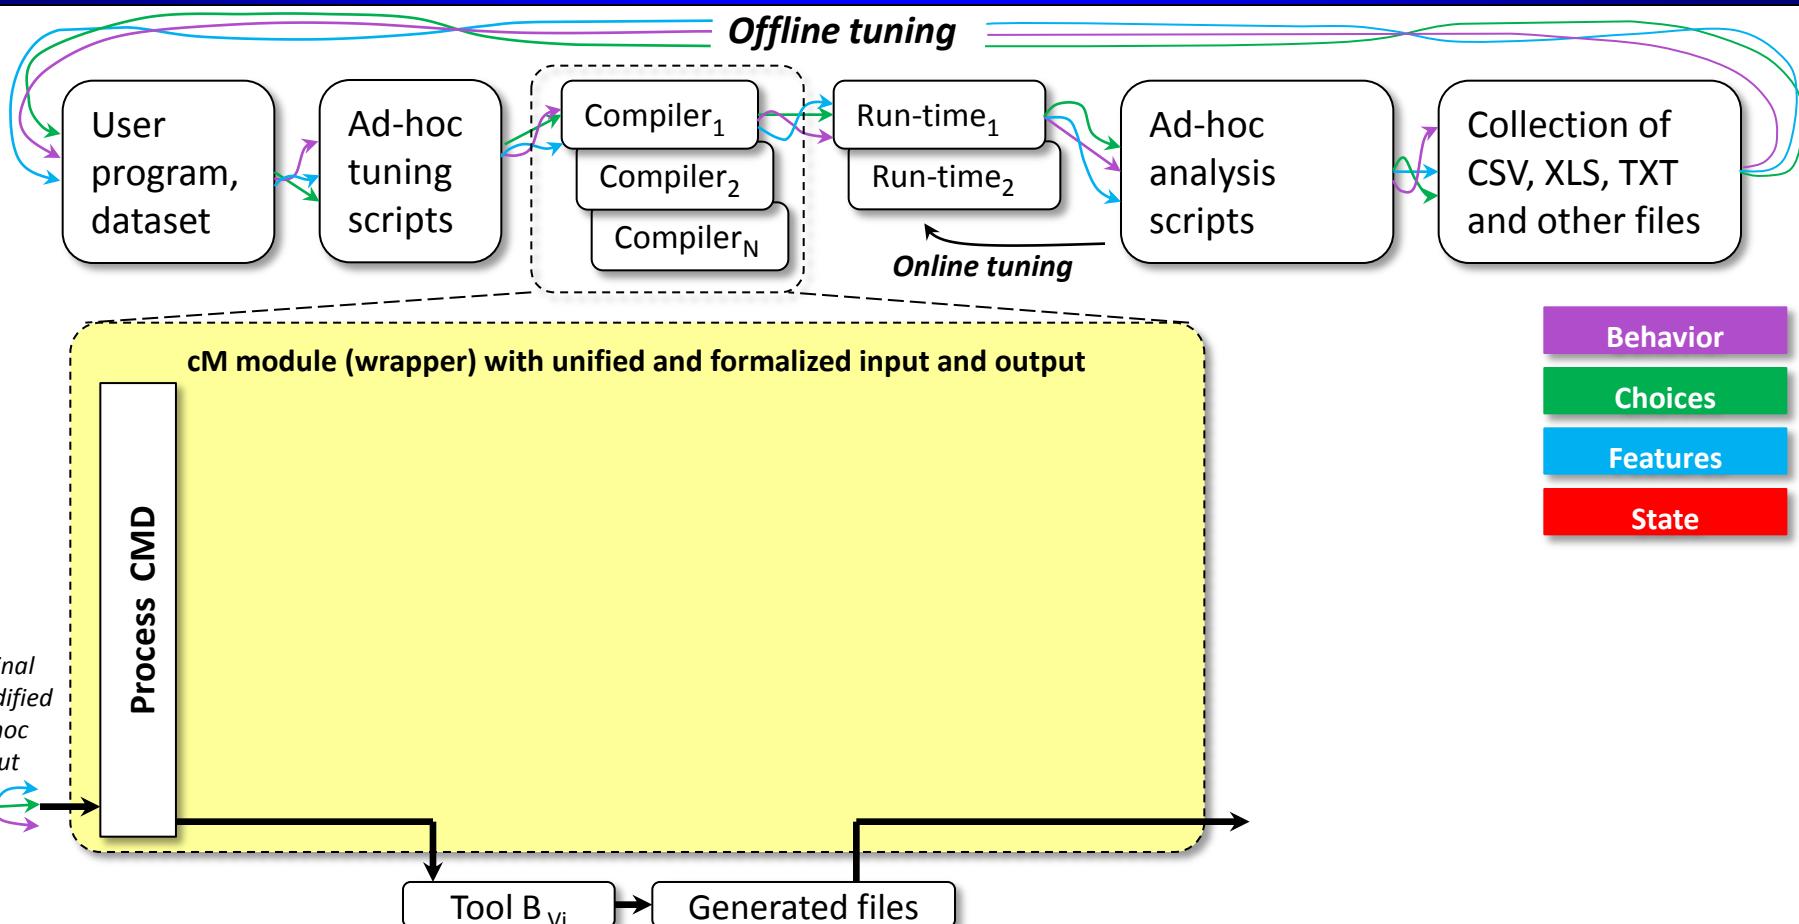


Hardwired experimental setups, very difficult to change, scale or share

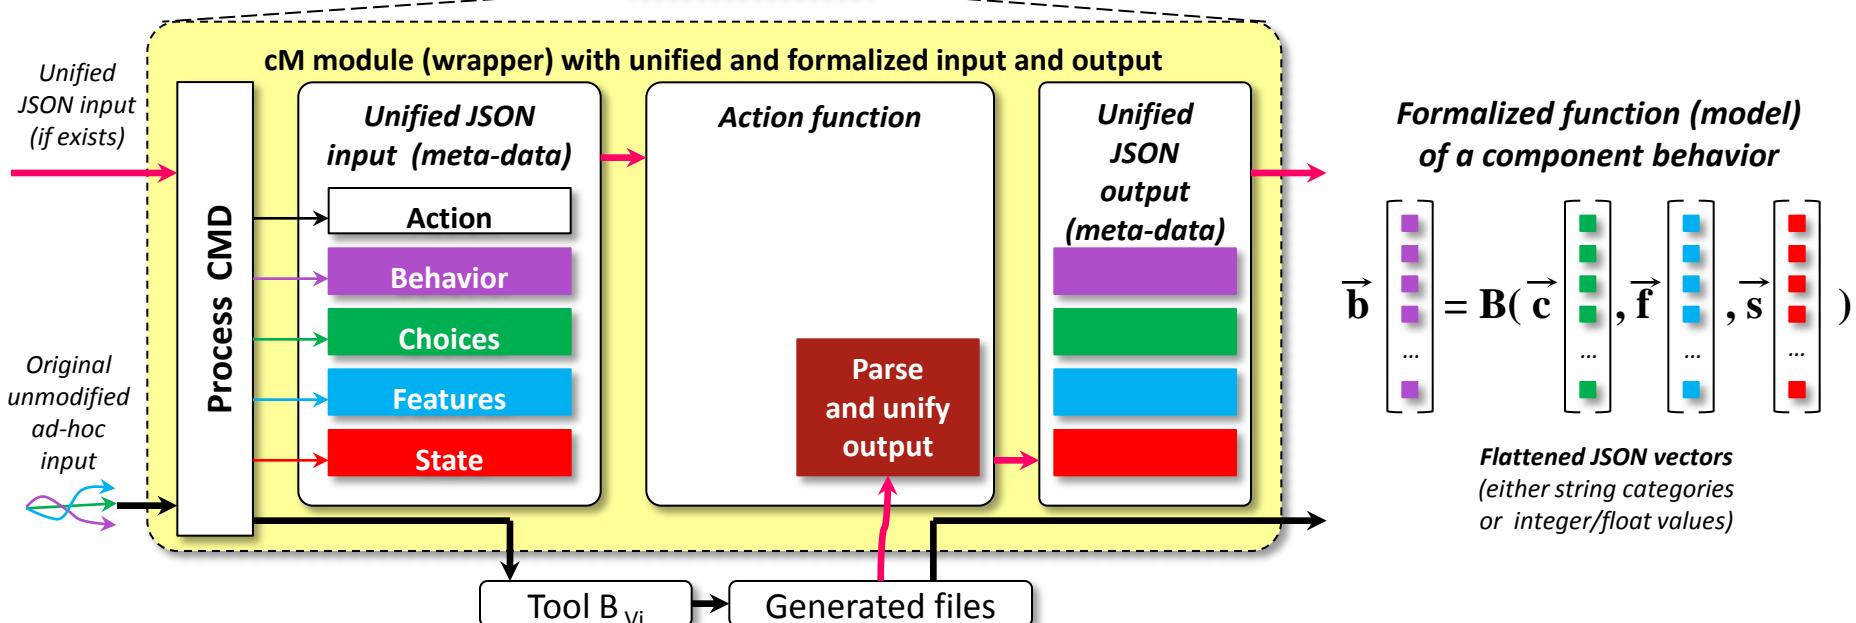
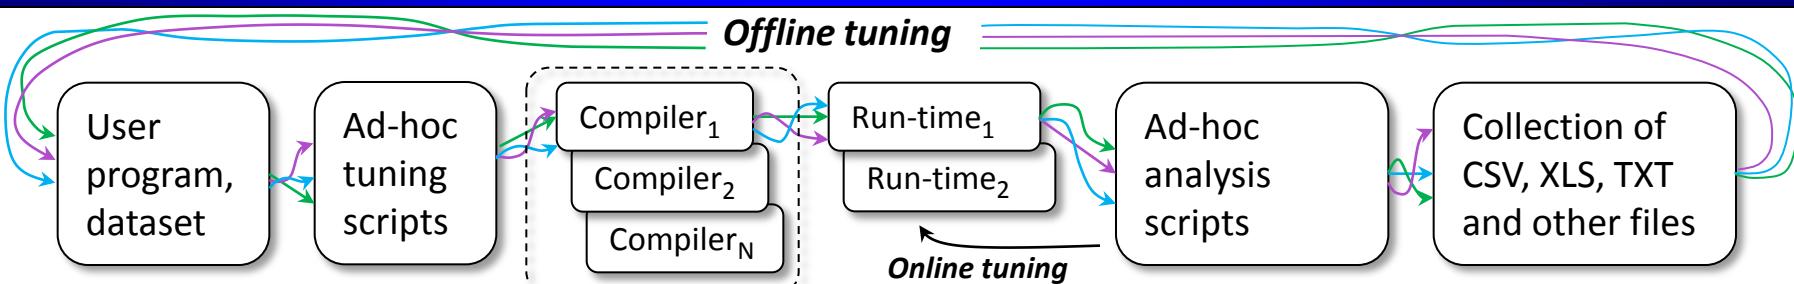


Meta description that should be exposed in the information flow for auto-tuning and machine learning

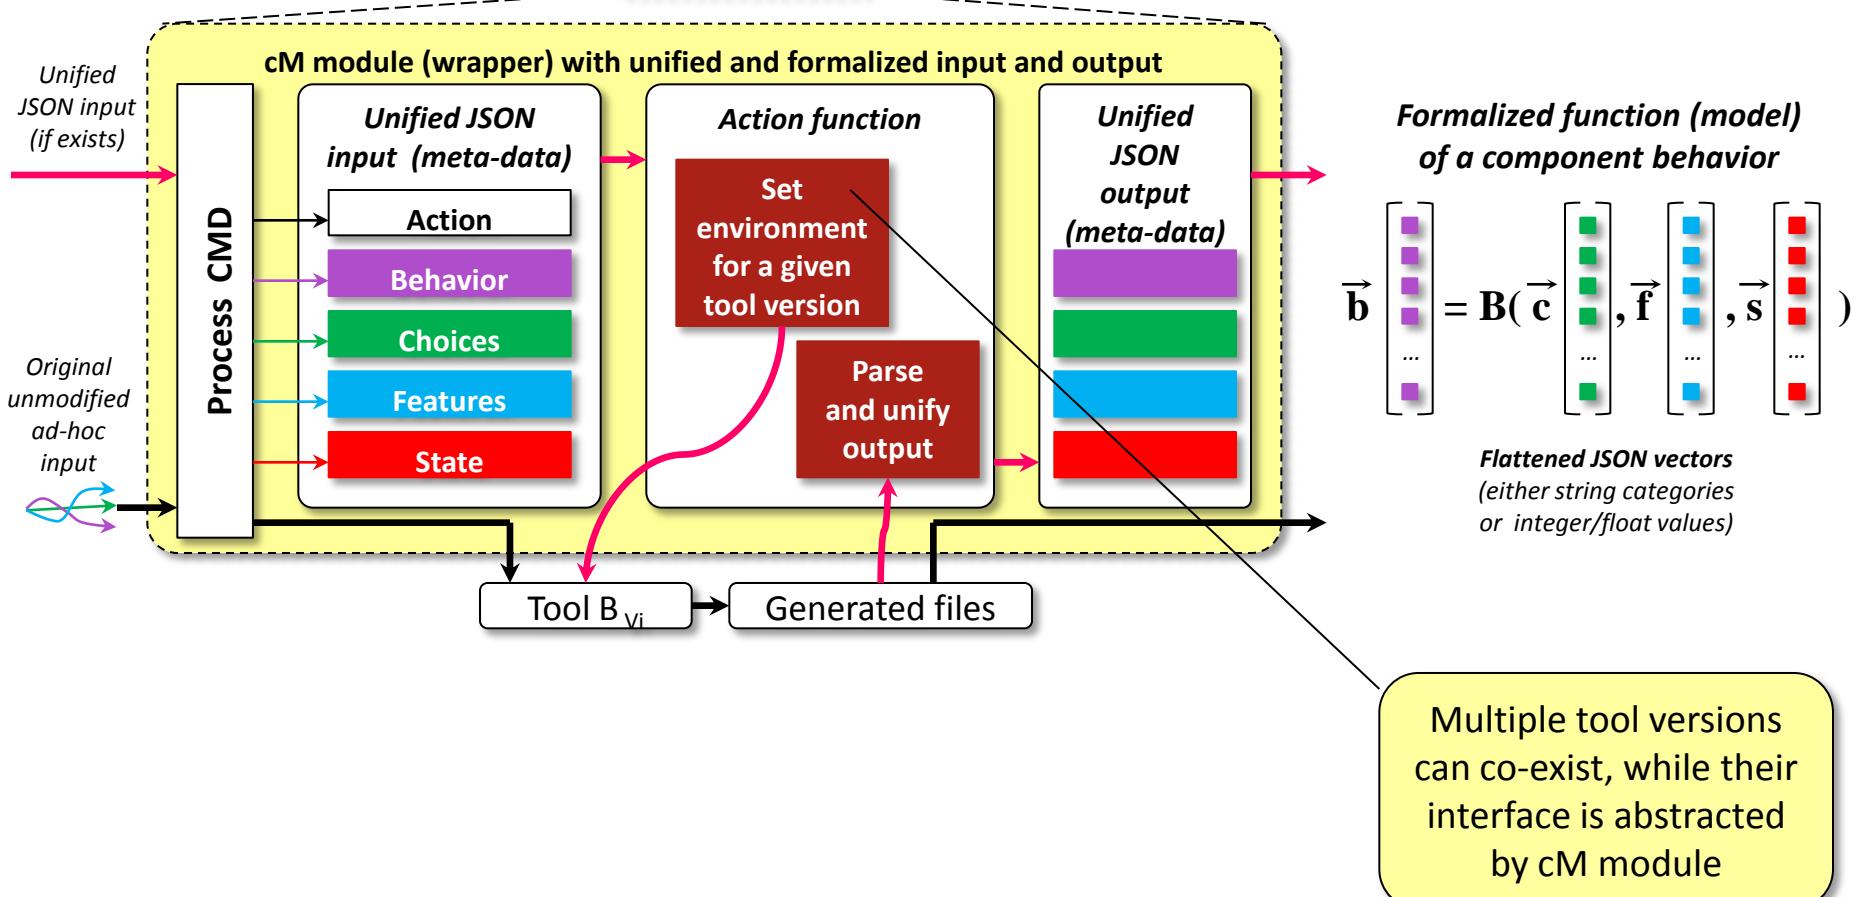
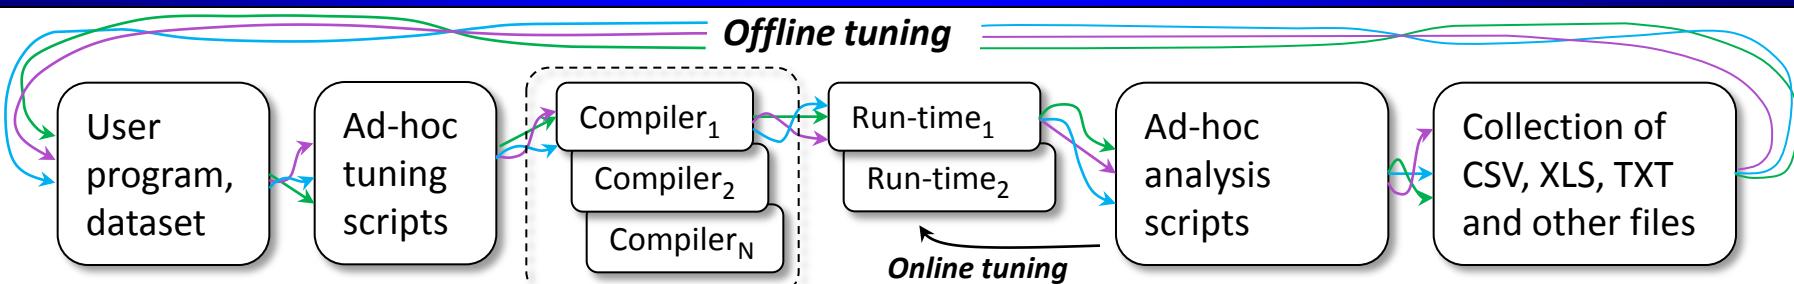
Convert ad-hoc experiments to cM python modules



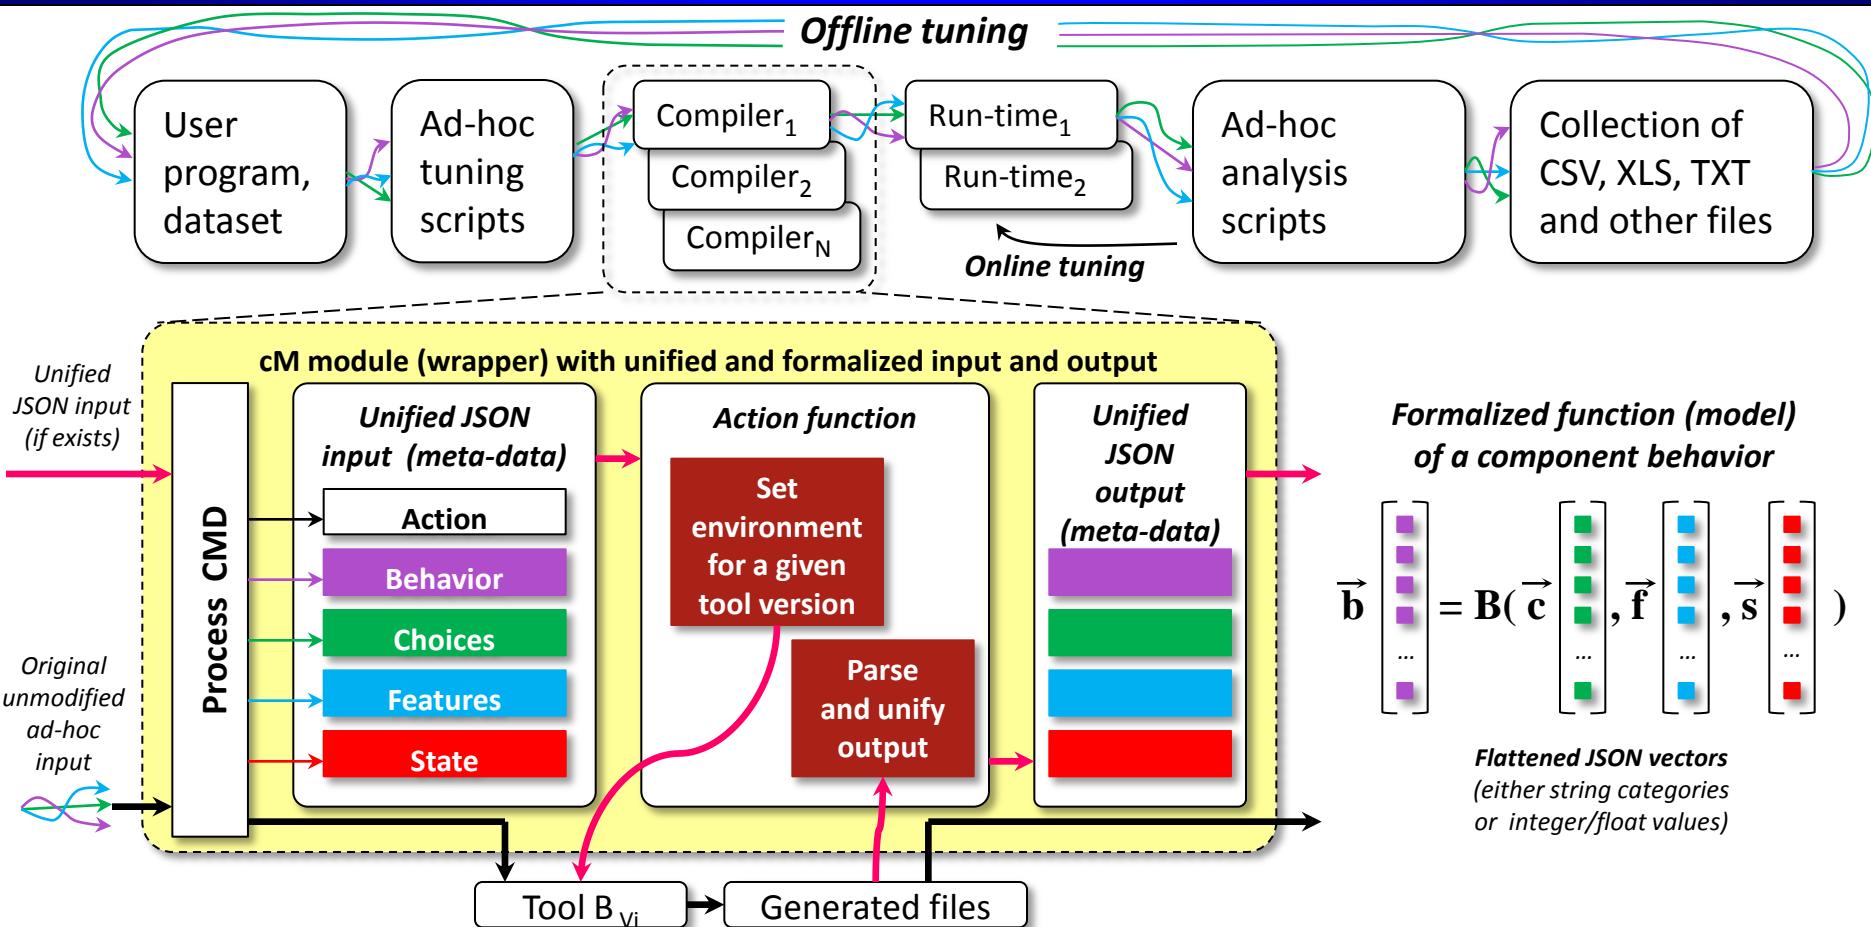
Convert ad-hoc experiments to cM python modules



Convert ad-hoc experiments to cM python modules



Convert ad-hoc experiments to cM python modules



cm [*module name*] [*action*] (param₁=value₁ param₂=value₂ ... -- unparsed command line)

cm **compiler build** -- icc -fast *.c

cm **code.source build** ct_compiler=icc13 ct_optimizations=-fast

cm **code run** os=android binary=./a.out dataset=image-crazy-scientist.pgm

Should be able to run on any OS (Windows, Linux, Android, MacOS, etc)!

Data abstraction in Collective Mind

cM module

compiler

JSON meta-description

GCC 4.4.4

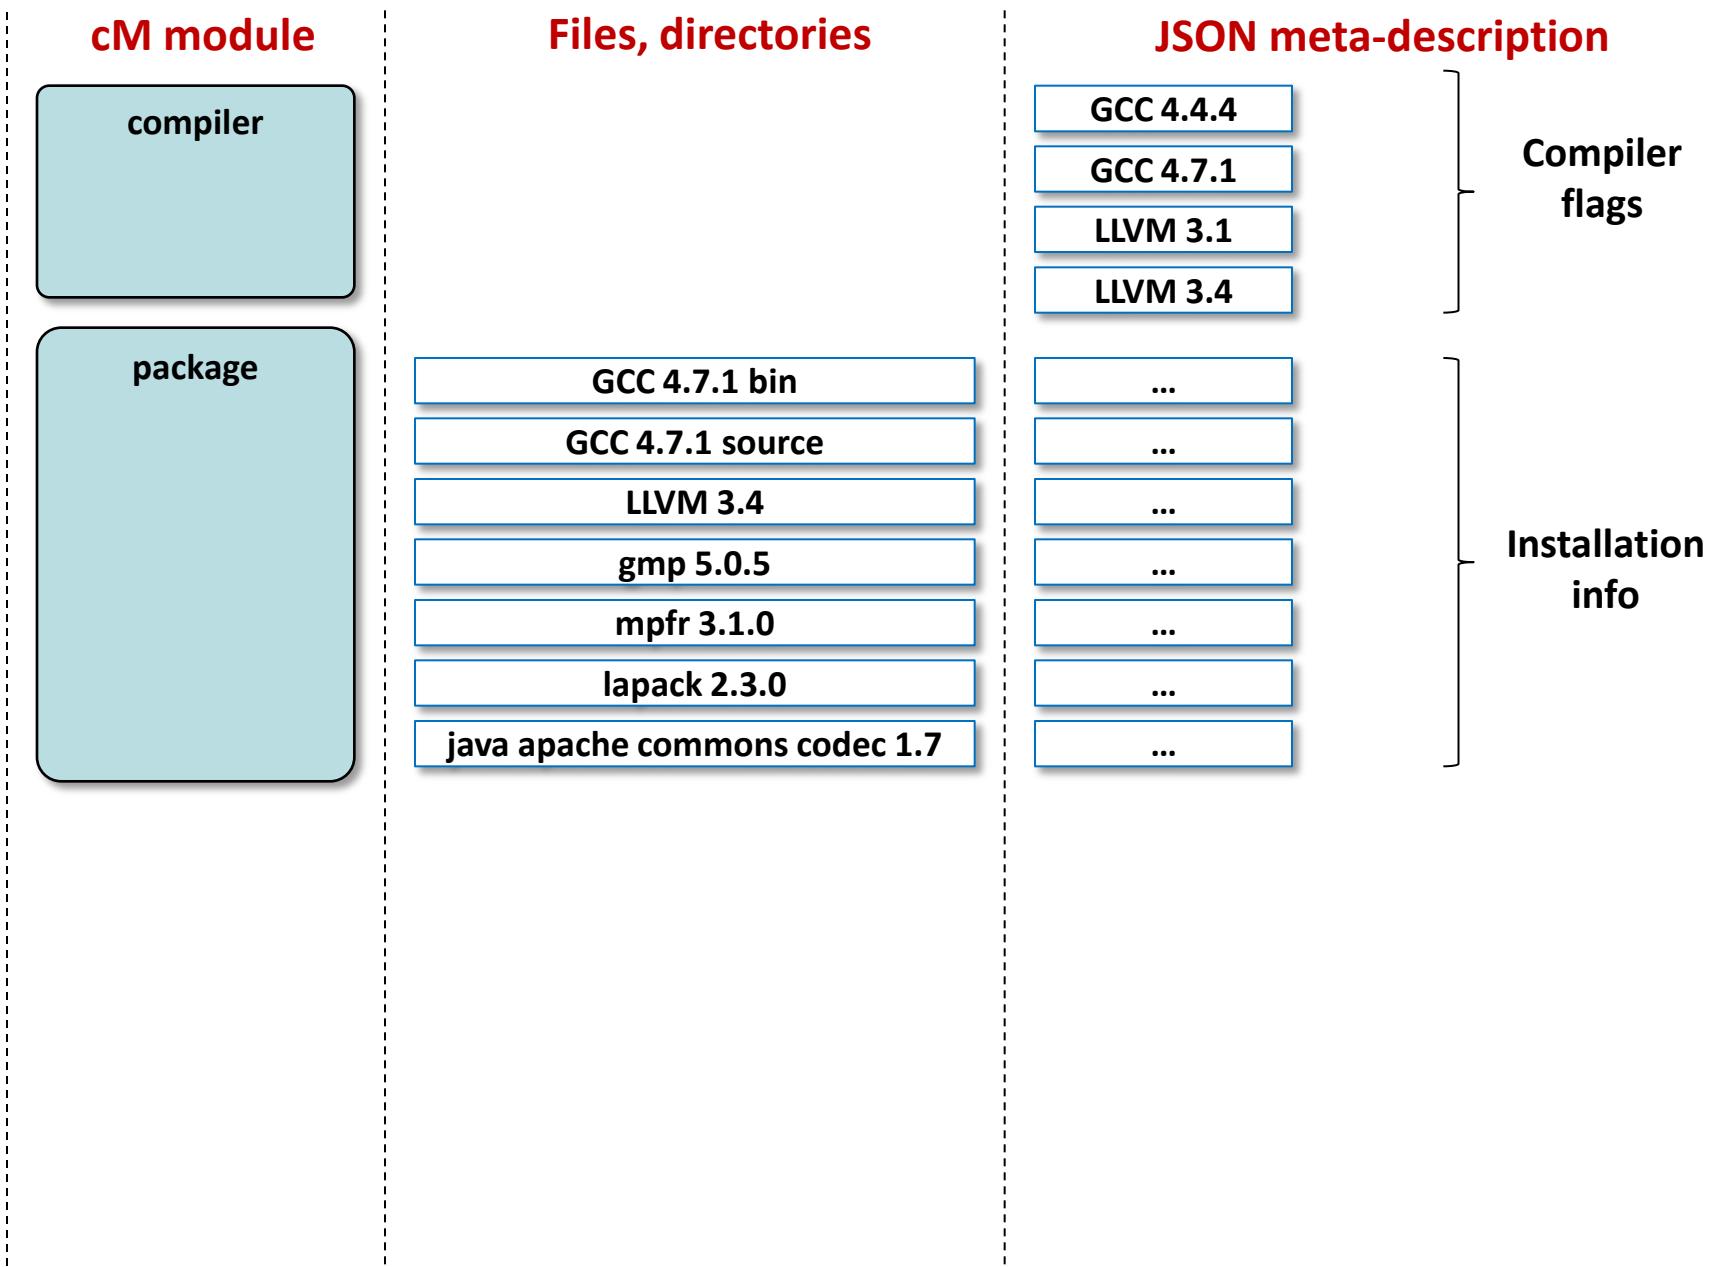
GCC 4.7.1

LLVM 3.1

LLVM 3.4

Compiler
flags

Data abstraction in Collective Mind



Data abstraction in Collective Mind

cM repository directory structure:

cM module

compiler

package

dataset

module

Files, directories

GCC 4.7.1 bin

GCC 4.7.1 source

LLVM 3.4

gmp 5.0.5

mpfr 3.1.0

lapack 2.3.0

java apache commons codec 1.7

image-jpeg-0001

bzip2-0006

txt-0012

compiler

package

dataset

JSON meta-description

GCC 4.4.4

GCC 4.7.1

LLVM 3.1

LLVM 3.4

...

...

...

...

...

...

...

...

...

...

...

...

...

Compiler flags

Installation info

Features

Actions

.cmr

/ module UOA

/ data UOA (UID or alias)

.cm / data.json

Data abstraction in Collective Mind

cM repository directory structure:

cM module

compiler

package

dataset

module

Files, directories

GCC 4.7.1 bin

JSON meta-description

GCC 4.4.4

GCC 4.7.1

LLVM 3.1

LLVM 3.4

...

Now can reference and find any data by CID
(similar to DOI):

<module UOA : data UOA>

.cmr

/ module UOA

/ data UOA (UID or alias)

.cm / data.json

Compiler flags

Installation info

Features

Actions

bzip2-0006

txt-0012

compiler

package

dataset

...

...

...

...

...

Data abstraction in Collective Mind

cM repository directory structure:

cM module

compiler

package

dataset

module

Files, directories

GCC 4.7.1 bin

GCC 4.7.1 source

LLVM 3.4

gmp 5.0.5

mpfr 3.1.0

lapack 2.3.0

java apache commons codec 1.7

image-jpeg-0001

bzip2-0006

txt-0012

compiler

package

dataset

JSON meta-description

GCC 4.4.4

GCC 4.7.1

LLVM 3.1

LLVM 3.4

...

...

...

...

...

...

...

...

...

...

...

...

...

...

Compiler flags

Installation info

Features

Actions

Dependencies between data and modules

.cmr

/ module UOA

/ data UOA (UID or alias)

.cm / data.json

Gradually adding specification (agile development)

cTuning experiment module data.json

{

"characteristics":{

 "execution times": ["10.3", "10.1", "13.3"],

 "code size": "131938", ...},

"choices":{

 "os": "linux", "os version": "2.6.32-5-amd64",

 "compiler": "gcc", "compiler version": "4.6.3",

 "compiler_flags": "-O3 -fno-if-conversion",

 "platform": {"processor": "intel xeon e5520",

 "l2": "8192", ...}, ...},

"features":{

 "semantic features": {"number_of_bb": "24", ...},

 "hardware counters": {"cpi": "1.4" ...}, ... }

"state":{

 "frequency": "2.27", ...}

}

cM flattened JSON key

##characteristics#execution_times@1

Gradually adding specification (agile development)

cTuning experiment module data.json

{

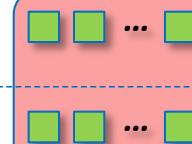
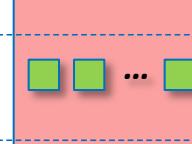
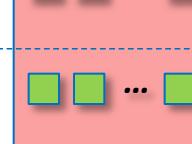
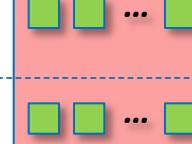
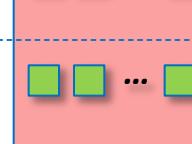
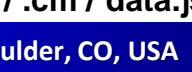
```
"characteristics":{  
    "execution times": ["10.3", "10.1", "13.3"],  
    "code size": "131938", ...},  
"choices":{  
    "os": "linux", "os version": "2.6.32-5-amd64",  
    "compiler": "gcc", "compiler version": "4.6.3",  
    "compiler_flags": "-O3 -fno-if-conversion",  
    "platform": {"processor": "intel xeon e5520",  
                "l2": "8192", ...}, ...},  
"features":{  
    "semantic features": {"number_of_bb": "24", ...},  
    "hardware counters": {"cpi": "1.4" ...}, ...}  
"state":{  
    "frequency": "2.27", ...}
```

cM flattened JSON key

##characteristics#execution_times@1

"flattened_json_key":{
 "type": "text"|"integer" | "float" | "dict" | "list"
 | "uid",
 "characteristic": "yes" | "no",
 "feature": "yes" | "no",
 "state": "yes" | "no",
 "has_choice": "yes" | "no",
 "choices": [list of strings if categorical
 choice],
 range",
 range",
 "explore_start": "start number if numerical
 "explore_stop": "stop number if numerical
 "explore_step": "step if numerical range",
 "can_be_omitted": "yes" | "no"
 ...

Gradually adding, extending and improving modules and data

Category	cM module	Module actions	All data	Meta-description
<i>Third-party tools, libraries</i>	package	<i>common*</i> , <i>install</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	 ... 
<i>High-level algorithms</i>	algorithm	<i>common*</i> , <i>transform</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	
<i>Applications, benchmarks, codelets, kernels</i>	code.source	<i>common*</i> , <i>build</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	
<i>Auto-tuning compilers</i>	ctuning.compiler	<i>common*</i> , <i>compile_program</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	
<i>Compiler machine learning plugins</i>	math.model	<i>common*</i> , <i>build</i> , <i>predict</i> , <i>fit</i> , <i>detectRepresentativePoints</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	
<i>Binaries and libraries</i>	code	<i>common*</i> , <i>run</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	
<i>Data sets</i>	dataset	<i>common*</i> , <i>create</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	
<i>Operating Systems</i>	os	<i>common*</i> , <i>detectHostFamily</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	
<i>Processors</i>	processor	<i>common*</i> , <i>detectHostProcessor</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	
<i>Statistical and data mining plugins</i>	math.statistics.r	<i>common*</i> , <i>analyze</i>	<input type="checkbox"/> <input type="checkbox"/> ... <input type="checkbox"/>	
<i>* add, list, view, copy, move, search</i>				
cM repository directory structure:	.cmr	/ module UOA (UID or alias)	/ data UOA	/ .cm / data.json

Transparently indexed by Elastic Search

cM interface: only one main function!

Simple and minimalistic high-level cM interface - **one function (!)**

(python dictionary) *output* = `cm_kernel.access` ((python dictionary) *input*)

Load data:

```
r = cm_kernel.access({'cm_run_module_uoa':'dataset',
                      'cm_action':'load',
                      'cm_data_uoa':'image-jpeg-0001'})

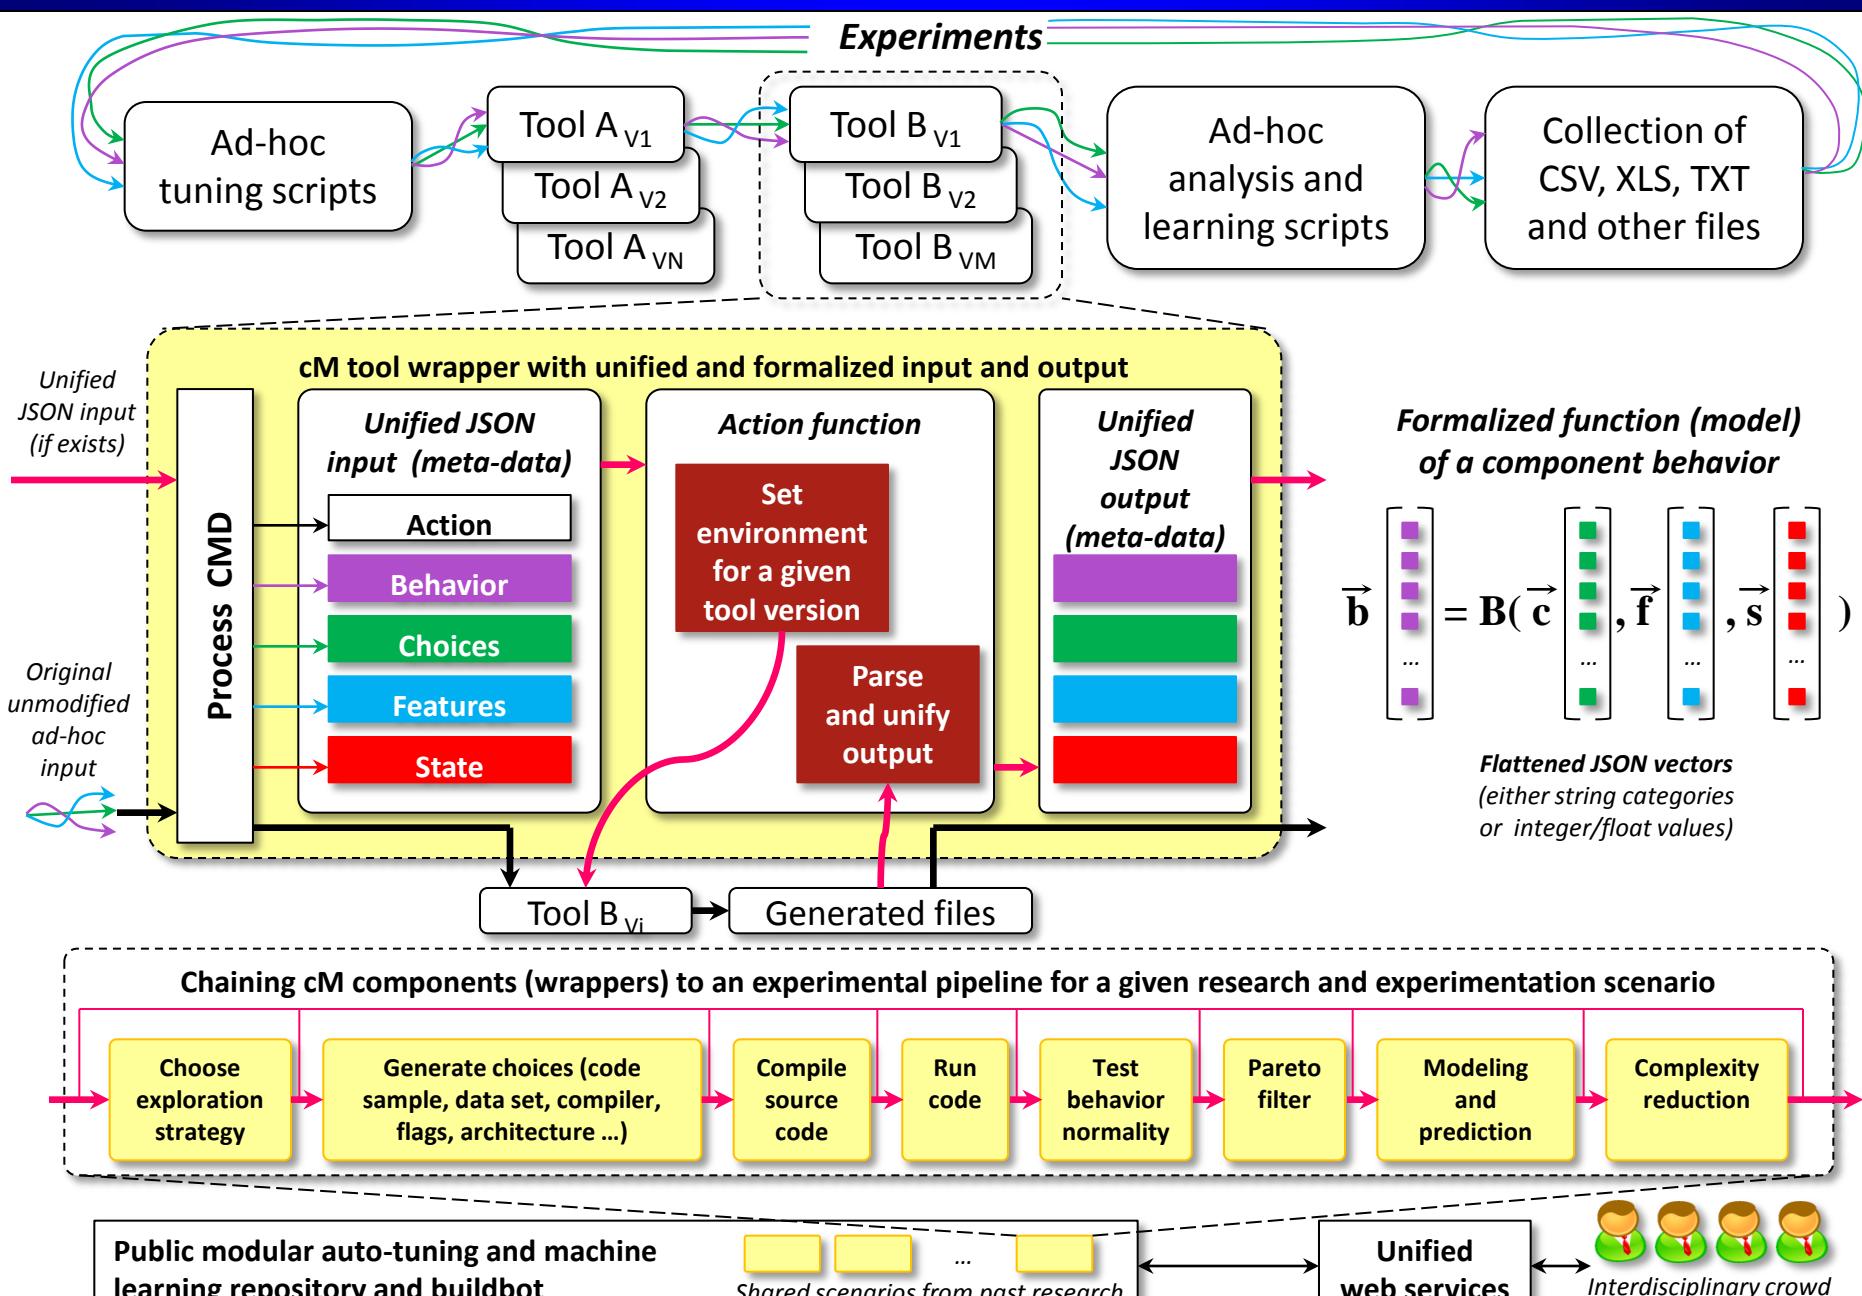
if r['cm_return']>0: return r
d=r['cm_data_obj']['cfg']
```

Call R machine learning module:

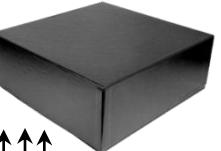
```
r = cm_kernel.access({'cm_run_module_uoa':'cm_math.model.r',
                      'cm_action':'build',
                      'model_name':'earth', ...})

if r['cm_return']>0: return r
```

Assembling, preserving, sharing and extending the whole pipeline as "LEGO"



Top-down decomposition and learning of computer systems



		Gradually expose some characteristics	Gradually expose some choices and features
→ Compile Program		time ...	compiler flags; pragmas ...
	<i>I now gradually convert to Collective Mind and validate past research techniques with the great help of volunteers (sadly such work is usually not funded, can't be easily published, and is often considered as a waste of time by academic community)</i>		
→ Run code	Run-time environment	time; CPI, power consumption ...	pinning/scheduling ...
	System	cost;	architecture; frequency; cache size...
	Data set	size; values; description ...	precision ...
→ Analyze profile		time; size ...	instrumentation; profiling ...

Start coarse-grain decomposition of a system (detect coarse-grain effects first). Add universal learning modules.

Growing, plugin-based cM pipeline for auto-tuning and learning

<http://c-mind.org/ctuning-pipeline>

•**Init pipeline**

- Detected system information
- Initialize parameters
- Prepare dataset

•**Clean program**

•**Prepare compiler flags**

- Use compiler profiling
- Use cTuning CC/MILEPOST GCC for fine-grain program analysis and tuning
- Use universal Alchemist plugin (with any OpenME-compatible compiler or tool)
- Use Alchemist plugin (currently for GCC)

•**Build program**

- Get objdump and md5sum (if supported)
- Use OpenME for fine-grain program analysis and online tuning (build & run)
- Use 'Intel VTune Amplifier' to collect hardware counters
- Use 'perf' to collect hardware counters
- Set frequency (in Unix, if supported)
- Get system state before execution

•**Run program**

- Check output for correctness (use dataset UID to save different outputs)
- Finish OpenME

- Misc info

•**Observed characteristics**

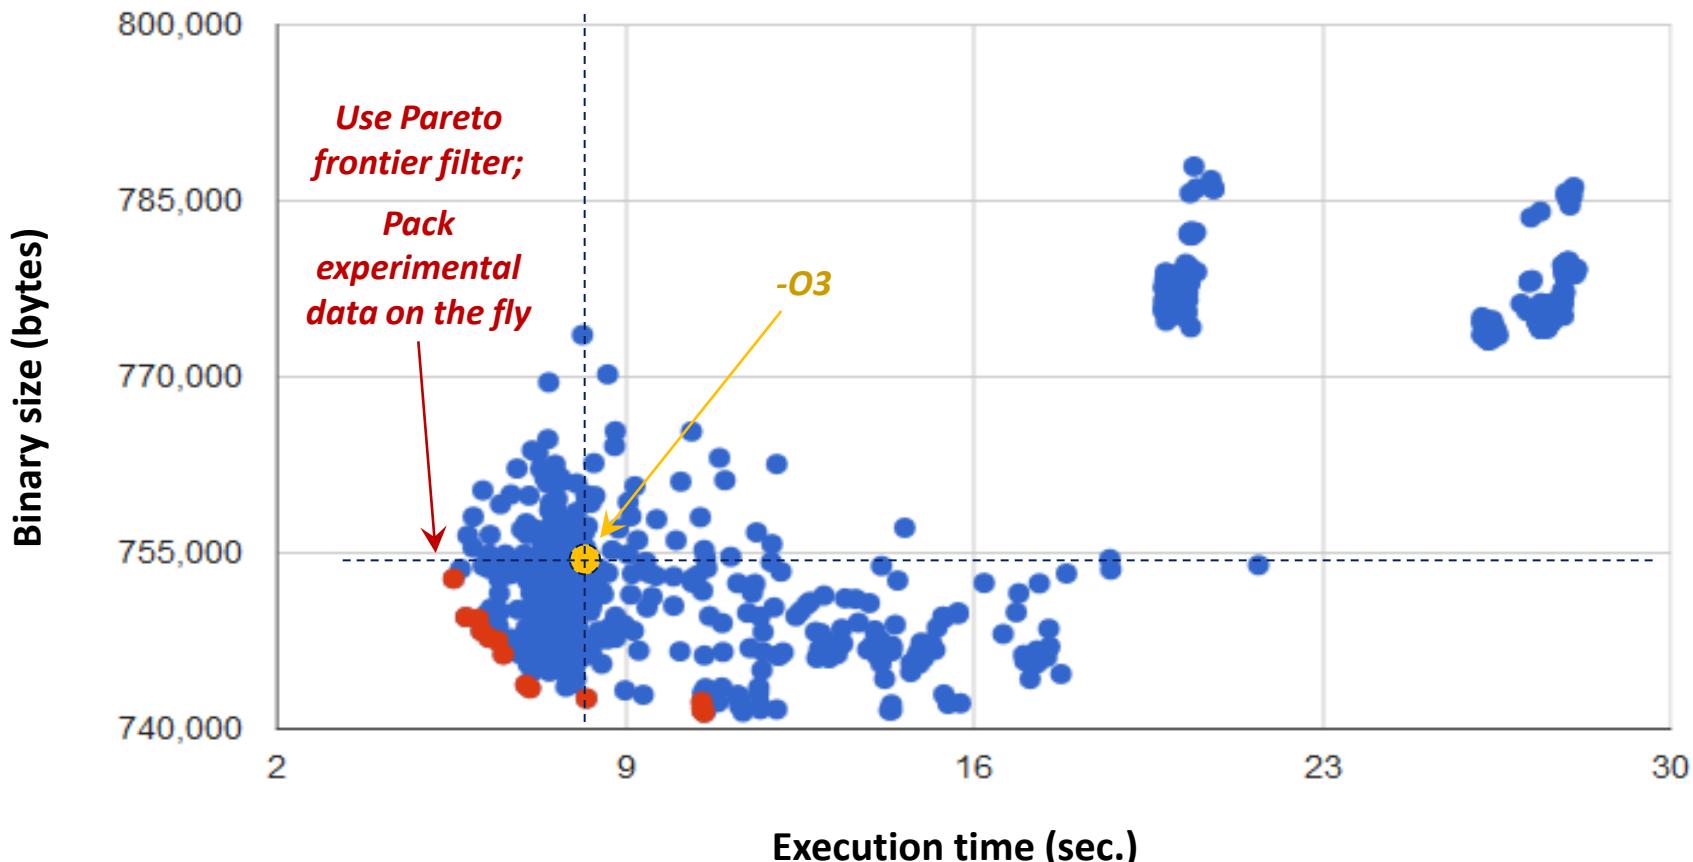
- Observed statistical characteristics

•**Finalize pipeline**

Our Collective Mind Buildbot and plugin-based auto-tuning pipeline supports the following shared benchmarks and codelets:

- Polybench - numerical kernels with exposed parameters of all matrices in cM
 - CPU: 28 prepared benchmarks
 - CUDA: 15 prepared benchmarks
 - OpenCL: 15 prepared benchmarks
- cBench - 23 benchmarks with 20 and 1000 datasets per benchmark
- Codelets - 44 codelets from embedded domain (provided by CAPS Entreprise)
- SPEC 2000/2006
- Description of 32-bit and 64-bit OS: Windows, Linux, Android
- Description of major compilers: GCC 4.x, LLVM 3.x, Open64/Pathscale 5.x, ICC 12.x
- Support for collection of hardware counters: perf, Intel vTune
- Support for frequency modification
- Validated on laptops, mobiles, tables, GRID/cloud - can work even from the USB key

Implemented scenario: multi-objective compiler auto-tuning using mobile phones



Program: *image corner detection*

Compiler: *Sourcery GCC for ARM v4.7.3*

System: *Samsung Galaxy Y*

Processor: *ARM v6, 830MHz*

OS: *Android OS v2.3.5*

Data set: *MiDataSet #1, image, 600x450x8b PGM, 263KB*

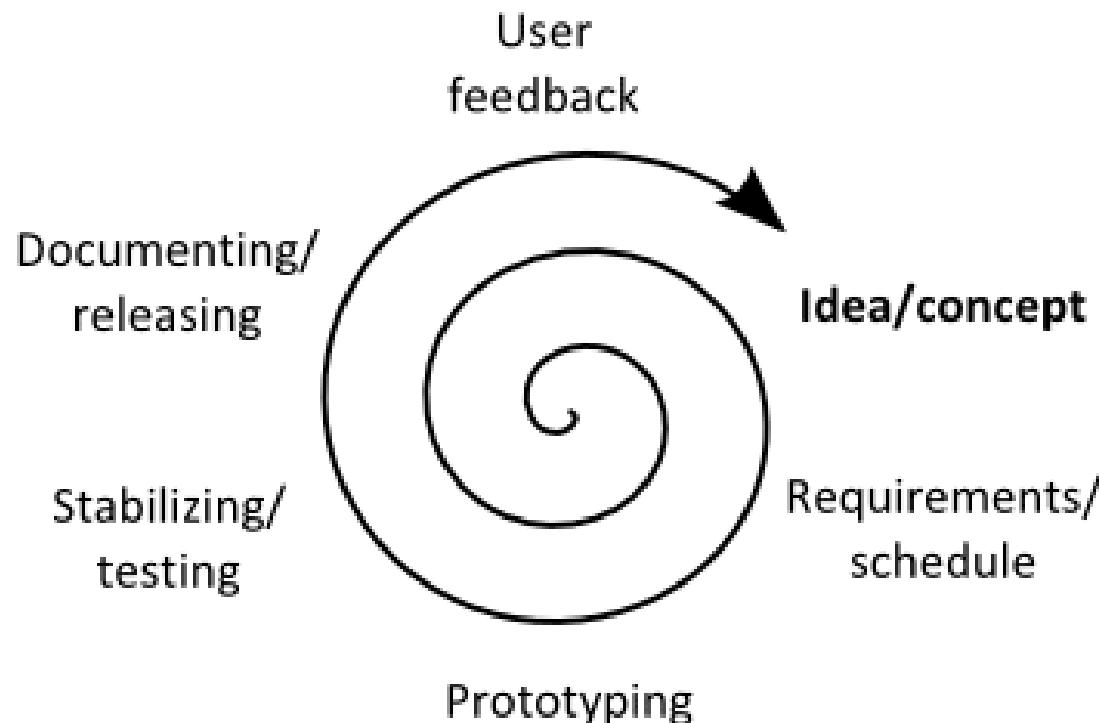
500 combinations of random flags -O3 -f(no-)FLAG

Powered by Collective Mind Node (Android Apps on Google Play)

Collective Mind agile research and development

From agile development to agile research!

Community shares, validates and improves benchmarks, data sets, tools, design and optimization choices, features, characteristics, models, ...



**Collective Mind experimental pipelines can now survive changes in the system!
Modules and data can evolve with the evolution of the technology!**

Reproducibility of experimental results

Reproducibility came as a side effect!

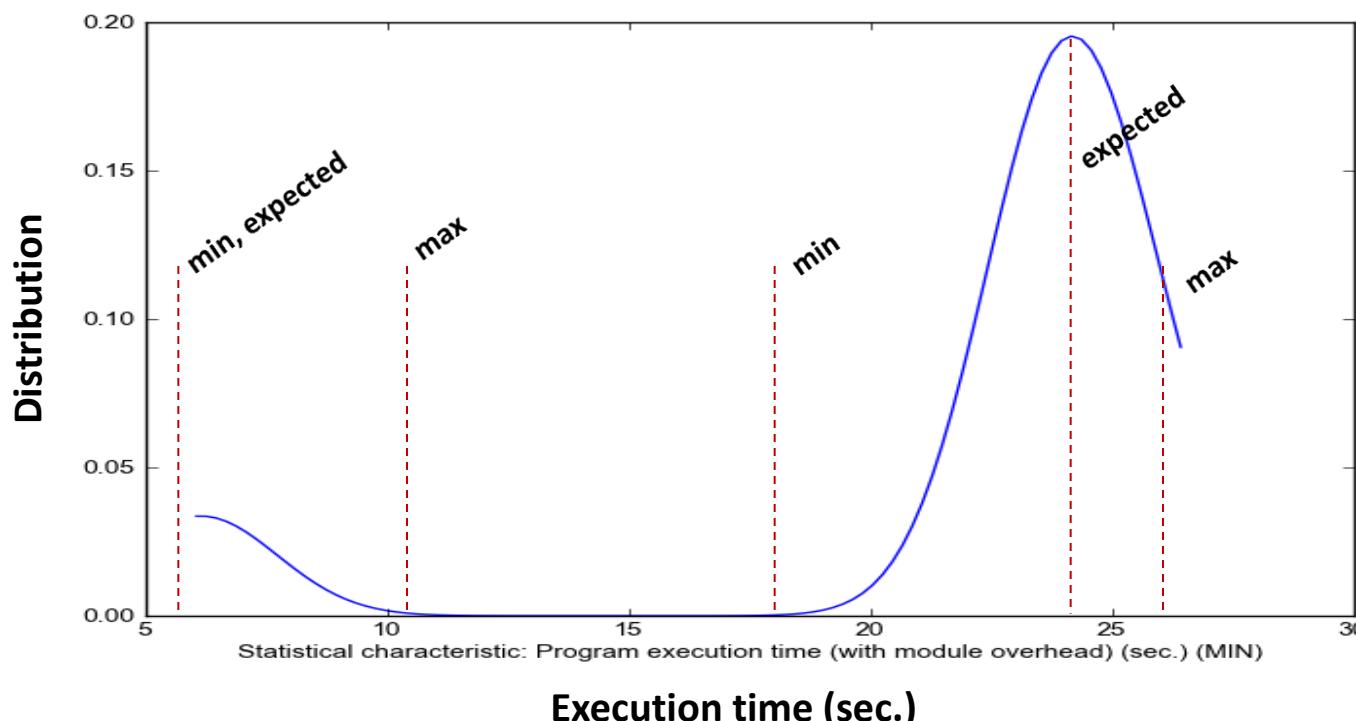
- Can preserve the whole experimental setup with all data and software dependencies
- Can perform statistical analysis (normality test) for characteristics
- Community can add missing features or improve machine learning models

Reproducibility of experimental results

Reproducibility came as a side effect!

- Can preserve the whole experimental setup with all data and software dependencies
- Can perform statistical analysis (normality test) for characteristics
- Community can add missing features or improve machine learning models

Unexpected behavior - expose to the community including domain specialists, explain, find missing feature and add to the system

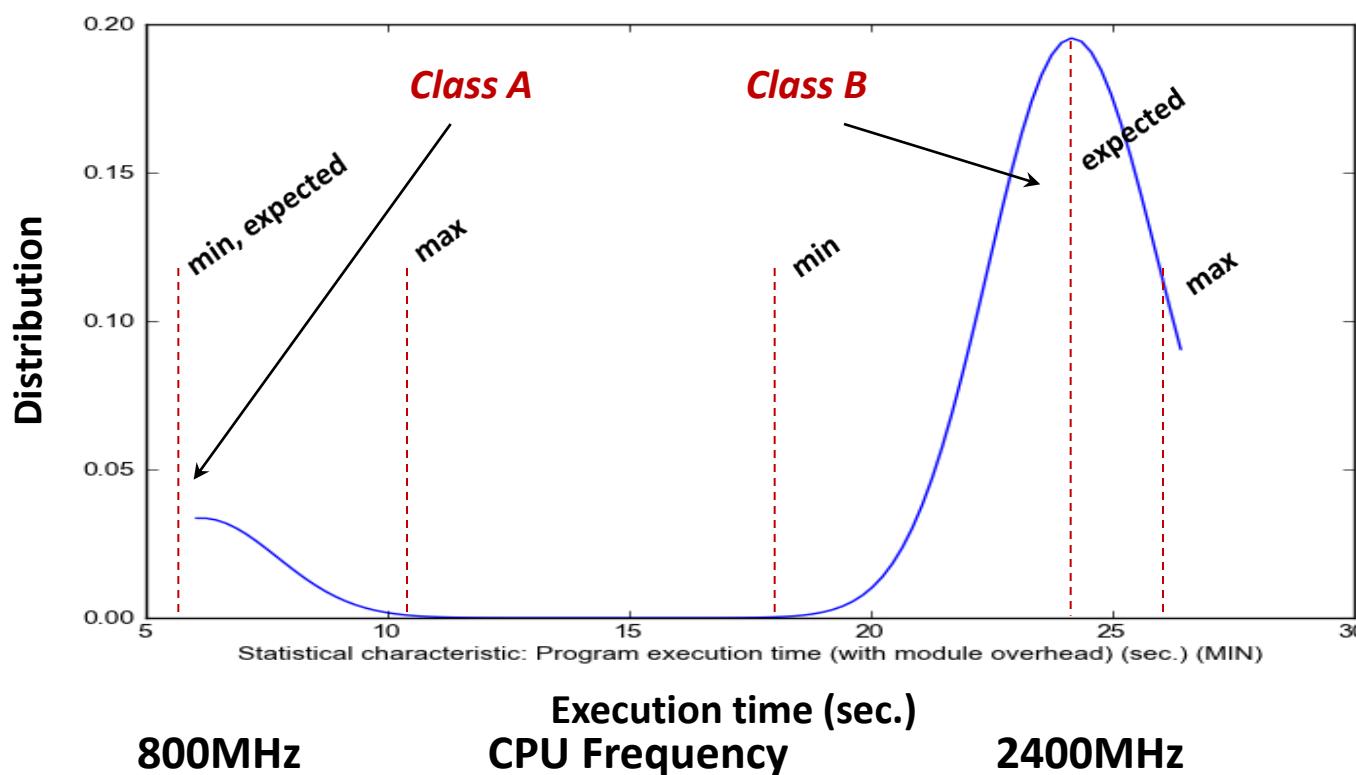


Reproducibility of experimental results

Reproducibility came as a side effect!

- Can preserve the whole experimental setup with all data and software dependencies
- Can perform statistical analysis (normality test) for characteristics
- Community can add missing features or improve machine learning models

Unexpected behavior - expose to the community including domain specialists, explain, find missing feature and add to the system

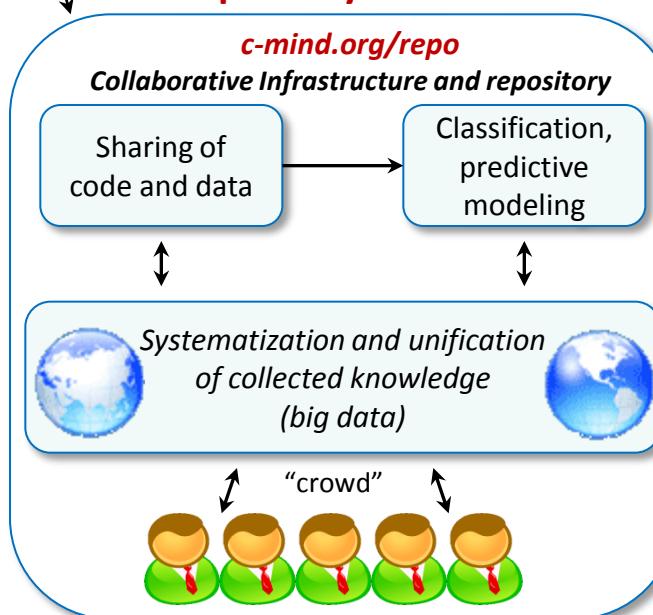
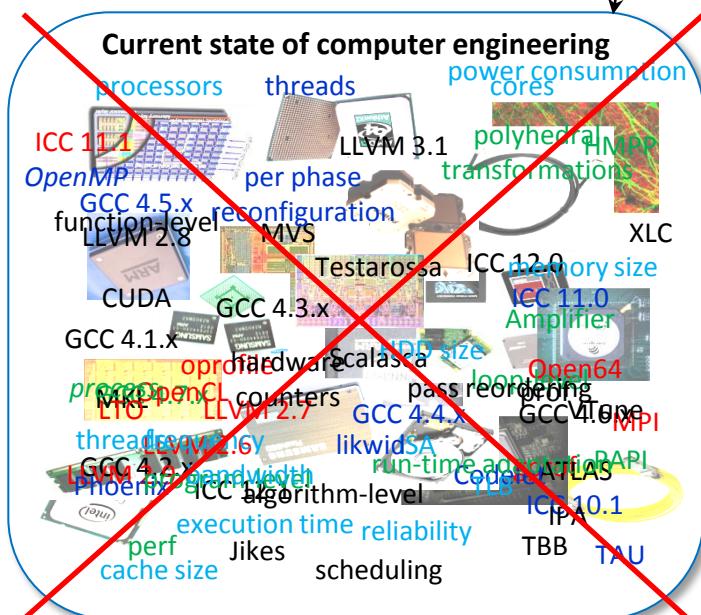


My dream: systematic, collaborative and reproducible computer engineering

- Quick, non-reproducible hack?
- Ad-hoc heuristic?
- Quick publication?
- No shared code and data?

- Prototype research idea
- Validate existing work
- Perform end-user task

- Share code, data with their meta-description and dependencies
- Systematize and classify collected optimization knowledge
- Develop and preserve the whole experimental pipeline
- Validate experimental results by the community
- Extrapolate collected knowledge to build faster, smaller, more power efficient and reliable computer systems



Current status and future work

- Pilot live repository for public curation of research material: <http://c-mind.org/repo>
- Infrastructure is available at SourceForge under standard BSD license: <http://c-mind.org>
- Example of crowdsourcing compiler flag auto-tuning using mobile phones: “**Collective Mind Node**” in Google Play Store
- Several publications under submission
- Raising funding to make cM more user friendly and add more research scenarios

Education	Academic research
<p>New publication model where research material and experimental results are shared, validated and reused by the community http://ctuning.org/reproducibility</p> <ul style="list-style-type: none">• ACM SIGPLAN TRUST 2014 @ PLDI 2014 http://c-mind.org/events/trust2014• Panel at ADAPT 2014 @ HiPEAC 2014 http://adapt-workshop.org (this year we attempted to reproduce some submitted articles)	<p></p> <ul style="list-style-type: none">• Systematizing, validating, sharing past research techniques on auto-tuning and machine learning through cM pipeline• Optimal feature and model selection• Finding representative benchmarks and data sets• Run-time adaptation and ML

Grigori Fursin, “Collective Mind: cleaning up the research and experimentation mess in computer engineering using crowdsourcing, big data and machine learning”, INRIA Tech. report No 00850880, August 2013

<http://hal.inria.fr/hal-00850880>

<http://arxiv.org/abs/1308.2410>

Acknowledgements

- Colleagues from NCAR (USA): ***Davide Del Vento et.al.***
- Colleagues from STMicroelectronics (France):
Christophe Guillone, Antoine Moynault, Christian Bertin
- Colleagues from ARM (UK): **Anton Lokhmotov**
- Colleagues from Intel (USA): ***David Kuck and David Wong***
- cTuning community:

<http://cTuning.org/lab/people>



- EU FP6, FP7 program and HiPEAC network of excellence
<http://www.hipeac.net>
- INRIA fellowship (2012-2016)

Thank you for your attention!

Questions and comments: grigori.fursin@inria.fr