

Scientific Filesystem

Vanessa Sochat

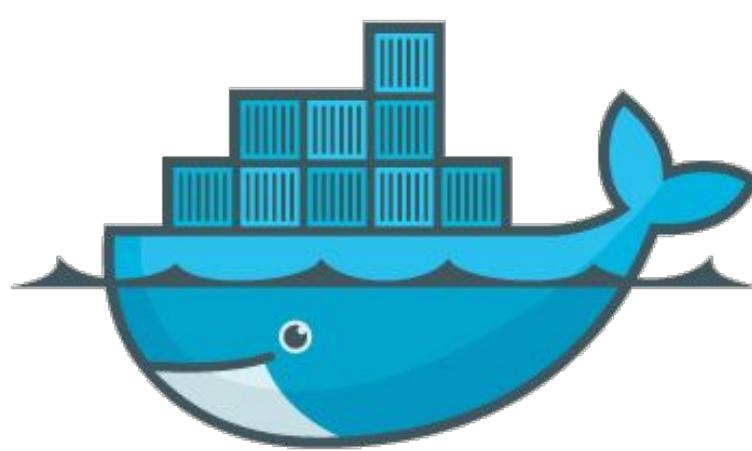
vsochat@stanford.edu



“

Welcome!





N
SHIFTER



Charliecloud

“

Let's all be friends!



Once upon a time...
There was a scientist



RESEARCH ARTICLE SUMMARY

PSYCHOLOGY

Estimating the reproducibility of psychological science

Open Science Collaboration*

INTRODUCTION: Reproducibility is a defining feature of science, but the extent to which it characterizes current research is unknown. Scientific claims should not gain credence because of the status or authority of their originator but by the replicability of their supporting evidence. Even research of exemplary quality may have irreproducible empirical findings because of random or systematic error.

RATIONALE: There is concern about the rate and predictors of reproducibility, but limited evidence. Potentially problematic practices in-

viously observed finding and is the means of establishing reproducibility of a finding with new data. We conducted a large-scale, collaborative effort to obtain an initial estimate of the reproducibility of psychological

substantial decline. Ninety-seven percent of original studies had significant results ($P < .05$). Thirty-six percent of replications had significant results; 47% of original effect sizes were in the 95% confidence interval of the replication effect size; 39% of effects were subjectively rated to have replicated the original result;

and if no bias in original results is assumed, combining original and replication results left 68% with statistically significant effects. Correlational tests suggest that replication success was better predicted by the strength of original evidence than by characteristics of the original and replication

ON OUR WEB SITE

Read the full article at <http://dx.doi.org/10.1126/science.aac4716>

NATURE | NEWS

Over half of psychology studies fail reproducibility test

Largest replication study to date casts doubt on many published positive results.

Monya Baker

27 August 2015



Foo



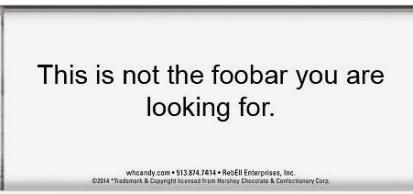
Bar

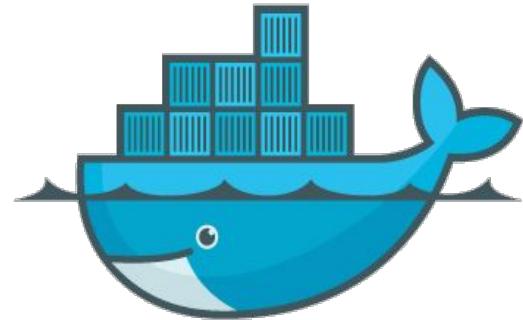


foO



baR

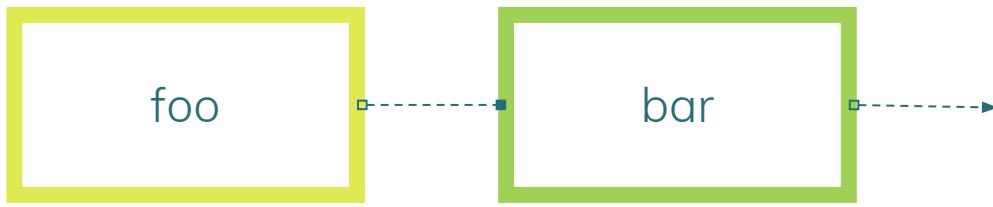




science +

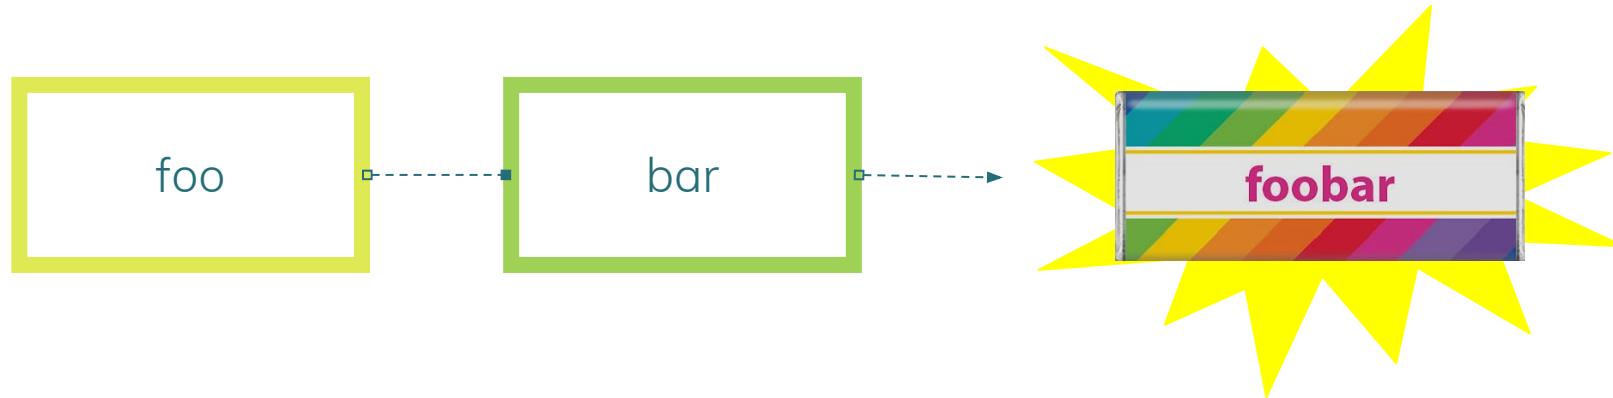


Pipeline foobar

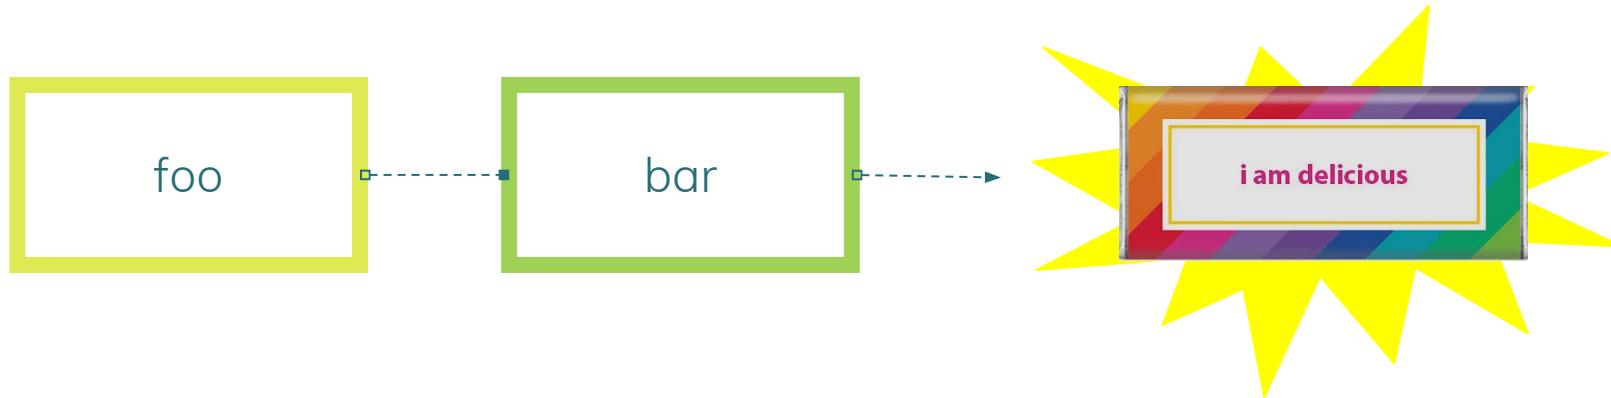


?

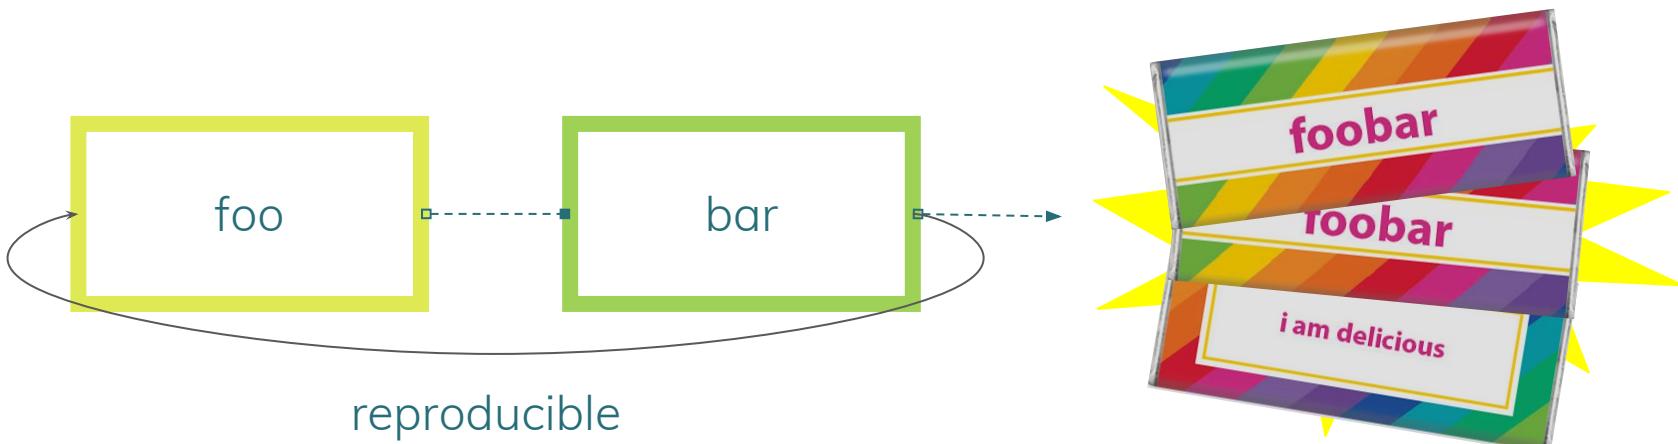
Pipeline foobar



Pipeline foobar



Pipeline foobar





Give it a try?



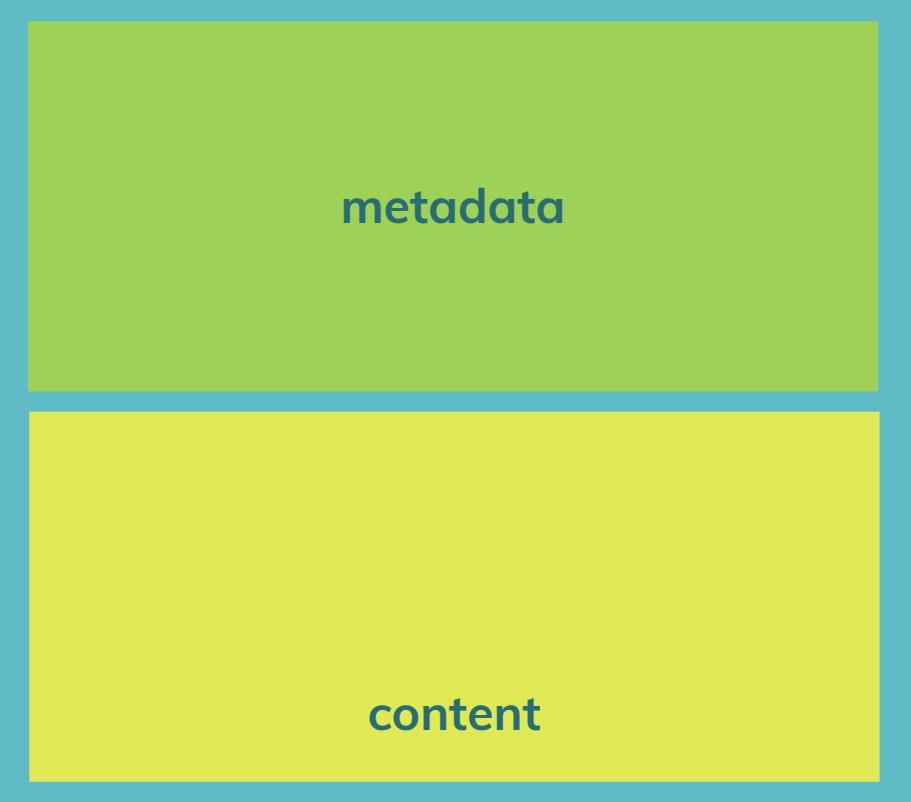


ENTRYPOINT



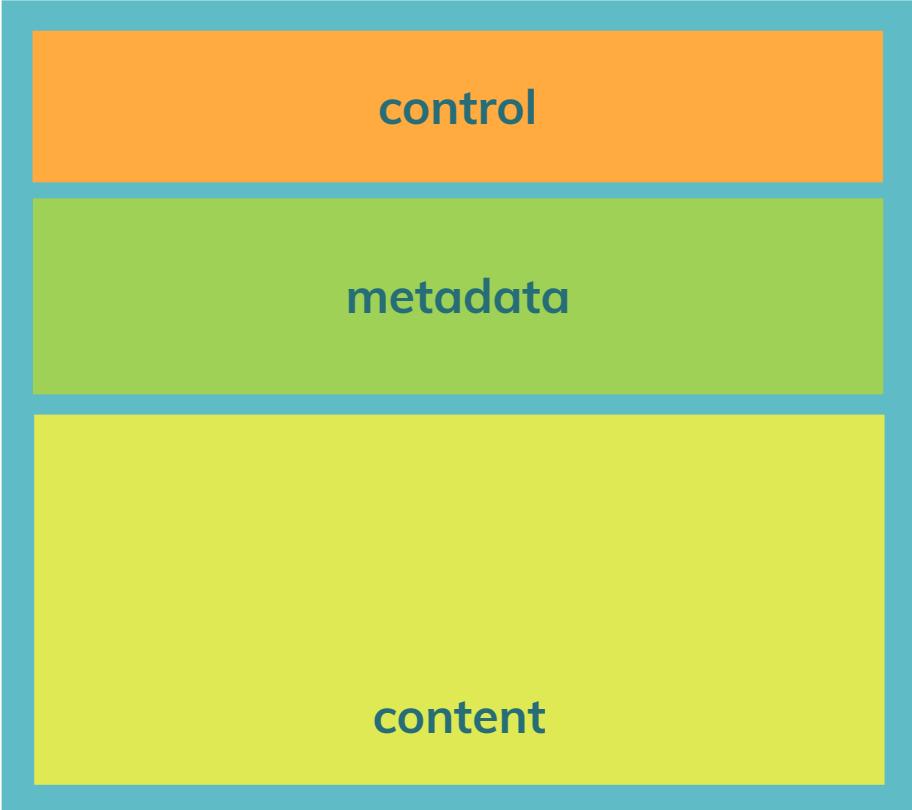
“

What about the human component
of reproducibility?



metadata

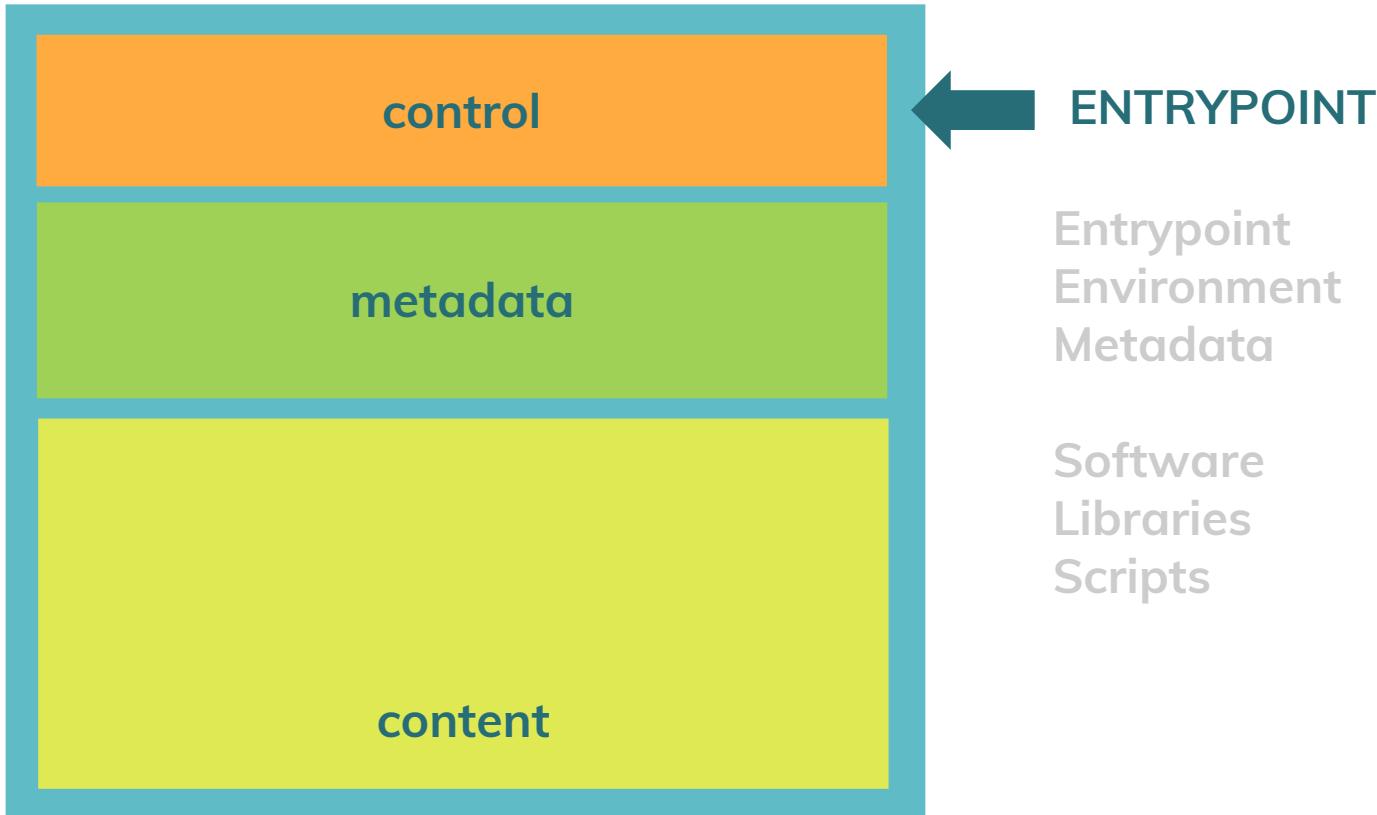
content



control

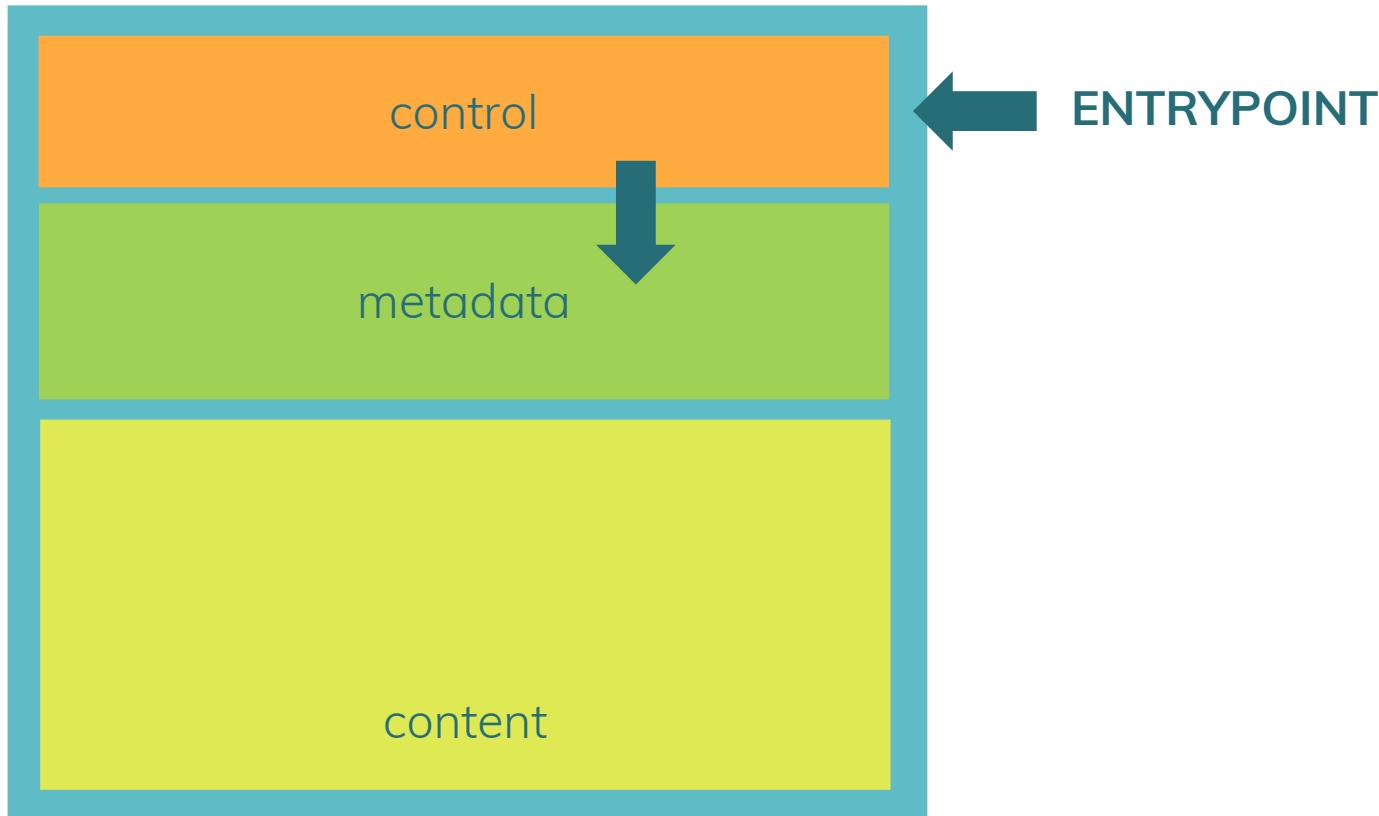
metadata

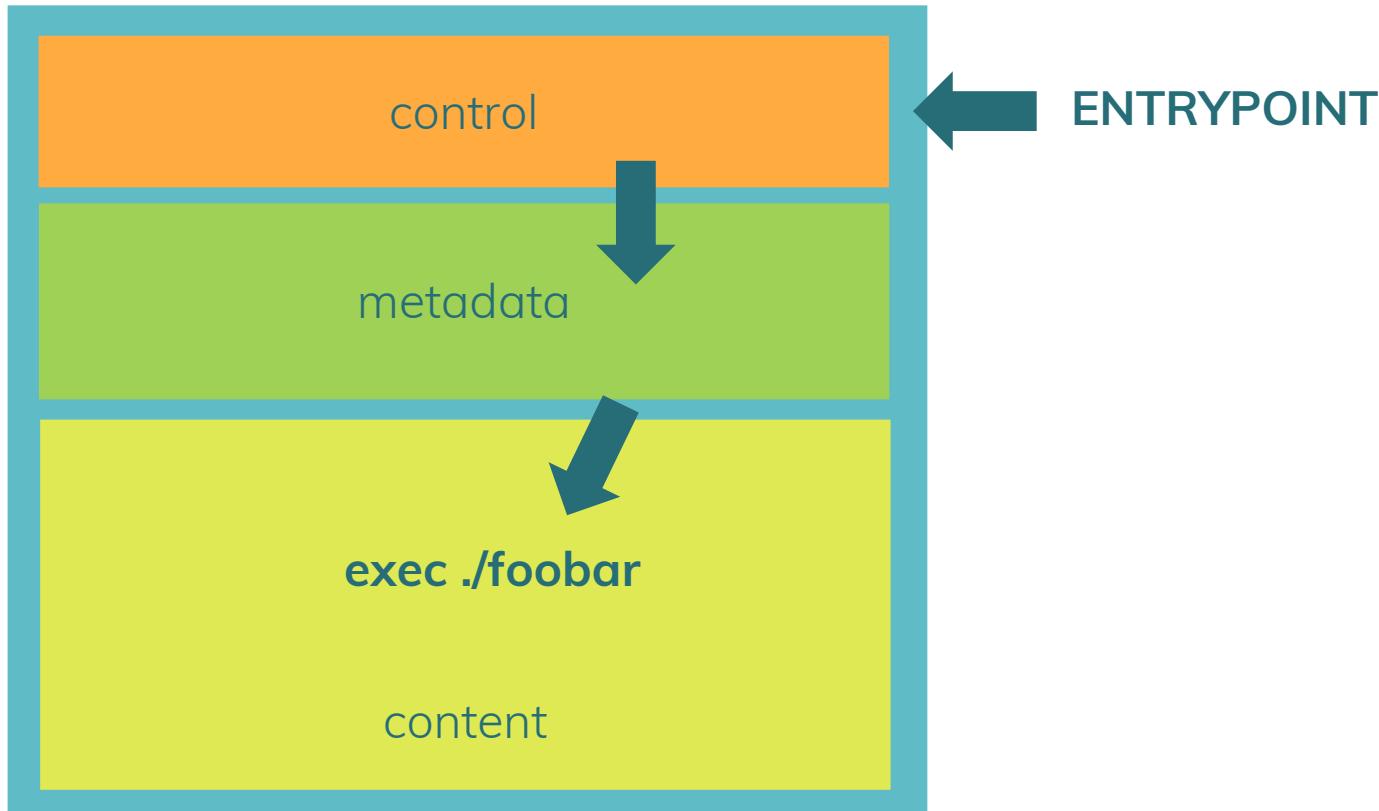
content

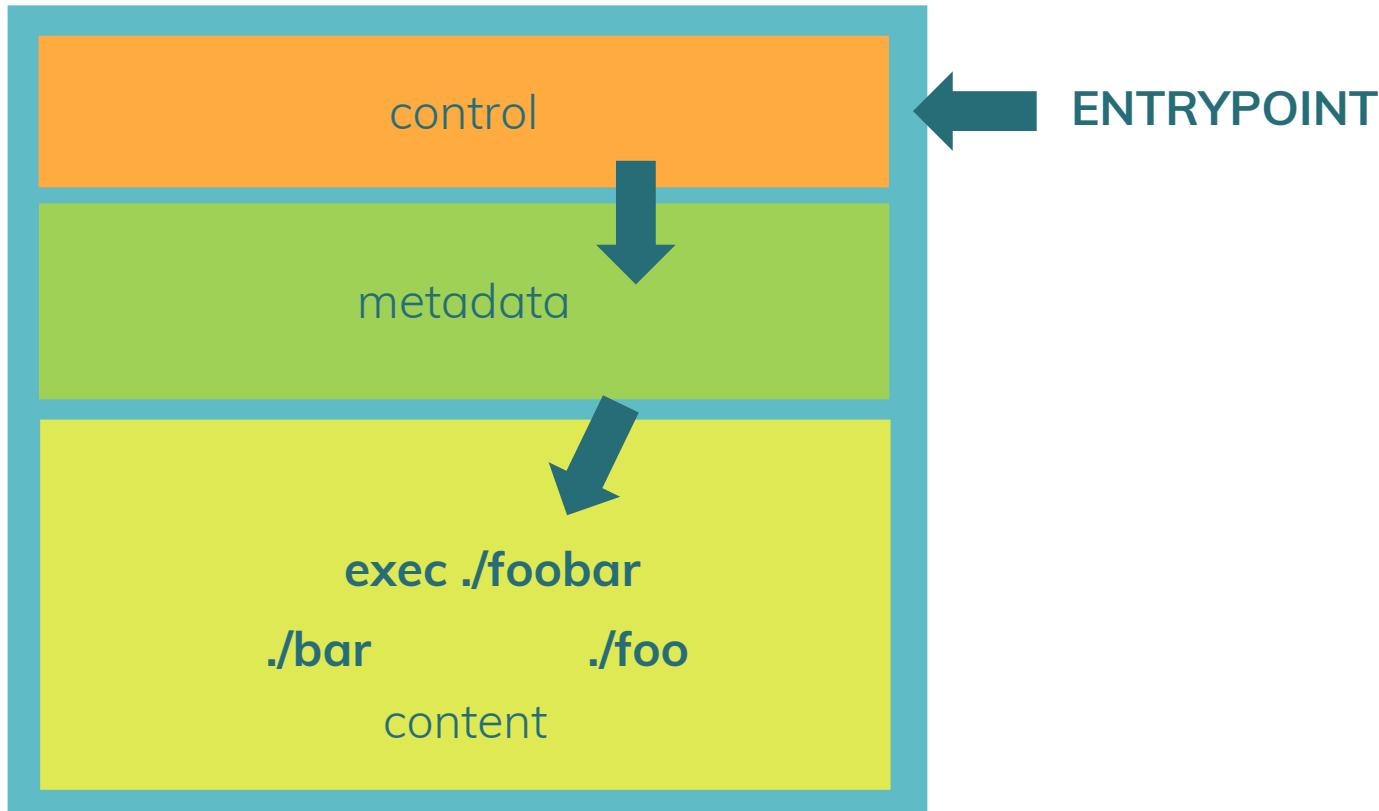


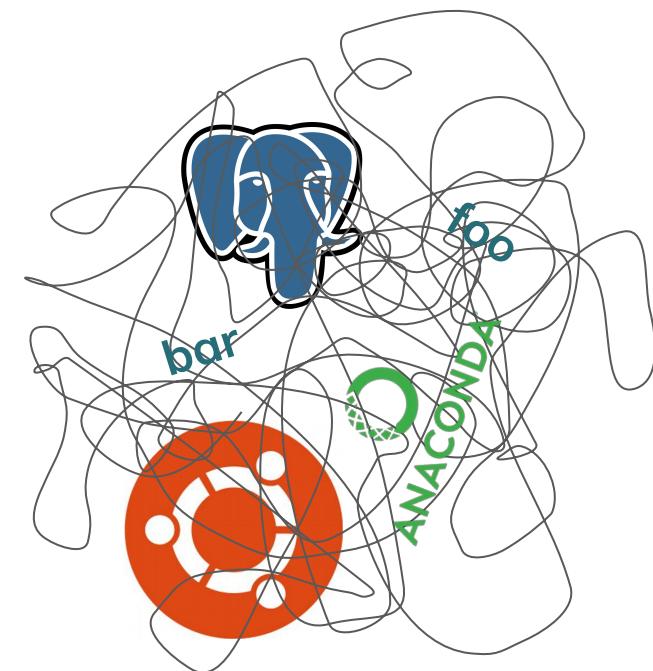
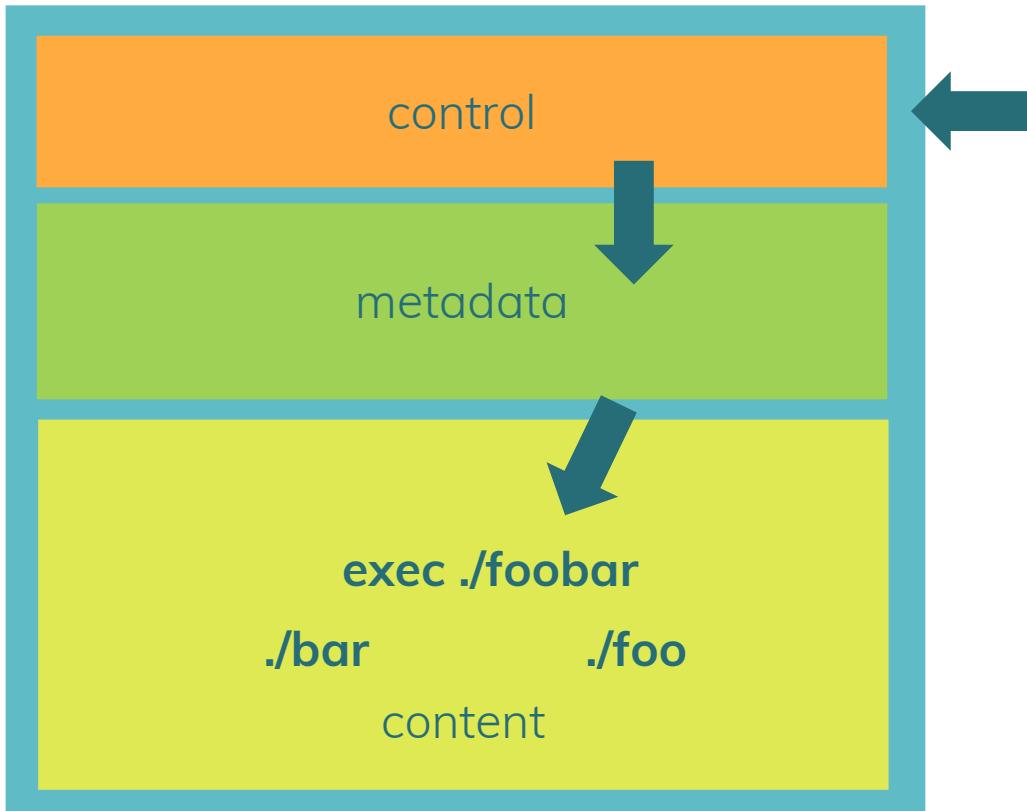
“

Where are foo and bar?









“

I just want to run bar.

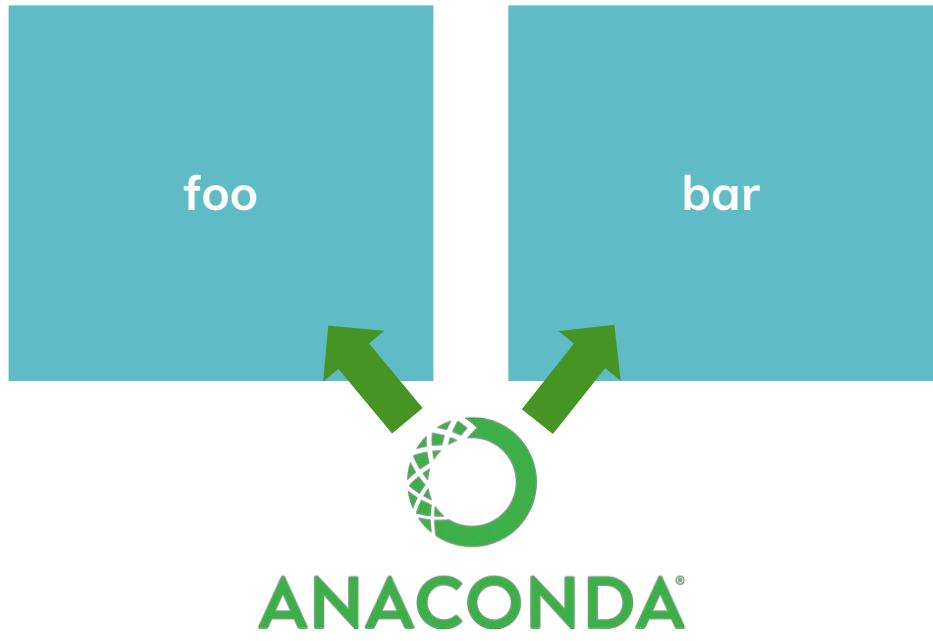
Solution 1: More containers



foo

bar

Solution 1: More containers



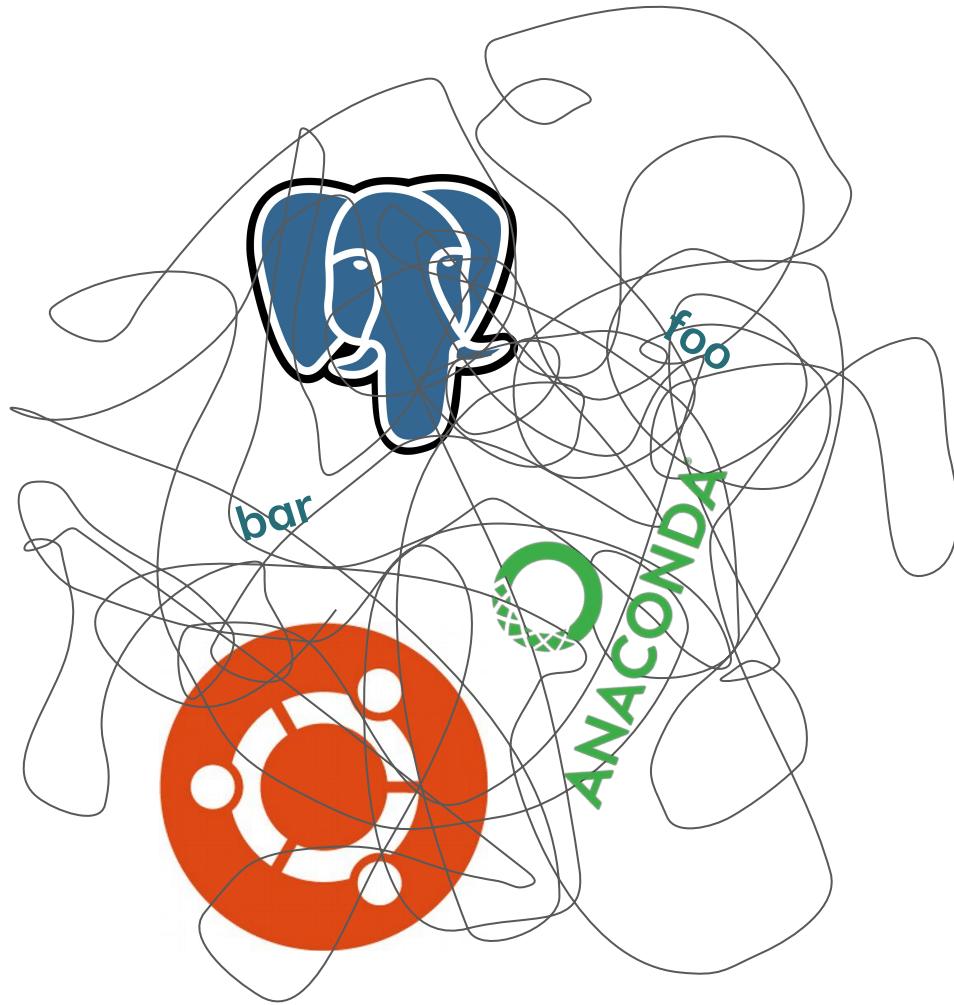
Solution 2: Custom entrypoint

```
If foo:  
    exec foo  
else If bar:  
    exec bar  
else:  
    exec foobar
```

Solution 2: Custom entrypoint

```
If foo:  
  exec foo  
else If bar:  
  exec bar  
else:  
  exec foobar
```

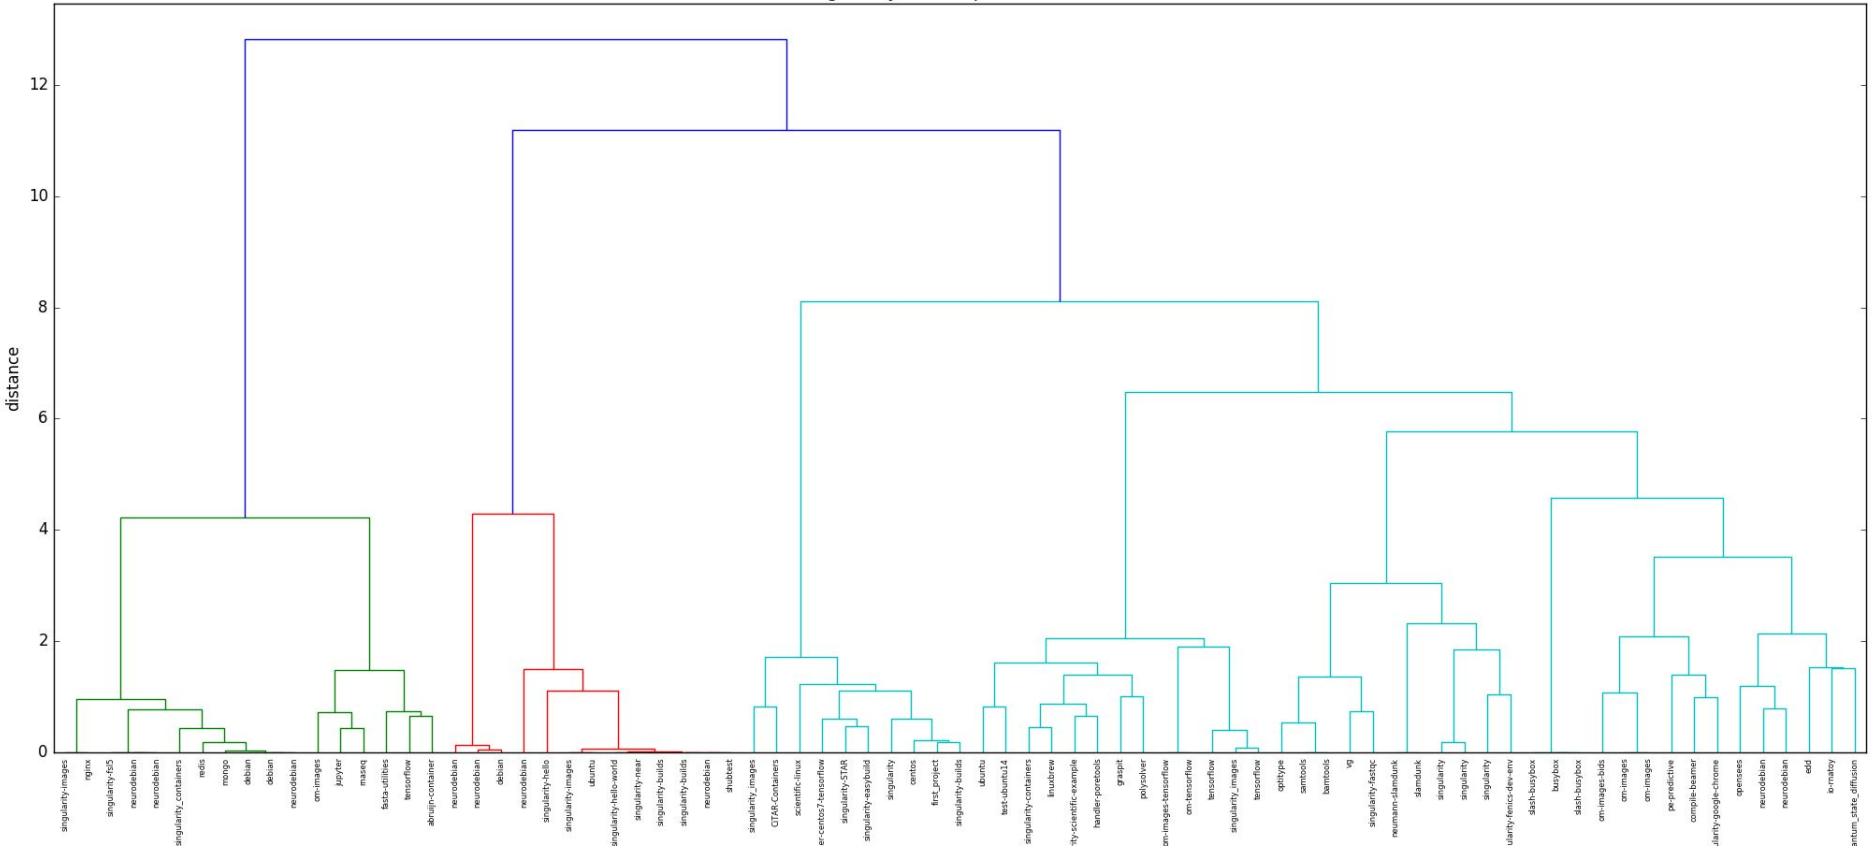
```
BESTAPP=FOO  
BESTAPP=BAR  
EXPORT BESTAPP
```

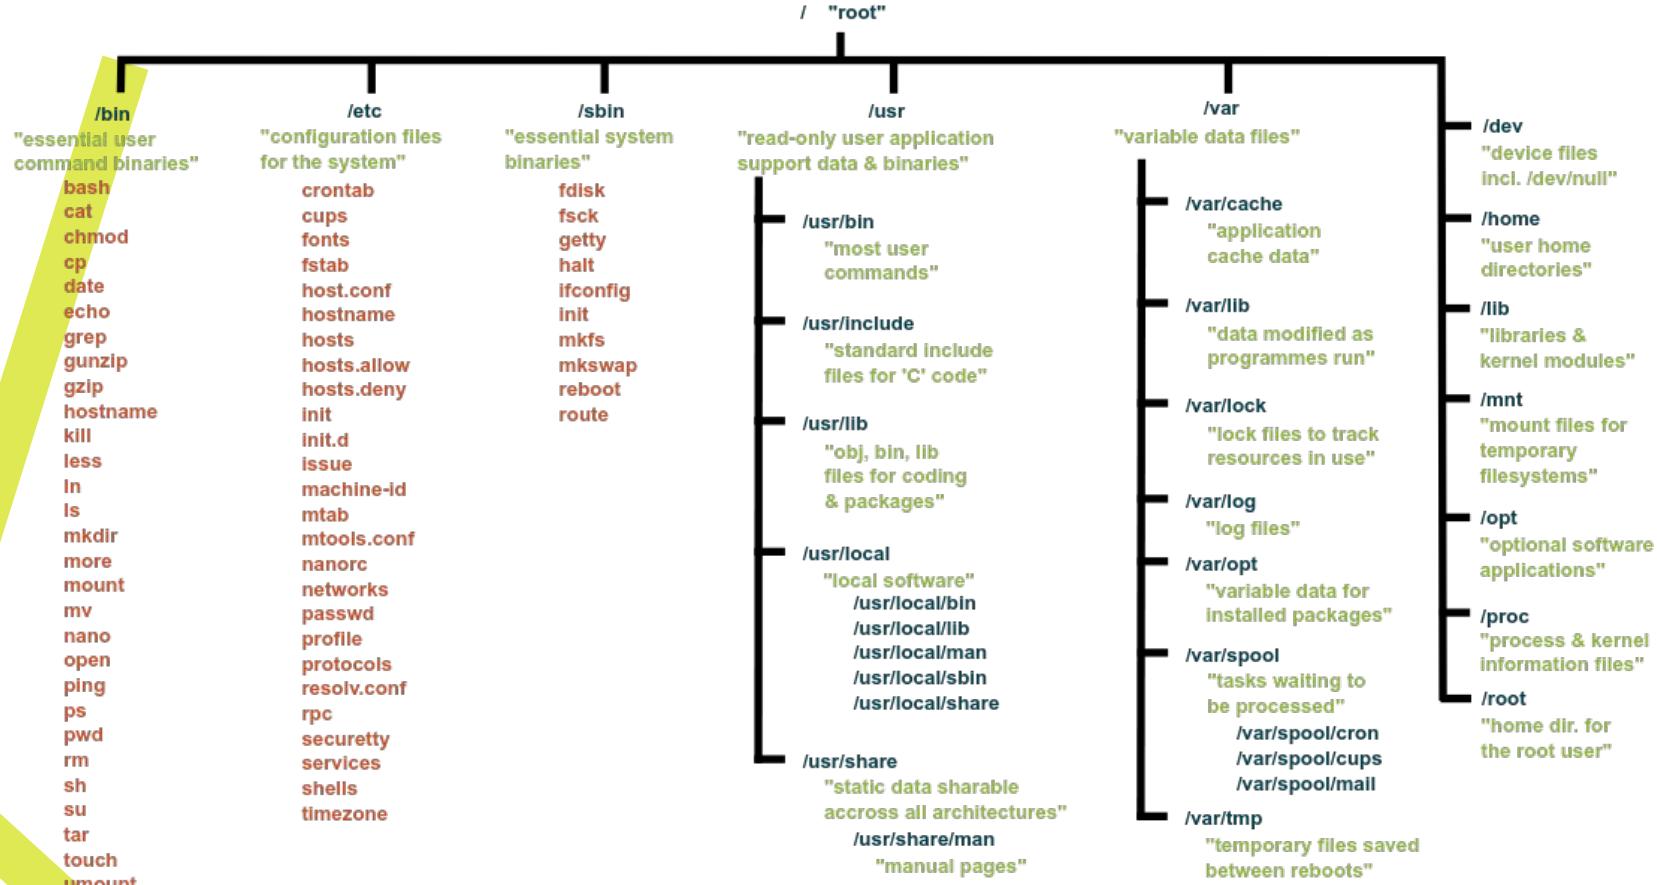


“

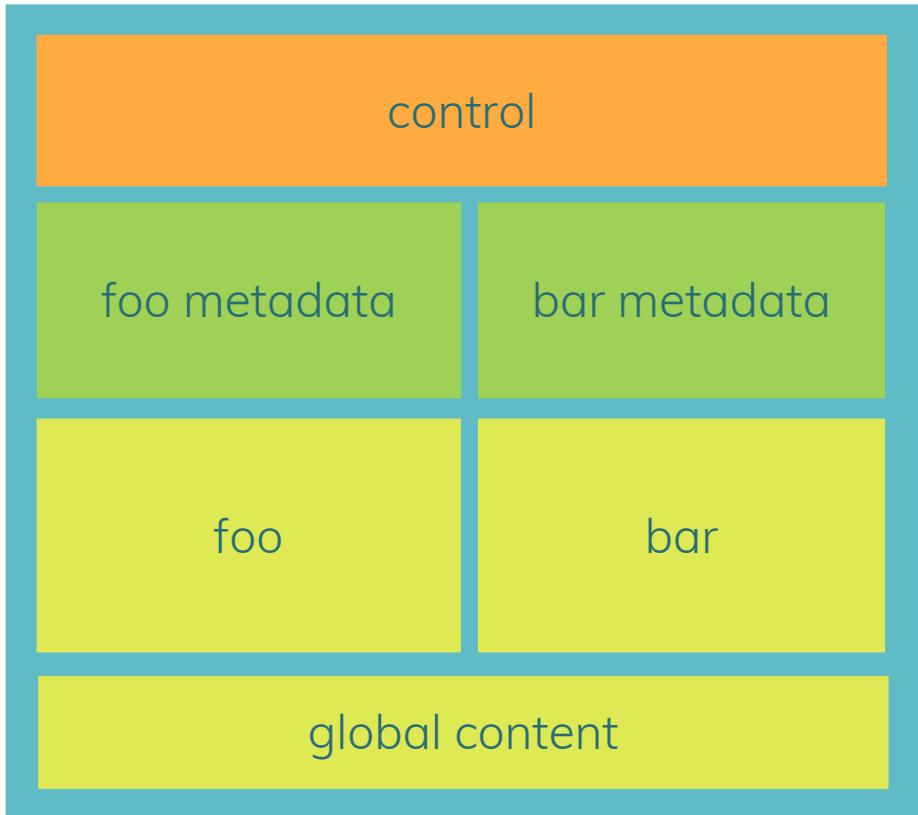
The Linux Filesystem Hierarchy was
not optimized for scientific applications.

Singularity Hub Replication Scores





Linux Filesystem Hierarchy



“

Let's make a Scientific Filesystem!

What is SCIF?

- The specification defines a **filesystem structure** and **environment variable** namespace that interact with **functions** to produce discoverable software applications.

Client

A controller for a SCIF, either for a developer or a user

Integration

A third party software or tool that understands the SCIF structure and interacts with all or some portion of it

Goals

- Enable users to easily generate predictable, discoverable applications
- A recipe format that maps to a SCIF, and vice versa
- Software to generate and manage SCIF

Non-Goals

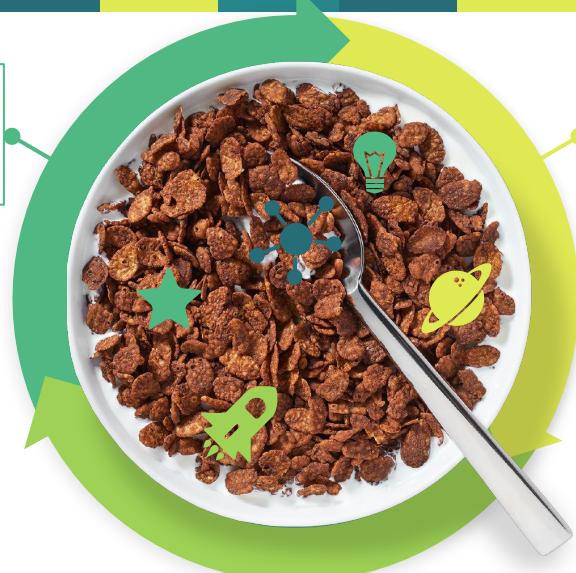
The Scientific Filesystem (SCIF) explicitly does not implement, define, or provide in v0.2:

- Package Management
- Workflow Management
- A mechanism for authentication and authorization or management of file permissions.
- Support or integration with non-linux operating systems.

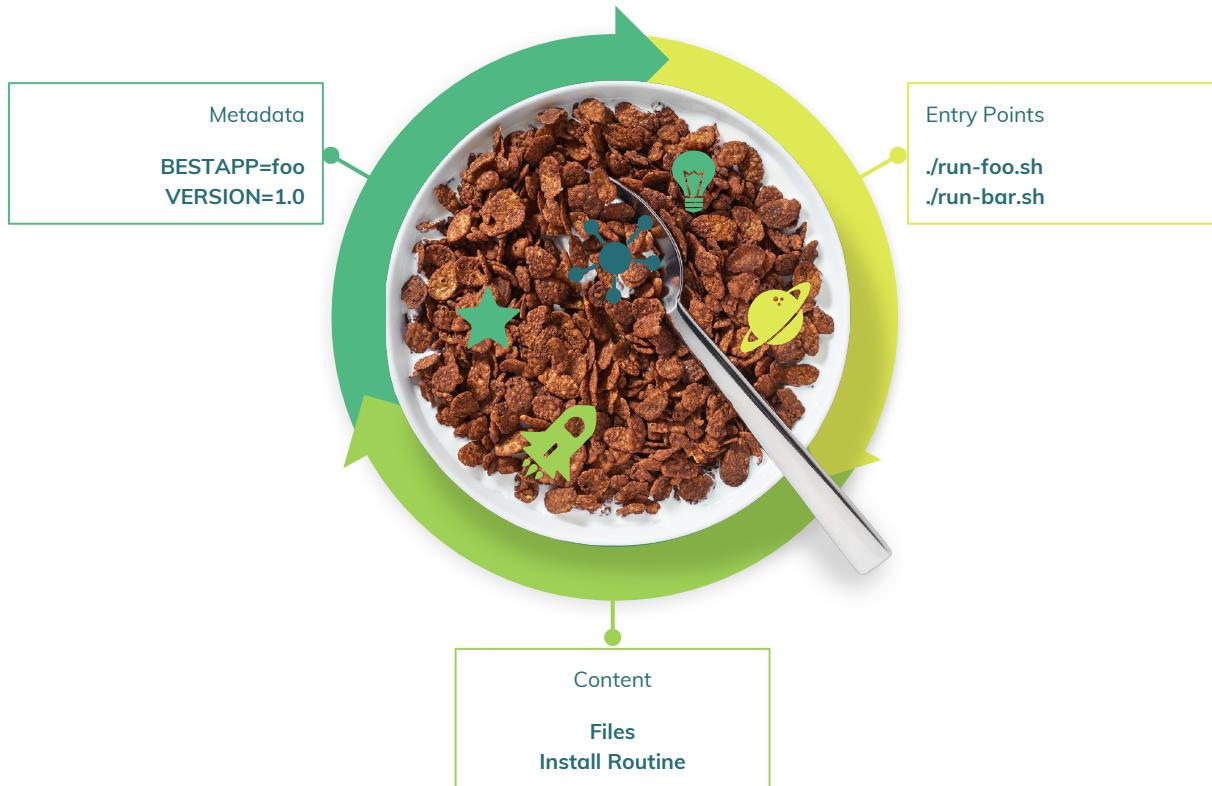
The SCIF Recipe

“
Metadata
BESTAPP=foo
VERSION=1.0

Entry Points
`/run-foo.sh`
`/run-bar.sh`



The SCIF Recipe



%app[action] [name]

```
%apprun hello-world  
/bin/bash hello-world.sh  
%appinstall hello-world  
echo "echo 'Hello World!'" >> $SCIF_APPBIN/hello-world.sh  
chmod u+x $SCIF_APPBIN/hello-world.sh  
%appenv hello-world  
THEBESTAPP=$SCIF_APPNAME  
export THEBESTAPP  
%applabels hello-world  
MAINTAINER Vanessasaur  
%apphelp hello-world  
This is an example "Hello World" application.  
You can install it to a Scientific Filesystem  
(scif) with the command:  
scif install hello-world.scif
```

<https://sci-f.github.io/tutorial-quick-start>

hello-world.scif

Developer Interaction

`recipe.scif`

`write recipe`

`scif install recipe.scif`

`install`

`scif run foo`

`interact`

User Interaction

`recipe.scif`

`write recipe`

`scif install recipe.scif`

`install`

`scif run foo`

`interact`

User Interaction

scif run foo

interact

./container ...

apps

run

exec

help

inspect

shell

pyshell

./scif apps

docker run **vanessa/scif:hw** apps

./scif-cli apps

`./scif apps`

`docker run vanessa/scif:hw apps`

`./scif-cli apps`

./scif run hello-world

docker run **vanessa/scif:hw** run hello-world-echo

./scif-cli run hello-world-echo

```
%apprun hello-world-echo  
exec echo "Hello World!"
```

./scif run hello-world

docker run **vanessa/scif:hw** run hello-world-echo

./scif-cli run hello-world-echo

```
%apprun hello-world-echo  
exec echo "Hello World!"
```

./scif exec hello-world

docker run **vanessa/scif:hw** exec hello-world echo [e]OMG

./scif-cli exec hello-world echo [e]OMG

%appenv hello-world

OMG=TACOS

export OMG

./scif exec hello-world

docker run **vanessa/scif:hw** exec hello-world echo [e]OMG

./scif-cli exec hello-world echo [e]OMG

%appenv hello-world

OMG=TACOS

export OMG

sticker	description	example
[e]	an environment variable prefix	[e]OMG converts to \$OMG
[pipe]	pipe (usually)	env [pipe] grep _SCIF
[out]	output direction (usually >)	cat input.txt [out] output.txt
[in]	input direction (usually <)	
[append]	append to a file (usually >>)	echo "pancakes" » recipe.txt

```
./scif inspect hello-world
```

./scif help hello-world

docker run **vanessa/scif:hw** help hello-world

./scif-cli help hello-world

%apphelp hello-world

This is an example "Hello World" application.

You can install it to a Scientific Filesystem
(scif) with the command:

scif install hello-world.scif

./scif shell hello-world

%appenv hello-world

OMG=TACOS

export OMG

```
./scif shell hello-world
```

`./scif pyshell`

```
--writable, -w      for relevant commands, if writable SCIF is needed

actions:
  actions for Scientific Filesystem

{version,pyshell,shell,preview,help,install,inspect,run,apps,dump,exec}
  scif actions
    version          show software version
    pyshell          Interactive python shell to scientific filesystem
    shell            shell to interact with scientific filesystem
    preview          preview changes to a filesystem
    help             look at help for an app, if it exists.
    install          install a recipe on the filesystem
    inspect          inspect an attribute for a recipe installation
    run              entrypoint to run a scientific filesystem
    apps             list apps installed
    dump             dump recipe
    exec             execute a command to a scientific filesystem
vanessa@vanessa-ThinkPad-T460s:~/Documents/Dropbox/Code/scif/examples/snakefile$ ./snakemake apps
  bwa
graphviz_create_dag
  samtools
  snakemake
vanessa@vanessa-ThinkPad-T460s:~/Documents/Dropbox/Code/scif/examples/snakefile$ sn
```



Architecture

The Scientific Filesystem has a default root on the host to intentionally be separate from standard linux folders, and to not interfere with likely existing folders (e.g., just /data or /apps).

```
/scif  
  /apps  
  /data
```

Organization

scif entrypoint

foo runtime

foo

```
/scif
  /apps
    /foo
      /bin
      /lib
      /scif
        runscript
        runscript.help
        labels.json
        environment.sh
```

Organization

scif entrypoint

foo runtime

foo

```
/scif
  /apps
    /foo
      /bin      ← $PATH
      /lib      ← $LD_LIBRARY_PATH
    /scif
      runscript
      runscript.help
      labels.json
      environment.sh
```

Interaction

```
./container apps ...  
  foo  
  bar
```

```
/scif  
  /apps  
    /bar  
    /foo
```

Interaction

```
./container exec bar [command]
```

```
/scif  
/apps  
/foo  
/bin  
/lib  
/scif  
runscript  
runscript.help  
labels.json  
environment.sh
```

Interaction

```
./container run bar [args]
```

```
/scif  
  /apps  
    /foo  
      /bin  
      /lib  
      /scif  
        runscript  
        runscript.help  
        labels.json  
        environment.sh
```

Interaction

```
./container shell bar $SCIF_SHELL
./container pyshell bar
```

/scif
/apps
/foo
/bin
/lib
/scif
runscript
runscript.help
labels.json
environment.sh

“

Environment Namespace

Environment Namespace

...

Variable	Default	Definition
SCIF_BASE	/scif	the root location for SCIF
SCIF_DATA	/scif/data	the root location for apps data
SCIF_APPS	/scif/apps	the root location for installed apps
SCIF_SHELL	/bin/bash	shell to use for “shell” command
SCIF_PYSHELL	ipython	interactive python shell for pyshell command
SCIF_ENTRYPOINT	/bin/bash	the command to run given no runscript or app defined
SCIF_ENTRYFOLDER	SCIF_BASE	the entry folder to run the entrypoint command
SCIF_MESSAGELEVEL	INFO	a client level of verbosity. Must be one of CRITICAL , ABORT , ERROR , WARNING , LOG , INFO , QUIET , VERBOSE , DEBUG

Table 1 During interaction and runtime of SCIF, the following environment variables must be defined.

```
./container run foo
```

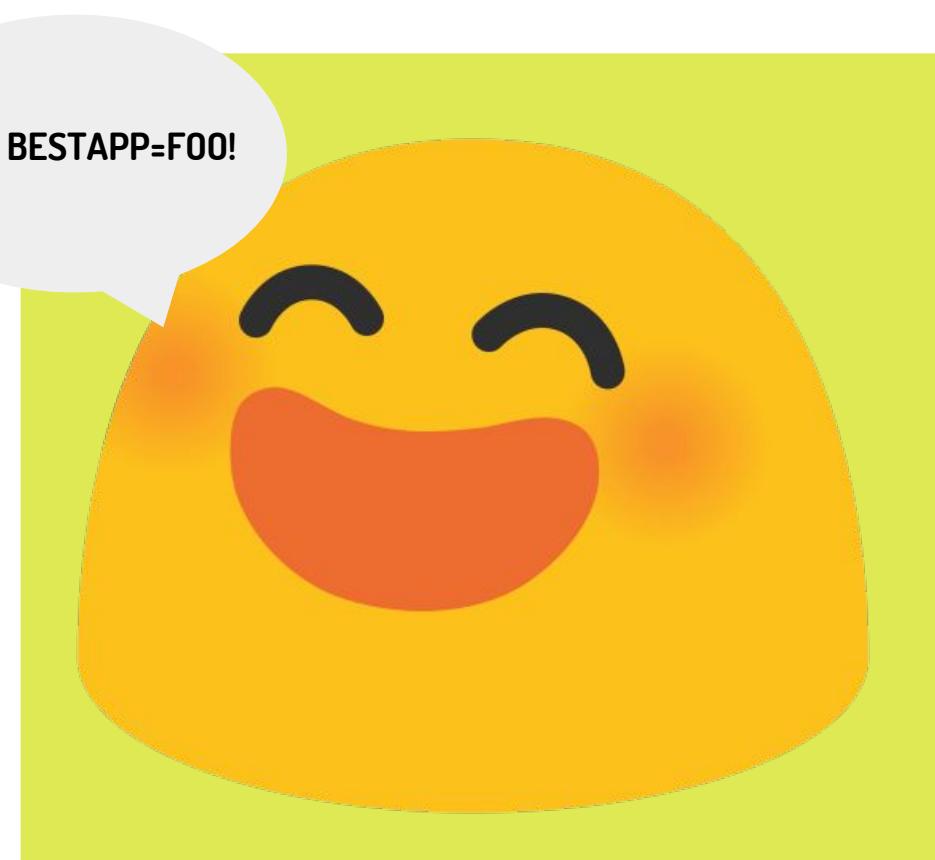


foo active



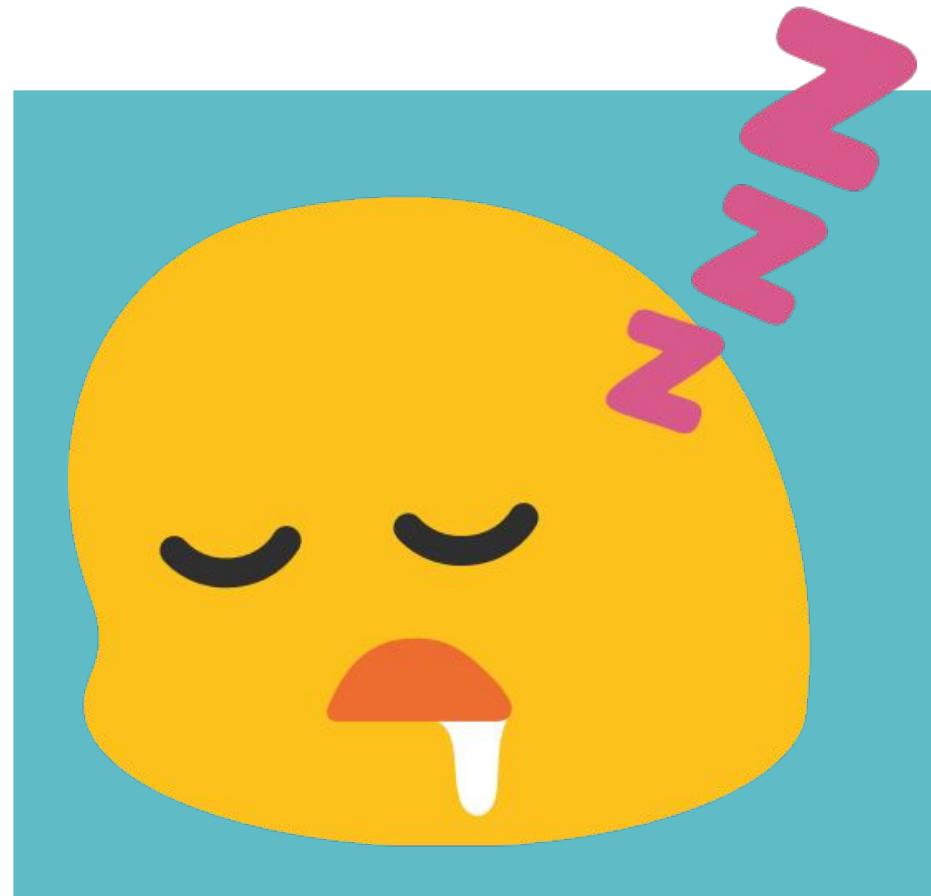
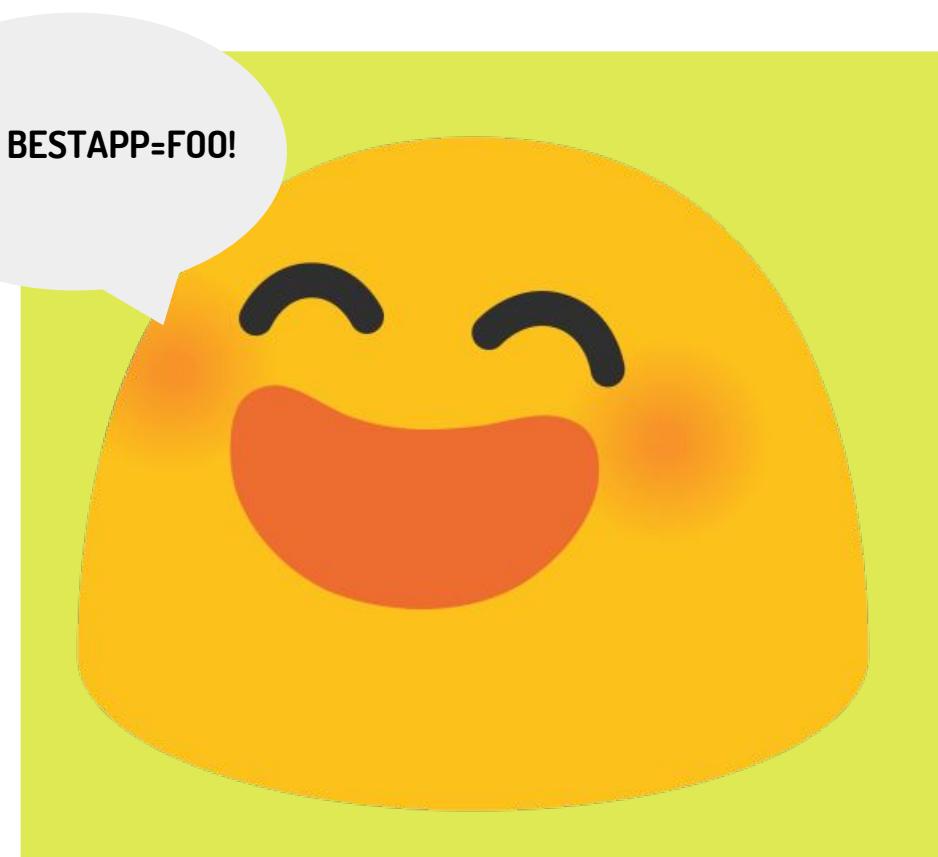
bar sleeping

./container run foo



bar sleeping

./container run foo



Active App Environment Namespace

...

Variable	Default	Definition
SCIF_APPNAME	example	the active software app
SCIF_APPDATA	/scif/data/example	the data root for the active software app
SCIF_APPROOT	/scif/apps/example	the install root for the active software app
SCIF_APPBIN	/scif/apps/example/bin	the app bin, which is automatically added to the path when active
SCIF_APPLIB	/scif/apps/example/lib	the app bin, which is automatically added to the path when active
SCIF_APPMETA	/scif/apps/example/scif	the metadata folder
SCIF_APPHELP	/scif/apps/example/scif/runscript.help	a text file with help to print for the user to the terminal
SCIF_APPRUN	/scif/apps/example/scif/runscript	the commands to run as the app entrypoint
SCIF_APPTEST	/scif/apps/example/scif/test	the commands to run to test the app
SCIF_APPLABELS	/scif/apps/example/scif/labels.json	a key:value json lookup dictionary of labels
SCIF_APPENV	/scif/apps/example/scif/environment.sh	a shell script to source for the software app environment

Sleeping App Environment Namespace

...

Variable	Default	Definition
SCIF_APPNAME_sleeper	sleeper	the inactive software app
SCIF_APPDATA_sleeper	/scif/data/sleeper	the data root for the inactive software app
SCIF_APPROOT_sleeper	/scif/apps/sleeper	the install root for the active software app
SCIF_APPBIN_sleeper	/scif/apps/sleeper/bin	the app bin, which is automatically added to the path when active
SCIF_APPLIB_sleeper	/scif/apps/sleeper/lib	the app bin, which is automatically added to the path when active
SCIF_APPMETA_sleeper	/scif/apps/sleeper/scif	the metadata folder
SCIF_APPHELP_sleeper	/scif/apps/sleeper/scif/runscript.help	a text file with help to print for the user to the terminal
SCIF_APPRUN_sleeper	/scif/apps/sleeper/scif/runscript	the commands to run as the app entrypoint
SCIF_APPTEST_sleeper	/scif/apps/sleeper/scif/test	the commands to run to test the app
SCIF_APPLABELS_sleeper	/scif/apps/sleeper/scif/labels.json	a key:value json lookup dictionary of labels
SCIF_APPENV_sleeper	/scif/apps/sleeper/scif/environment.sh	a shell script to source for the software app environment

“

What can I do with a SCIF?

- expose multiple container entrypoints
- package metrics/evaluation alongside the analysis
- SCIF apps as different working environments, runtime
- SCIF apps for different runtime contexts

PAPER

The Scientific Filesystem (SCIF)

Vanessa Sochat^{1,*†}

¹Stanford Research Computing Center and ²Stanford University School of Medicine

* vsochat@stanford.edu

Abstract

Background. Here we present the Scientific Filesystem (SCIF), an organizational format that supports exposure of executables and metadata for discoverability of scientific applications. The format includes a known filesystem structure, a definition for a set of environment variables describing it, and functions for generation of the variables and interaction with the libraries, metadata, and executables located within. SCIF makes it easy to expose metadata, multiple environments, installation steps, files, and entrypoints to render scientific applications consistent, modular, and discoverable. A SCIF can be installed on a traditional host or in a container technology such as Docker or Singularity. We will start by reviewing the background and rationale for the Scientific Filesystem, followed by an overview of the specification, and the different levels of internal modules ("apps") that the organizational format affords. Finally, we will conclude with some examples of how the SCIF can be used to facilitate reproducibility and portability of scientific applications.

“

What are we going to do today?

What are we going to do today?

Build Containers

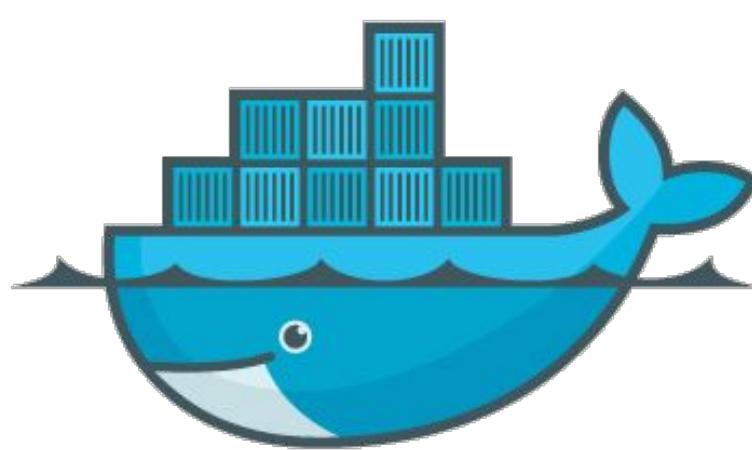
Across different container technologies, all with the same scientific filesystem (SCIF)

Run a Workflow

The same snakemake workflow made modular by SCIF applications

Evaluate

How do the different container technologies compare?



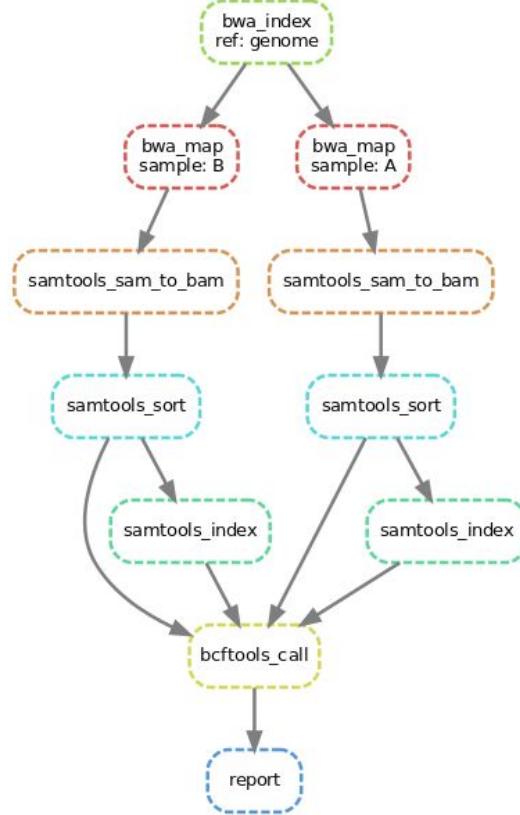
N
SHIFTER



Charliecloud

Scientific Pipeline

- Felix Bartusch
- variant calling
- Snakemake



“

What level of modularity?

- Level 1: Container Modularity

- One entrypoint
- Run the entire thing, or don't

`./container`

- Level 2: Pipeline Modularity

- Multiple entrypoints
- Correspond to pipeline steps
- Step → collection of commands

`./container apps`
`mapping`
`sorting`

`./container run mapping`



- Level 3: Developer Modularity

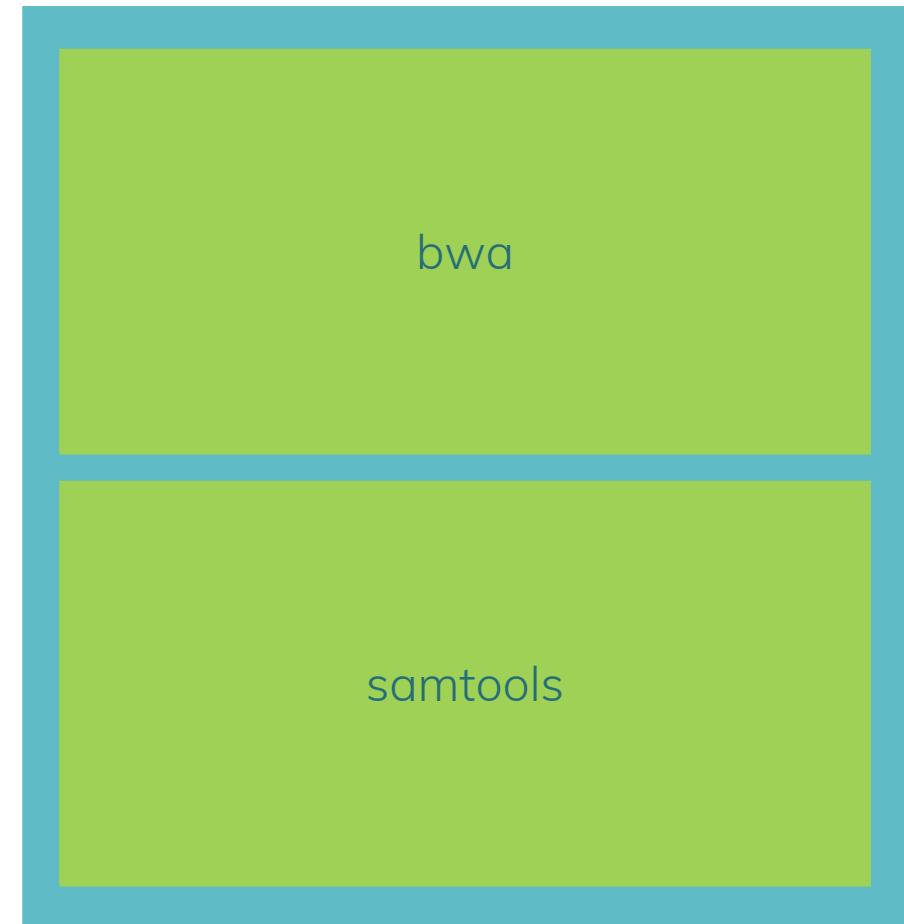
- Expose lowest level executables
- Ideal for building pipelines
- Step → core tools

./container apps

bwa

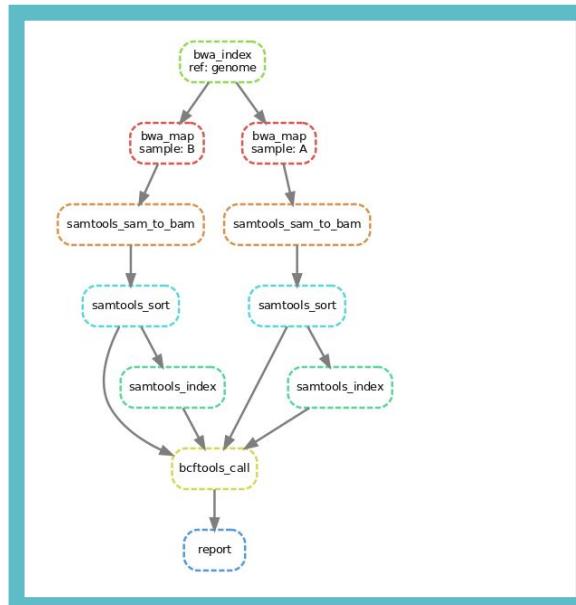
samtools

./container run samtools



A Scientific Filesystem Design for Variant Calling

What level of modularity should be exposed to interact with the container?



bwa

samtools

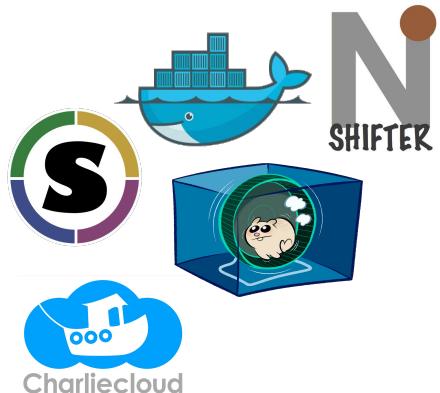
graphviz_create_dag

snakemake

./container apps
bwa
samtools
graphviz_create_dag
snakemake

What did I do with the container?

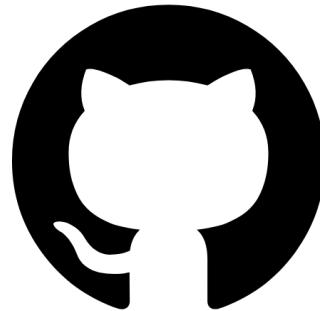
Build Containers



Run a Workflow



Evaluate



<https://github.com/sci-f/snakefile.scif>

Container Comparisons

What steps did we take across container technologies?

0. install the software

1. build the container

2. run components of the workflow

3. run the entire workflow

What did we learn?

What did I learn about user interaction?

Good software without documentation and tutorial for several use cases is akin to a treasure chest on a remote planet.

Container Comparisons

What steps did we take across container technologies?

0. install the software

Installation

- Docker
- Singularity
- Charliecloud
- runc
- Shifter

Container Comparisons

What steps did we take across container technologies?

0. install the software

1. Documentation is the first interaction

- a. Finding it
- b. Is it up to date?
- c. Is there too much detail?
- d. Is there a clear avenue to get more help?



Container Comparisons

What steps did we take across container technologies?

0. install the software

1. build the container

Container Comparisons

What steps did we take across container technologies?

0. install the software

1. build the container

1. Examples are needed for learning

- a. Does it show me all the things I might need?
 - i. binds
 - ii. options
 - iii. permissions



Container Comparisons

What steps did we take across container technologies?

0. install the software

1. build the container

2. run components of the workflow

Container Comparisons

What steps did we take across container technologies?

0. install the software

1. build the container

2. run components of the workflow

1. Establish boundaries
 - a. show inside vs. outside



Container Comparisons

What steps did we take across container technologies?

0. install the software

1. build the container

2. run components of the workflow

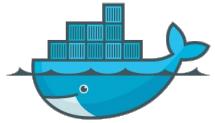
3. run the entire workflow

“

Run the entire workflow !



```
runc --root /tmp/runc run --bundle /tmp/runc snakemake.runc
```



```
docker run -v $PWD/data:/scif/data:z -v $PWD:/scif/data/snakefile:z -it  
vanessa/snakefile.scif run snakefile all
```



```
singularity run --bind data:/scif/data snakefile.simg run snakefile all
```



Charliecloud

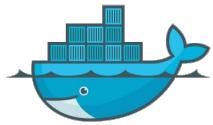
```
ch-run --cd $PWD -b data:/scif/data /var/tmp/snakefile.ch -- /opt/conda/bin/scif run  
snakefile all
```



```
shifter --image=vanessa/snakefile.scif --volume `pwd`/data:/scif/data --workdir /scif/data  
/opt/conda/bin/scif run snakefile all
```



```
runc --root /tmp/runc run --bundle /tmp/runc snakmake.runc
```



```
docker run -v $PWD/data:/scif/data:z -v $PWD:/scif/data/snakefile:z -it  
vanessa/snakefile.scif run snakefile all
```



```
singularity run --bind data:/scif/data snakefile.simg run snakefile all
```



```
ch-run --cd $PWD -b data:/scif/data /var/tmp/snakefile.ch -- /opt/conda/bin/scif run  
snakefile all
```

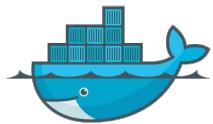


```
shifter --image=vanessa/snakefile.scif --volume `pwd`/data:/scif/data --workdir /scif/data  
/opt/conda/bin/scif run snakefile all
```

executable



runc --root /tmp/runc run --bundle /tmp/runc snakemake.runc



docker run -v \$PWD/data:/scif/data -v \$PWD:/scif/data/snakefile -it vanessa/snakefile.scif run snakefile all



singularity run --bind data:/scif/data snakefile.simg run snakefile all



Charliecloud

ch-run --cd \$PWD -b data:/scif/data /var/tmp/snakefile.ch -- /opt/conda/bin/scif run snakefile all



shifter --image=vanessa/snakefile.scif --volume `pwd`/data:/scif/data --workdir /scif/data /opt/conda/bin/scif run snakefile all

executable + action



```
runc --root /tmp/runc run --bundle /tmp/runc snakmake.runc
```



```
docker run -v $PWD/data:/scif/data -v $PWD:/scif/data/snakefile -it  
vanessa/snakefile:scif run snakefile all
```



```
singularity run --bind data:/scif/data snakefile.simg run snakefile all
```



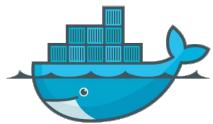
```
ch-run --cd $PWD -b data:/scif/data /var/tmp/snakefile.ch -- /opt/conda/bin/scif run  
snakefile all
```



```
shifter --image=vanessa/snakefile:scif --volume `pwd`/data:/scif/data --workdir /scif/data  
/opt/conda/bin/scif run snakefile all
```



```
runc --root /tmp/runc run --bundle /tmp/runc snakemake.runc
```



```
docker run -v $PWD/data:/scif/data -v $PWD:/scif/data/snakefile -it  
vanessa/snakefile.scif run snakefile all
```



```
singularity run --bind data:/scif/data snakefile.simg run snakefile all
```



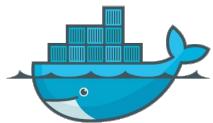
```
ch-run --cd $PWD -b data:/scif/data /var/tmp/snakefile.ch -- /opt/conda/bin/scif run  
snakefile all
```



```
shifter --image=vanessa/snakefile.scif --volume `pwd`/data:/scif/data --workdir  
/scif/data /opt/conda/bin/scif run snakefile all
```



```
runc --root /tmp/runc run --bundle /tmp/runc snakmake.runc
```



```
docker run -v $PWD/data:/scif/data -v $PWD:/scif/data/snakefile -it  
vanessa/snakefile.scif run snakefile all
```



```
singularity run --bind data:/scif/data snakefile.simg run snakefile all
```



```
ch-run --cd $PWD -b data:/scif/data /var/tmp/snakefile.ch -- /opt/conda/bin/scif run  
snakefile all
```



```
shifter --image=vanessa/snakefile.scif --volume `pwd`/data:/scif/data --workdir  
/scif/data /opt/conda/bin/scif run snakefile all
```



executable action options container options command



<https://github.com/sci-f/snakefile.scif>



“

Are the outputs comparable?

Results are the same

```
$ diff results/docker/calls.vcf results/singularity/calls.vcf
28c28
< ##bcftools_callCommand=call -mv -; Date=Sat Mar 10 01:07:35 2018
---
> ##bcftools_callCommand=call -mv -; Date=Fri Mar  9 20:07:47 2018

$ diff results/docker/calls.vcf results/runc/calls.vcf
28c28
< ##bcftools_callCommand=call -mv -; Date=Sat Mar 10 01:07:35 2018
---
> ##bcftools_callCommand=call -mv -; Date=Sat Mar 10 01:08:05 2018

$ diff results/docker/calls.vcf results/shifter/calls.vcf
28c28
< ##bcftools_callCommand=call -mv -; Date=Sat Mar 10 01:07:35 2018
---
> ##bcftools_callCommand=call -mv -; Date=Mon Mar 12 20:45:21 2018

$ diff results/docker/calls.vcf results/charliecloud/calls.vcf
28c28
< ##bcftools_callCommand=call -mv -; Date=Sat Mar 10 01:07:35 2018
---
> ##bcftools_callCommand=call -mv -; Date=Sat Mar 10 01:07:57 2018
```



vs



What did we learn?

Do the technologies produce the same result?

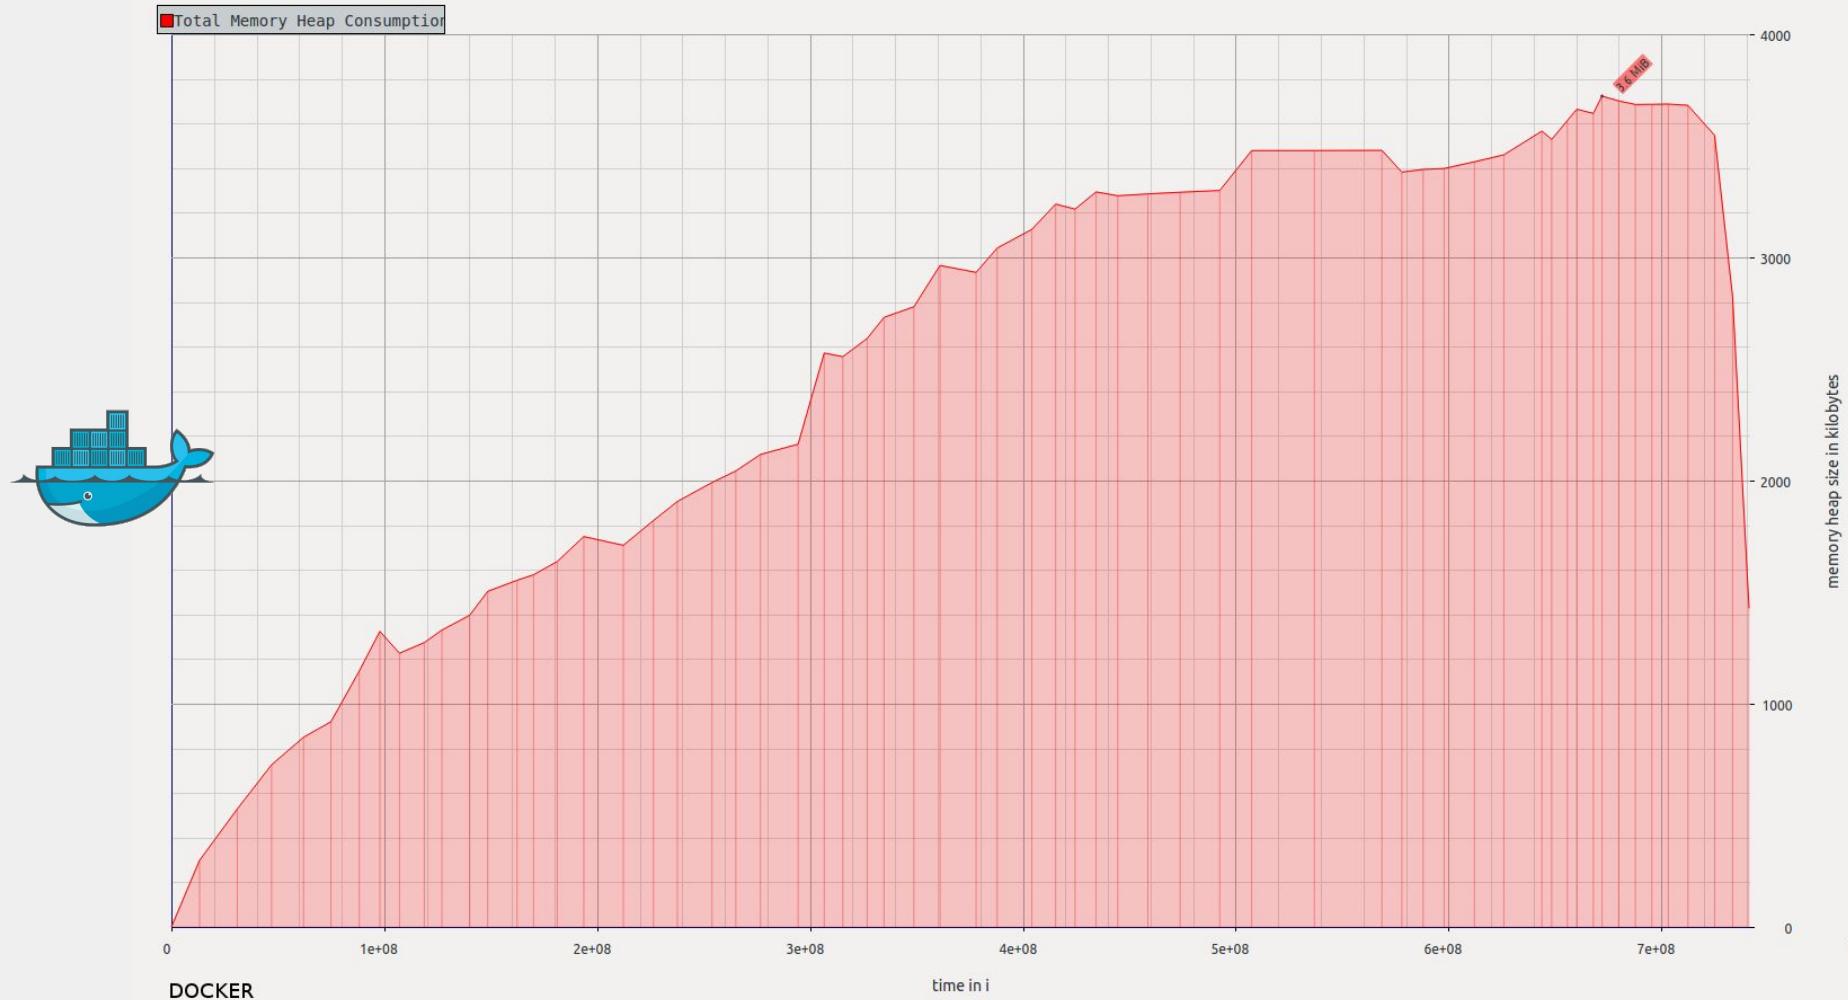
From the perspective of the scientist, we have comparable resource usage and produce the same output.

“

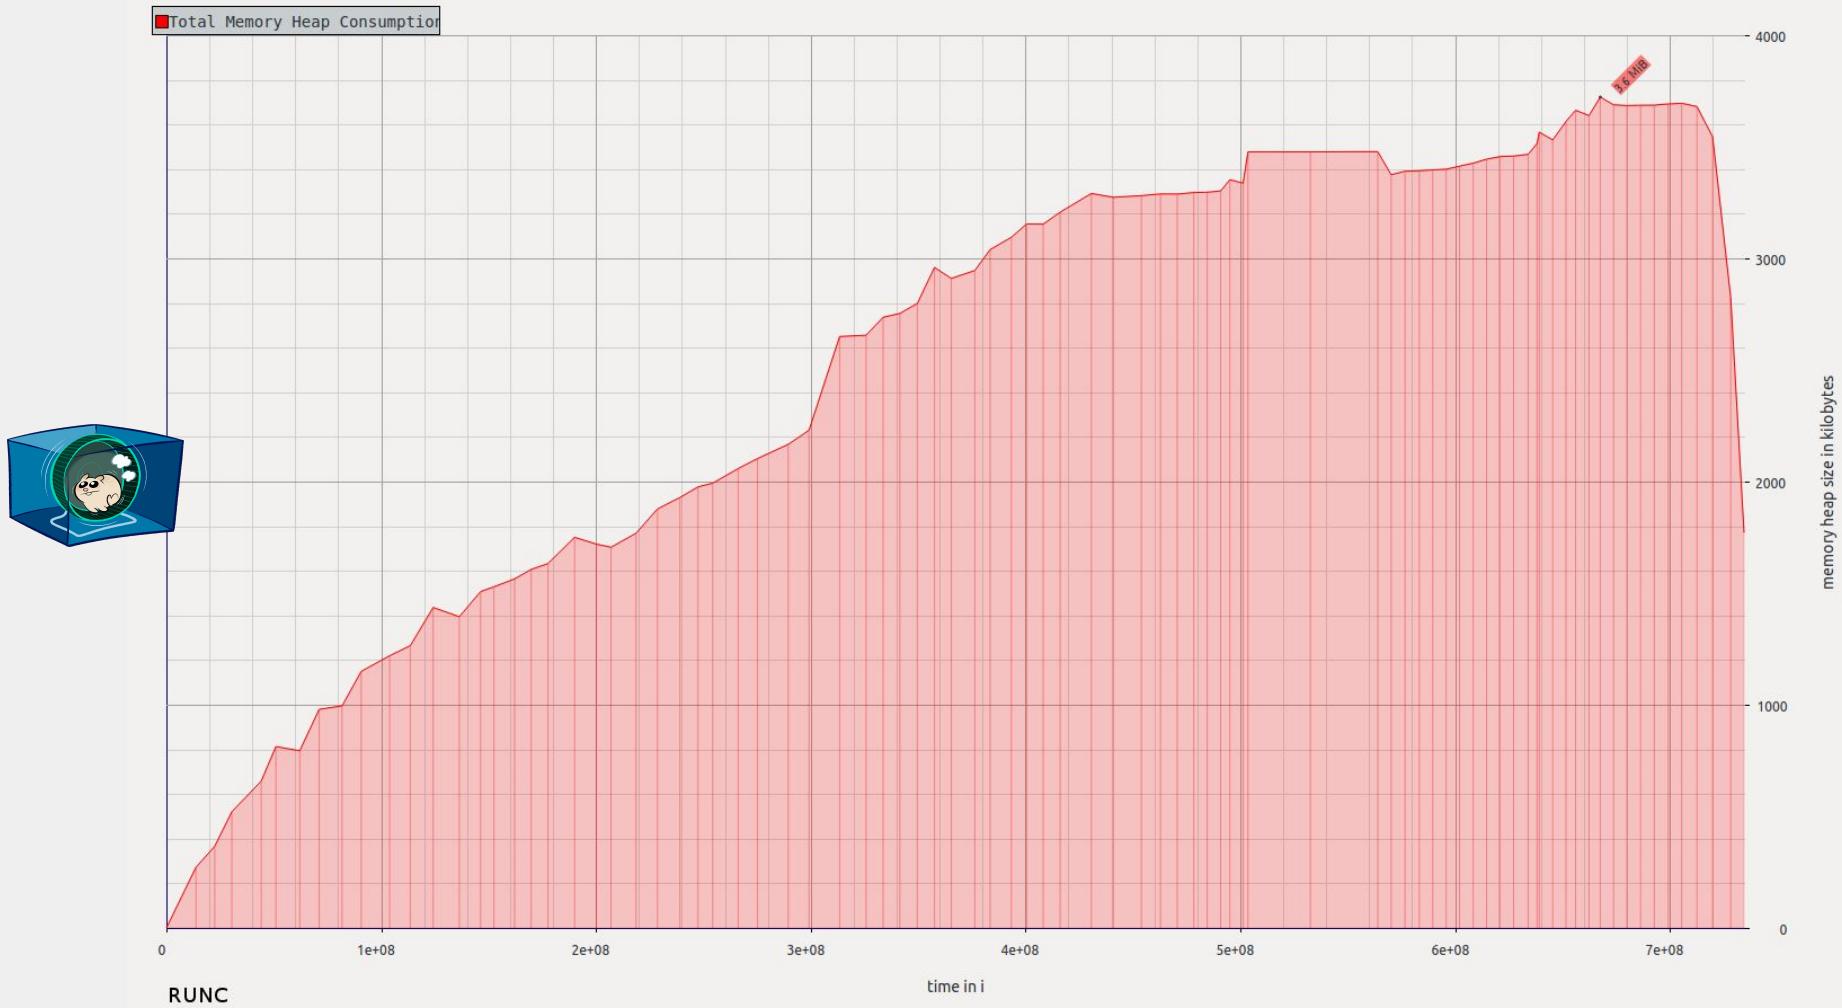
What about resource usage?

```
singularity run --bind data/:/scif/data snakemake.simg run valgrind  
singularity run --bind data/:/scif/data snakemake.simg run timeit
```

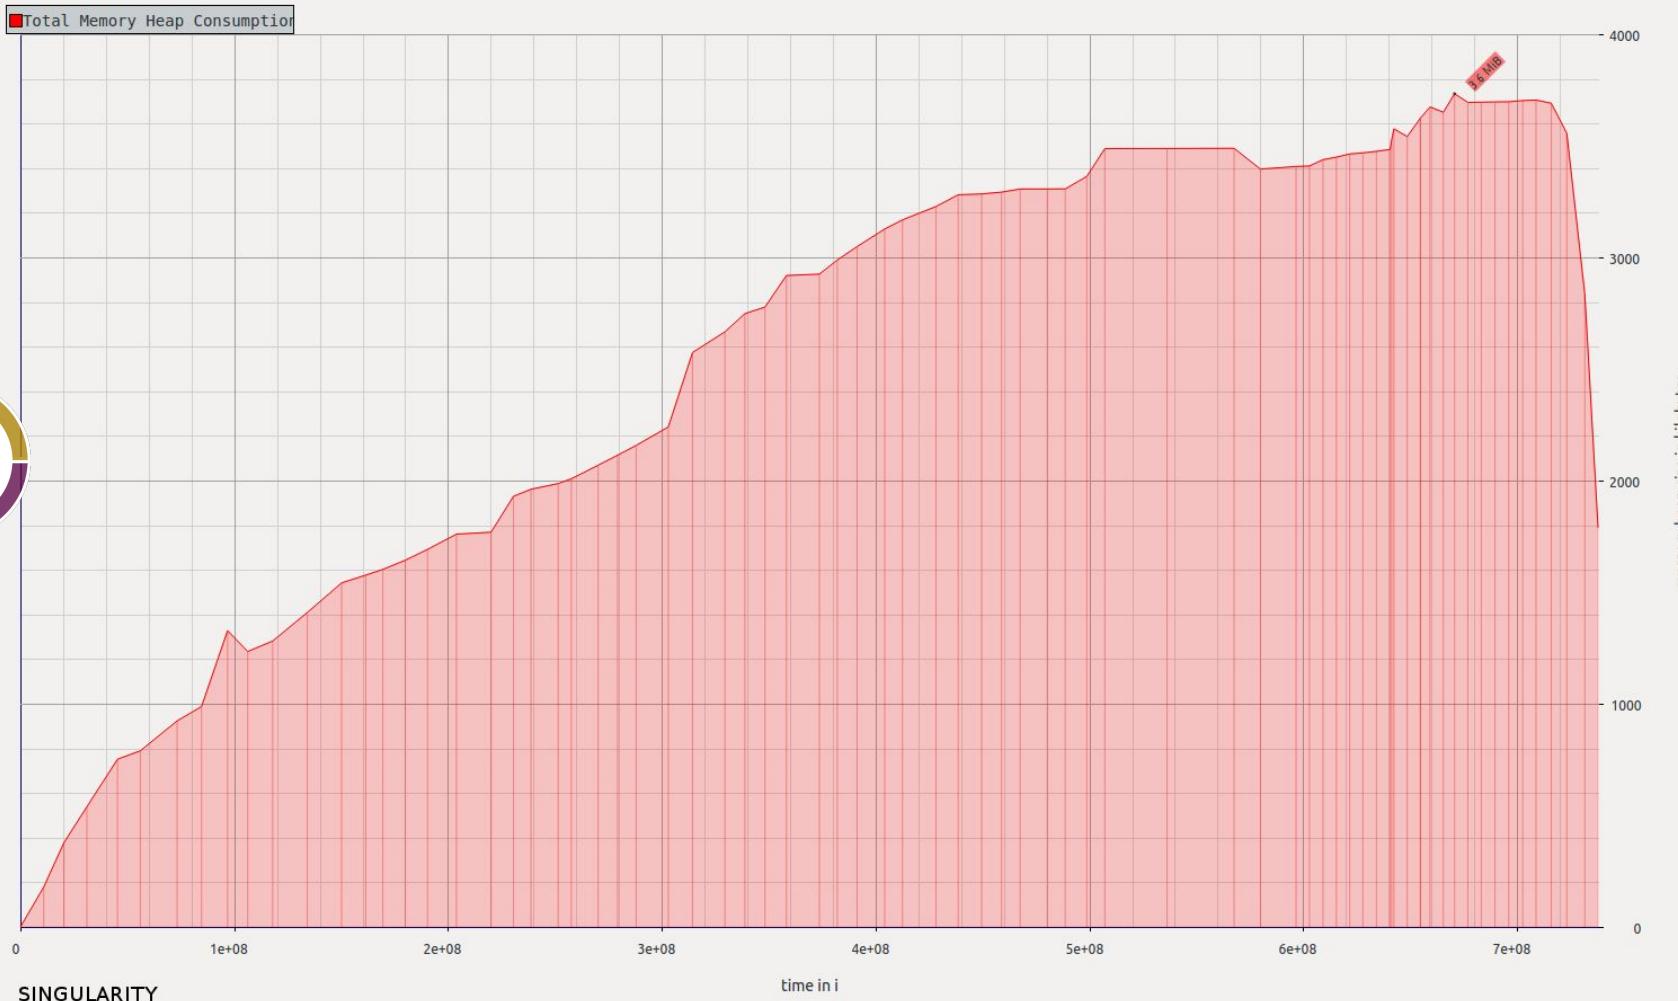
Memory consumption of /opt/conda/bin/snakefile
Peak of 3.6 MiB at snapshot #53



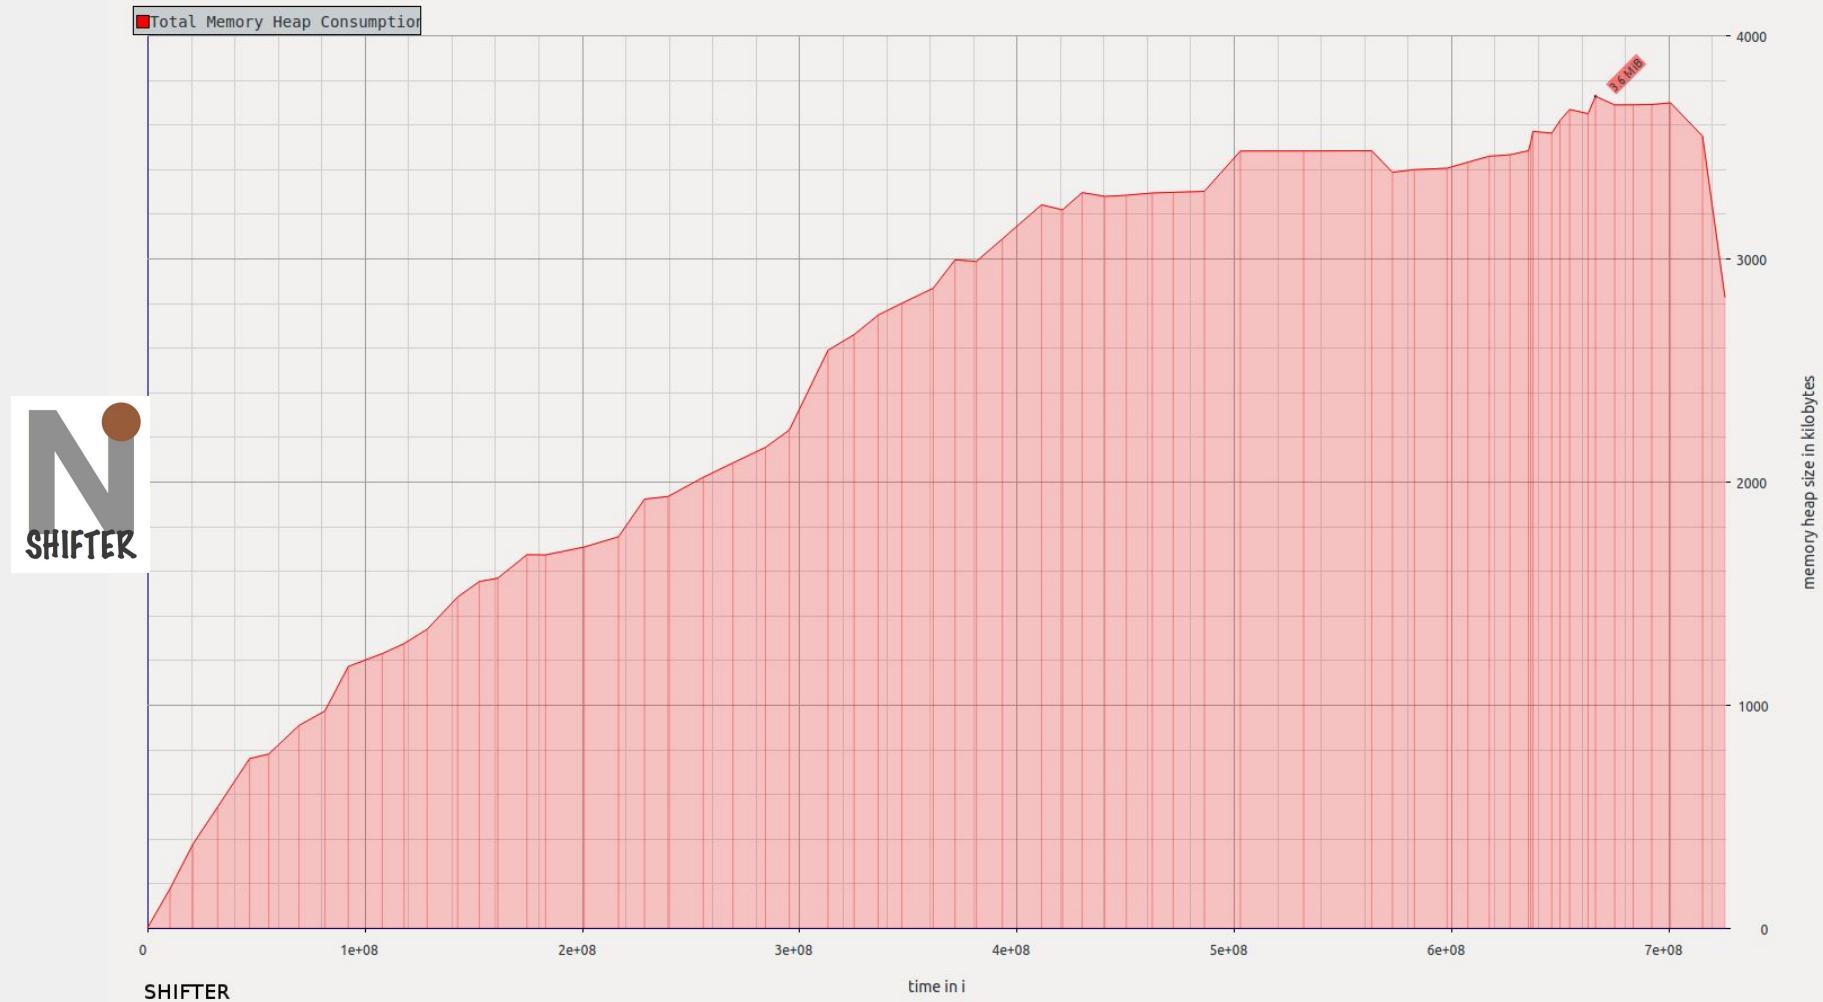
Memory consumption of /opt/conda/bin/snakefile
Peak of 3.6 MiB at snapshot #74



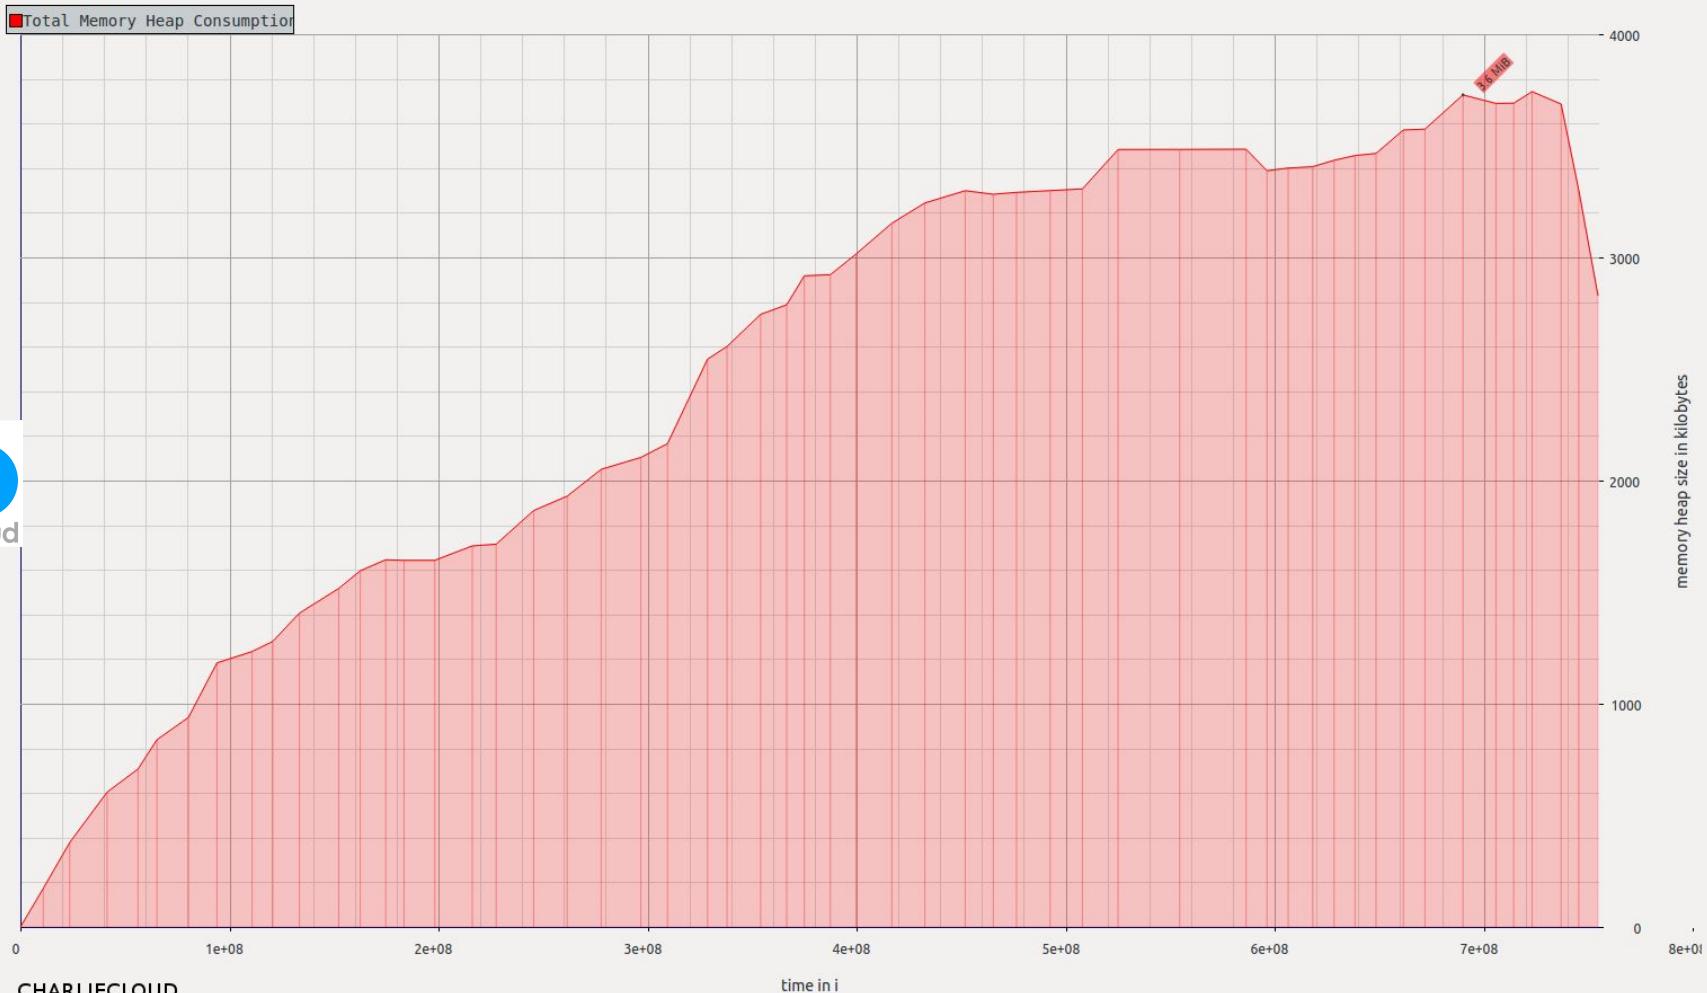
Memory consumption of /opt/conda/bin/snakefile
Peak of 3.6 MiB at snapshot #64



Memory consumption of /opt/conda/bin/snakefile
Peak of 3.6 MiB at snapshot #57

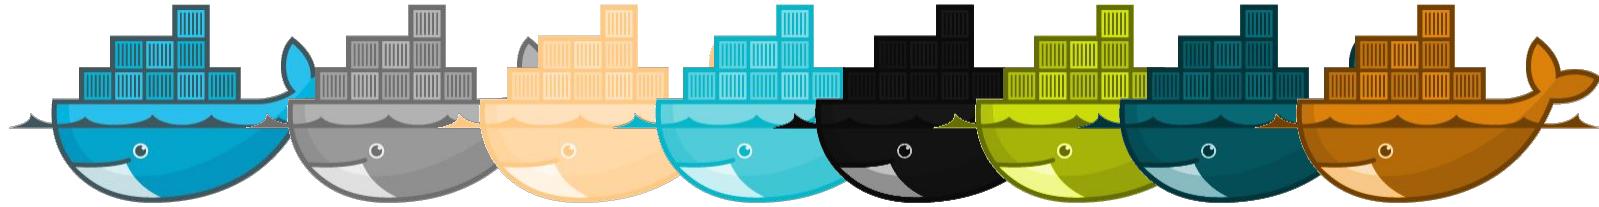


Memory consumption of /opt/conda/bin/snakefile
Peak of 3.6 MiB at snapshot #48



Charliecloud

...but how do my choices of the content, and host, influence my result?



Does Heap Consumption Vary Across Machine Type?

```
vanessa@vanessa-ThinkPad-T460s:~/Documents/Dropbox/Code/scif/examples/snake&gt;make.scif/results/cloud$ sregistry build --config config-n1-highcpu-2.json https://www.github.com/sci-f/snake&gt;make.scif
[client|google-compute] [database|sqlite:///home/vanessa/.singularity/sregistry.db]
[bucket][sregistry-vanessa]
INSTANCE purple-lentil-2448
sci-f-snake&gt;make.scif-builder https://vsocz.github.io/builders/cloud/google/compute/ubuntu/metrics/securebuild-2.4.3
Robot Logger: http://35.203.190.60
Allow a few minutes for web server install, beepboop!
vanessa@vanessa-ThinkPad-T460s:~/Documents/Dropbox/Code/scif/examples/snake&gt;make.scif/results/cloud$ sregistry build --config config-n1-highcpu-4.json https://www.github.com/sci-f/snake&gt;make.scif
[client|google-compute] [database|sqlite:///home/vanessa/.singularity/sregistry.db]
[bucket][sregistry-vanessa]
INSTANCE placid-signal-8619
sci-f-snake&gt;make.scif-builder https://vsocz.github.io/builders/cloud/google/compute/ubuntu/metrics/securebuild-2.4.3
Robot Logger: http://104.198.10.164
Allow a few minutes for web server install, beepboop!
vanessa@vanessa-ThinkPad-T460s:~/Documents/Dropbox/Code/scif/examples/snake&gt;make.scif/results/cloud$ sregistry build --config config-n1-highcpu-8.json https://www.github.com/sci-f/snake&gt;make.scif
[client|google-compute] [database|sqlite:///home/vanessa/.singularity/sregistry.db]
[bucket][sregistry-vanessa]
INSTANCE gloopy-nalgas-6369
sci-f-snake&gt;make.scif-builder https://vsocz.github.io/builders/cloud/google/compute/ubuntu/metrics/securebuild-2.4.3
Robot Logger: http://35.230.72.13
Allow a few minutes for web server install, beepboop!
vanessa@vanessa-ThinkPad-T460s:~/Documents/Dropbox/Code/scif/examples/snake&gt;make.scif/results/cloud$ sregistry build --config config-n1-highcpu-16.json https://www.github.com/sci-f/snake&gt;make.scif
[client|google-compute] [database|sqlite:///home/vanessa/.singularity/sregistry.db]
[bucket][sregistry-vanessa]
INSTANCE bricky-despacito-9882
sci-f-snake&gt;make.scif-builder https://vsocz.github.io/builders/cloud/google/compute/ubuntu/metrics/securebuild-2.4.3
Robot Logger: http://35.230.90.118
Allow a few minutes for web server install, beepboop!
vanessa@vanessa-ThinkPad-T460s:~/Documents/Dropbox/Code/scif/examples/snake&gt;make.scif/results/cloud$ sregistry build --config config-n1-highcpu-32.json https://www.github.com/sci-f/snake&gt;make.scif
[client|google-compute] [database|sqlite:///home/vanessa/.singularity/sregistry.db]
[bucket][sregistry-vanessa]
INSTANCE expressive-pedo-7623
sci-f-snake&gt;make.scif-builder https://vsocz.github.io/builders/cloud/google/compute/ubuntu/metrics/securebuild-2.4.3
Robot Logger: http://35.230.21.109
Allow a few minutes for web server install, beepboop!
vanessa@vanessa-ThinkPad-T460s:~/Documents/Dropbox/Code/scif/examples/snake&gt;make.scif/results/cloud$ sregistry build --config config-n1-highcpu-64.json https://www.github.com/sci-f/snake&gt;make.scif
[client|google-compute] [database|sqlite:///home/vanessa/.singularity/sregistry.db]
[bucket][sregistry-vanessa]
INSTANCE butterscotch-house-3317
sci-f-snake&gt;make.scif-builder https://vsocz.github.io/builders/cloud/google/compute/ubuntu/metrics/securebuild-2.4.3
Robot Logger: http://35.197.117.86
Allow a few minutes for web server install, beepboop!
vanessa@vanessa-ThinkPad-T460s:~/Documents/Dropbox/Code/scif/examples/snake&gt;make.scif/results/cloud$ □
```



The Builders

Welcome to the Builders documentation! [builder-manifest][static-demo].

What are the Builders?

Hi friends! Welcome to the builders documentation! The Builder robots are served faithfully by the [Singularity Global Client](#) so you can run your own builds on different clouds of your choosing. This documentation is intended for **developers** of builders, and not using the builders, see the [Singularity Global Client](#). The builder that you use (e.g., Google Compute) will logically map to the client there!

```
> SINGULARITY BUILDER ROBOT LOGGER, REPORTING FOR DUTY!
> INSTALLING SINGULARITY DEPENDENCIES
PREPARING LOGGING...
LOGS AVAILABLE AT HTTP://10.138.0.6/
# SINGULARITY
SINGULARITY_REPO: HTTPS://GITHUB.COM/CCLERGET/SINGULARITY.GIT
THE SINGULARITY REPOSITORY BEING CLONED BY THE BUILDER.
SINGULARITY_BRANCH: FEATURE-SQUASHBUILD-SECBUILD-2.4.3
THE BRANCH OF THE REPOSITORY BEING USED.
SINGULARITY_COMMIT:
IF DEFINED, A PARTICULAR COMMIT TO CHECKOUT.
# SETTINGS
SREGISTRY_USER_REPO: HTTPS://GITHUB.COM/EILON-S/SHERLOCK_VEP
YOUR REPOSITORY WE ARE BUILDING FROM!
```

Local machine

sregistry build <https://www.github.com/vsoch/hello-world>

build template

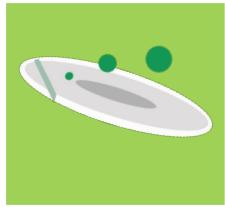


cloud builder

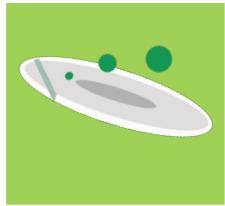


cloud storage

sregistry pull sregistry pull vsoch/hello-world:latest...



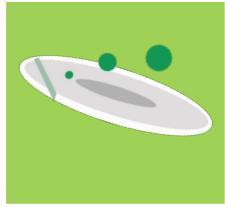
n1-standard-1



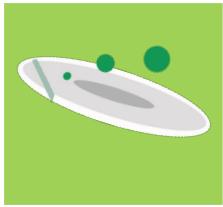
n1-standard-2



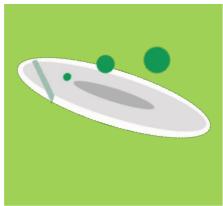
n1-standard-4



n1-standard-8



n1-standard-1



n1-standard-2



n1-standard-4



n1-standard-8

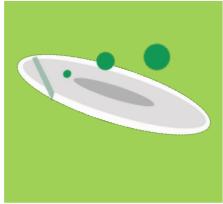




n1-standard-1



n1-standard-2



n1-standard-4

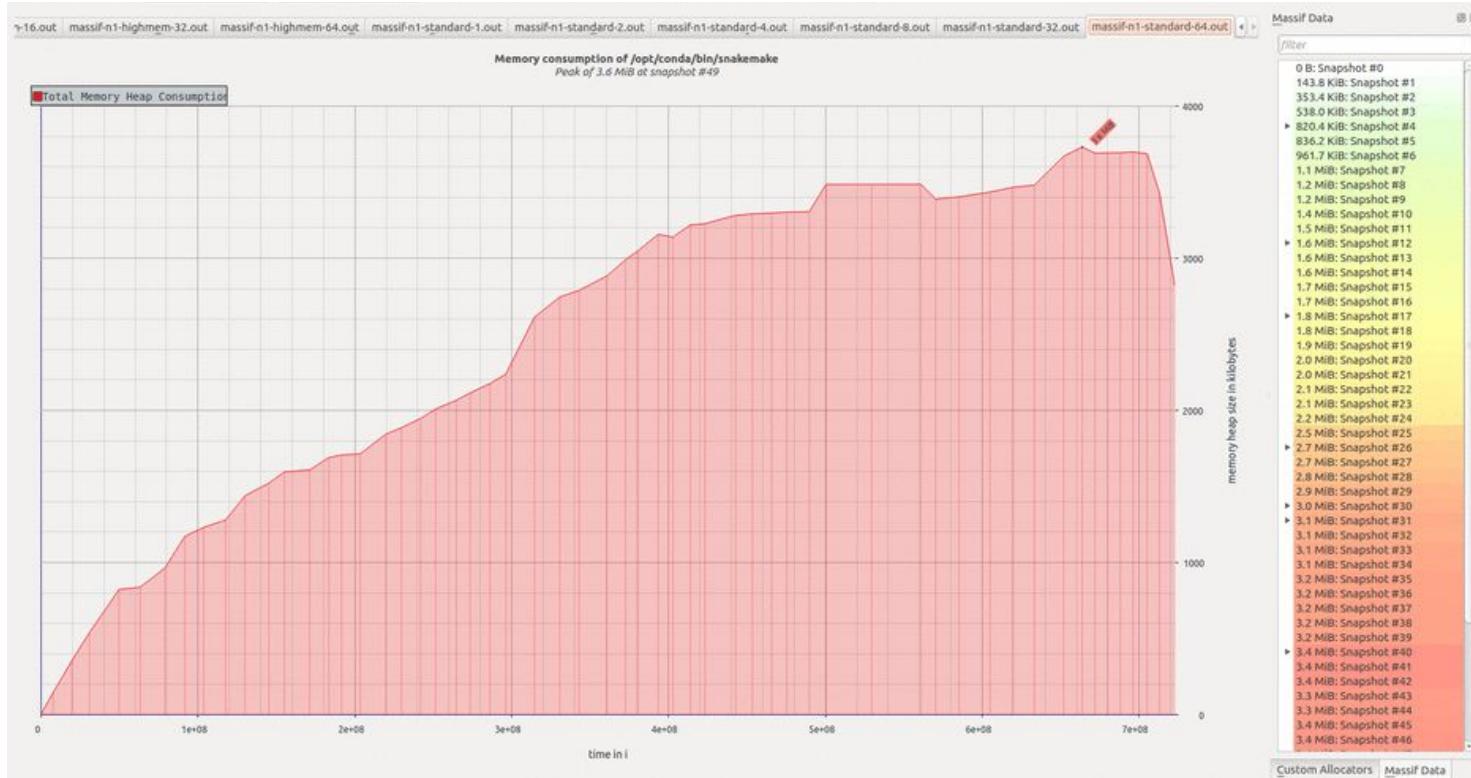


n1-standard-8



Does Heap Consumption Vary Across Machine Type?

Variant calling memory usage is comparable from 1 to 64 cores



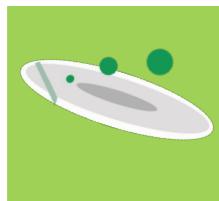


What is the optimal design?

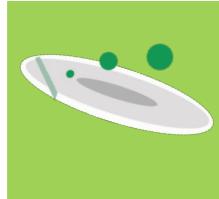
tensorflow 17.0



tensorflow 17.1



theano 1.0.1



theano 1.0.0



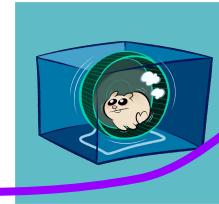
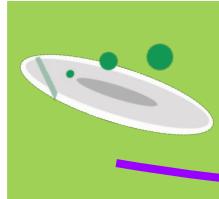
tensorflow 17.0



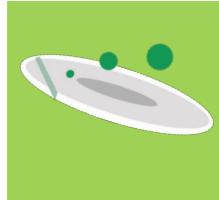
tensorflow 17.1



theano 1.0.1



theano 1.0.0

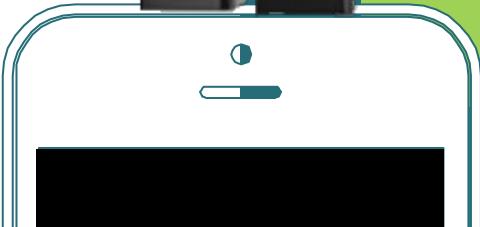


“

TLDR



It does the same thing!



Containerization seems to produce a reproducible workflow.

We need to make tools that are usable across technologies.

Transparency, ease of use, and ability to contribute are important.

It's our job to encourage introspection and evaluation of variation.



Once upon a time...
There was a scientist

There was a
reproducibility crisis





...but friends were there to help :)



WHY CAN'T WE BE

FRIENDS?

<https://sci-f.github.io>

PRESENTER: VANESSASAURUS

THIS IS THE END

vsochat@stanford.edu