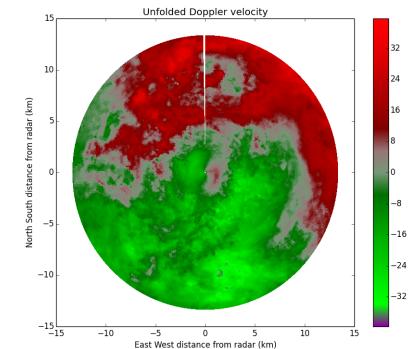
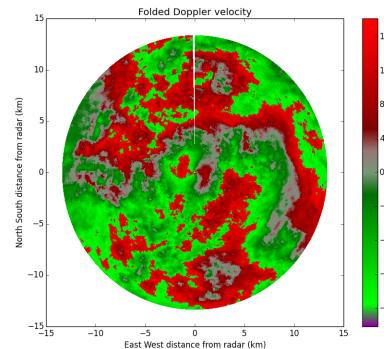
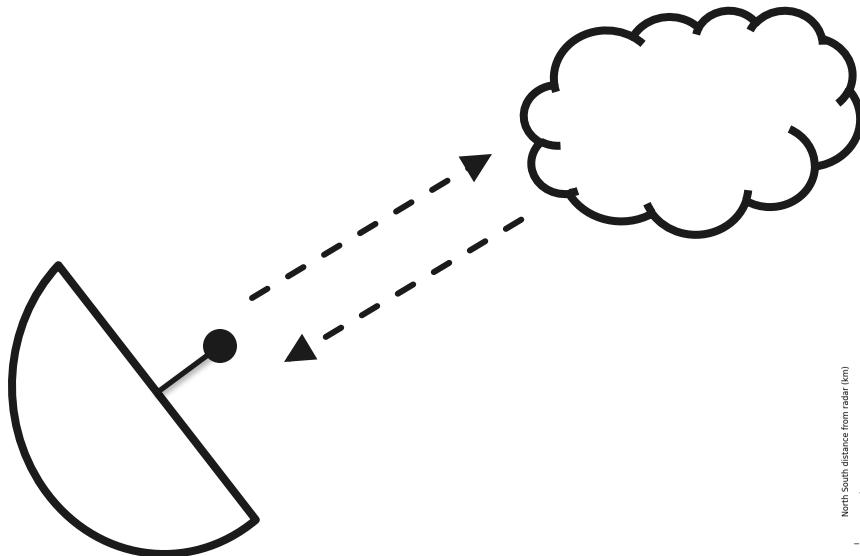


Profiling Python code to improve memory usage and execution time

Jonathan Helmus, Argonne National Laboratory



The ARM Climate Research Facility

- The U.S. Department of Energy's Atmospheric Radiation Measurement (ARM) Climate Research facility provides *in situ* and remote sensing data with a mission to **improve climate and earth systems models**.
- The program operates a **large number of instruments** at three **fixed sites** as well as two **mobile facilities** that can be deployed in support of field campaigns around the globe.
- This instrumentation includes a number of scanning **cloud and precipitation radars** which were acquired with funding from the American Recovery Act.
- The program is the process of preparing **two new scanning radars** for deployment in 2015.

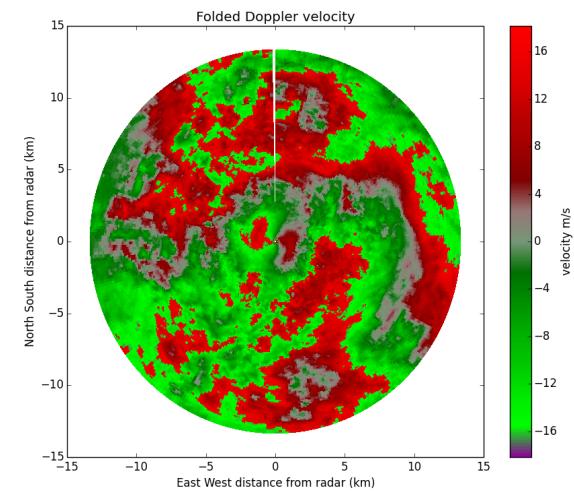
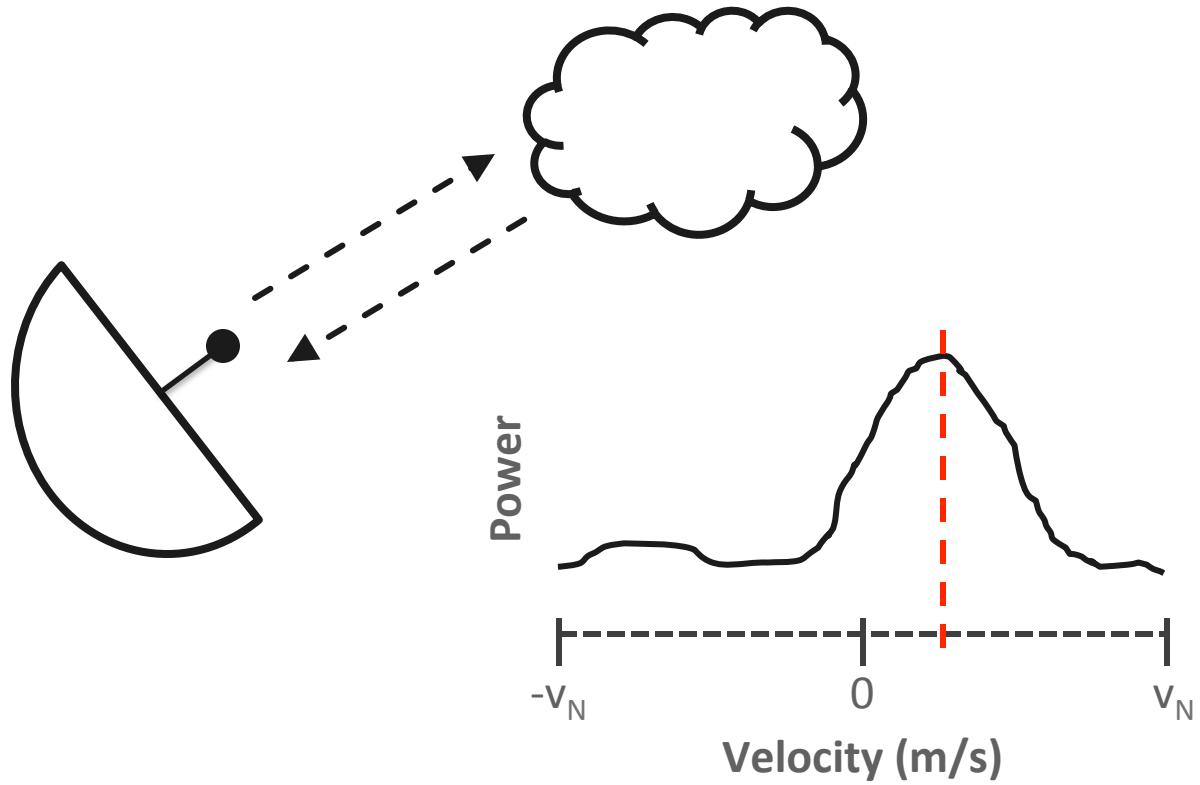


The Python ARM Radar Toolkit: Py-ART

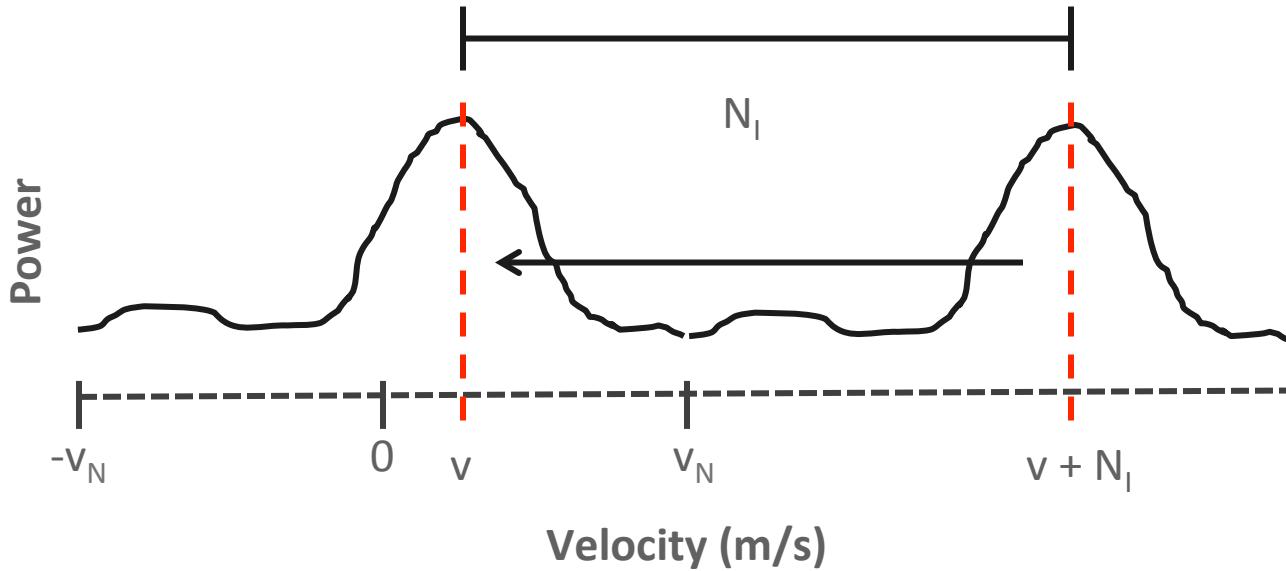
- Py-ART is a module for visualizing, correcting and analyzing **weather radar data** using packages from the scientific Python stack.
- Development began to address the needs of the **ARM** program with the acquisition of **multiple scanning cloud and precipitation radars**.
- The project has since been expanded to work with a **variety of weather radars**, including NEXRAD and TDWR radars, and a wide user base including radar researchers, weather enthusiast and climate modelers.
- Available on GitHub as **open source software** under a BSD license, arm-doe.github.io/pyart/.
- Conda packages are available at binstar.org/jjhelmus for Windows and OS X.



Doppler Velocity Measurements by Radar

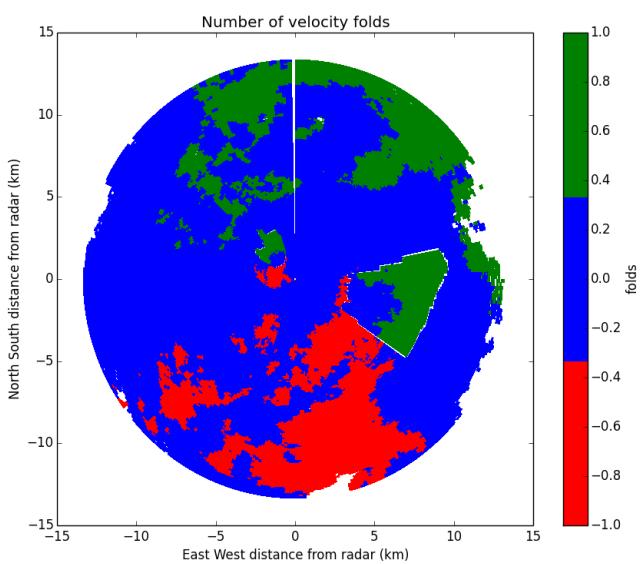
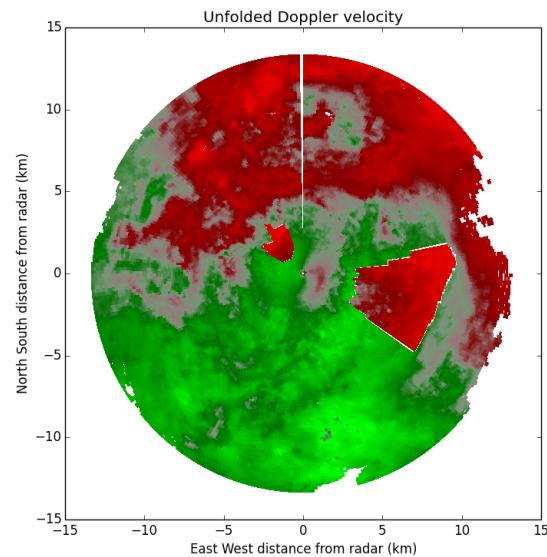
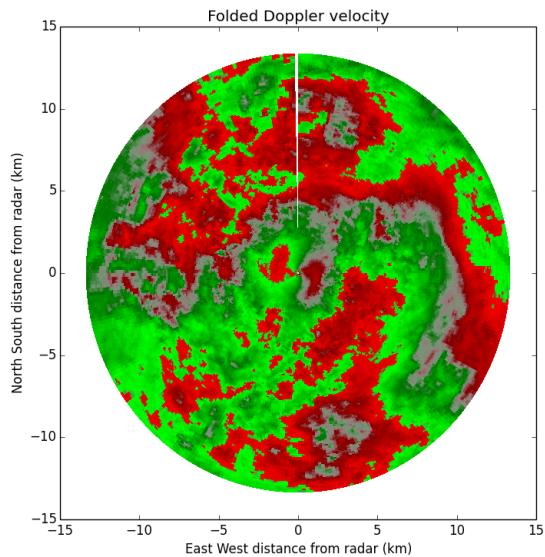


Doppler Velocity Aliasing



$$v' = v + n \times N_I \quad \text{where } n = 0, \pm 1, \dots$$

Four-Dimensional Doppler Dealiasing



A Real-Time Four-Dimensional Doppler Dealiasing Scheme

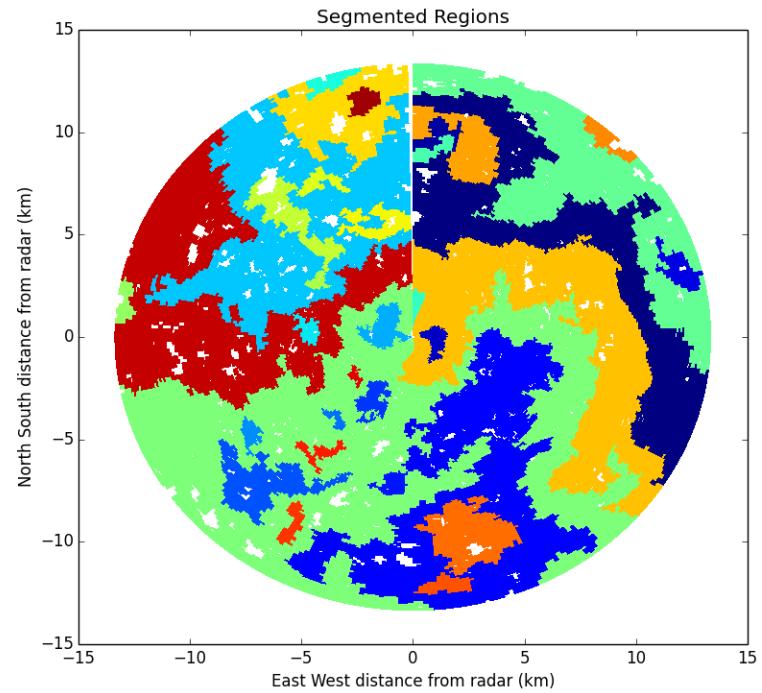
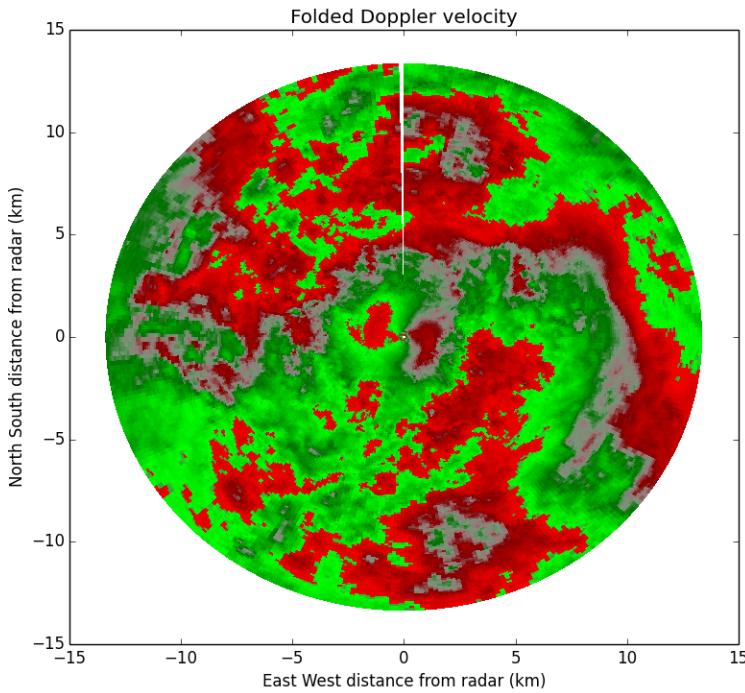
CURTIS N. JAMES* AND ROBERT A. HOUZE JR.

Department of Atmospheric Sciences, University of Washington, Seattle, Washington

J Tech, Vol 18, 2001, pg. 1674.

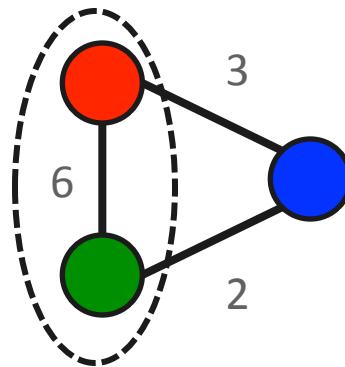
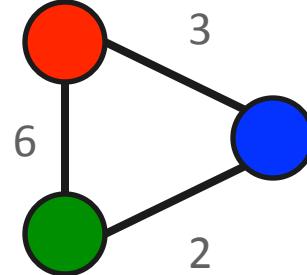
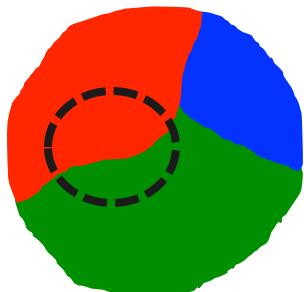
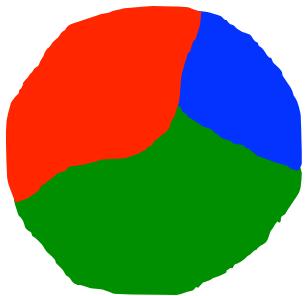
Profiling Python code, Jonathan Helmus, SEA Conference 2015, Boulder, CO

Region Based Dealiasing : Algorithm Design I

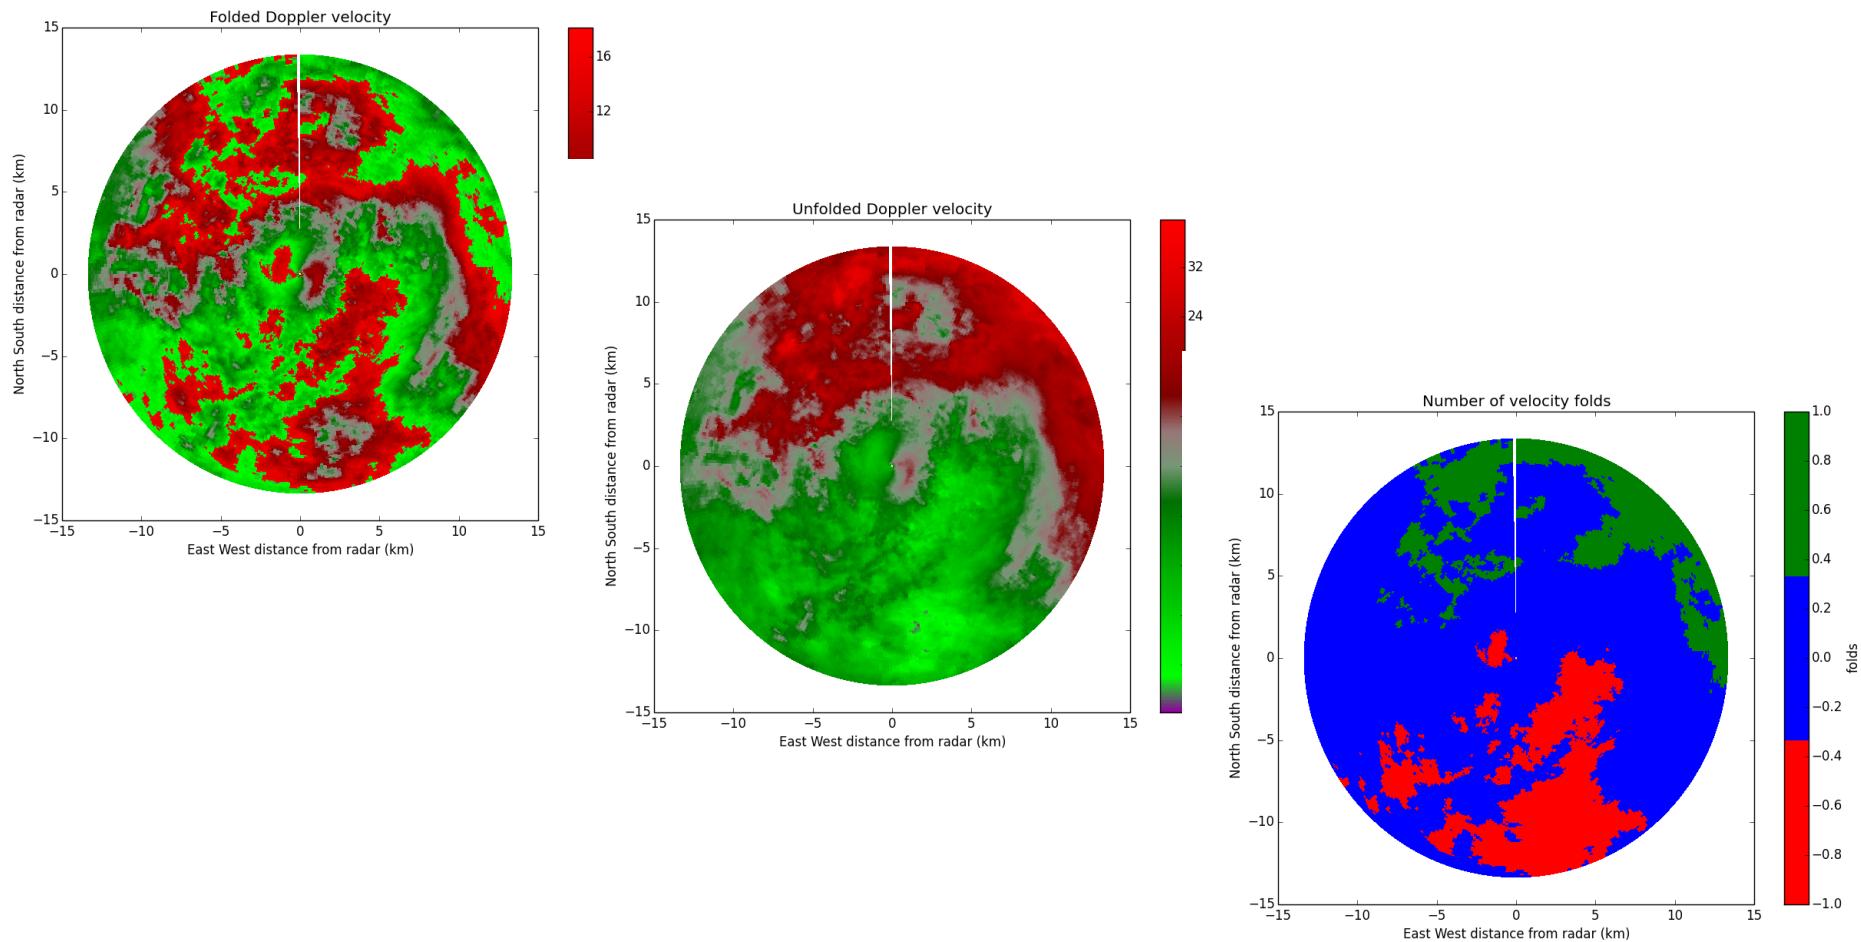


- Segmentation into regions based upon folded velocities.
- Uses `scipy.ndimage` module.

Region Based Dealiasing : Algorithm Design II



Region Based Dealiasing : Initial results



Memory usage reduction with *memory_profiler*

https://pypi.python.org/pypi/memory_profiler

```
import dealias_lib
import numpy as np
import pyart

@profile
def dealias(original_data, nyquist):
    dealias_data = original_data.copy()
    labels, nfeatures = dealias_lib.find_segments(original_data)
    segment_sizes = dealias_lib.segment_sizes(labels, nfeatures)
    edge_sum, edge_count = dealias_lib.edge_sum_and_count(
        labels, nfeatures, original_data)

    upper_indices = np.triu_indices_from(edge_count, 1)
    for i in xrange(4):
        n = np.argmax(edge_count[upper_indices])
        x, y = upper_indices[0][n], upper_indices[1][n]
        if edge_count[x,y] == 0:
            break
        dealias_lib.combine_segments(
            x, y, labels, segment_sizes, edge_sum,
            edge_count, nyquist, dealias_data, step=None)

    return

if __name__ == "__main__":
    radar = pyart.io.read('105235.mdv')
    original_data = radar.fields['velocity']['data'].copy()
    original_data = original_data[:100]
    nyquist = radar.instrument_parameters['nyquist_velocity']['data'][0]
    dealias(original_data, nyquist)
```



Memory usage reduction with *memory_profiler*

```
$ python -m memory_profiler script.py
Filename: script.py

Line #      Mem usage      Increment   Line Contents
=====
6    627.469 MiB      0.000 MiB   @profile
7                                def dealias(original_data, nyquist):
8    627.844 MiB      0.375 MiB   dealias_data = original_data.copy()
9    516.098 MiB     -111.746 MiB  labels, nfeatures = dealias_lib.find_segments(original_data)
10   493.902 MiB     -22.195 MiB  segment_sizes = dealias_lib.segment_sizes(labels, nfeatures)
11   493.902 MiB      0.000 MiB   edge_sum, edge_count = dealias_lib.edge_sum_and_count(
12   776.426 MiB     282.523 MiB   labels, nfeatures, original_data)
13
14  2264.855 MiB  1488.430 MiB   upper_indices = np.triu_indices_from(edge_count, 1)
15 2916.219 MiB   651.363 MiB   for i in xrange(4):
16 2854.086 MiB   -62.133 MiB   n = np.argmax(edge_count[upper_indices])
17 2854.086 MiB      0.000 MiB   x, y = upper_indices[0][n], upper_indices[1][n]
18 2854.086 MiB      0.000 MiB   if edge_count[x,y] == 0:
19                                break
20 2854.086 MiB      0.000 MiB   dealias_lib.combine_segments(
21 2854.086 MiB      0.000 MiB   x, y, labels, segment_sizes, edge_sum,
22 2916.219 MiB     62.133 MiB   edge_count, nyquist, dealias_data, step=None)
23
24 2916.219 MiB      0.000 MiB   return
```

Memory usage reduction with *memory_profiler*

Line #	Mem usage	Increment	Line Contents
6	627.492 MiB	0.000 MiB	@profile
7			def dealias(original_data, nyquist):
8	627.867 MiB	0.375 MiB	dealias_data = original_data.copy()
9	516.406 MiB	-111.461 MiB	labels, nfeatures = dealias_lib.find_segments(original_data)
10	494.211 MiB	-22.195 MiB	segment_sizes = dealias_lib.segment_sizes(labels, nfeatures)
11	494.211 MiB	0.000 MiB	edge_sum, edge_count = dealias_lib.edge_sum_and_count(
12	311.348 MiB	-182.863 MiB	labels, nfeatures, original_data)
13	302.684 MiB	-8.664 MiB	edge_count.setdiag(0)
14			
15	313.363 MiB	10.680 MiB	for i in xrange(4):
16	310.547 MiB	-2.816 MiB	argmax_idx = edge_count.data.argmax()
17	310.547 MiB	0.000 MiB	x = np.nonzero(
18	310.547 MiB	0.000 MiB	edge_count.indptr <= argmax_idx)[0][-1]
19	310.547 MiB	0.000 MiB	y = edge_count.indices[argmax_idx]
20	310.547 MiB	0.000 MiB	if x > y:
21			x,y = y,x
22	310.547 MiB	0.000 MiB	if edge_count[x,y] == 0:
23			break
24	310.547 MiB	0.000 MiB	dealias_lib.combine_segments(
25	310.547 MiB	0.000 MiB	x, y, labels, segment_sizes, edge_sum,
26	313.363 MiB	2.816 MiB	edge_count, nyquist, dealias_data, step=None)
27			
28	313.363 MiB	0.000 MiB	return

Run time optimization with *cProfile*

```
import dealias_lib
import numpy as np
import pyart

def dealias(original_data, nyquist):
    ...
    ...
    ...
    return

if __name__ == "__main__":
    radar = pyart.io.read('single_folded_sweep.nc')
    original_data = radar.fields['velocity']['data'].copy()
    nyquist = radar.instrument_parameters['nyquist_velocity']['data'][0]

    # profile
    import cProfile
    cProfile.run("dealias(original_data, nyquist)",
                "Profile.prof")

    # print out stats
    import pstats
    s = pstats.Stats("Profile.prof")
    s.strip_dirs().sort_stats("time").print_stats(20)
```

Run time optimization with *cProfile*

```
$ python profile_script.py
Thu Jan  1 21:59:03 2015      Profile.prof

    7344341 function calls (7342262 primitive calls) in 75.729 seconds

Ordered by: internal time
List reduced from 192 to 20 due to restriction <20>

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
   693    48.491    0.070    48.491    0.070 {scipy.sparse._sparsetools.csr_sample_values}
  2772     5.209    0.002    15.494    0.000 compressed.py:702(_insert_many)
 494518     4.260    0.000    4.260    0.000 {numpy.core.multiarray.concatenate}
 483430     4.173    0.000    6.971    0.000 arraysetops.py:93(unique)
   6920     2.386    0.000    2.386    0.000 {scipy.sparse._sparsetools.get_csr_submatrix}
  55388     1.642    0.000    1.642    0.000 {method 'reduce' of 'numpy.ufunc' objects}
   8304     1.425    0.000    1.425    0.000 {scipy.sparse._sparsetools.csr_sample_offsets}
   8304     1.301    0.000   18.600    0.002 compressed.py:649(_set_many)
  483430     0.564    0.000    0.564    0.000 {method 'flatten' of 'numpy.ndarray' objects}
  486202     0.540    0.000    0.540    0.000 {method 'argsort' of 'numpy.ndarray' objects}
  17998     0.437    0.000    0.938    0.000 compressed.py:126(check_format)
   42919     0.362    0.000    0.362    0.000 {method 'astype' of 'numpy.ndarray' objects}
    692     0.291    0.000   24.949    0.036 dealias_lib.py:156(combine_segments)
  282452     0.279    0.000    0.279    0.000 {numpy.core.multiarray.array}
   44308     0.225    0.000    1.317    0.000 sputils.py:132(get_index_dtype)
  16608     0.224    0.000    0.590    0.000 stride_tricks.py:36(broadcast_arrays)
  47080     0.206    0.000    0.206    0.000 getlimits.py:244(__init__)
   22144     0.181    0.000    0.283    0.000 stride_tricks.py:22(as_strided)
      1     0.155    0.155    0.214    0.214 dealias_lib.py:97(edge_sum_and_count)
 1966244     0.153    0.000    0.153    0.000 {method 'append' of 'list' objects}
```

Run time optimization with *cProfile*

```
import dealias_lib
from trackers import RegionTracker, EdgeTracker
import pyart

def dealias(original_data, nyquist):

    labels, nfeatures = dealias_lib.find_segments(original_data)
    segment_sizes = dealias_lib.segment_sizes(labels, nfeatures)
    edge_sum, edge_count = dealias_lib.edge_sum_and_count(
        labels, nfeatures, original_data)

    region_tracker = RegionTracker(segment_sizes)
    edge_tracker = EdgeTracker(edge_sum, edge_count, nyquist)
    for i in xrange(80000):
        if dealias_lib.combine_segments(region_tracker, edge_tracker):
            break

    dealias_data = original_data.copy()
    for i in range(nfeatures+1):
        nowrap = region_tracker.unwrap_number[i]
        dealias_data[labels == i] += nowrap * nyquist
    return dealias_data

if __name__ == "__main__":
    radar = pyart.io.read('single_folded_sweep.nc')
    ...
```

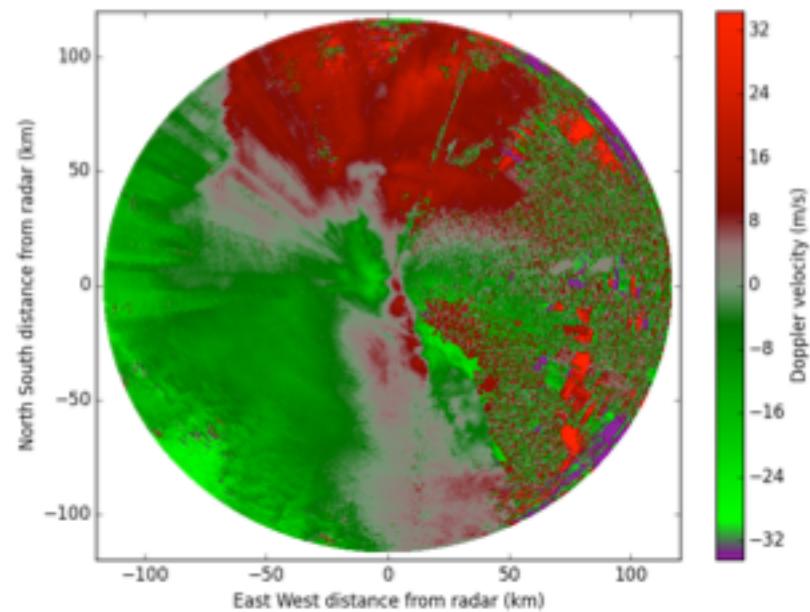
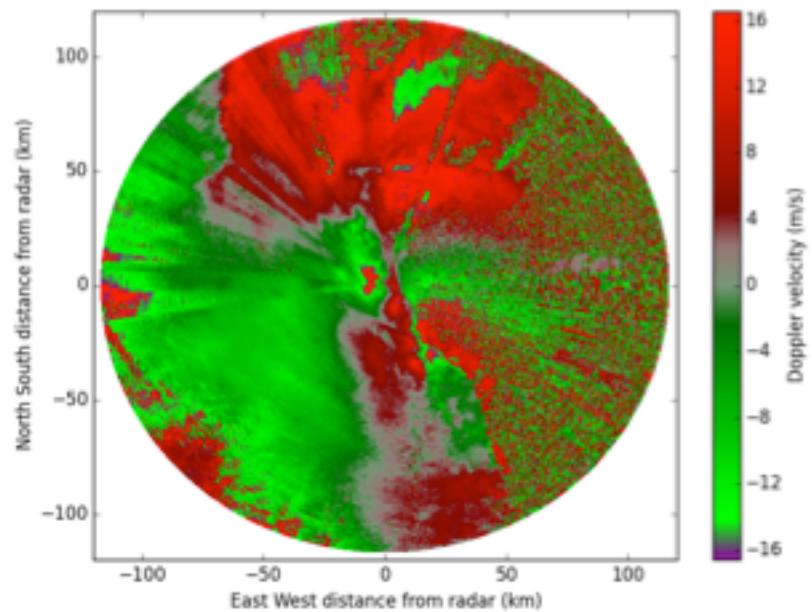
Run time optimization with *cProfile*

```
$ python profile_script.py
Thu Jan  1 22:04:25 2015      Profile.prof
    199063 function calls in 0.626 seconds

Ordered by: internal time
List reduced from 124 to 20 due to restriction <20>

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
      1    0.129    0.129    0.626    0.626 profile_script_2.py:6(dealias)
  2505    0.117    0.000    0.144    0.000 compressed.py:772(_get_single_element)
      1    0.112    0.112    0.147    0.147 dealias_lib.py:97(edge_sum_and_count)
  1670    0.060    0.000    0.060    0.000 {method 'remove' of 'list' objects}
  5013    0.038    0.000    0.045    0.000 sputils.py:188(isintlike)
 54001    0.025    0.000    0.033    0.000 index_tricks.py:490(__next__)
  2526    0.014    0.000    0.014    0.000 {method 'reduce' of 'numpy.ufunc' objects}
  2505    0.012    0.000    0.031    0.000 sputils.py:244(_unpack_index)
   693    0.010    0.000    0.012    0.000 trackers.py:206(pop_edge)
 27583    0.009    0.000    0.009    0.000 {isinstance}
   692    0.009    0.000    0.073    0.000 trackers.py:102(merge_nodes)
 54001    0.008    0.000    0.008    0.000 {next}
  2505    0.008    0.000    0.012    0.000 sputils.py:310(_check_boolean)
   304    0.008    0.000    0.008    0.000 trackers.py:193(unwrap_node)
      1    0.007    0.007    0.234    0.234 trackers.py:58(__init__)
  2505    0.007    0.000    0.226    0.000 csr.py:197(__getitem__)
  2505    0.005    0.000    0.005    0.000 {method 'compress' of 'numpy.ndarray' objects}
   304    0.005    0.000    0.005    0.000 trackers.py:42(unwrap_node)
   693    0.004    0.000    0.107    0.000 dealias_lib.py:159(combine_segments)
  2505    0.004    0.000    0.004    0.000 sputils.py:272(_check_ellipsis)
```

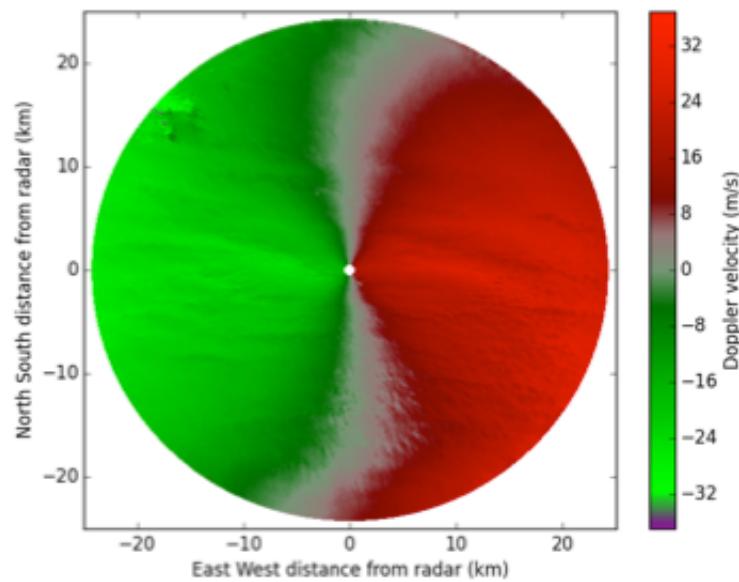
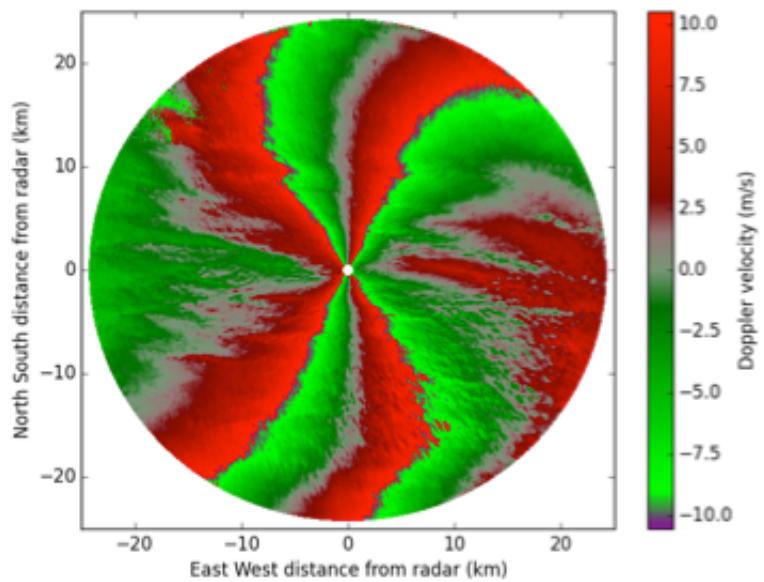
Sample results: ARM CSAPR



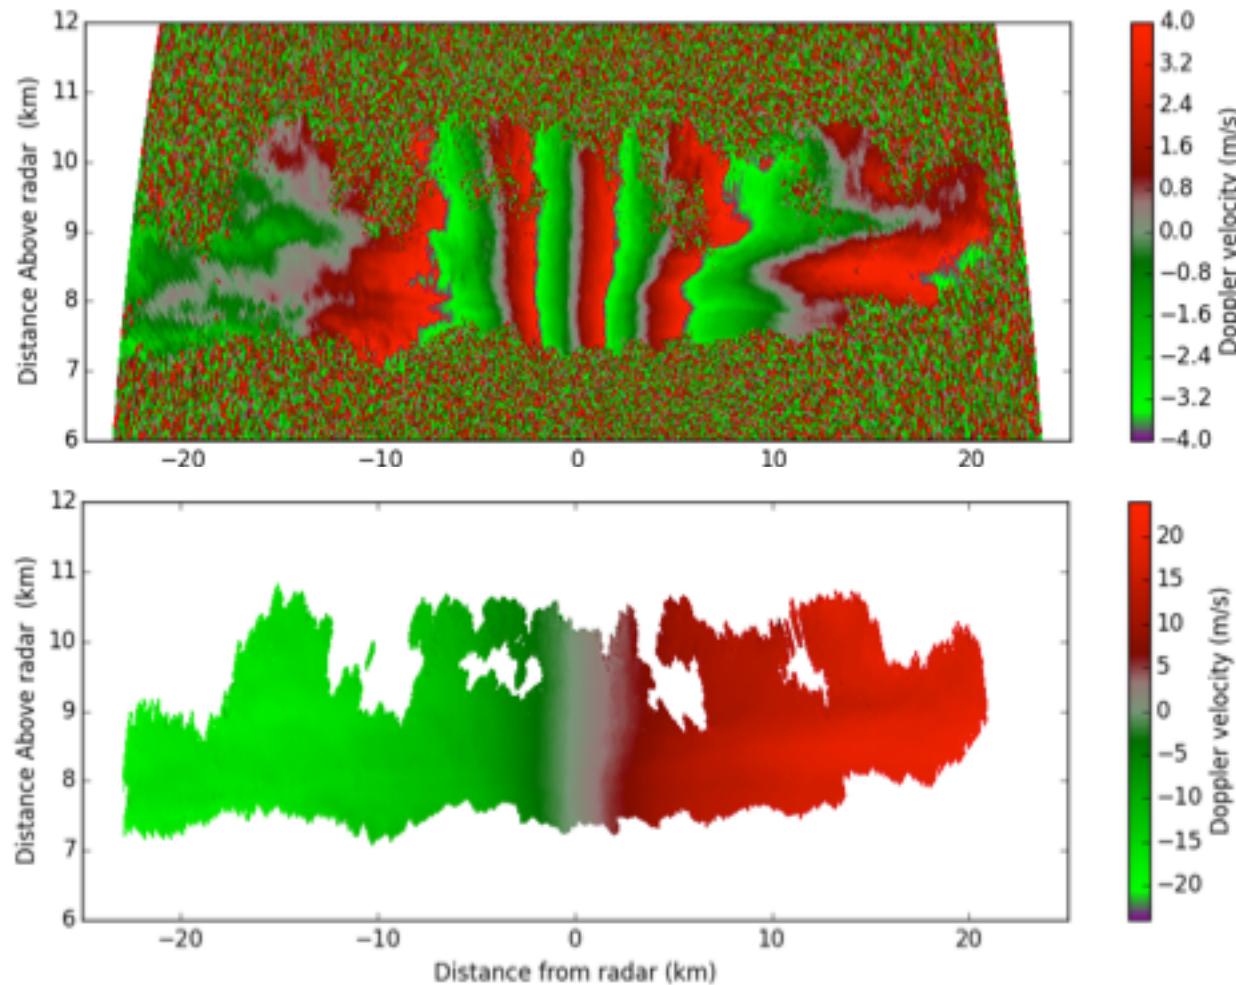
Time required to dealias a full radar volume: ~1.5 minutes



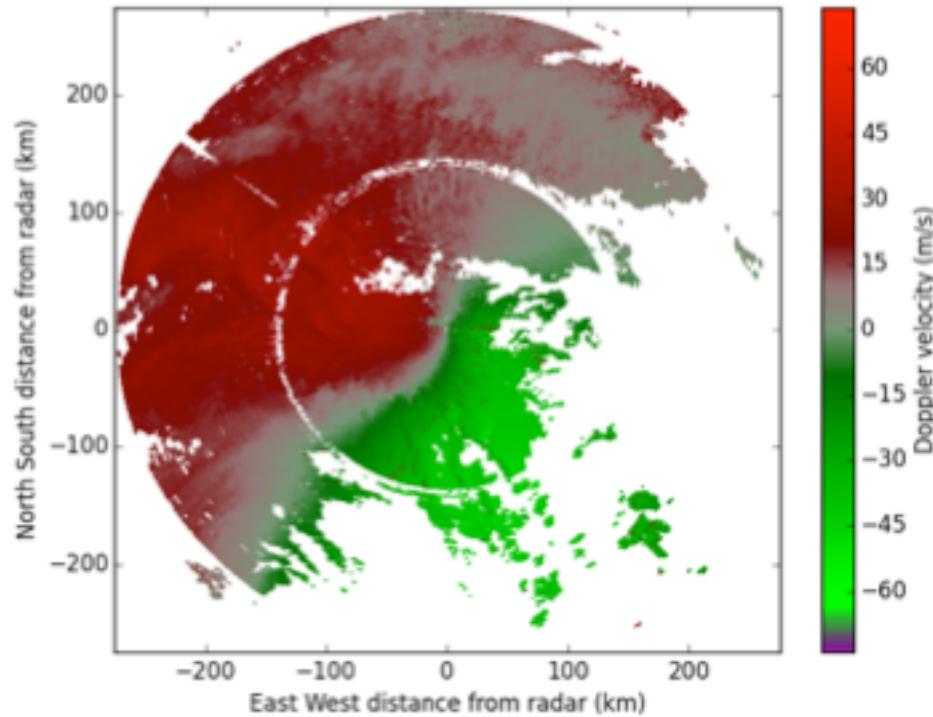
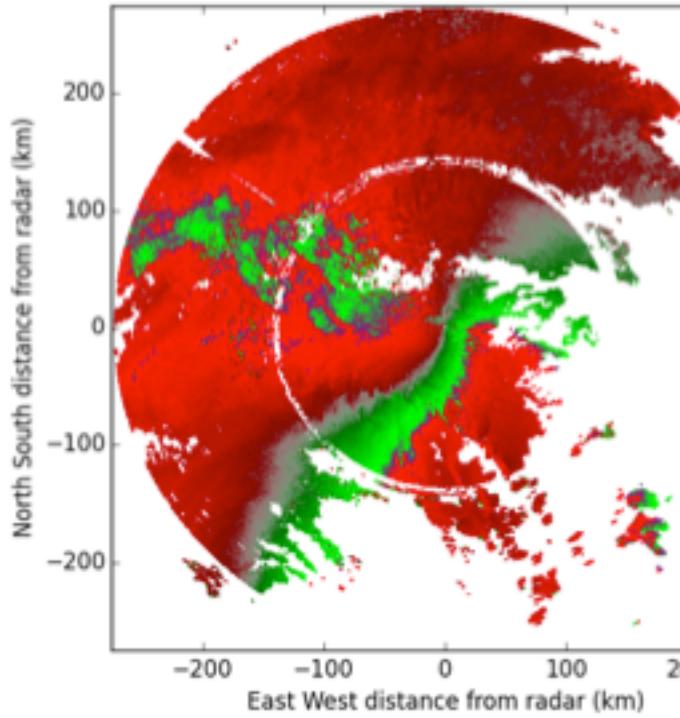
Sample results: ARM KaSACR



Sample results: ARM WSACR



Sample results: NEXRAD



Run time optimization with *line_profiler*

https://pypi.python.org/pypi/line_profiler/

Source code

```
@profile
def dealias_region_based(
    radar, interval_splits=3, interval_limits=None,
    skip_between_rays=2, skip_along_ray=2, centered=True,
    nyquist_vel=None, gatefilter=None, rays_wrap_around=None,
    keep_original=True, vel_field=None, corr_vel_field=None, **kwargs):
    """
        Dealias Doppler velocities using a region based algorithm.
    ...

```

Script

```
import pyart
radar = pyart.io.read('105235.mdv')
pyart.correct.dealias_region_based(radar, keep_original=False,
    skip_along_ray=100, skip_between_rays=100)
```

Run time optimization with *line_profiler*

```
$ kernprof.py -l -v script.py
```

```
Timer unit: 1e-06 s
```

```
File: /home/jhelmus/dev/pyart/pyart/correct/region_dealias.py
```

```
Function: dealias_region_based at line 57
```

```
Total time: 46.1975 s
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
57					@profile
58					def dealias_region_based(
...					
156					# extract sweep data
157	17	3116	183.3	0.0	sdata = vdata[sweep_slice].copy() # is a copy needed here?
158	17	38	2.2	0.0	scorr = data[sweep_slice]
159	17	28	1.6	0.0	sfilter = gfilter[sweep_slice]
160					
161					# find regions in original data
162	17	278528	16384.0	0.6	labels, nfeatures = _find_regions(sdata, sfilter, interval_limits)
163	17	40	2.4	0.0	edge_sum, edge_count, region_sizes = _edge_sum_and_count(
164	17	21	1.2	0.0	labels, nfeatures, sdata, rays_wrap_around, skip_between_rays,
165	17	163290	9605.3	0.4	skip_along_ray)
166					
167					# find the number of folds in the regions
168	17	57276	3369.2	0.1	region_tracker = RegionTracker(region_sizes)
169	17	33134917	1949112.8	71.7	edge_tracker = _EdgeTracker(edge_sum, edge_count, nyquist_interval)
170	48314	56051	1.2	0.1	while True:
171	48314	6810804	141.0	14.7	if _combine_regions(region_tracker, edge_tracker):
172	17	25	1.5	0.0	break

Run time optimization with *line_profiler*

```
$ kernprof.py -l -v script.py
```

Total time: 13.4539 s

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
...					
171					
172					# find the number of folds in the regions
173	17	57249	3367.6	0.4	region_tracker = _RegionTracker(region_sizes)
174	17	30	1.8	0.0	edge_tracker = _EdgeTracker(indices, edge_count, velos,
175	17	625889	36817.0	4.7	nyquist_interval, nfeatures+1)
176	48314	57803	1.2	0.4	while True:
177	48314	6622303	137.1	49.2	if _combine_regions(region_tracker, edge_tracker):
178	17	24	1.4	0.0	break
179					
180					# center sweep if requested, determine a global sweep unfold number
181					# so that the average number of gate folds is zero.
182	17	17	1.0	0.0	if centered:
183	17	329	19.4	0.0	gates_dealiased = region_sizes.sum()
184	17	26	1.5	0.0	total_folds = np.sum(
185	17	564	33.2	0.0	region_sizes * region_tracker.unwrap_number[1:])
186	17	71	4.2	0.0	sweep_offset = int(round(float(total_folds) / gates_dealiased))
187	17	22	1.3	0.0	if sweep_offset != 0:
188	3	22	7.3	0.0	region_tracker.unwrap_number -= sweep_offset
189					
190					# dealias the data using the fold numbers
191					# start from label 1 to skip masked region
192	48334	58536	1.2	0.4	for i in range(1, nfeatures+1):
193	48317	68051	1.4	0.5	nwrap = region_tracker.unwrap_number[i]
194	48317	182673	3.8	1.4	if nwrap != 0:
195	16432	5252213	319.6	39.0	scorr[labels == i] += nwrap * nyquist_interval

Run time optimization with *line_profiler*

```
$ kernprof.py -l -v script.py
```

Total time: 8.20599 s

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
...					
171					
172					# find the number of folds in the regions
173	17	60067	3533.4	0.7	region_tracker = _RegionTracker(region_sizes)
174	17	32	1.9	0.0	edge_tracker = _EdgeTracker(indices, edge_count, velos,
175	17	640319	37665.8	7.8	nyquist_interval, nfeatures+1)
176	48314	60980	1.3	0.7	while True:
177	48314	6872862	142.3	83.8	if _combine_regions(region_tracker, edge_tracker):
178	17	21	1.2	0.0	break
179					
180					# center sweep if requested, determine a global sweep unfold number
181					# so that the average number of gate folds is zero.
182	17	21	1.2	0.0	if centered:
183	17	376	22.1	0.0	gates_dealiased = region_sizes.sum()
184	17	29	1.7	0.0	total_folds = np.sum(
185	17	687	40.4	0.0	region_sizes * region_tracker.unwrap_number[1:])
186	17	76	4.5	0.0	sweep_offset = int(round(float(total_folds) / gates_dealiased))
187	17	25	1.5	0.0	if sweep_offset != 0:
188	3	23	7.7	0.0	region_tracker.unwrap_number -= sweep_offset
189					
190					# dealias the data using the fold numbers
191	17	14154	832.6	0.2	nwrap = np.take(region_tracker.unwrap_number, labels)
192	17	13471	792.4	0.2	scorr += nwrap * nyquist_interval

Run time optimization using Cython

```
$ kernprof.py -l -v script.py
```

```
Total time: 153.644 s
```

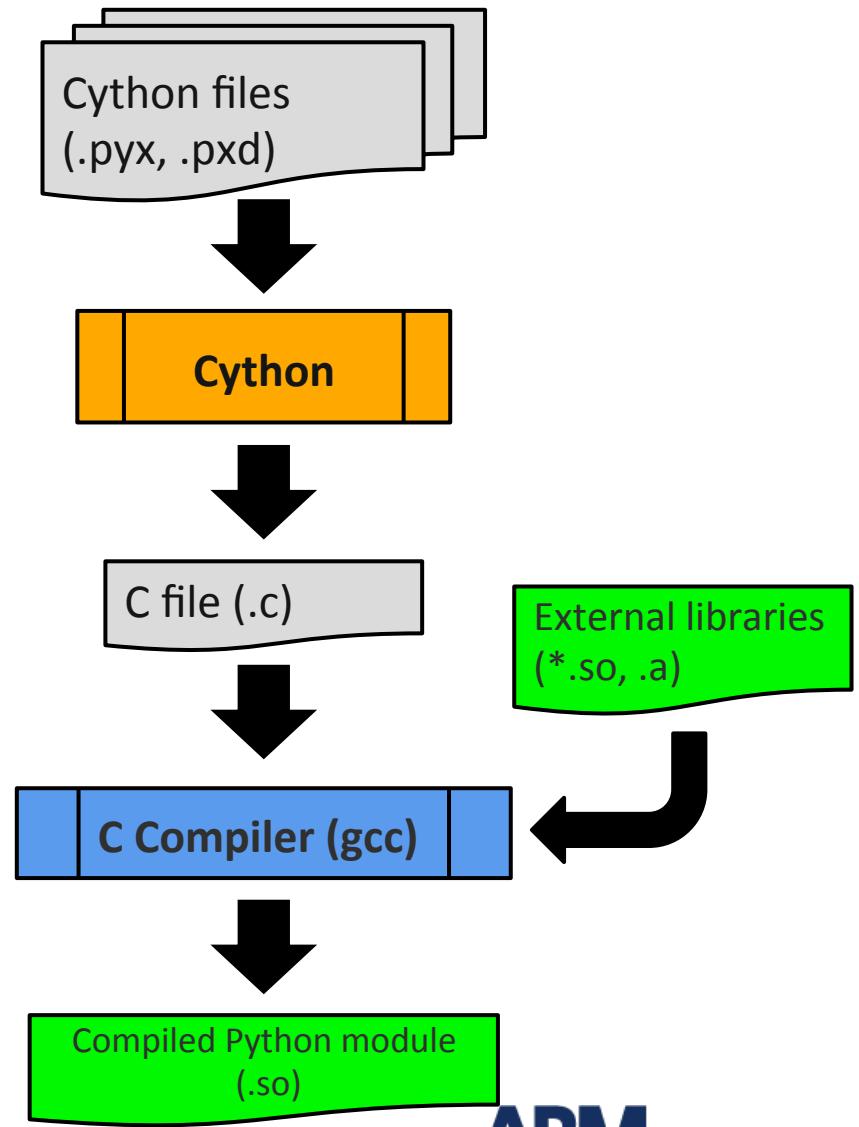
Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
159					# extract sweep data
160	17	3170	186.5	0.0	sdata = vdata[sweep_slice].copy() # is a copy needed here?
161	17	38	2.2	0.0	scorr = data[sweep_slice]
162	17	28	1.6	0.0	sfilter = gfilter[sweep_slice]
163					
164					# find regions in original data
165	17	284246	16720.4	0.2	labels, nfeatures = _find_regions(sdata, sfilter, interval_limits)
166	17	43	2.5	0.0	edge_sum, edge_count, region_sizes = _edge_sum_and_count(
167	17	19	1.1	0.0	labels, nfeatures, sdata, rays_wrap_around, skip_between_rays,
168	17	107524384	6324963.8	70.0	skip_along_ray)
169					
170					# find the number of folds in the regions
171	17	58564	3444.9	0.0	region_tracker = _RegionTracker(region_sizes)
172	17	33254618	1956154.0	21.6	edge_tracker = _EdgeTracker(edge_sum, edge_count, nyquist_interval)
173	48314	56613	1.2	0.0	while True:
174	48314	6727992	139.3	4.4	if _combine_regions(region_tracker, edge_tracker):
175	17	19	1.1	0.0	break

Run time optimization using Cython

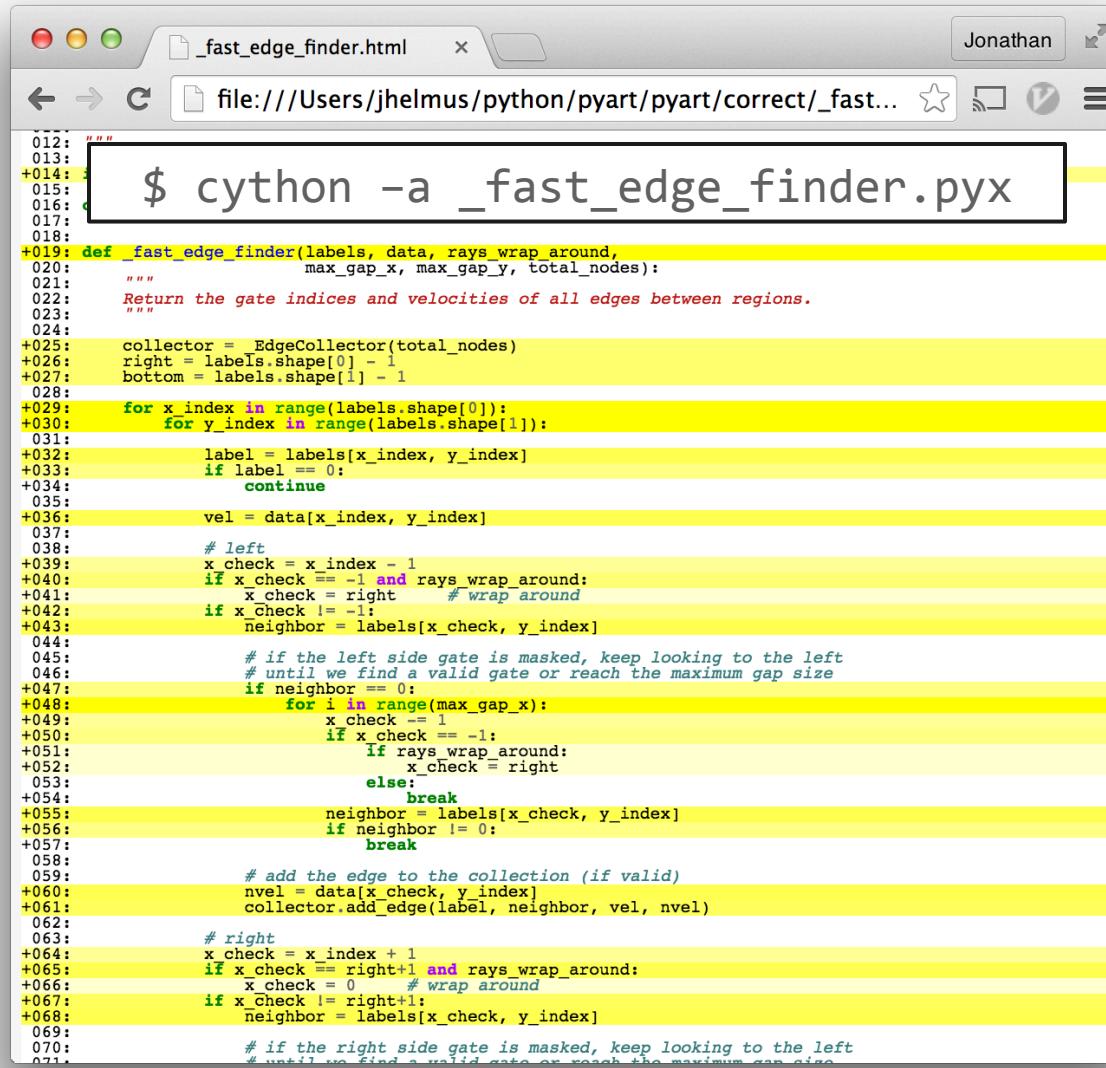
Python code, especially when it contains nested loops, can be slow. One method to speed up the run time is to use Cython.

Cython

- Python to C code translator.
- Generates a Python extension module.
- Can be used to speed up Python code by adding static type information.
- Also can be used to interact with C/C++ function and classes in external libraries.



Optimizing Cython code

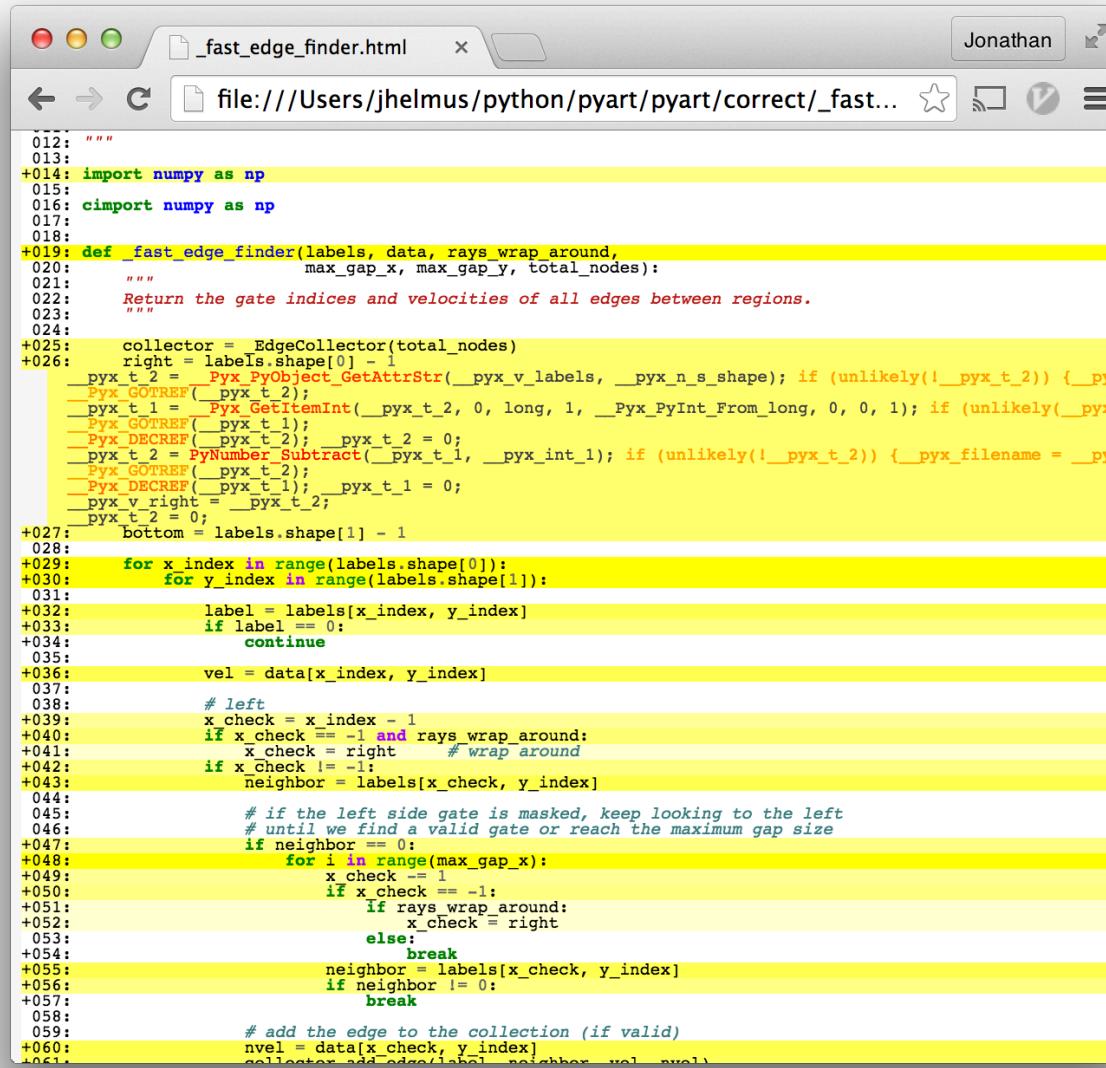


The screenshot shows a web browser window titled '_fast_edge_finder.html' with a URL starting 'file:///Users/jhelmus/python/pyart/pyart/correct/_fast...'. The browser interface includes standard controls like back, forward, and search. A tab labeled 'Jonathan' is visible in the top right. The main content area displays a terminal window with the following text:

```
$ cython -a _fast_edge_finder.pyx
```

```
012: """
013:
014: +
015: def _fast_edge_finder(labels, data, rays_wrap_around,
016:                       max_gap_x, max_gap_y, total_nodes):
017:     """
018:
019:     """Return the gate indices and velocities of all edges between regions.
020:
021:     collector = EdgeCollector(total_nodes)
022:     right = labels.shape[0] - 1
023:     bottom = labels.shape[1] - 1
024:
025:     for x_index in range(labels.shape[0]):
026:         for y_index in range(labels.shape[1]):
027:             label = labels[x_index, y_index]
028:             if label == 0:
029:                 continue
030:
031:             vel = data[x_index, y_index]
032:             # left
033:             x_check = x_index - 1
034:             if x_check == -1 and rays_wrap_around:
035:                 x_check = right # wrap around
036:             if x_check != -1:
037:                 neighbor = labels[x_check, y_index]
038:
039:                 # if the left side gate is masked, keep looking to the left
040:                 # until we find a valid gate or reach the maximum gap size
041:                 if neighbor == 0:
042:                     for i in range(max_gap_x):
043:                         x_check -= 1
044:                         if x_check == -1:
045:                             if rays_wrap_around:
046:                                 x_check = right
047:                             else:
048:                                 break
049:                         neighbor = labels[x_check, y_index]
050:                         if neighbor != 0:
051:                             break
052:
053:             # add the edge to the collection (if valid)
054:             nvel = data[x_check, y_index]
055:             collector.add_edge(label, neighbor, vel, nvel)
056:
057:
058:
059:             # right
060:             x_check = x_index + 1
061:             if x_check == right+1 and rays_wrap_around:
062:                 x_check = 0 # wrap around
063:             if x_check != right+1:
064:                 neighbor = labels[x_check, y_index]
065:
066:
067:
068:
069:
070:             # if the right side gate is masked, keep looking to the left
071:             # until we find a valid gate or reach the maximum gap size
```

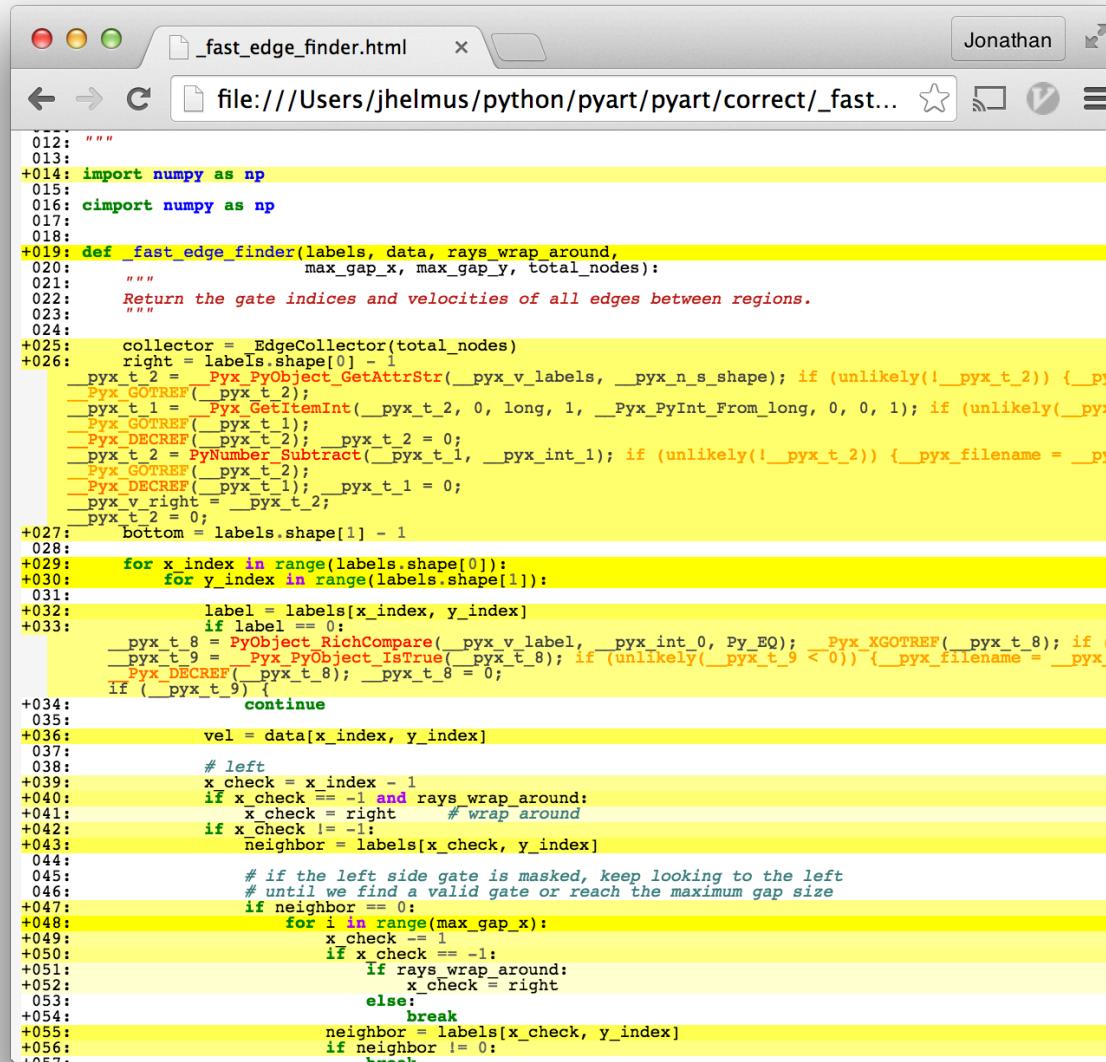
Optimizing Cython code



The screenshot shows a web browser window titled '_fast_edge_finder.html' with the URL 'file:///Users/jhelmus/python/pyart/pyart/correct/_fast...'. The page content displays a block of Cython code. The code is annotated with numerous orange and yellow highlights, primarily around memory management and performance-critical sections. The code itself is a function named '_fast_edge_finder' that takes 'labels', 'data', 'rays_wrap_around', 'max_gap_x', 'max_gap_y', and 'total_nodes' as parameters. It uses a C-style collector to find edges between regions and handles boundary conditions like wrap-around.

```
012: """
013:
014: import numpy as np
015:
016: cimport numpy as np
017:
018:
019: def _fast_edge_finder(labels, data, rays_wrap_around,
020:                         max_gap_x, max_gap_y, total_nodes):
021:     """
022:         Return the gate indices and velocities of all edges between regions.
023:
024:
025:         collector = EdgeCollector(total_nodes)
026:         right = labels.shape[0] - 1
027:         pyx_t_2 = __Pyx_PyObject_GetAttrStr(__pyx_v_labels, __pyx_n_s_shape); if (unlikely(!__pyx_t_2)) {__pyx_t_2 = Pyx_GOTREF(__pyx_t_2);
028:         pyx_t_1 = __Pyx_GetItemInt(__pyx_t_2, 0, long, 1, __Pyx_PyInt_From_long, 0, 0, 1); if (unlikely(__pyx_t_1 == Pyx_DECREF(__pyx_t_2));
029:         pyx_t_2 = PyNumber_Subtract(__pyx_t_1, __pyx_int_1); if (unlikely(!__pyx_t_2)) {__pyx_filename = __pyx_t_2 = Pyx_DECREF(__pyx_t_2);
030:         pyx_t_1 = __Pyx_PyObject_SetItemStr(__pyx_t_2, __pyx_n_s_right, __pyx_t_1); if (unlikely(!__pyx_t_1)) {__pyx_t_1 = Pyx_DECREF(__pyx_t_2);
031:         bottom = labels.shape[1] - 1
032:         for x_index in range(labels.shape[0]):
033:             for y_index in range(labels.shape[1]):
034:                 label = labels[x_index, y_index]
035:                 if label == 0:
036:                     continue
037:                 vel = data[x_index, y_index]
038:                 # left
039:                 x_check = x_index - 1
040:                 if x_check == -1 and rays_wrap_around:
041:                     x_check = right #wrap around
042:                 if x_check != -1:
043:                     neighbor = labels[x_check, y_index]
044:                     # if the left side gate is masked, keep looking to the left
045:                     # until we find a valid gate or reach the maximum gap size
046:                     if neighbor == 0:
047:                         for i in range(max_gap_x):
048:                             x_check -= 1
049:                             if x_check == -1:
050:                                 If rays_wrap_around:
051:                                     x_check = right
052:                                 else:
053:                                     break
054:                                 neighbor = labels[x_check, y_index]
055:                                 if neighbor != 0:
056:                                     break
057:
058:                                 # add the edge to the collection (if valid)
059:                                 nvel = data[x_check, y_index]
060:                                 collector.add_edge(label, neighbor, vel, nvel)
```

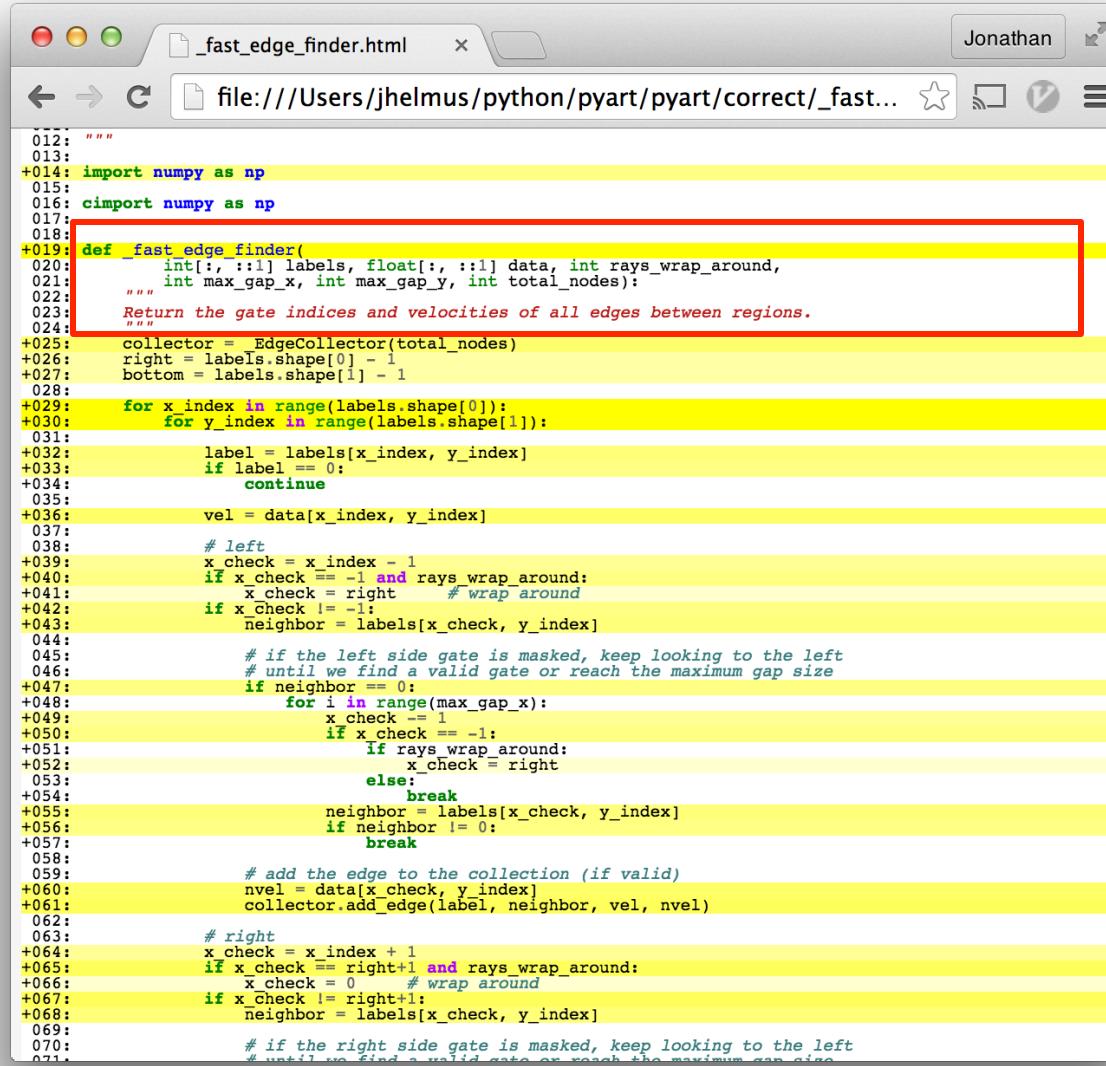
Optimizing Cython code



The screenshot shows a web browser window titled '_fast_edge_finder.html'. The URL in the address bar is 'file:///Users/jhelmus/python/pyart/pyart/correct/_fast...'. The browser interface includes standard controls like back, forward, and search, and a tab labeled 'Jonathan'.

```
012: """
013:
+014: import numpy as np
015:
016: cimport numpy as np
017:
018:
+019: def _fast_edge_finder(labels, data, rays_wrap_around,
020:                         max_gap_x, max_gap_y, total_nodes):
021:     """
022:         Return the gate indices and velocities of all edges between regions.
023:     """
024:
+025:     collector = EdgeCollector(total_nodes)
+026:     right = labels.shape[0] - 1
+027:     pyx_t_2 = Pyx_PyObject_GetAttrStr(_pyx_v_labels, _pyx_n_s_shape); if (unlikely(!_pyx_t_2)) {__pyx_t_2 = Pyx_GOTREF(_pyx_t_2);
+028:     pyx_t_1 = Pyx_GetItemInt(_pyx_t_2, 0, long, 1, __Pyx_PyInt_From_long, 0, 0, 1); if (unlikely(_pyx_t_1 == Pyx_DECREF(_pyx_t_2)); __pyx_t_2 = 0;
+029:     pyx_t_2 = PyNumber_Subtract(_pyx_t_1, __pyx_int_1); if (unlikely(!_pyx_t_2)) {__pyx_filename = __pyx_t_1;
+030:     Pyx_DECREF(_pyx_t_2);
+031:     pyx_t_1 = Pyx_DECREF(_pyx_t_1);
+032:     pyx_v_right = __pyx_t_2;
+033:     pyx_t_2 = 0;
+034:     bottom = labels.shape[1] - 1
+035:     for x_index in range(labels.shape[0]):
+036:         for y_index in range(labels.shape[1]):
+037:             label = labels[x_index, y_index]
+038:             if label == 0:
+039:                 pyx_t_8 = PyObject_RichCompare(_pyx_v_label, __pyx_int_0, Py_EQ); if (unlikely(_pyx_t_8 < 0)) {__pyx_t_8 = Pyx_XGOTREF(_pyx_t_8);
+040:                 pyx_t_9 = Pyx_PyObject_IsTrue(_pyx_t_8); if (unlikely(_pyx_t_9 < 0)) {__pyx_filename = __pyx_t_8;
+041:                 Pyx_DECREF(_pyx_t_8); __pyx_t_8 = 0;
+042:                 if (_pyx_t_9) {
+043:                     continue
+044:                 vel = data[x_index, y_index]
+045:                 # left
+046:                 x_check = x_index - 1
+047:                 if x_check == -1 and rays_wrap_around:
+048:                     if neighbor == 0:
+049:                         for i in range(max_gap_x):
+050:                             x_check -= 1
+051:                             if x_check == -1:
+052:                                 if rays_wrap_around:
+053:                                     x_check = right
+054:                                 else:
+055:                                     break
+056:                             neighbor = labels[x_check, y_index]
+057:                             if neighbor != 0:
```

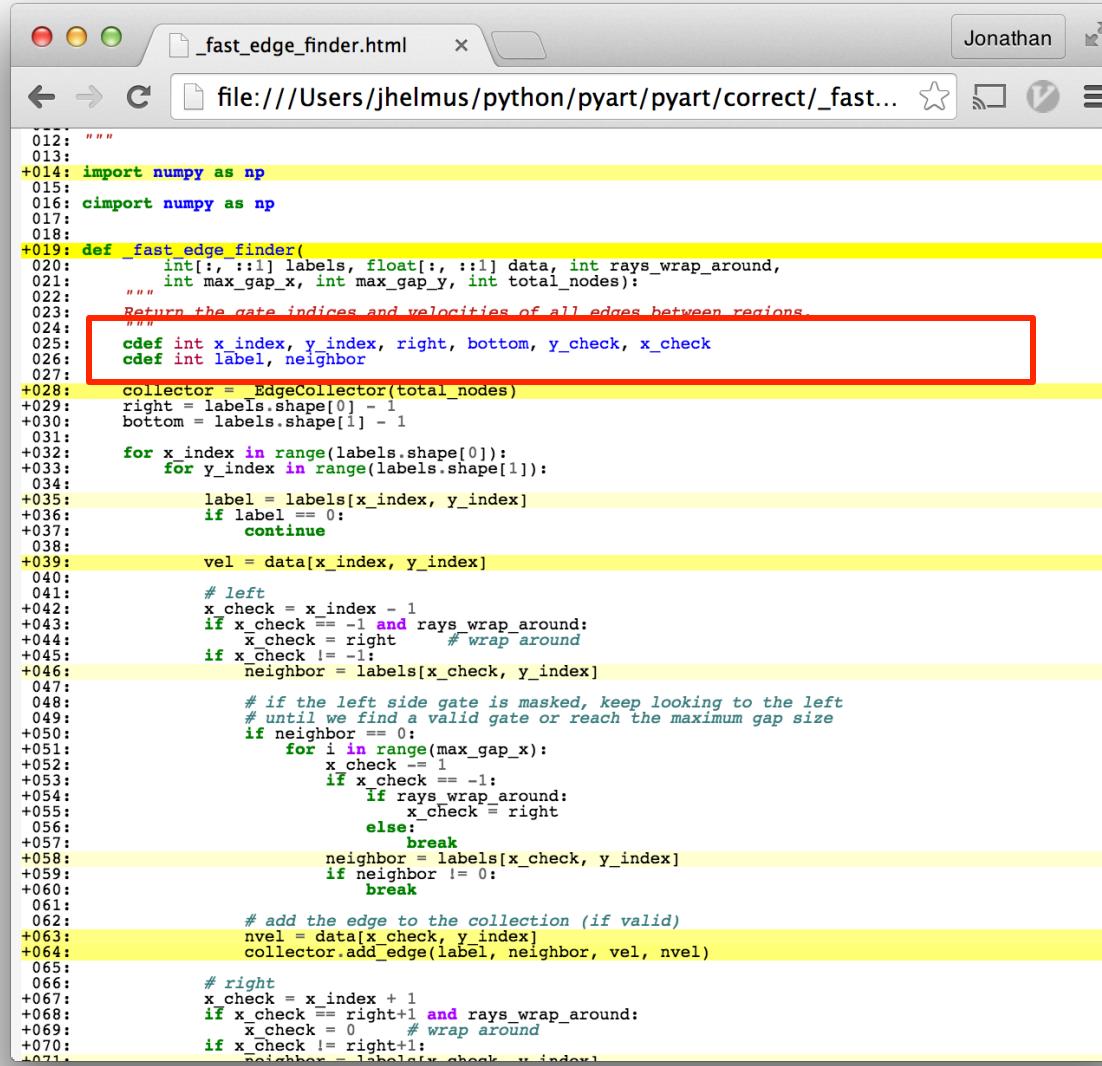
Optimizing Cython code



The screenshot shows a web browser window with the title '_fast_edge_finder.html'. The URL is 'file:///Users/jhelmus/python/pyart/pyart/correct/_fast...'. The page content is a Python script with syntax highlighting. A red box highlights the following code block:

```
+014: import numpy as np
+015:
+016: cimport numpy as np
+017:
+018:
+019: def _fast_edge_finder(
+020:     int[:, ::1] labels, float[:, ::1] data, int rays_wrap_around,
+021:     int max_gap_x, int max_gap_y, int total_nodes):
+022:     """
+023:     Return the gate indices and velocities of all edges between regions.
+024:
+025:     collector = EdgeCollector(total_nodes)
+026:     right = labels.shape[0] - 1
+027:     bottom = labels.shape[1] - 1
+028:
+029:     for x_index in range(labels.shape[0]):
+030:         for y_index in range(labels.shape[1]):
+031:
+032:             label = labels[x_index, y_index]
+033:             if label == 0:
+034:                 continue
+035:
+036:             vel = data[x_index, y_index]
+037:
+038:             # left
+039:             x_check = x_index - 1
+040:             if x_check == -1 and rays_wrap_around:
+041:                 x_check = right # wrap around
+042:             if x_check != -1:
+043:                 neighbor = labels[x_check, y_index]
+044:
+045:                 # if the left side gate is masked, keep looking to the left
+046:                 # until we find a valid gate or reach the maximum gap size
+047:                 if neighbor == 0:
+048:                     for i in range(max_gap_x):
+049:                         x_check -= 1
+050:                         if x_check == -1:
+051:                             If rays_wrap_around:
+052:                                 x_check = right
+053:                             else:
+054:                                 break
+055:                         neighbor = labels[x_check, y_index]
+056:                         if neighbor != 0:
+057:                             break
+058:
+059:                 # add the edge to the collection (if valid)
+060:                 nvel = data[x_check, y_index]
+061:                 collector.add_edge(label, neighbor, vel, nvel)
+062:
+063:             # right
+064:             x_check = x_index + 1
+065:             if x_check == right+1 and rays_wrap_around:
+066:                 x_check = 0 # wrap around
+067:             if x_check != right+1:
+068:                 neighbor = labels[x_check, y_index]
+069:
+070:                 # if the right side gate is masked, keep looking to the left
+071:                 # until we find a valid gate or reach the maximum gap size
```

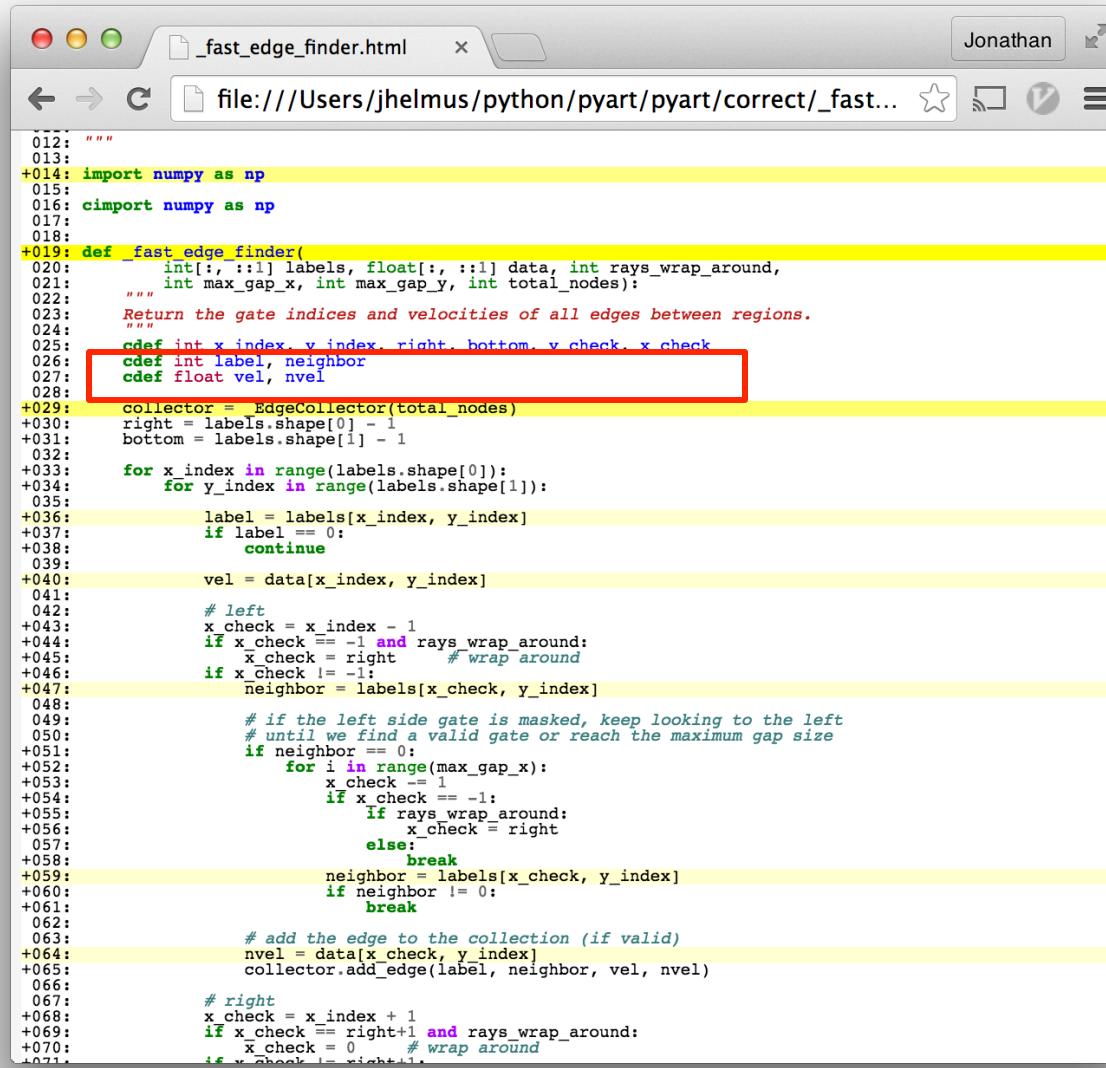
Optimizing Cython code



The screenshot shows a web browser window with the title '_fast_edge_finder.html'. The URL in the address bar is 'file:///Users/jhelmus/python/pyart/pyart/correct/_fast...'. The page content displays a block of Python code with syntax highlighting. A red rectangular box highlights a specific section of the code, which includes several C-style type annotations ('cdef') and variable declarations ('label', 'neighbor'). The code is a Cython implementation of an edge finder function, likely for a scientific application involving grid data.

```
+012: """
+013:
+014: import numpy as np
+015:
+016: cimport numpy as np
+017:
+018:
+019: def _fast_edge_finder(
+020:     int[:, ::1] labels, float[:, ::1] data, int rays_wrap_around,
+021:     int max_gap_x, int max_gap_y, int total_nodes):
+022:
+023:     """
+024:     Return the gate indices and velocities of all edges between regions.
+025:
+026:     cdef int x_index, y_index, right, bottom, y_check, x_check
+027:
+028:     collector = EdgeCollector(total_nodes)
+029:     right = labels.shape[0] - 1
+030:     bottom = labels.shape[1] - 1
+031:
+032:     for x_index in range(labels.shape[0]):
+033:         for y_index in range(labels.shape[1]):
+034:
+035:             label = labels[x_index, y_index]
+036:             if label == 0:
+037:                 continue
+038:
+039:             vel = data[x_index, y_index]
+040:
+041:             # left
+042:             x_check = x_index - 1
+043:             if x_check == -1 and rays_wrap_around:
+044:                 x_check = right  # wrap around
+045:             if x_check != -1:
+046:                 neighbor = labels[x_check, y_index]
+047:
+048:                 # if the left side gate is masked, keep looking to the left
+049:                 # until we find a valid gate or reach the maximum gap size
+050:                 if neighbor == 0:
+051:                     for i in range(max_gap_x):
+052:                         x_check = 1
+053:                         if x_check == -1:
+054:                             if rays_wrap_around:
+055:                                 x_check = right
+056:                             else:
+057:                                 break
+058:                         neighbor = labels[x_check, y_index]
+059:                         if neighbor != 0:
+060:                             break
+061:
+062:                         # add the edge to the collection (if valid)
+063:                         nvel = data[x_check, y_index]
+064:                         collector.add_edge(label, neighbor, vel, nvel)
+065:
+066:             # right
+067:             x_check = x_index + 1
+068:             if x_check == right+1 and rays_wrap_around:
+069:                 x_check = 0  # wrap around
+070:             if x_check != right+1:
+071:                 neighbor = labels[x_check, y_index]
```

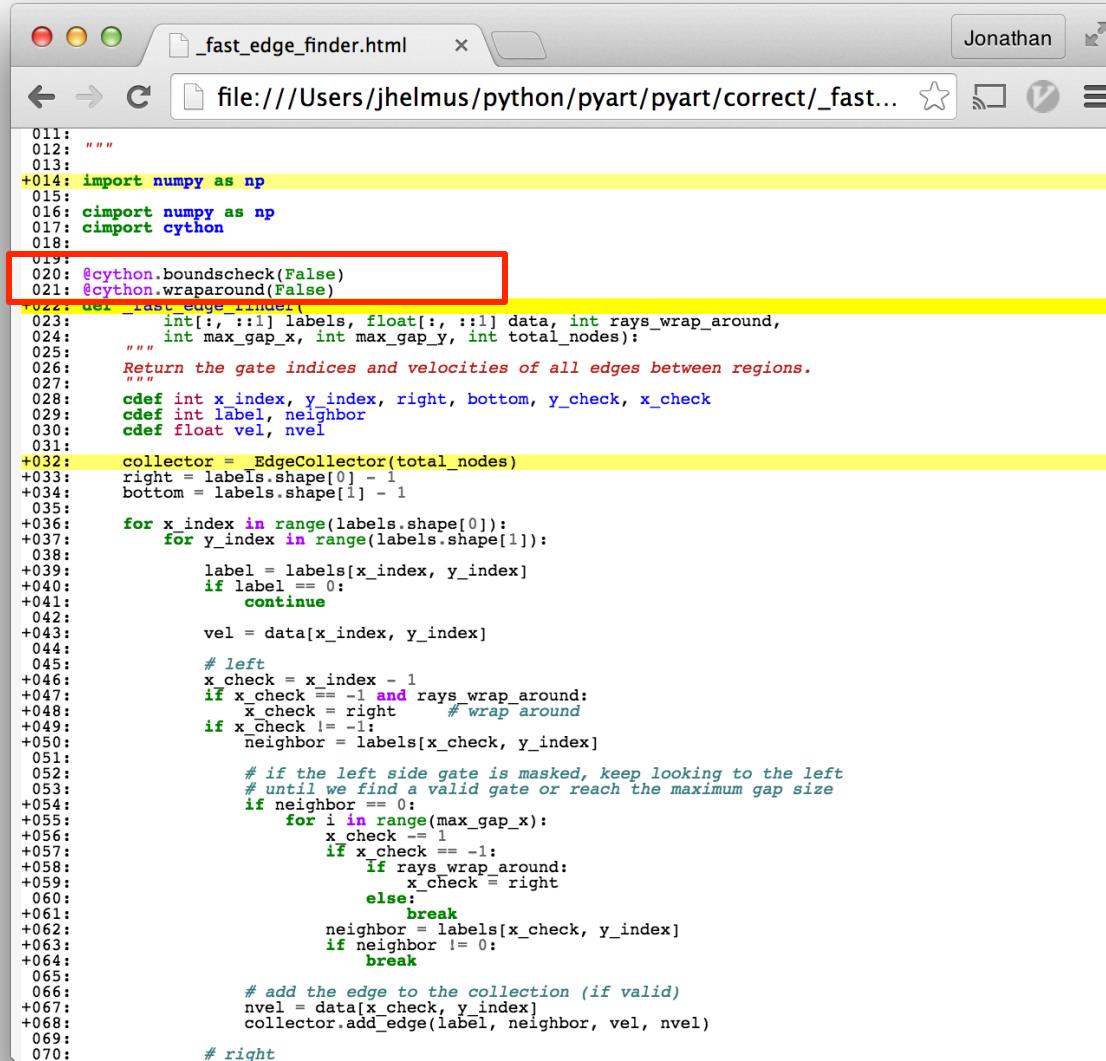
Optimizing Cython code



The screenshot shows a web browser window titled '_fast_edge_finder.html'. The URL in the address bar is 'file:///Users/jhelmus/python/pyart/pyart/correct/_fast...'. The browser interface includes standard controls like back, forward, and search, and a tab labeled 'Jonathan'.

```
012: """
013:
+014: import numpy as np
015:
016: cimport numpy as np
017:
018:
+019: def _fast_edge_finder(
020:     int[:, ::1] labels, float[:, ::1] data, int rays_wrap_around,
021:     int max_gap_x, int max_gap_y, int total_nodes):
022: """
023:     Return the gate indices and velocities of all edges between regions.
024:
025:     cdef int x_index, y_index, right, bottom, v_check, x_check
026:     cdef int label, neighbor
027:     cdef float vel, nvel
028:
+029:     collector = EdgeCollector(total_nodes)
+030:     right = labels.shape[0] - 1
+031:     bottom = labels.shape[1] - 1
+032:
+033:     for x_index in range(labels.shape[0]):
+034:         for y_index in range(labels.shape[1]):
+035:
+036:             label = labels[x_index, y_index]
+037:             if label == 0:
+038:                 continue
+039:
+040:             vel = data[x_index, y_index]
+041:
+042:             # left
+043:             x_check = x_index - 1
+044:             if x_check == -1 and rays_wrap_around:
+045:                 x_check = right      # wrap around
+046:             if x_check != -1:
+047:                 neighbor = labels[x_check, y_index]
+048:
+049:                 # if the left side gate is masked, keep looking to the left
+050:                 # until we find a valid gate or reach the maximum gap size
+051:                 if neighbor == 0:
+052:                     for i in range(max_gap_x):
+053:                         x_check -= 1
+054:                         if x_check == -1:
+055:                             if rays_wrap_around:
+056:                                 x_check = right
+057:                             else:
+058:                                 break
+059:                         neighbor = labels[x_check, y_index]
+060:                         if neighbor != 0:
+061:                             break
+062:
+063:                         # add the edge to the collection (if valid)
+064:                         nvel = data[x_check, y_index]
+065:                         collector.add_edge(label, neighbor, vel, nvel)
+066:
+067:             # right
+068:             x_check = x_index + 1
+069:             if x_check == right+1 and rays_wrap_around:
+070:                 x_check = 0          # wrap around
+071:                 if v_check != right+1:
```

Optimizing Cython code



The screenshot shows a web browser window titled '_fast_edge_finder.html' with the URL 'file:///Users/jhelmus/python/pyart/pyart/correct/_fast...'. The code is written in Cython and is annotated with line numbers. A red box highlights the first few lines of the code:

```
+014: import numpy as np
+015:
+016: cimport numpy as np
+017: cimport cython
+018:
+019: #cython.boundscheck(False)
+020: #cython.wraparound(False)
+021:
+022: def __fast_edge_finder(
+023:     int[:, ::1] labels, float[:, ::1] data, int rays_wrap_around,
+024:     int max_gap_x, int max_gap_y, int total_nodes):
+025:
+026:     """Return the gate indices and velocities of all edges between regions.
+027:
+028:     cdef int x_index, y_index, right, bottom, y_check, x_check
+029:     cdef int label, neighbor
+030:     cdef float vel, nvel
+031:
+032:     collector = EdgeCollector(total_nodes)
+033:     right = labels.shape[0] - 1
+034:     bottom = labels.shape[1] - 1
+035:
+036:     for x_index in range(labels.shape[0]):
+037:         for y_index in range(labels.shape[1]):
+038:
+039:             label = labels[x_index, y_index]
+040:             if label == 0:
+041:                 continue
+042:
+043:             vel = data[x_index, y_index]
+044:
+045:             # left
+046:             x_check = x_index - 1
+047:             if x_check == -1 and rays_wrap_around:
+048:                 x_check = right      # wrap around
+049:             if x_check != -1:
+050:                 neighbor = labels[x_check, y_index]
+051:
+052:                 # if the left side gate is masked, keep looking to the left
+053:                 # until we find a valid gate or reach the maximum gap size
+054:                 if neighbor == 0:
+055:                     for i in range(max_gap_x):
+056:                         x_check -= 1
+057:                         if x_check == -1:
+058:                             if rays_wrap_around:
+059:                                 x_check = right
+060:                             else:
+061:                                 break
+062:                         neighbor = labels[x_check, y_index]
+063:                         if neighbor != 0:
+064:                             break
+065:
+066:                 # add the edge to the collection (if valid)
+067:                 nvel = data[x_check, y_index]
+068:                 collector.add_edge(label, neighbor, vel, nvel)
+069:
+070:             # right
```

Run time optimization using Cython

```
$ kernprof.py -l -v script.py
```

Total time: 155.869 s

Line #	Hits	Time	Per Hit	% Time	Line Contents
--------	------	------	---------	--------	---------------

=====

```
...
164                                # find regions in original data
165    17      347781  20457.7      0.2
166    17        46      2.7      0.0
167    17        20      1.2      0.0
168    17  107326070  6313298.2     68.9
```

Total time: 49.5179 s

Line #	Hits	Time	Per Hit	% Time	Line Contents
--------	------	------	---------	--------	---------------

=====

```
...
160                                # find regions in original data
161    17      348933  20525.5      0.7
162    17        46      2.7      0.0
163    17        22      1.3      0.0
164    17      234857  13815.1      0.5
```

107 seconds vs. 0.234 seconds, x450 performance improvement



Conclusions

- We have developed and implemented, in Python, a novel new algorithm for Doppler velocity dealiasing based upon unfolding regions of similar velocities.
- The ***memory_profiler*** module provides a list of memory usage within a Python function and can be used as a guide to reduce memory usage.
- The ***cProfile*** and ***line_profiler*** modules provide statistics on execution time, either by function or line-by-line. This information can be used to optimize runtime.
- ***Cython*** is a Python to C translator that can be used to speed up the execution time of Python code by adding static type information.

Modification	Run time
Original	94 sec
Cython edge finding	35.7 sec
Edge tracker refactor	14.0 sec
Vectorized unwrapping	8.8 sec