HIP Performance Portability & GPUS

Dr. Joseph Schoonover

CEO, Cloud-HPC Systems Architect
Fluid Numerics, LLC
Boulder, CO

joe@fluidnumerics.com

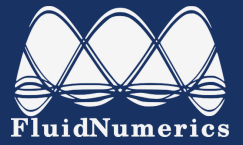

HPC GPU Programming Ecosystem

Compiler →

API ↓

	GCC	Intel	PGI	XL	CLANG
OpenCL					
OpenMP					
OpenACC					
CUDA					
HIP					

FluidNumerics

HPC GPU Programming Ecosystem

Architecture →

API ↓

	Nvidia	AMD	Arm	Intel
OpenCL				
OpenMP				
HIP				
OpenACC				
CUDA				

FluidNumerics

We don’t always know the what the next system will be.

Cloud platforms are delivering more GPU options

Rewriting code can be expensive

How proactive can we be ?

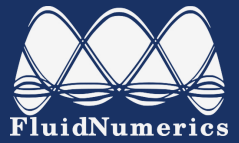

ROCm

Radeon **O**pen **Co**m**p**ute

https://rocm.github.io

Open Source platform for programming AMD and Nvidia GPUs

- OpenCL
- LLVM based C/C++ compilers (Branch off clang)
  - hcc
- GPU offloading wrapper
  - hipcc
- Accelerated Math and Deep Learning Libraries
- Multi-GPU support
  - Peer-to-peer multi-GPU operation with RDMA
- Profilers and Debuggers
- System monitoring tools
  - rocm-smi

FluidNumerics

ROCm

Radeon **O**pen **Co**m**m**pute

https://rocm.github.io

Debian and RPM repositories available

```
$ sudo {yum|apt-get} install rocm-dkms
```

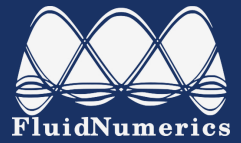

HIP

Heterogeneous-compute **I**nterface for **P**ortability
https://github.com/ROCm-Developer-Tools/HIP

Interface for compiling C/C++ for AMD and Nvidia GPUs

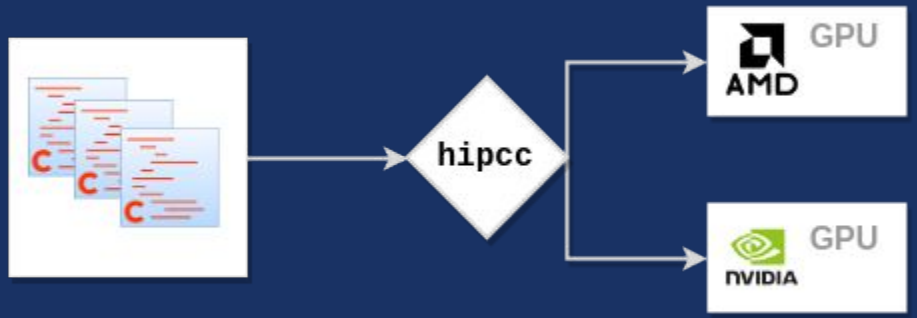

FluidNumerics

HIP
API comparisons with CUDA

CUDA

```
cudaMalloc((void**)&nodes_dev, N*sizeof(float) );
```

HIP

```
hipMalloc((void**)&nodes_dev, N*sizeof(float) );
```

FluidNumerics

HIP

API comparisons with CUDA

CUDA

```
dim3 threadsPerBlock(nthreads,nthreads,nthreads);
dim3 blocks(n_elements);
GPUKernel<<<blocks,threadsPerBlock>>>( input );
```

HIP

```
dim3 threadsPerBlock(nthreads,nthreads,nthreads);
dim3 blocks(n_elements);
hipLaunchKernelGGL(GPUKernel, dim3(blocks), dim3(threadsPerBlock), 0, 0, input);
```

FluidNumerics

HIP

Compilation comparisons with CUDA

CUDA

```
$ nvcc source_code.cu
```

HIP

```
$ hipcc source_code.cu
```

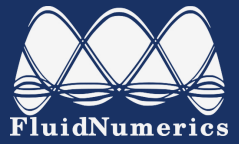

HIP

HIPify existing CUDA Code

FluidNumerics

Spectral Element Divergence (3-D)

Cost & Performance Study

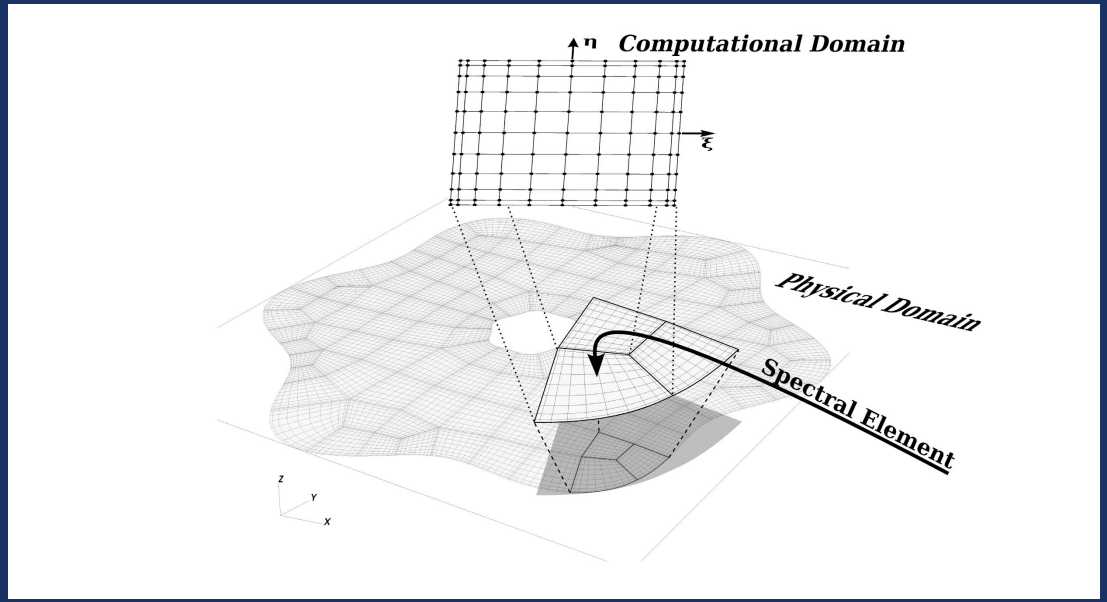

FluidNumerics

Spectral Element Divergence (3-D)

Cost & Performance Study

Serial Code

```
for( iel=0; iel<n_elements; iel++ ){
  for( k=0; k<N; k++ ){
    for( j=0; j<N; j++ ){
      for( i=0; i<N; i++ ){
        df[SP_3D_INDEX(i,j,k,iel,N)] = 0.0;
        for( ii=0; ii<=N; ii++){
          df[3D_INDEX(i,j,k,iel,N)]=df[3D_INDEX(i,j,k,iel,N)] +
            D[INDEX(ii,i,N)]*f[VEC_3D_INDEX(0,ii,j,k,iel,N)] +
            D[INDEX(ii,j,N)]*f[VEC_3D_INDEX(1,i,ii,k,iel,N)] +
            D[INDEX(ii,k,N)]*f[VEC_3D_INDEX(2,i,j,ii,iel,N)] ;

        }
      }
    }
  }
}
```

FluidNumerics

Spectral Element Divergence (3-D)

Cost & Performance Study

CUDA Kernel

```
int iel = blockidx.x;
int i   = threadidx.x;
int j   = threadidx.y;
int k   = threadidx.z;

df[SP_3D_INDEX(i,j,k,iel,N)] = 0.0;
for( ii=0; ii<=N; ii++){
  df[3D_INDEX(i,j,k,iel,N)]=df[3D_INDEX(i,j,k,iel,N)] +
    D[INDEX(ii,i,N)]*f[VEC_3D_INDEX(0,ii,j,k,iel,N)] +
    D[INDEX(ii,j,N)]*f[VEC_3D_INDEX(1,i,ii,k,iel,N)] +
    D[INDEX(ii,k,N)]*f[VEC_3D_INDEX(2,i,j,ii,iel,N)] ;
 }
```

FluidNumerics

Spectral Element Divergence (3-D)

Cost & Performance Study

- **API's Compared**
  - CUDA - nvcc compiled
  - HIP - hipified CUDA, hipcc compiled
- **Compute Platforms**
  - **Google Cloud Platform**
    - **Nvidia Tesla P100**
    - **Nvidia Tesla V100**
  - **GPUEater (gpueater.com)**
    - **Radeon Vega 64**
    - **Radeon Vega Frontier Edition**

FluidNumerics

Cost/hour vs. GPU

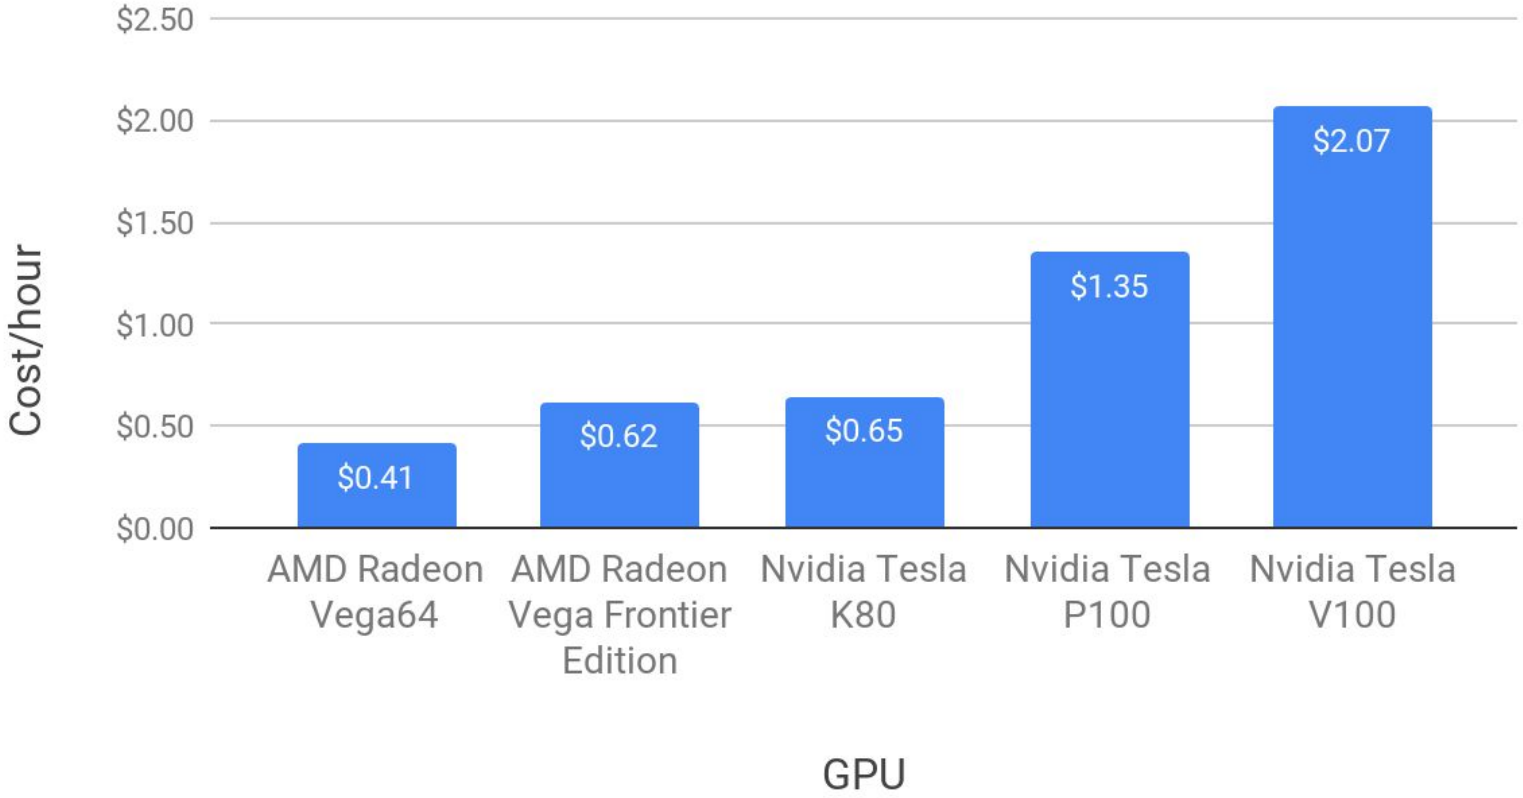

FluidNumerics

Divergence Runtimes : hipcc and nvcc

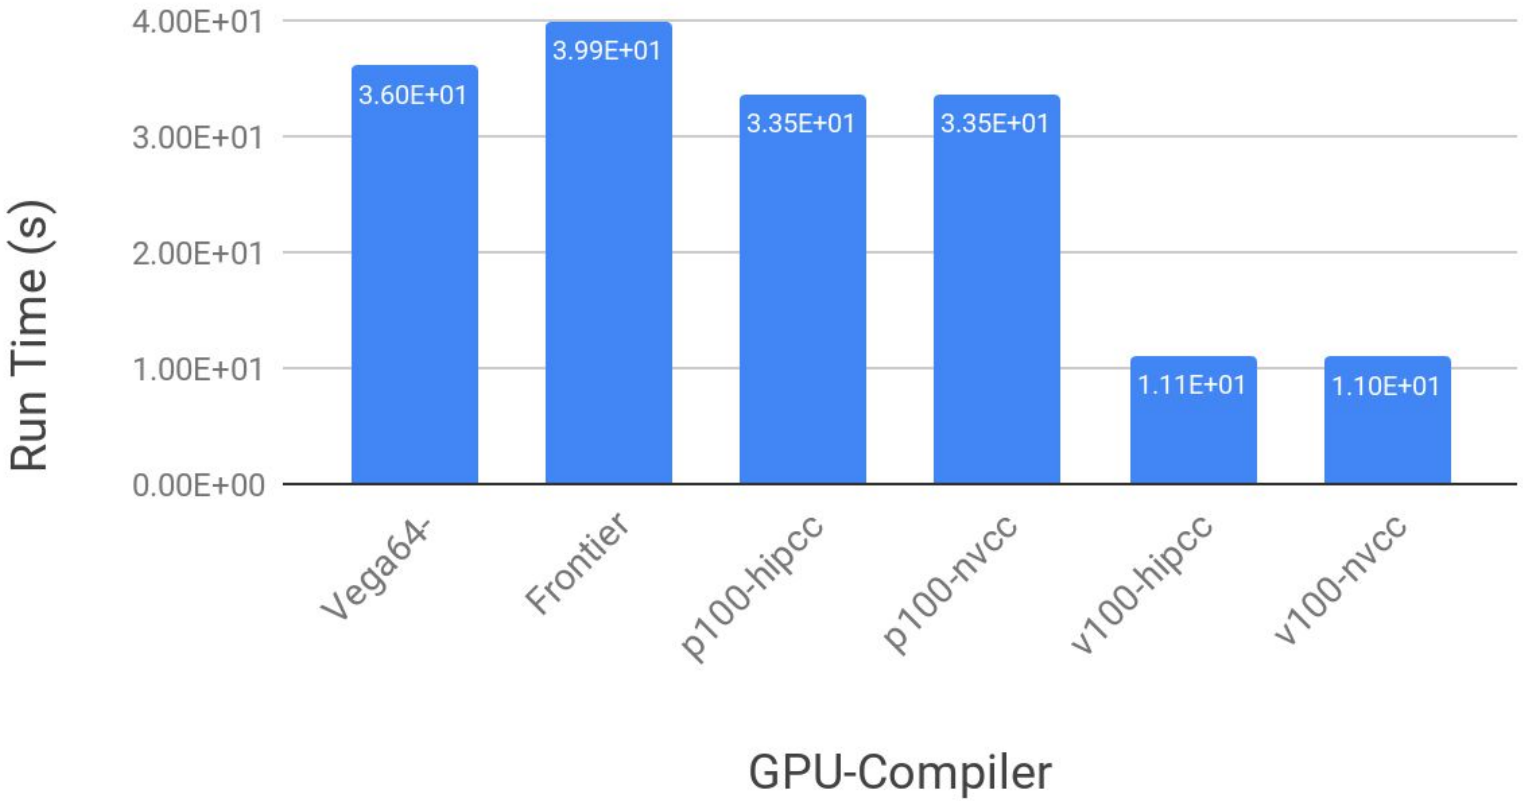

FluidNumerics

Cloud Costs : hipcc and nvcc

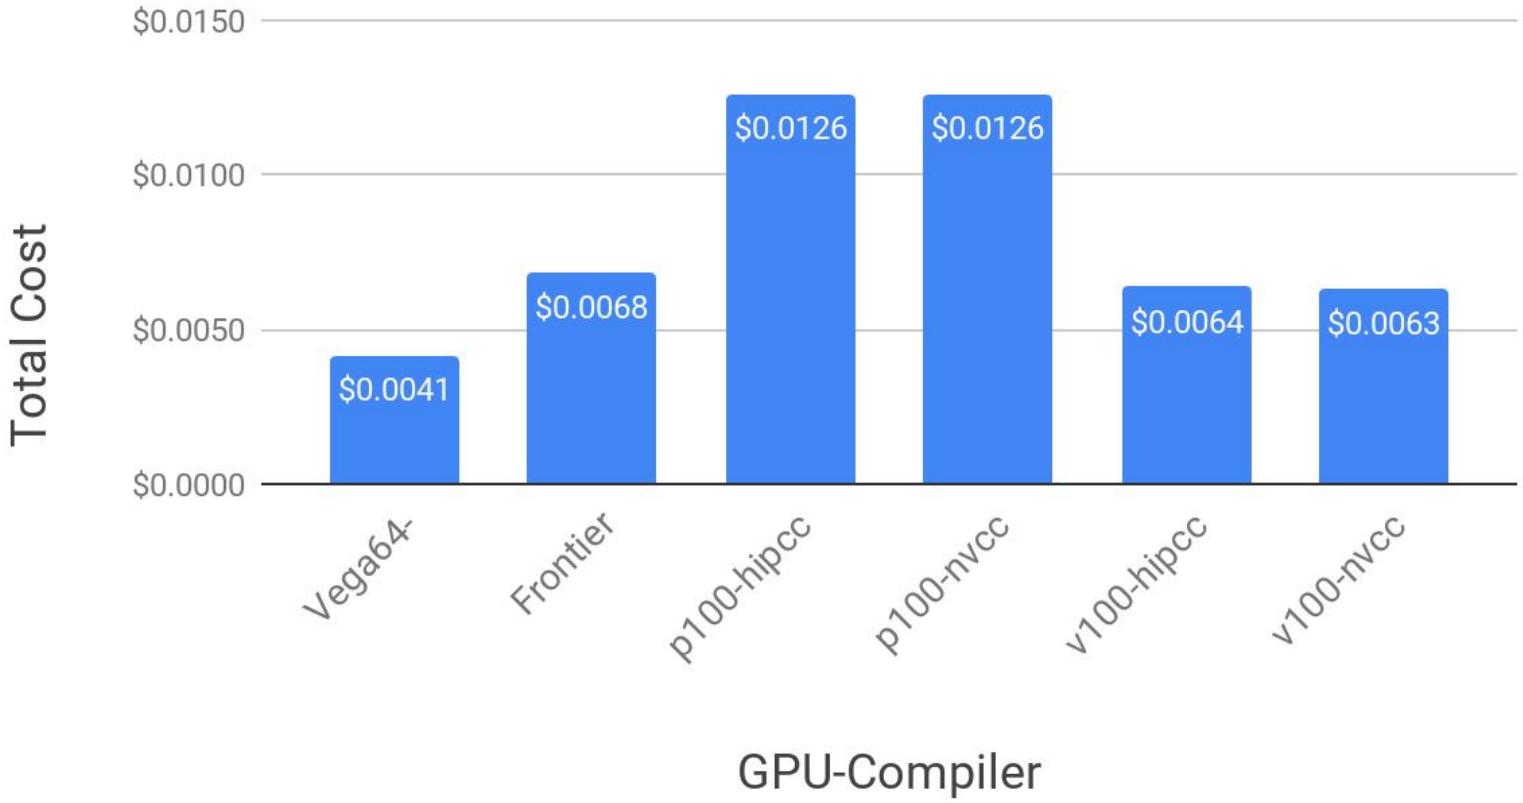

FluidNumerics

Spectral Element Divergence (3-D)

Cost & Performance Study

- **Total cost for cost & performance study was less than $5 in compute resources.**
- **HIP kernels give consistent performance with CUDA kernels on identical platforms**
- **V100's result in shortest time-to-solution**
- **Radeon Vega64 gives P100-like performance at the least expensive cost**
- **Hardware matters - Kernel performance varies across hardware**

FluidNumerics

Further Questions

Contact : joe@fluidnumerics.com

https://fluidnumerics.com

FluidNumerics