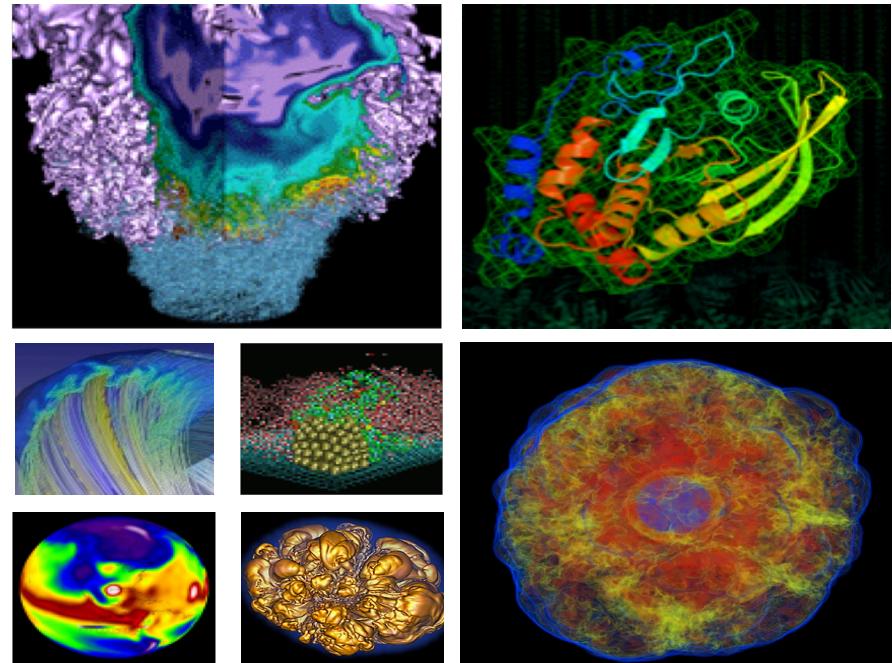


Experience supporting Containers with Shifter at NERSC



Shane Canon

Data & Analytics Group, NERSC

Agenda

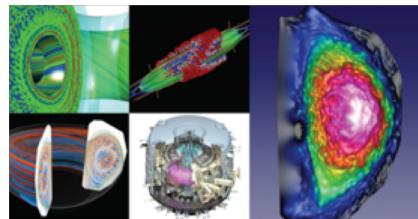
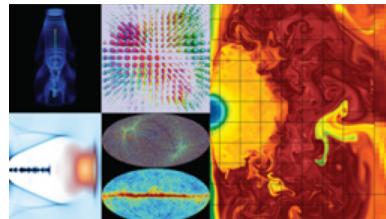
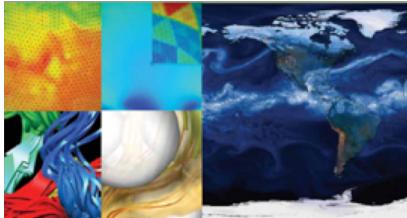


- Motivations for containers at NERSC
- Shifter: Goals and Design
- Experience supporting containers at NERSC

Why Containers at NERSC



- NERSC deploys advanced HPC and data systems for the broad Office of Science community
- Approximately 6000 users and 750 projects
- Growing number of users around Analyzing Experimental and Observational Data, "Big Data" Analytics, and Machine Learning
- Shift towards converged systems that support traditional modeling and simulation workloads plus new models



The Struggles

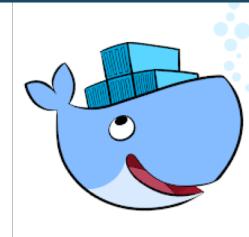
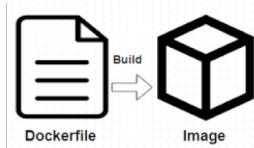


- My software doesn't build on this system...
- I'm missing dependencies...
- I need version 1.3.2 but this system has version 1.0.2..
- I need to re-run the exact same thing 12 months from now...
- I want to run this exact same thing somewhere else...
- I want my collaborators to have the same exact software as me...
- I've heard about these Containers, can I just run that?
- Can I run docker on this HPC system?

Docker Basics



- Build images that captures applications requirements.
- Manually commit or use a recipe file.
- Push an image to DockerCloud, a hosted registry, or a private Docker Registry.
- Share Images
- Use Docker Engine to pull images down and execute a container from the image.



Why not just run Docker

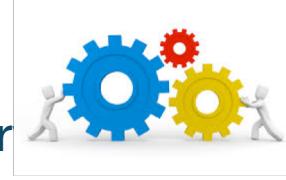


- Security: Docker currently uses an all or nothing security model. Users would effectively have system privileges

```
> docker run -it -v /:/mnt --rm busybox
```



- System Architecture: Docker assumes local disk
- Integration: Docker doesn't play nice with batch systems.
- System Requirements: Docker typically requires a very modern kernel
- Complexity: Running real Docker would add new layers of complexity



Long History with Custom Environments



- ~2003: ChrootOS(CHOS) – System to enable projects to have customized environments on a shared cluster (PDSF). Integrated with login and batch and used a custom kernel module to provide a seamless experience.
- ~2014: MyDock – Thin wrapper around Docker to allow users to securely run Docker on a shared cluster (Jesup/Carver).
- ~2015: Shifter: “User Defined Environments”

Shifter (HPC) versus Spin (Services)



Shifter

Runs processes as the user
Runs on the HPC systems

Best for:

Simulation or analysis runs
Need to run at scale
Need to read/write a lot of data

Spin

Runs with stock Docker and Rancher
Runs on dedicated hardware

Best for:

Running services or processes that
need to run “indefinitely”
Services that need to be externally
accessible

Solution: Shifter



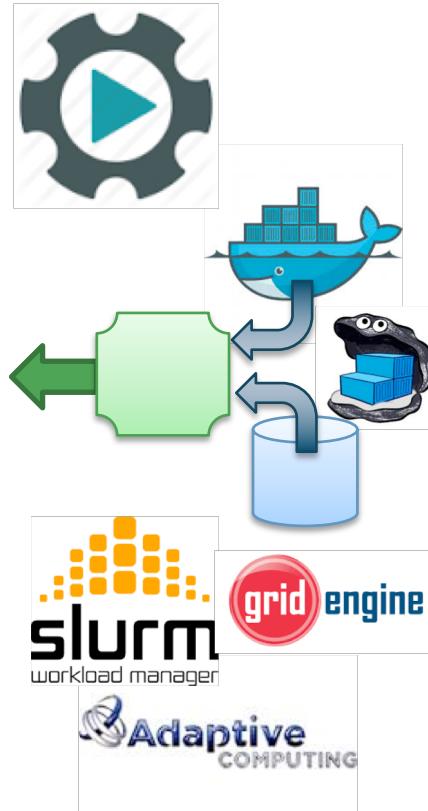
- Design Goals:
 - User independence: Require no administrator assistance to launch an application inside an image
 - Shared resource availability (e.g., file systems and network interfaces)
 - Leverages or integrates with public image repos (i.e. DockerHub)
 - Seamless user experience
 - Robust and secure implementation
- Hosted at GitHub:
 - <https://github.com/nersc/shifter>



Shifter Components



- Shifter Image Gateway
 - Imports and converts images from DockerHub and Private Registries
- Shifter Runtime
 - Instantiates images securely on compute resources
- Work Load Manager Integration
 - Integrates Shifter with WLM
 - Implementations include Slurm, UGE, and Torque/Moab



Usage



- Use shifterimg pull to pull images from a registry
 - Only do this once or after an update

```
> shifterimg pull ubuntu:14.04
```

- Use shifter command to run a container with an image

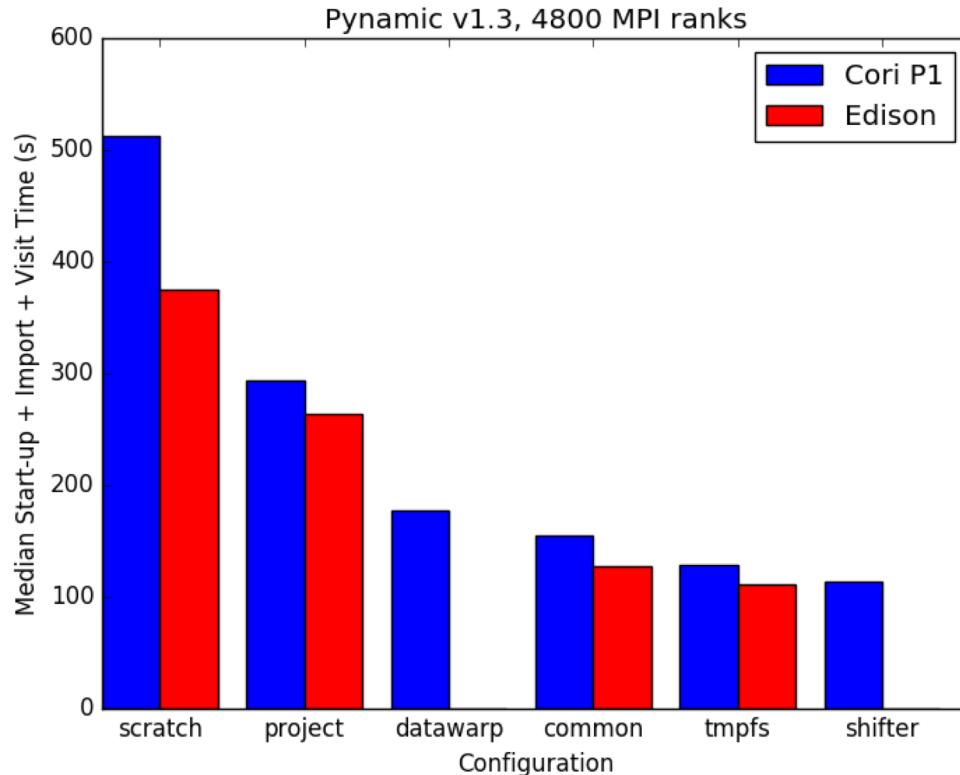
```
> shifter --image=ubuntu:14.04 bash
$ lsb_release -a
No LSB modules are available.
Distributor ID:      Ubuntu
Description:         Ubuntu 14.04.5 LTS
Release:            14.04
Codename:           trusty
```

Containers and Science



- Productivity
 - Pick the OS that works best for your app and use the system package manager to install dependencies.
- Reusability and Collaboration
 - Share images across a project to avoid rebuilds and avoid mistakes
- Reproducibility
 - Everything you need to redo a scientific analysis can be in the image (apps, libraries, environment setup, scripts)
- Portability
 - Can easily run on different resources (of the same architecture)

Shifter accelerates many Apps



Shifter and MPI



- In Image
 - Add required libraries directly into image.
 - Users would have to maintain libraries and rebuild images after an upgrade.
- Managed Base Image (Golden Images)
 - User builds off of a managed image that has required libraries.
 - Images are built or provided as part of a system upgrade.
 - Constrained OS choices and a rebuild is still required.
- Volume Mounting
 - Applications built using ABI compatibility.
 - Appropriate libraries are volume mounted at run time.
 - No rebuild required, but may not work for all cases.

Running an MPI Job – Building Image



```
FROM nersc/mpi-ubuntu:14.04  
  
ADD . /app  
RUN cd /app && \  
    mpicc -o hello helloworld.c
```

Dockerfile

```
> docker build -t scanon/hello .  
> docker push scanon/hello
```

Running an MPI Job – Submit and run



```
#!/bin/sh
#SBATCH --image=scanon/hello
srun -np 10 shifter /app/hello
```

submit.sl

> sbatch submit.sl

How does Shifter differ from Docker?



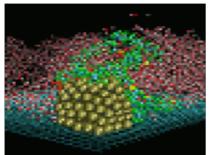
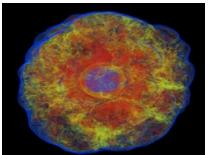
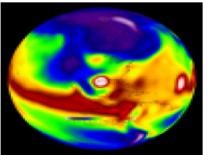
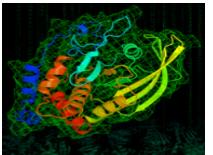
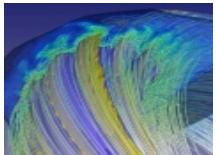
Most Noticeable

- Image read-only on the Computational Platform
- User runs as the user in the container – not root
- Image modified at container construction time (e.g. additional mounts)

Less Noticeable:

- Shifter only uses mount namespaces, not network or process namespaces
- Shifter does not use cgroups directly (integrated with the Workload Manager)
- Shifter uses individual compressed filesystem files to store images, not the Docker graph (slows down iterative updates)
- Shifter starts some additional services (e.g. sshd in container space)

Other Tips and Tricks



Volume Mounts



- Volume Mounts provide a way to map external paths into container paths.
- This allows paths in the container to be abstracted so it can be portable across different systems.
- All runtimes support volume mounts but the syntax may vary.
- Basic syntax is:
-volume <external path>:<container path>

Using Volume Mounts



```
canon@cori06:~> ls $SCRATCH/myjob  
config data.in  
  
canon@cori06:~> shifter --image=ubuntu \  
--volume=$SCRATCH/myjob:/data bash  
~$ ls /data/  
config data.in
```

PerNode Write Cache (Shifter)



- PerNodeWrite extends the volume concept to create temporary writeable space that aren't shared across nodes.
- These spaces are ephemeral (removed on exit)
- These are node local and the size can be adjusted
- Performs like a local disk but is more flexible
- Basic syntax is

```
--volume <external path>:<container path>:perNodeCache=size=XXG
```

Using PerNode Mounts



```
canon@cori06:~> shifter --image=ubuntu \
    --volume=$SCRATCH:/scratch:perNodeCache=size=100G /bin/bash
~$ df -h /scratch/
Filesystem      Size  Used Avail Use% Mounted on
/dev/loop4       100G   33M  100G   1% /scratch
~$ dd if=/dev/zero bs=1k count=10M of=/scratch/output
10737418240 bytes (11 GB, 10 GiB) copied, 22.2795 s, 482 MB/s
~$ ls -lh /scratch/output
-rw-r--r-- 1 canon canon 10G Nov  9 23:38 /scratch/output
~$ exit
canon@cori06:~> shifter --image=ubuntu \
    --volume=$SCRATCH:/scratch:perNodeCache=size=100G /bin/bash
~$ ls -l /scratch
total 0
```

Multi-Stage Builds



- Added in Docker 17.05
- Allows a build to progress through stages
- Files can be copied from a stage to later stages
- Useful for splitting images between build and run-time to keep image sizes small
- Can be used to make public images that make use of commercial compilers

Dockerfile – Multistage build



```
FROM centos:7 as build
RUN yum -y install gcc make
ADD code.c /src/code.c
RUN gcc -o /src/mycode /src/code.c

FROM centos:7
COPY --from=build /src/mycode /usr/bin/mycode
```

Multistage Builds with Intel Compilers



- We have a model for compiling applications using Intel Compilers
- Requires tunneling a connection to NERSC's license server
- Uses multistage build so resulting runtime image doesn't contain any licensed software
- Images are hosted in NERSC's registry which requires a NERSC login
- Documented on NERSC's web site
- NERSC Postdoc Jonathan Madsen figured this all out



Jonathan Madsen

Example Dockerfile



```
##### Stage 1 #####
FROM registry.services.nersc.gov/nersc/intel_cxx_mpi-devel as builder
# ... build your code with "devel" image variant ...
# ... recommended to use a common install prefix, such as "/opt/local", e.g.
#

ENV CC /opt/intel/bin/icc
ENV CXX /opt/intel/bin/icpc

RUN cd ${HOME}/package_a && \
    ./configure --prefix=/opt/local && \
    make install -j8 && \
    cd ${HOME}/package_b && \
    mkdir build_package_b && \
    cd build_package_b && \
    cmake -DCMAKE_INSTALL_PREFIX=/opt/local .. && \
    make install -j8

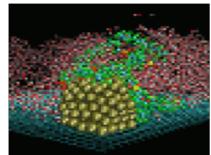
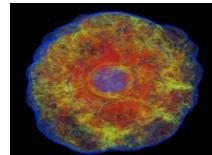
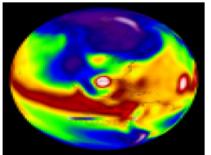
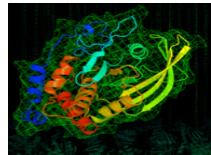
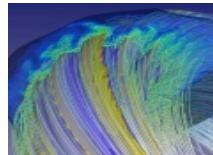
#####
# Stage 2 #####
# ... don't have to clean above, just copy over installation
FROM registry.services.nersc.gov/nersc/intel_cxx_mpi_runtime
COPY --from=builder /opt/local/ /opt/local
RUN echo '/opt/local/lib' > /etc/ld.so.conf.d/package-libs.conf && \
    ldconfig
```

Shane's Advice



- Use a Dockerfile - Leverages the build caching and is the most portable
- Use a multi-stage build – Helps with size and licensing issues
- Create base images – Base images that are common across a suite of applications
- Version images – Helps with reproducibility
- Leverage ABI – Best supported in MPICH but hopefully this improves
- Use Volume Mounts – Separate data from the application. Use volume mounts for data
- Run as non-root – Design your images to run as non-root (most portable and safe)

Use Case Example and Summary



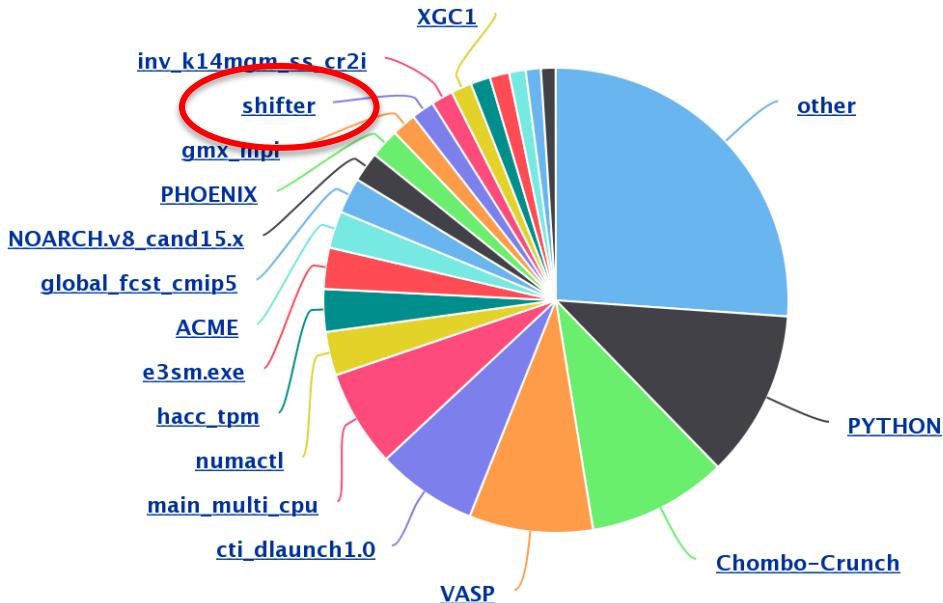
Stats on Shifter Use at NERSC



- 10M+ Image Lookups
- 1900+ Unique Image tags
- 700+ Unique Users
- Still a small fraction of NERSC overall use (~3%)

Cori Machine Hours Breakdown by Binary Names

Processes as a percentage of 100% of total machine hours (5374916 hours).



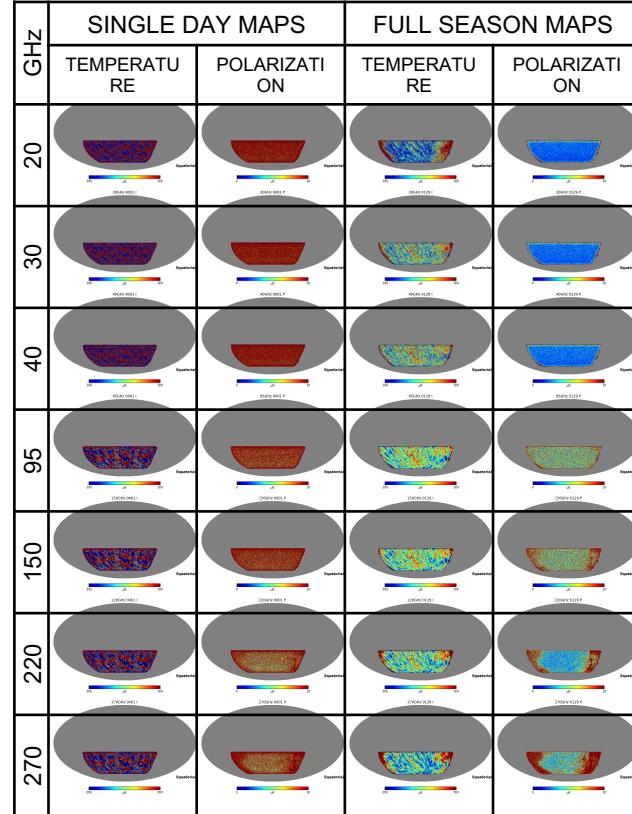
Highcharts.com

Shifter in Action



Measuring the Composition of the Universe

- **CMB – S4** - Ambitious collection of telescopes to measure the remnants of the Big Bang with unprecedented precision
- Simulated 50,000 instances of telescope using **600,000 cores** on Cori KNL nodes.
- Why Shifter?
 - Python wrapped code needs to start at scale
 - Without Shifter application timed out during startup
- Uses Intel Compilers



Courtesy of Ted Kisner



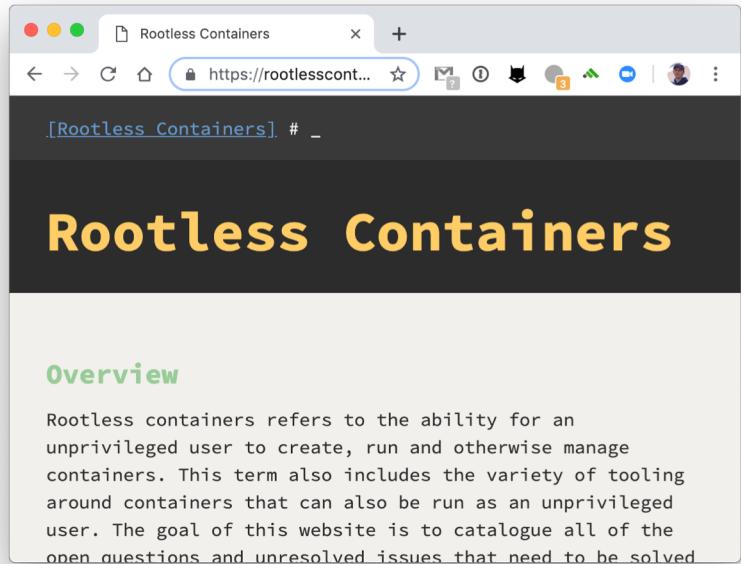
Office of
Science



What's next?



- Continued training and development best practices (in partnership with ECP projects)
- Preparing for Perlmutter (GPU support) – Leverage work done by CSCS
- Exploring rootless/nosuid for Shifter
- Exploring emerging rootless support in the general docker/container sphere (Docker, podman)



A screenshot of a web browser window titled "Rootless Containers". The URL in the address bar is <https://rootlesscont...>. The page content includes a large title "Rootless Containers" and a section titled "Overview". The "Overview" section contains text about rootless containers and their ability to be managed by unprivileged users.

Rootless containers refers to the ability for an unprivileged user to create, run and otherwise manage containers. This term also includes the variety of tooling around containers that can also be run as an unprivileged user. The goal of this website is to catalogue all of the open questions and unresolved issues that need to be solved



Summary



Containers are great

- ✓ Productivity – Get exactly what you need for your application
- ✓ Portable – Run the same software on different resources (assuming architectural compatibility)
- ✓ Sharable – Collaborators can run the same code as you with less chance of problems
- ✓ Reproducible – Run the same image later
- ✓ Performant – Can actually speed up applications in some cases



The background image shows a modern, multi-story building with a glass and steel facade. The sky is a gradient from blue to orange, suggesting sunset or sunrise. In the foreground, there's a red rectangular sign with white text that reads "WE'RE HIRING".

Questions...

Shane Canon: scanon@lbl.gov

Upcoming Events:

- June 11-13 - Jupyter Facilities Workshop
- June 20 – HPCW at ISC-HP
- Nov - SC19: Stay tuned

This work was supported by the Director, Office of Science, Office of Advanced Scientific Computing Research of the U.S. Department of Energy under Contract No.

DE-AC02-05CH11231.

WE'RE
HIRING