

Interactive Performance Analysis with Vampir

UCAR Software Engineering Assembly
in Boulder/CO, 2013-04-03

Andreas Knüpfer, Thomas William
TU Dresden, Germany

- Introduction
- Vampir displays
- GPGPU support
- Scalability
- Performance Problems
- Conclusions & Preview on Tutorial

Why should I bother with performance analysis?

- Efficient usage of expensive and limited resources
- Scalability to achieve next bigger simulation
- Well, why are you here after all?

How should I do it?

- Optimization phase, just like the testing phase
- Use tools, avoid do-it-yourself-with-printf, really!

What methods are available?

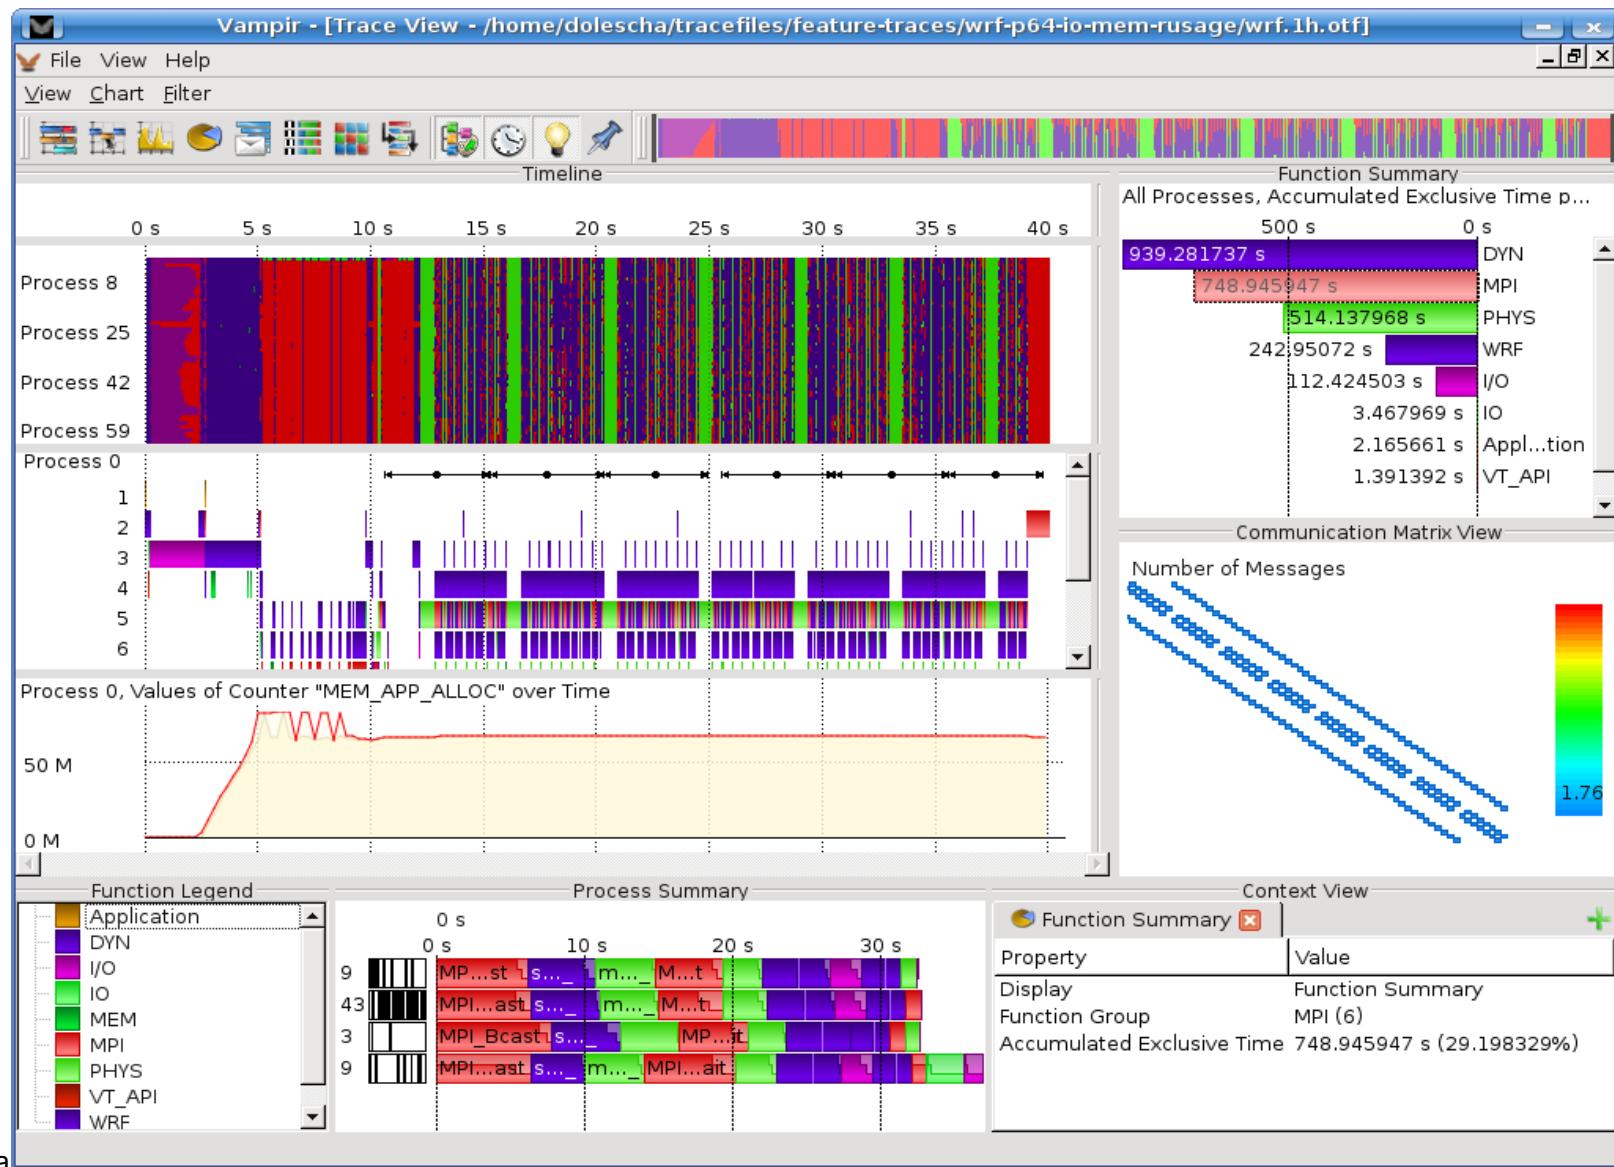
- Profiling
- Event Tracing

- Based on recorded event traces
 - Sorted collection of run-time events per process
- Event types with time stamp and process:
 - Enter given function/subroutine
 - Send or receive message of N bytes to/from P
 - Hardware performance counter values
 - More ...
- Visualize dynamic run-time behavior
 - Various timeline displays
 - Various statistics displays
- Interactive browsing/zooming/selection
 - Zoom in time and processes
 - Feasible for very large and highly parallel traces

- Introduction
- Vampir displays
- GPGPU support
- Scalability
- Performance Problems
- Conclusions & Preview on Tutorial

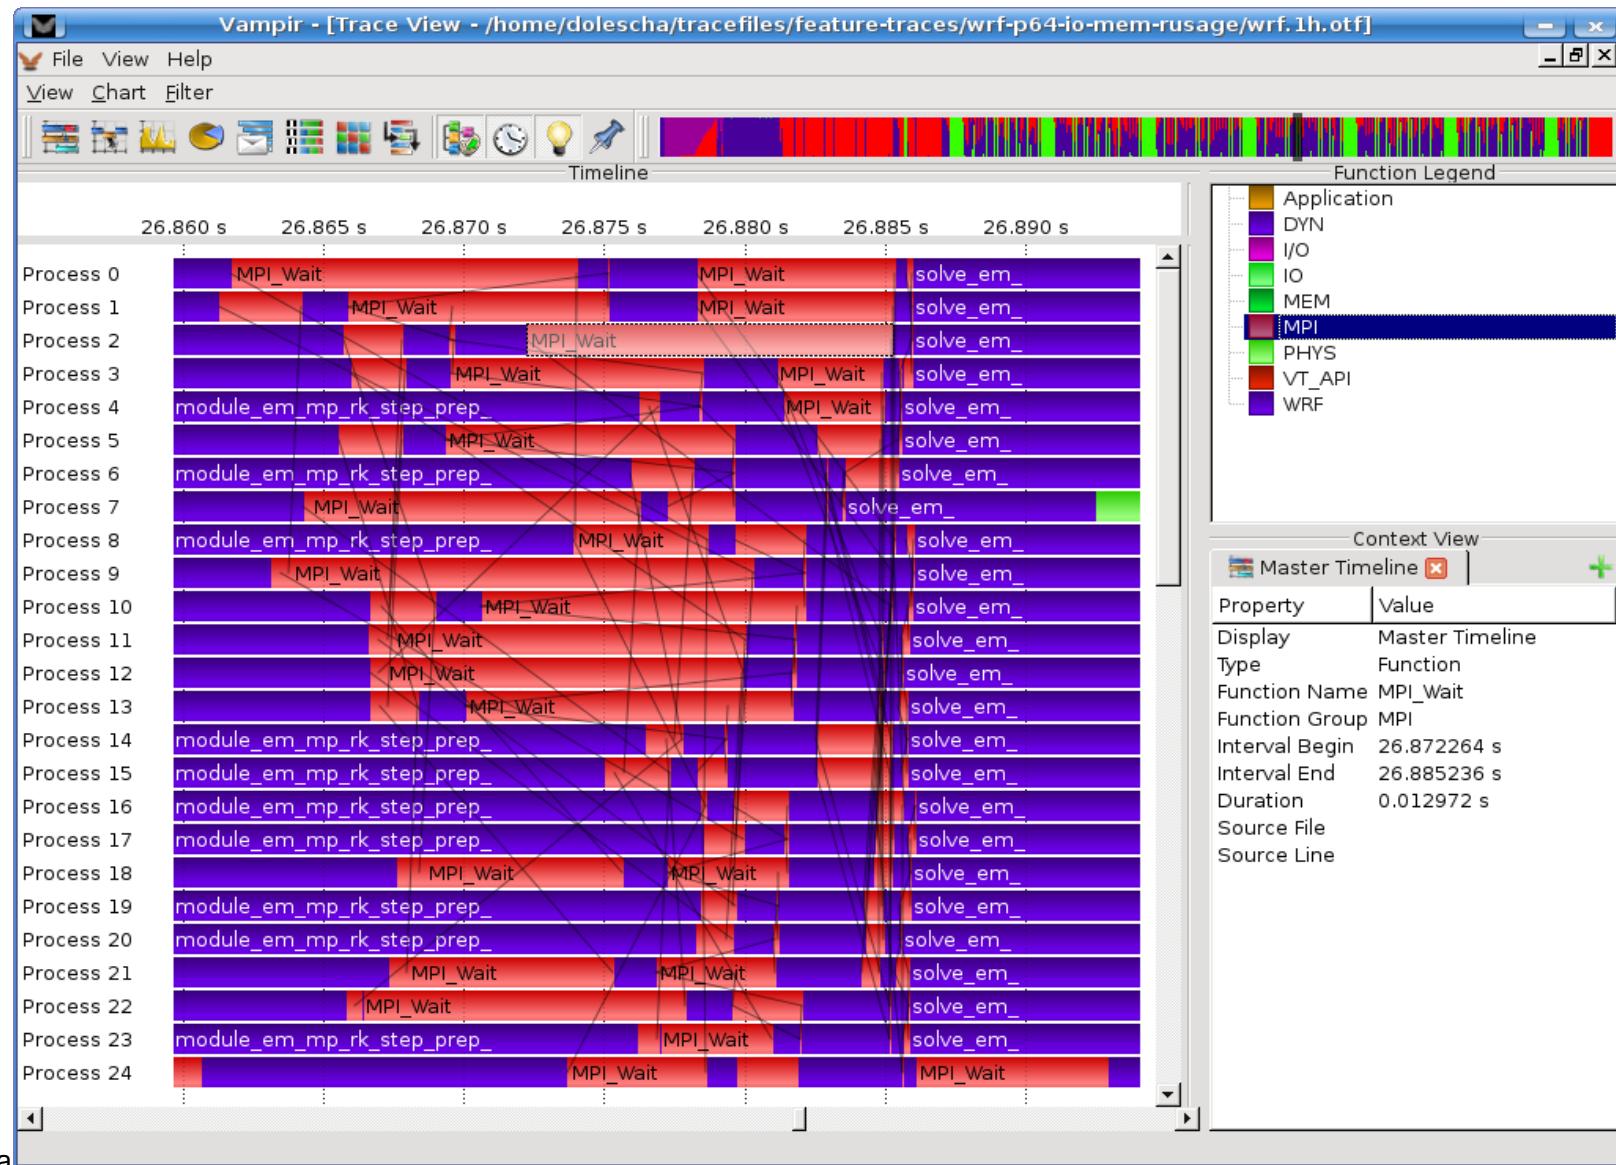
Vampir Display Overview

VI-HPS



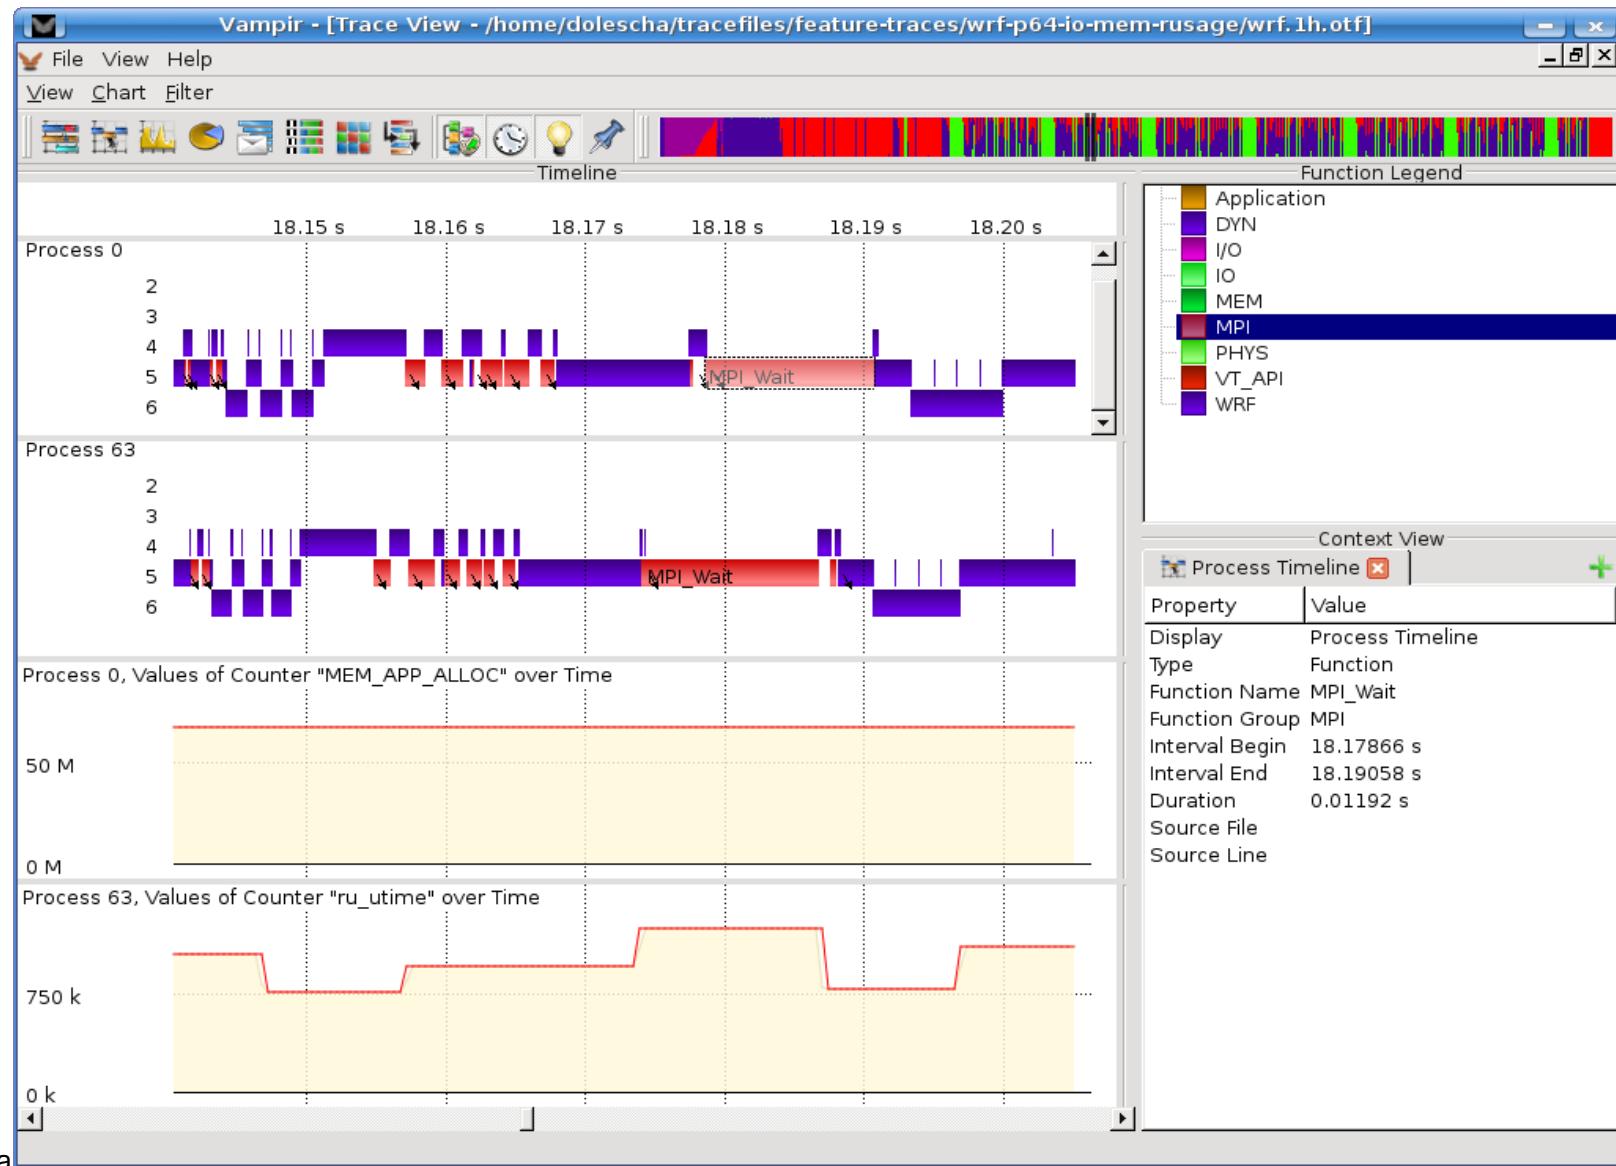
Master Timeline

VI-HPS



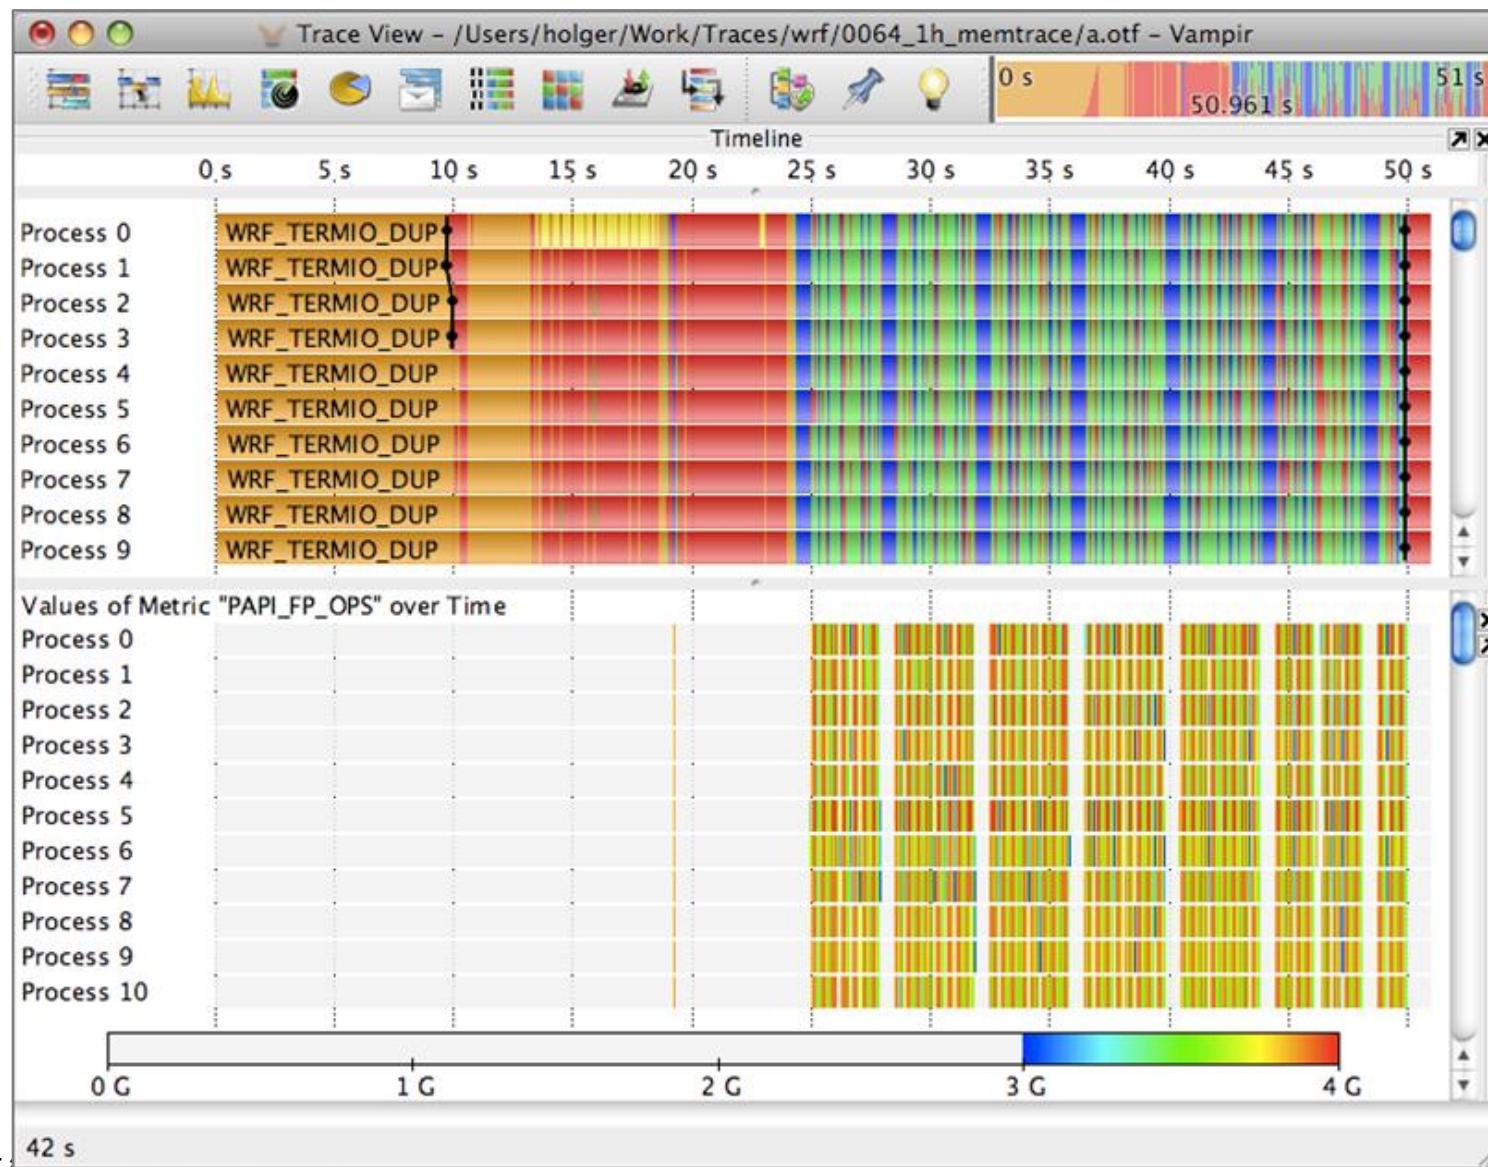
Process and Counter Timelines

VI-HPS



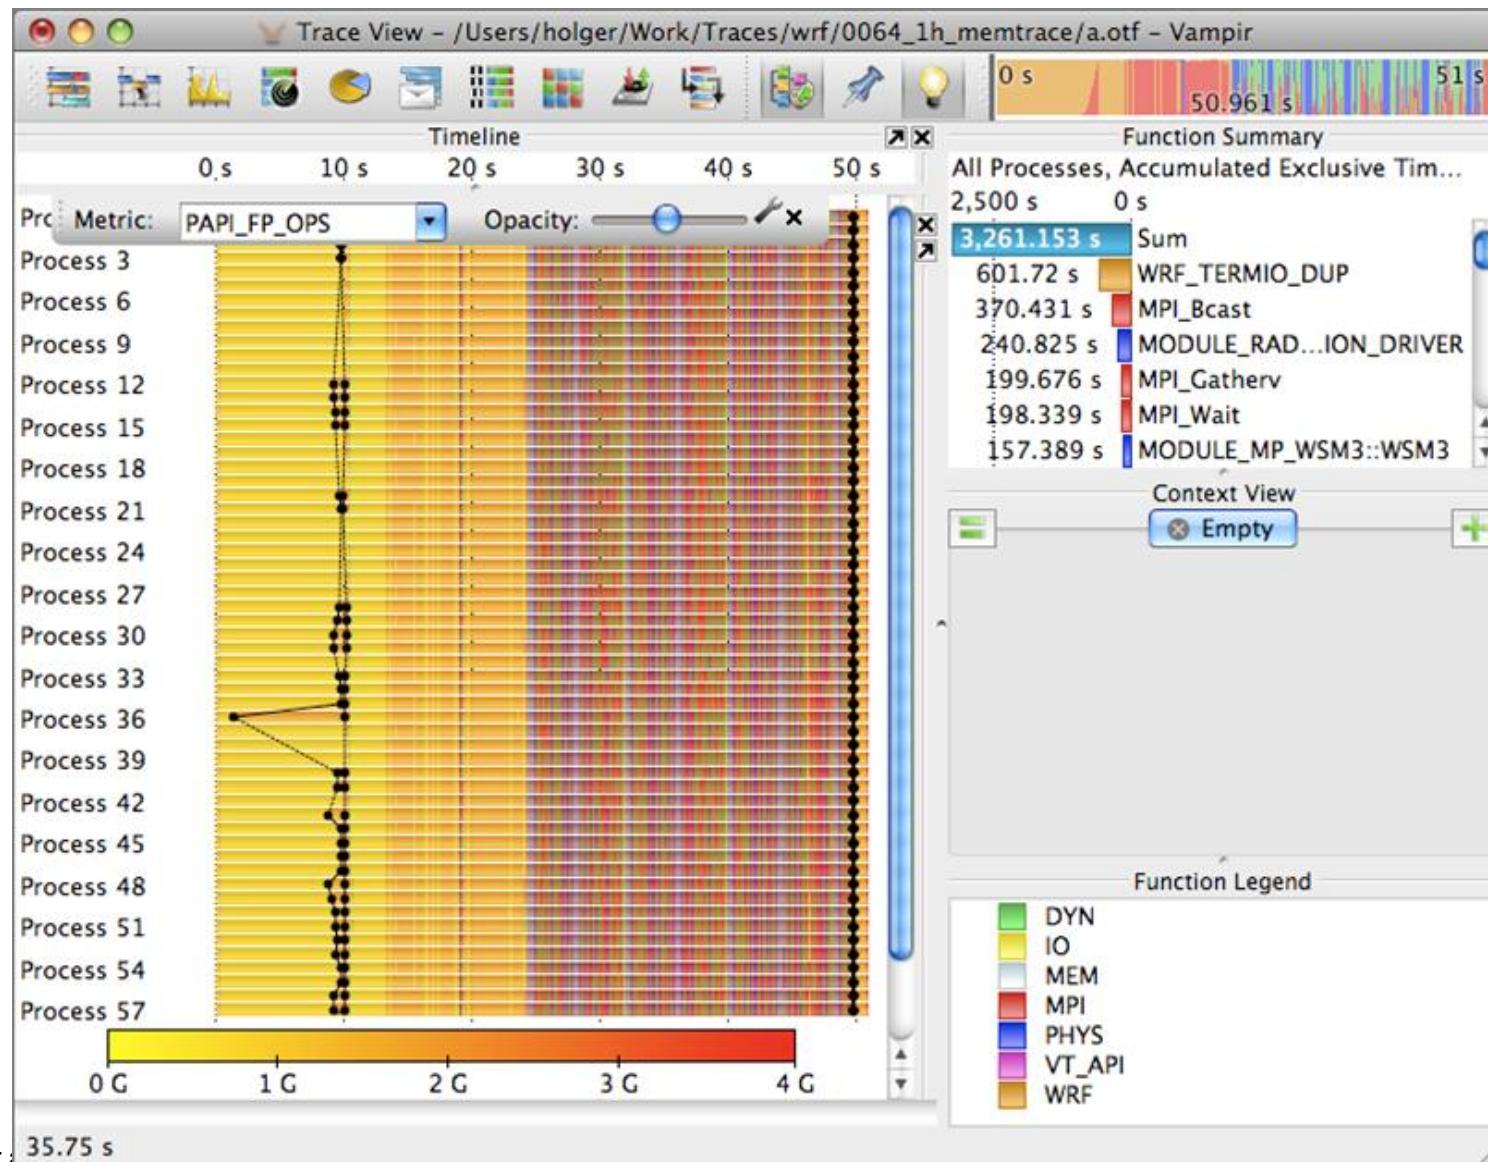
Performance Radar

VI-HPS

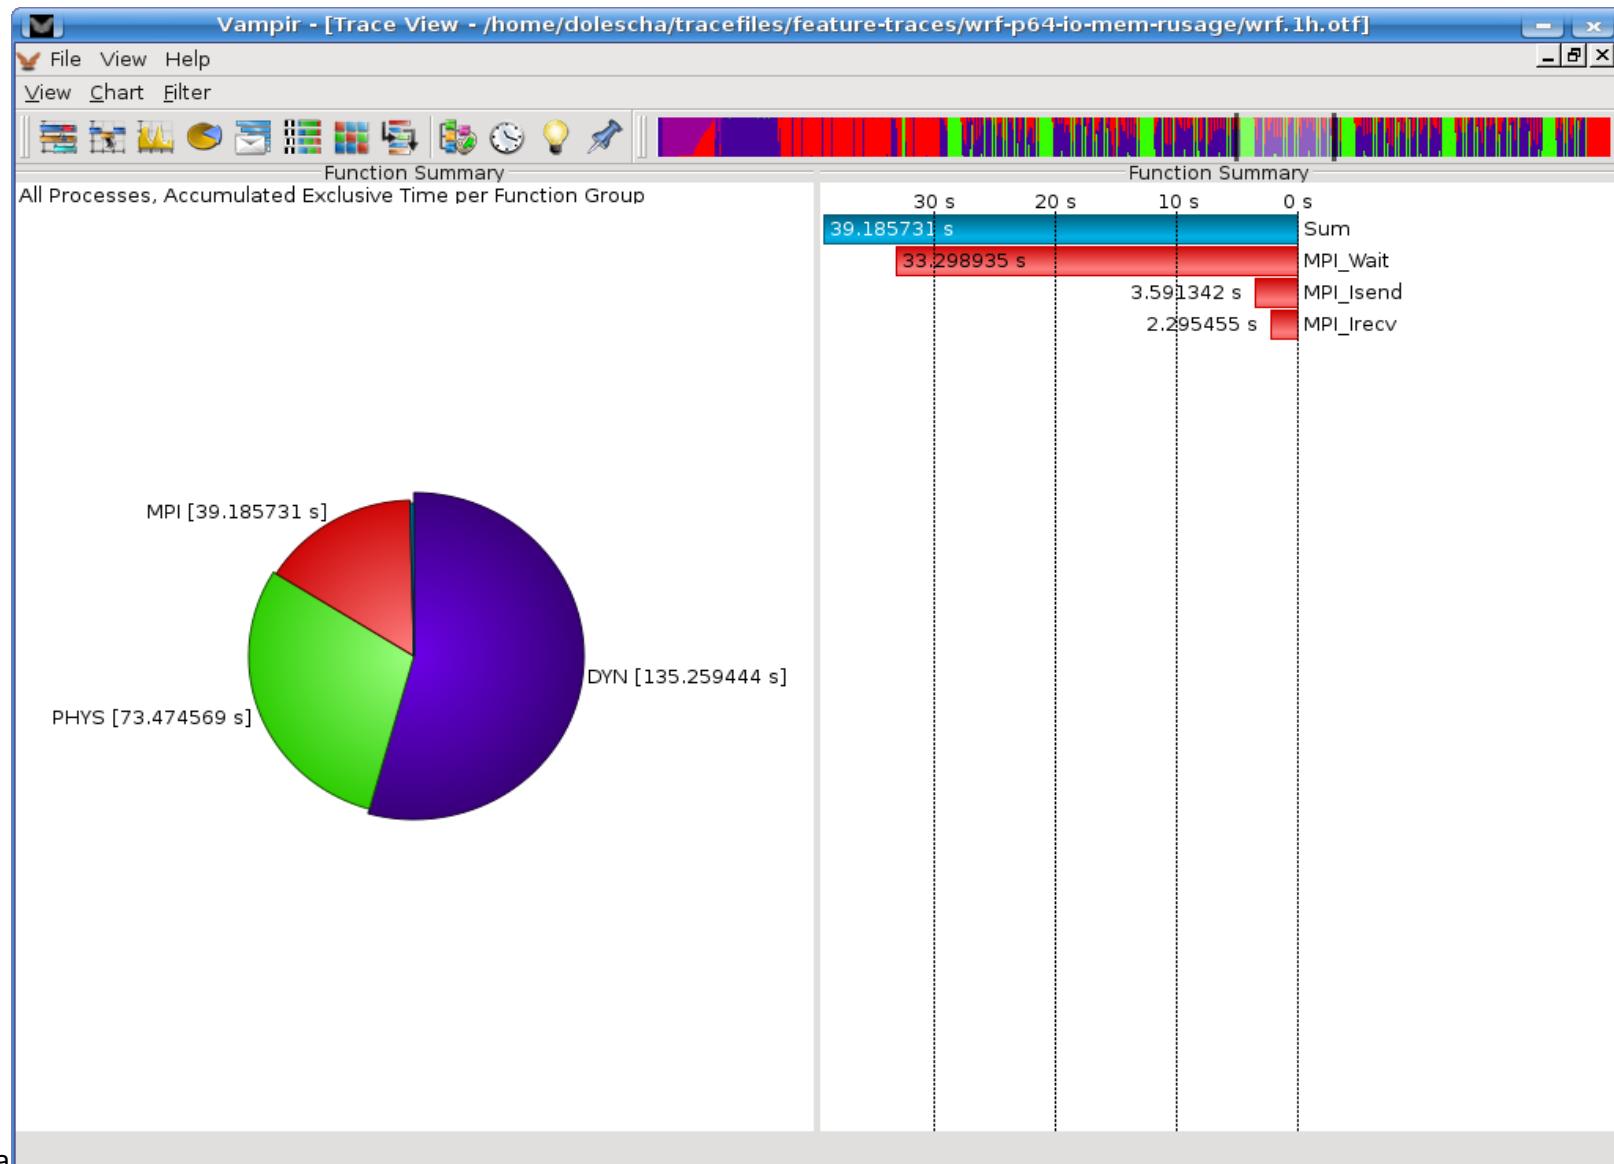


Performance Radar in Overlay Mode

VI-HPS

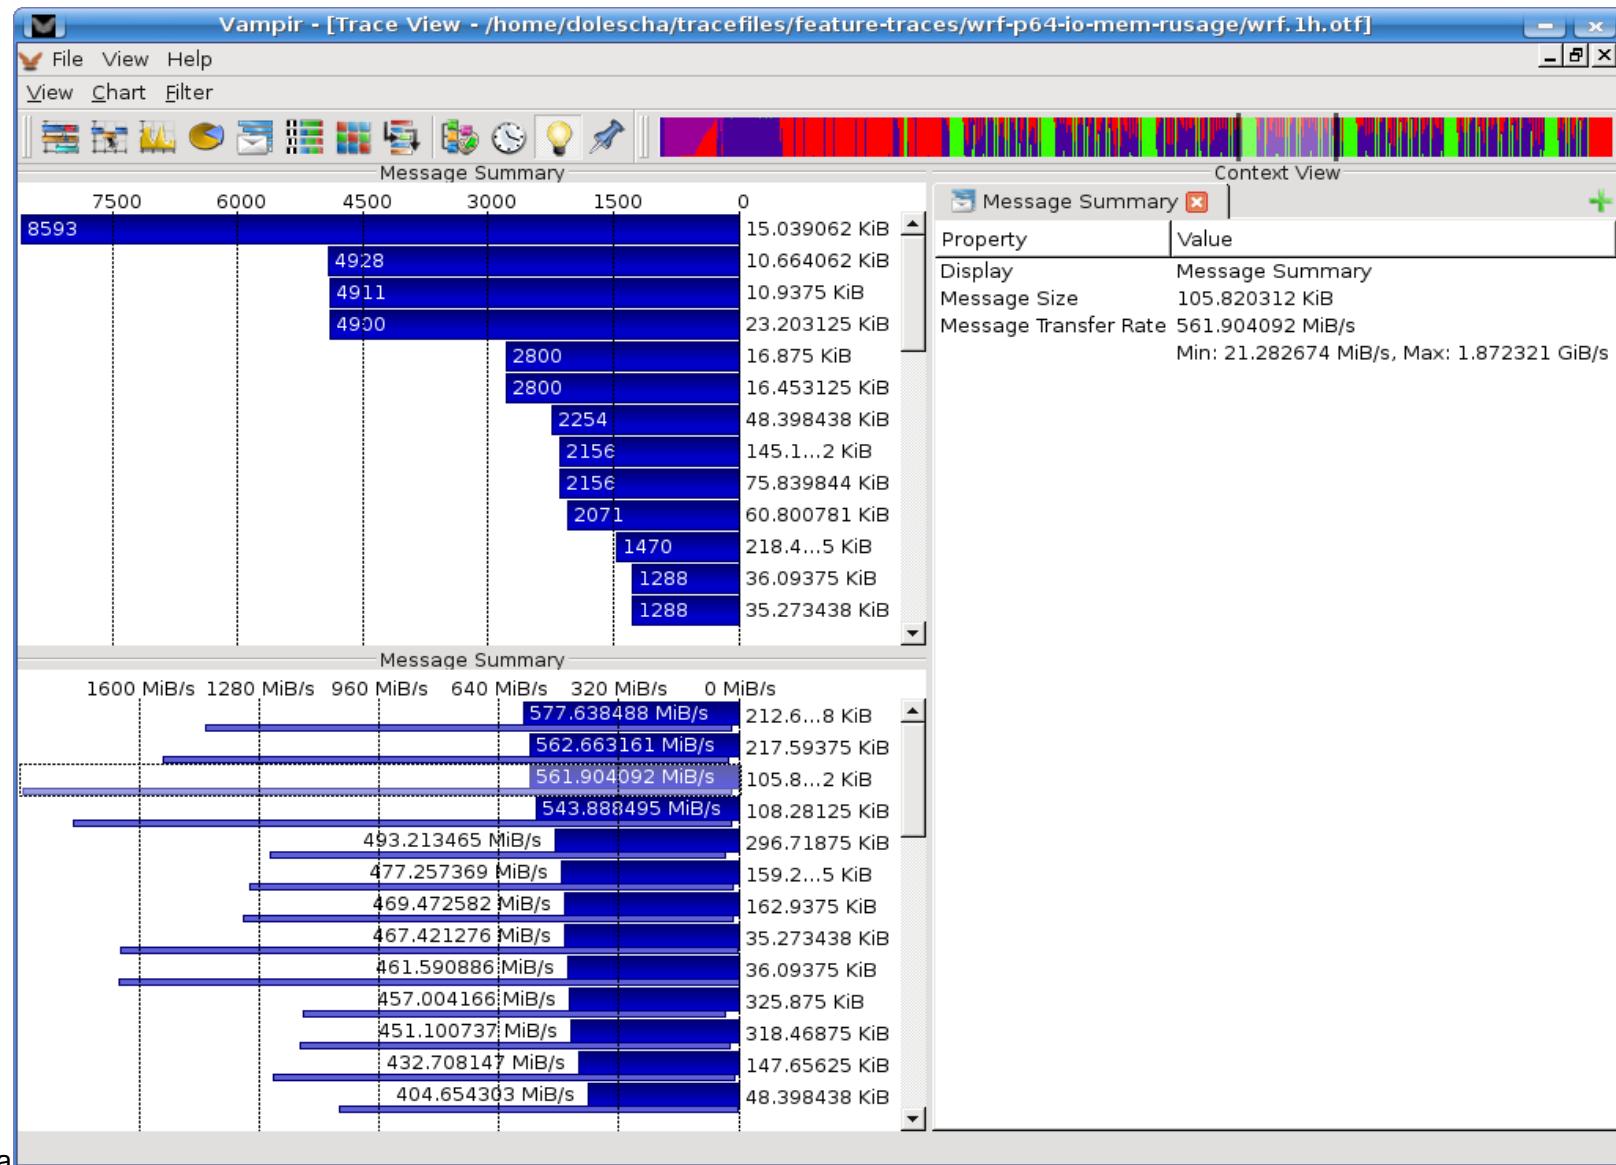


Function Summary



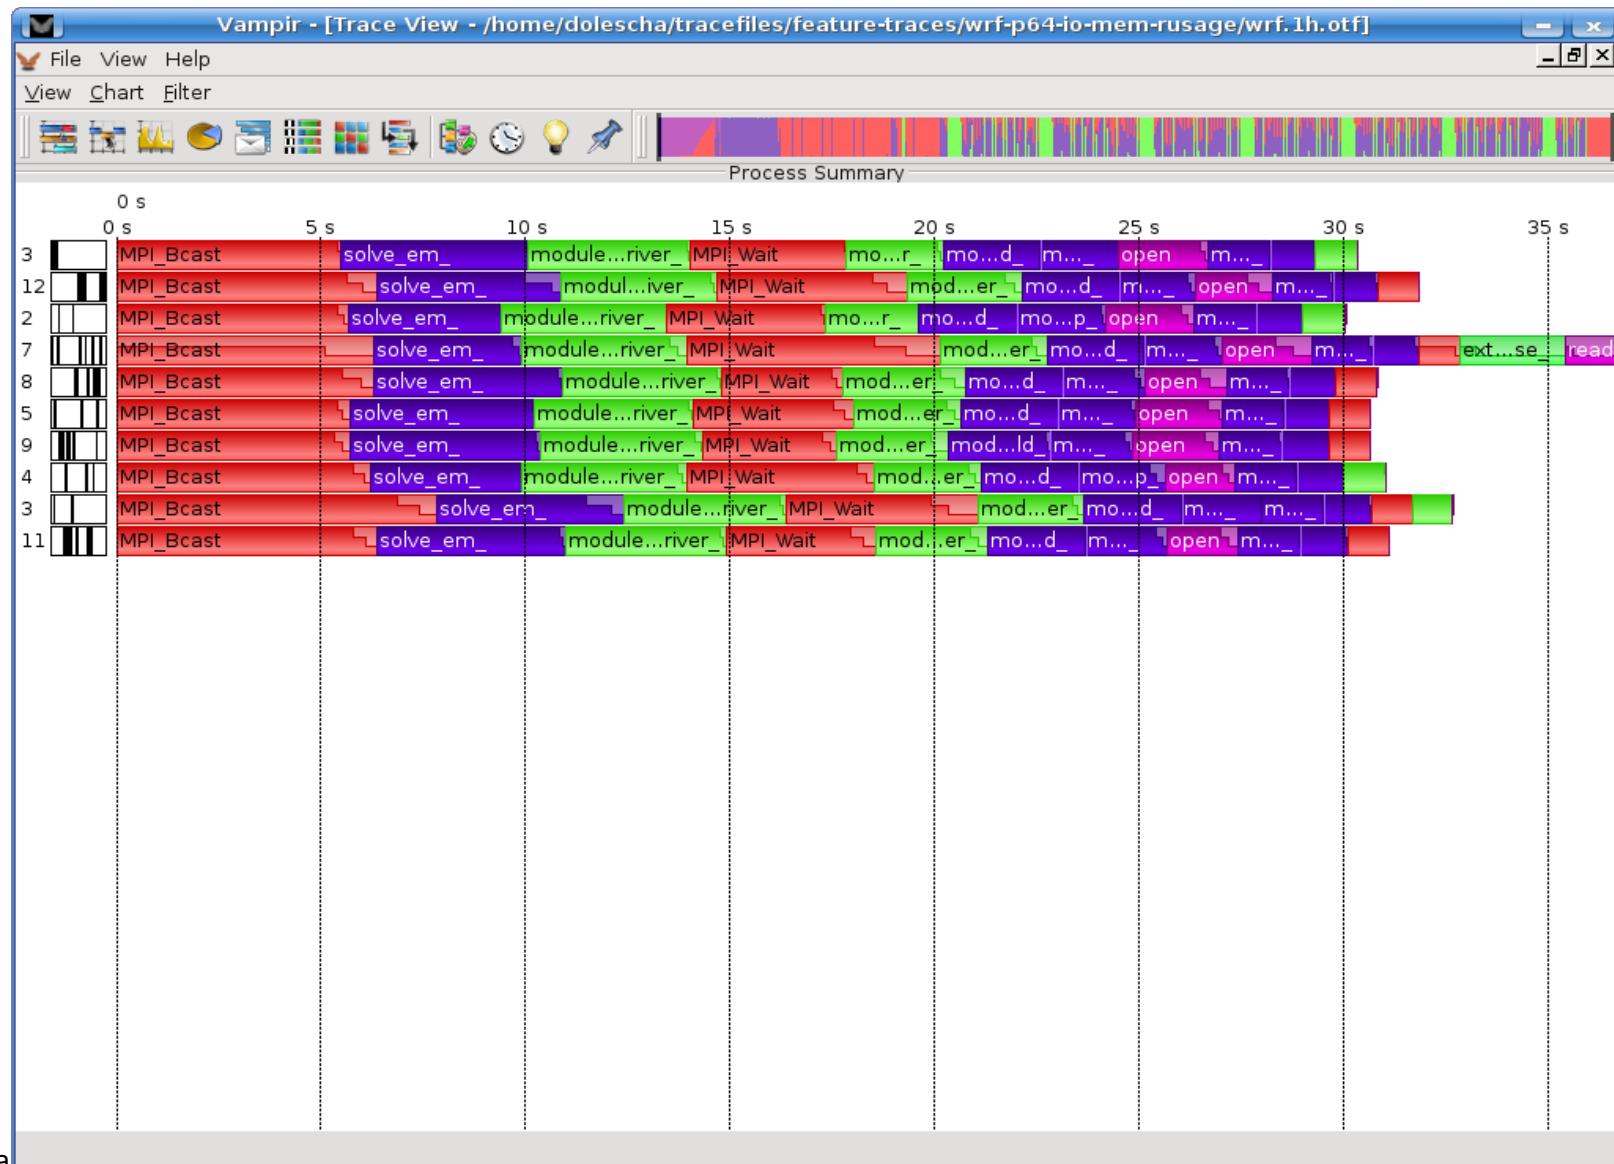
Message Summary

VI-HPS



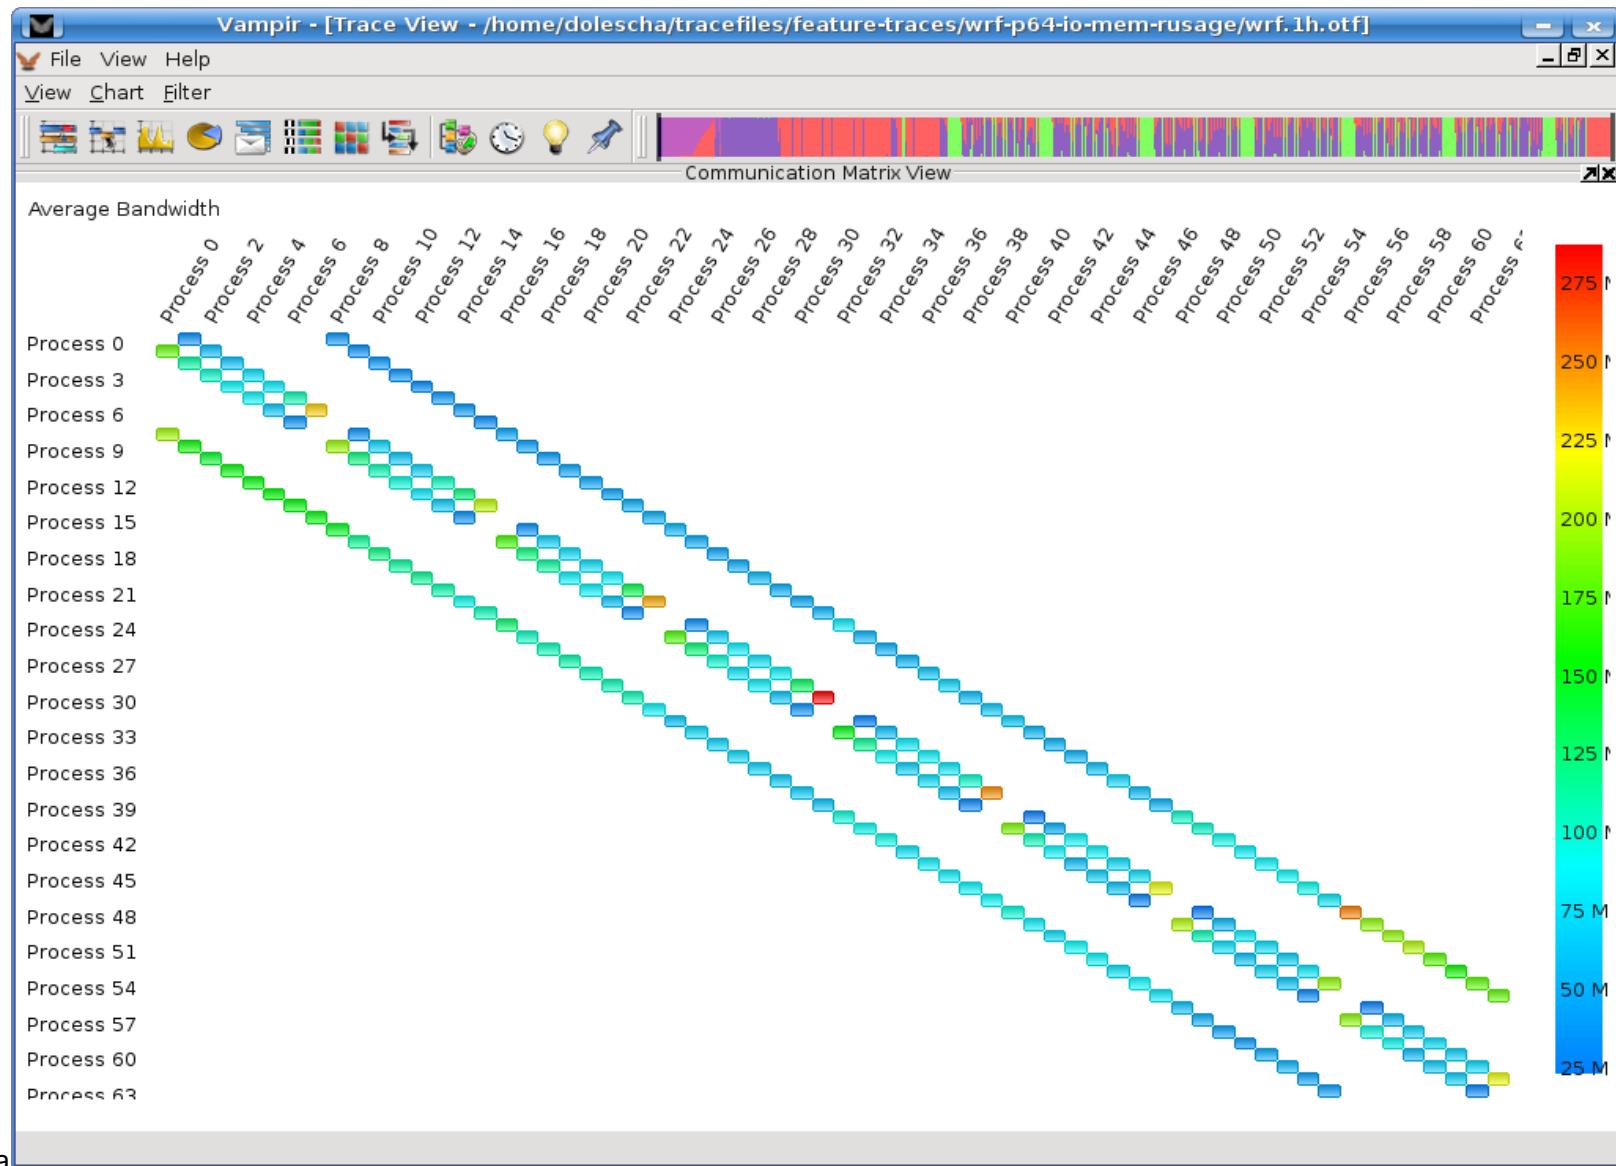
Process Summary

VI-HPS



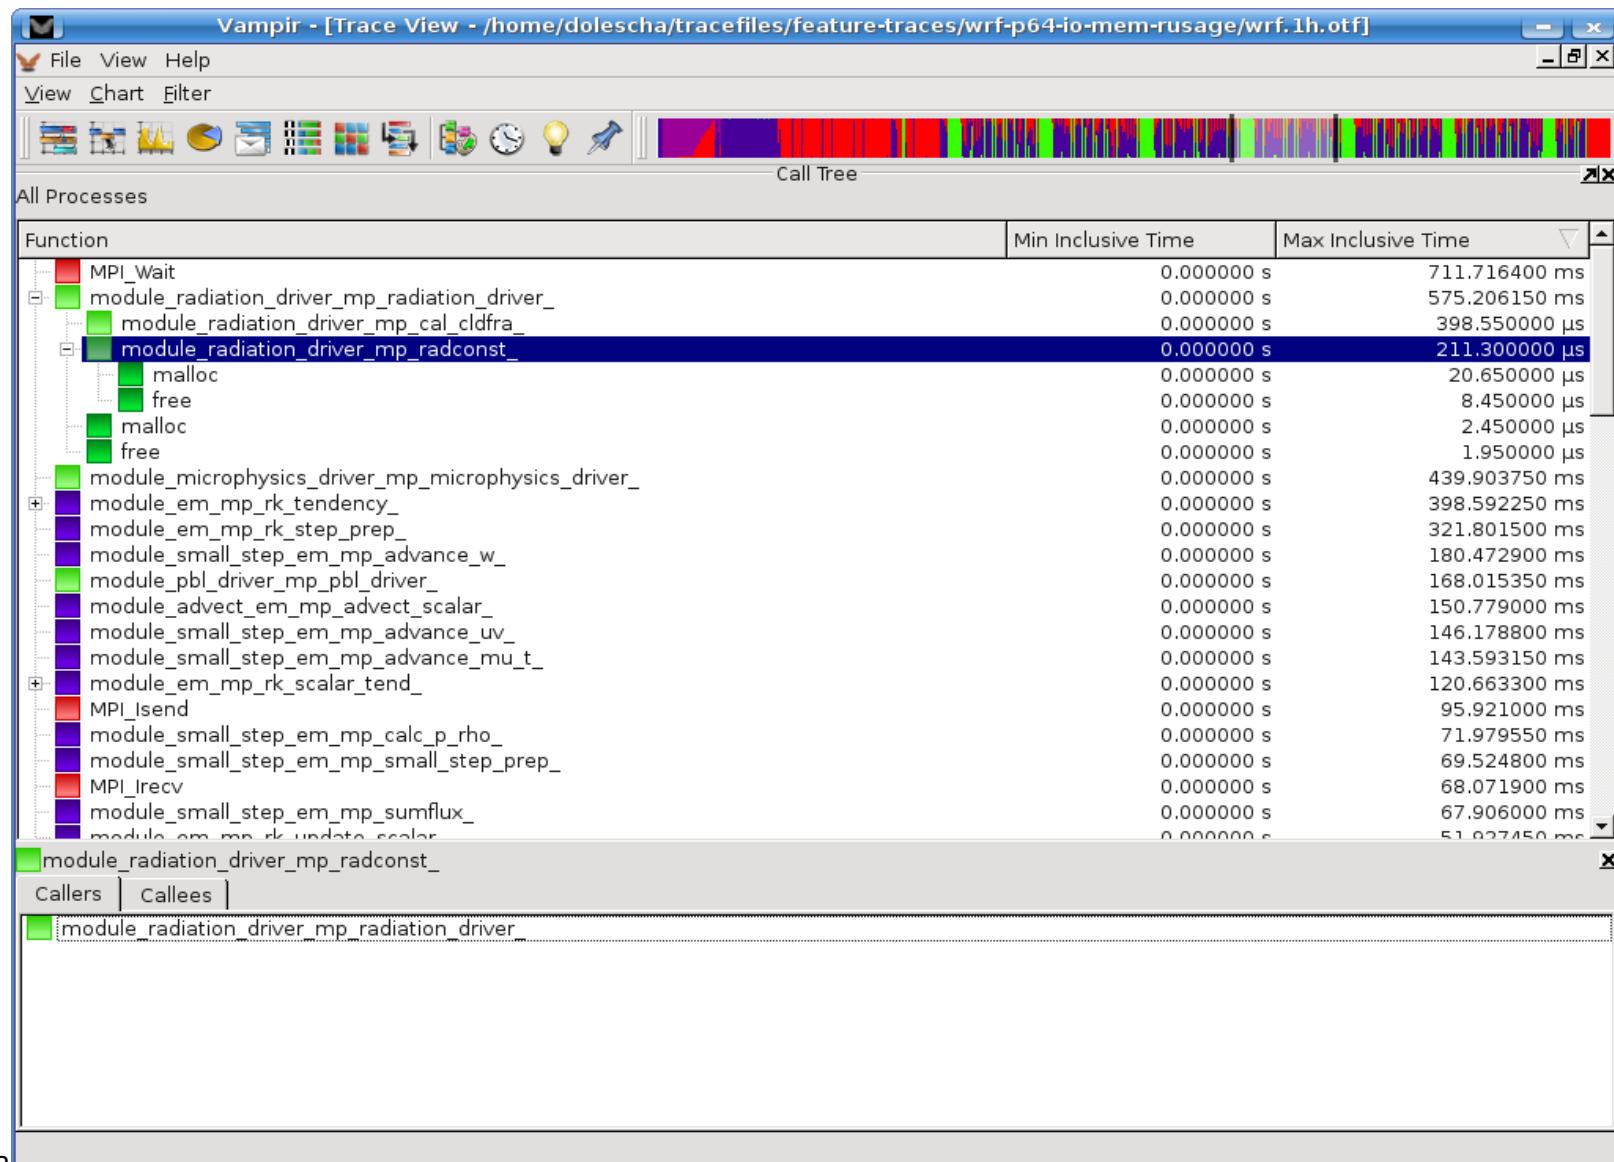
Communication Matrix

VI-HPS



Call Tree

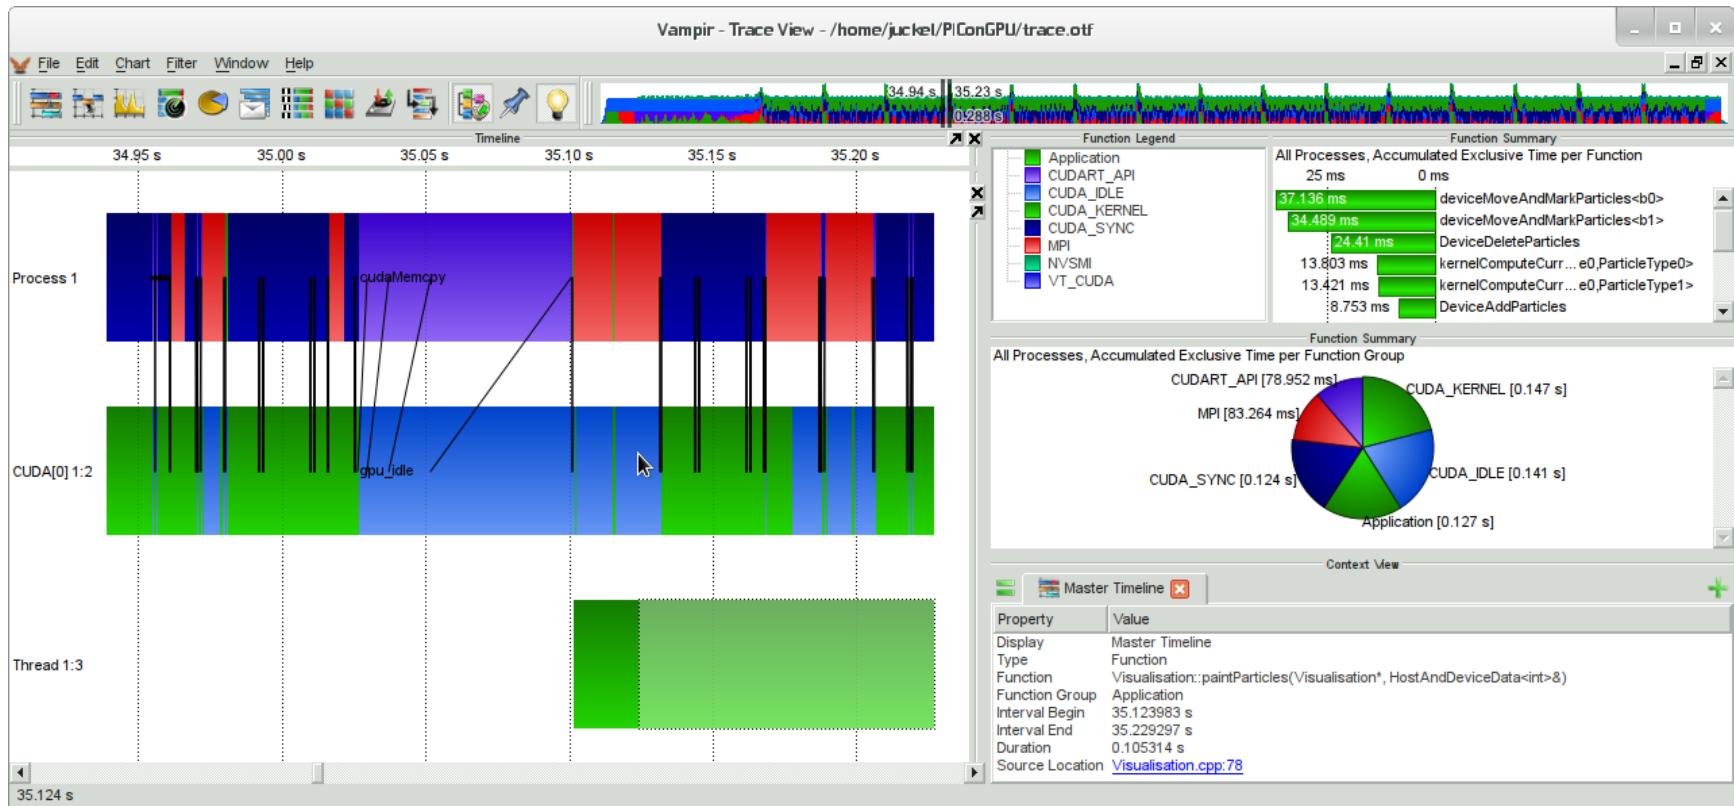
VI-HPS



- Introduction
- Vampir displays
- GPGPU support
- Scalability
- Performance Problems
- Conclusions & Preview on Tutorial

GPU Tracing for CUDA

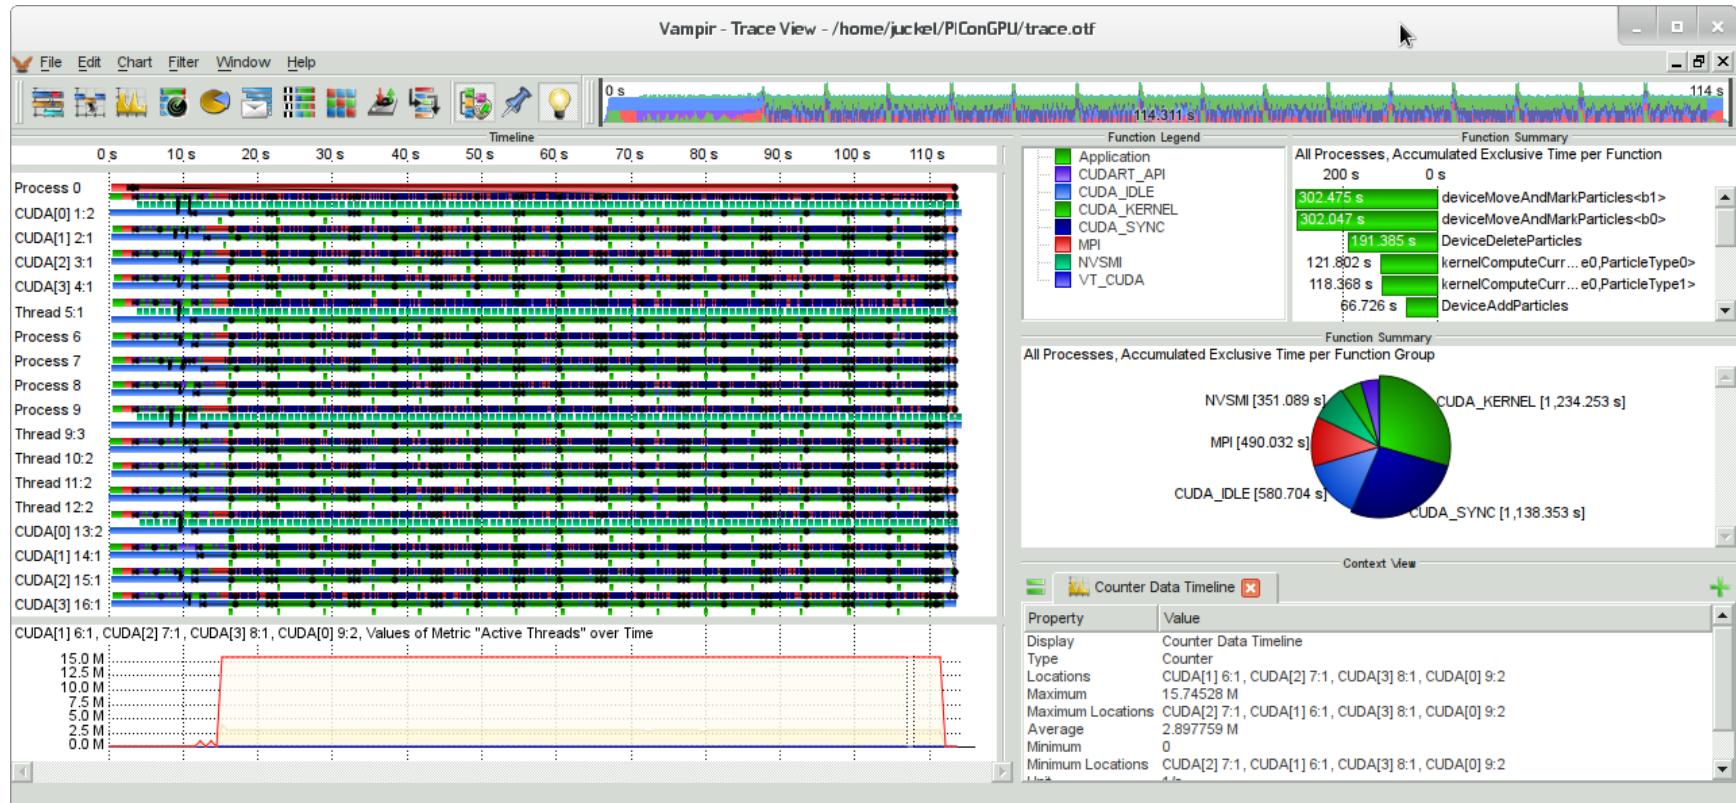
VI-HPS



- Support for GPGPU computing with CUDA (OpenCL support in Score-P soon)
- Shows Kernel invocations and idle time (individual DMA transfers shown with Score-P soon)

GPU Tracing and MPI

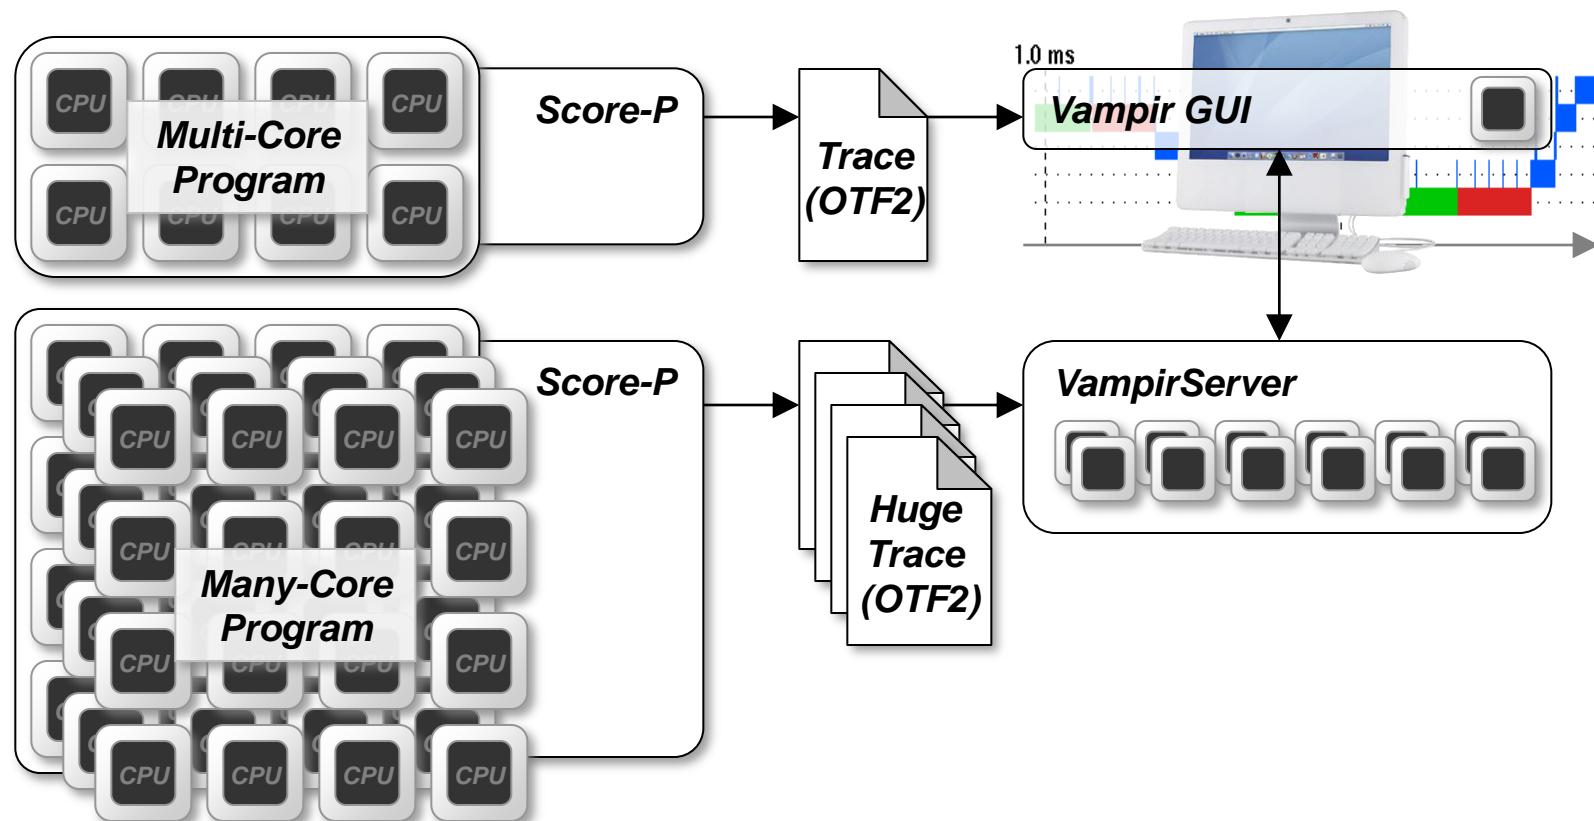
VI-HPS



- Also for hybrid MPI+CUDA applications for large scale GPU computing beyond a single compute node

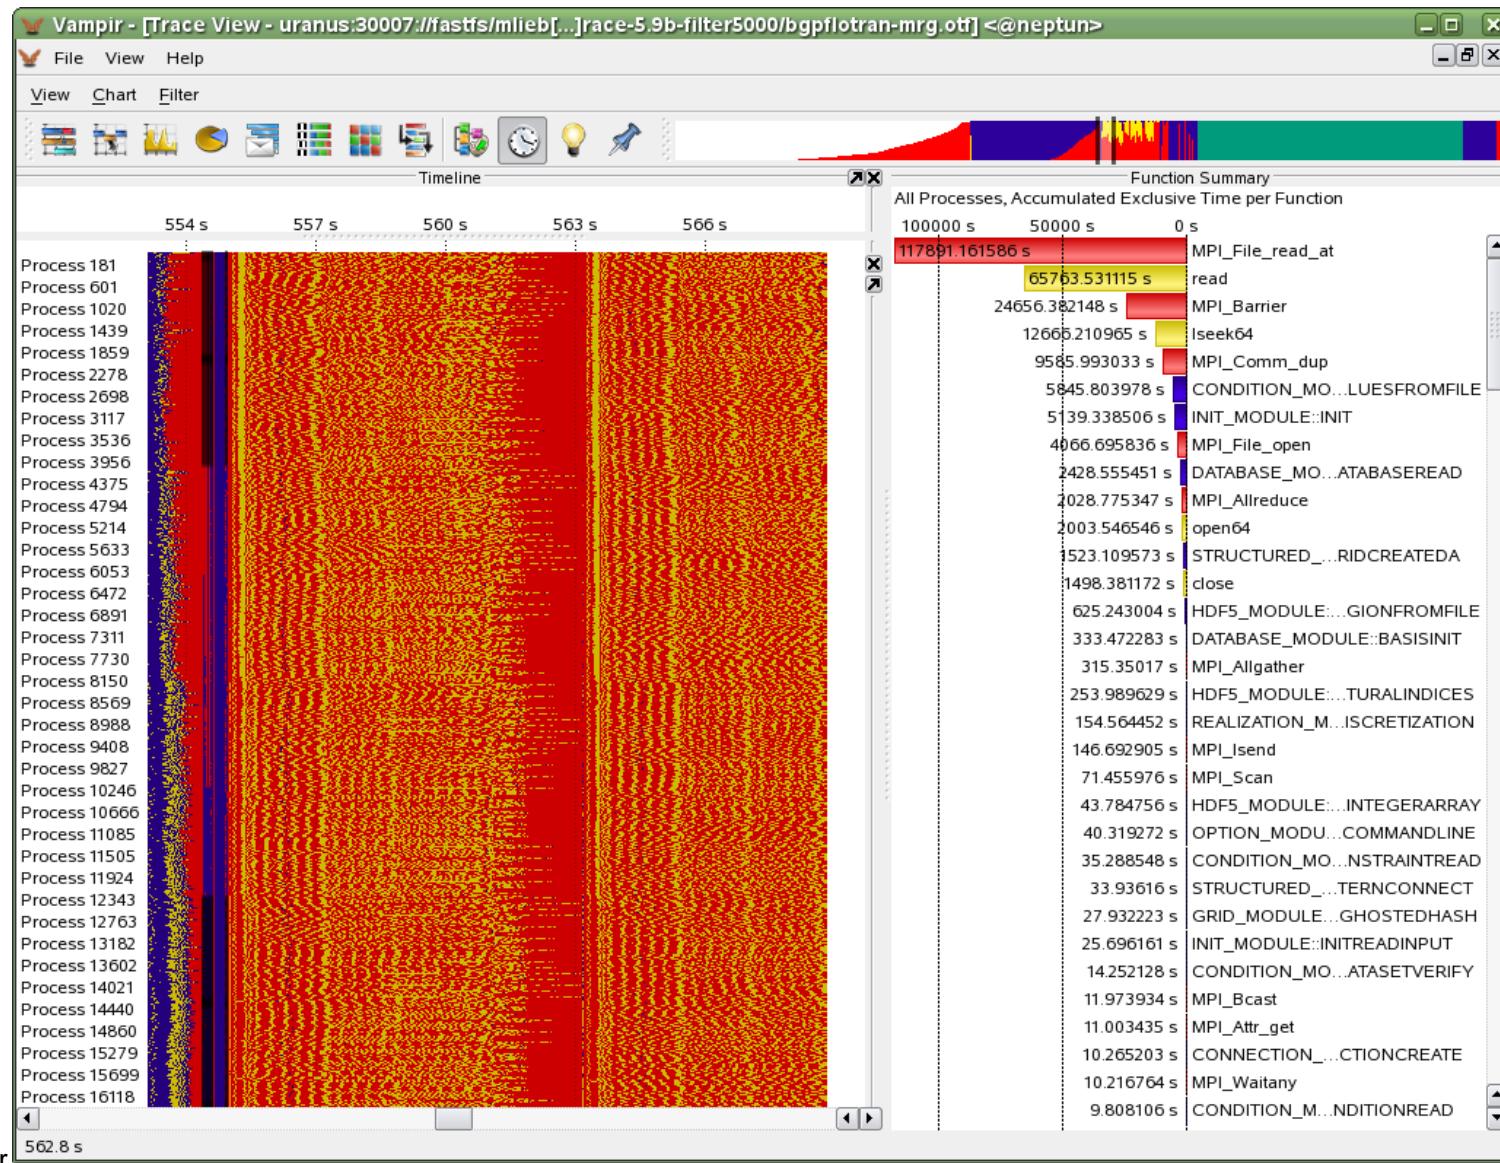
- Introduction
- Vampir displays
- GPGPU support
- Scalability
- Performance Problems
- Conclusions & Preview on Tutorial

Scalable parallel Vampir visualization architecture



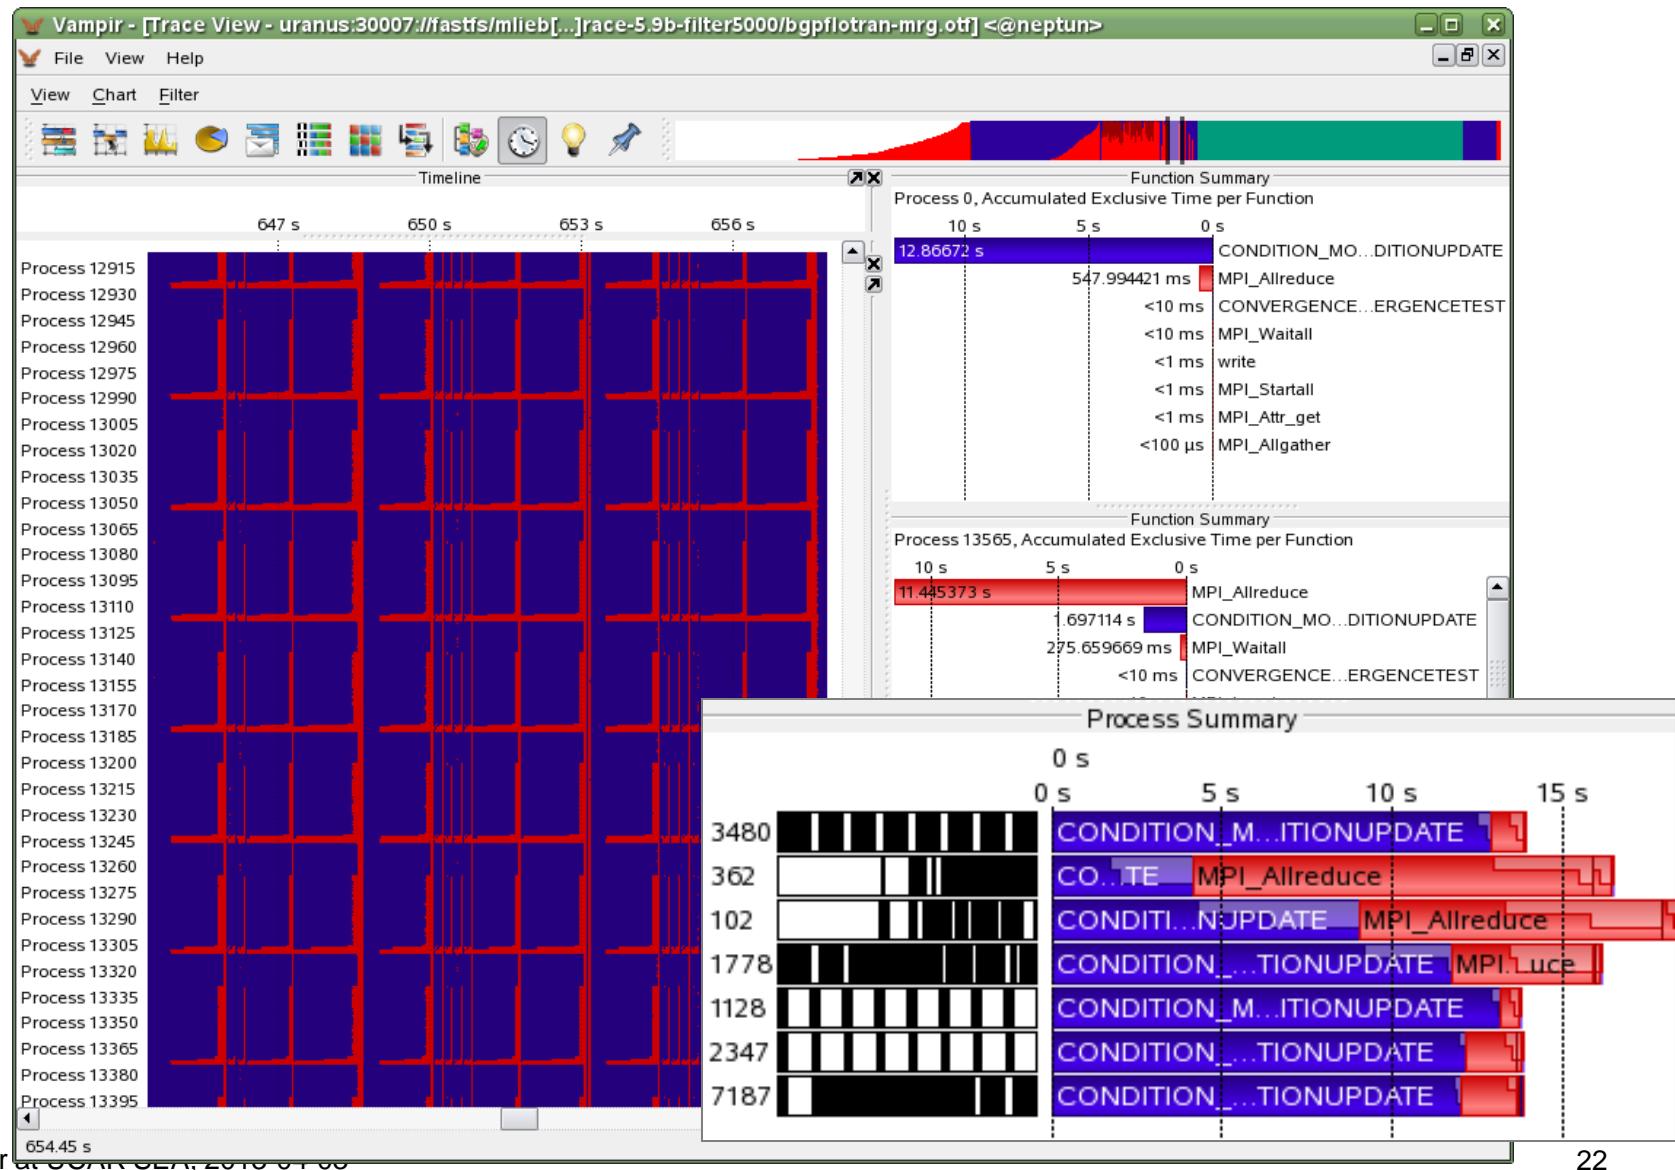
Vampir's Scalability: Full View

VI-HPS



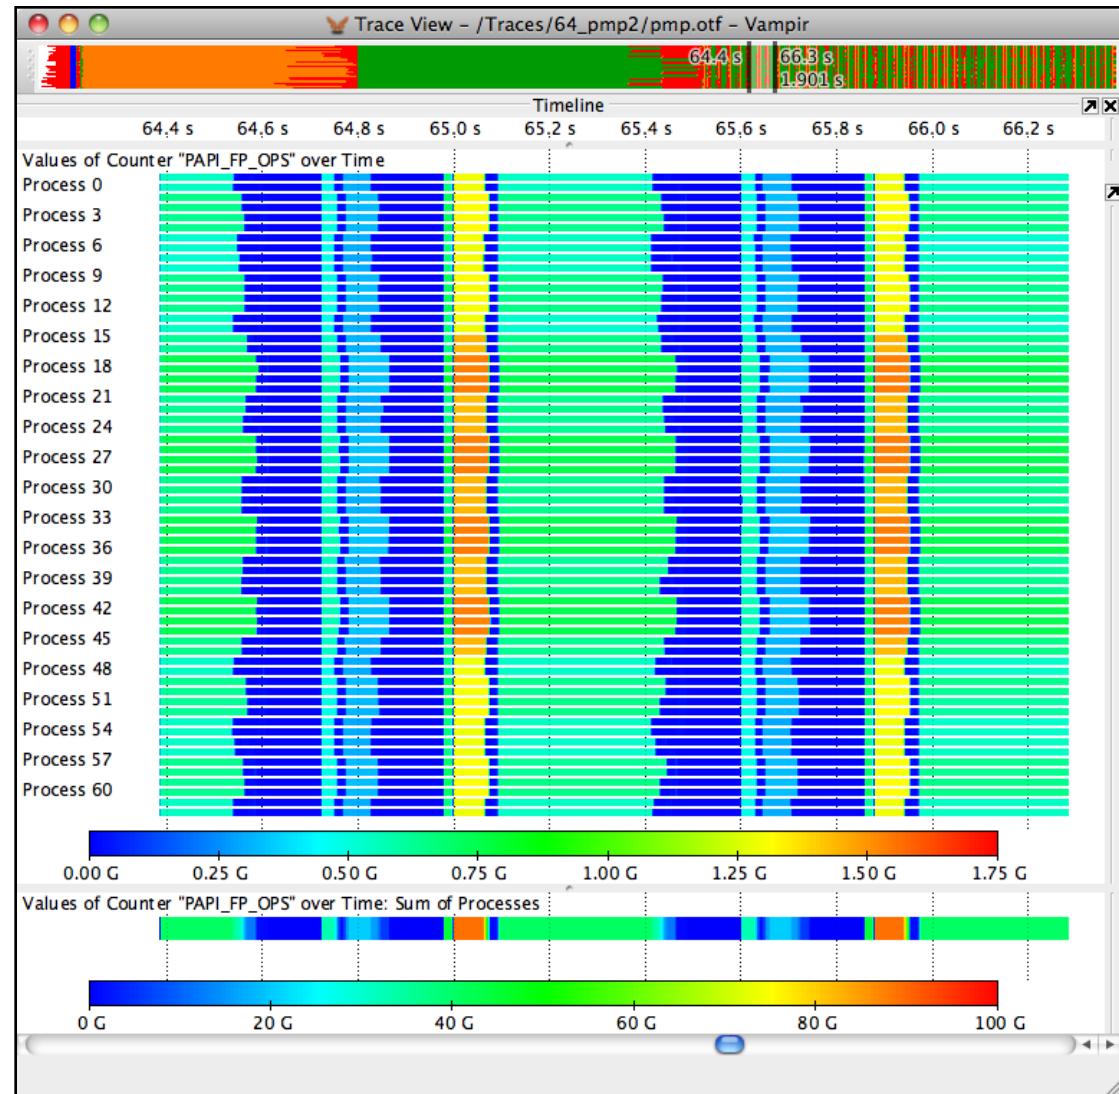
Vampir's Scalability: Process Clustering

VI-HPS



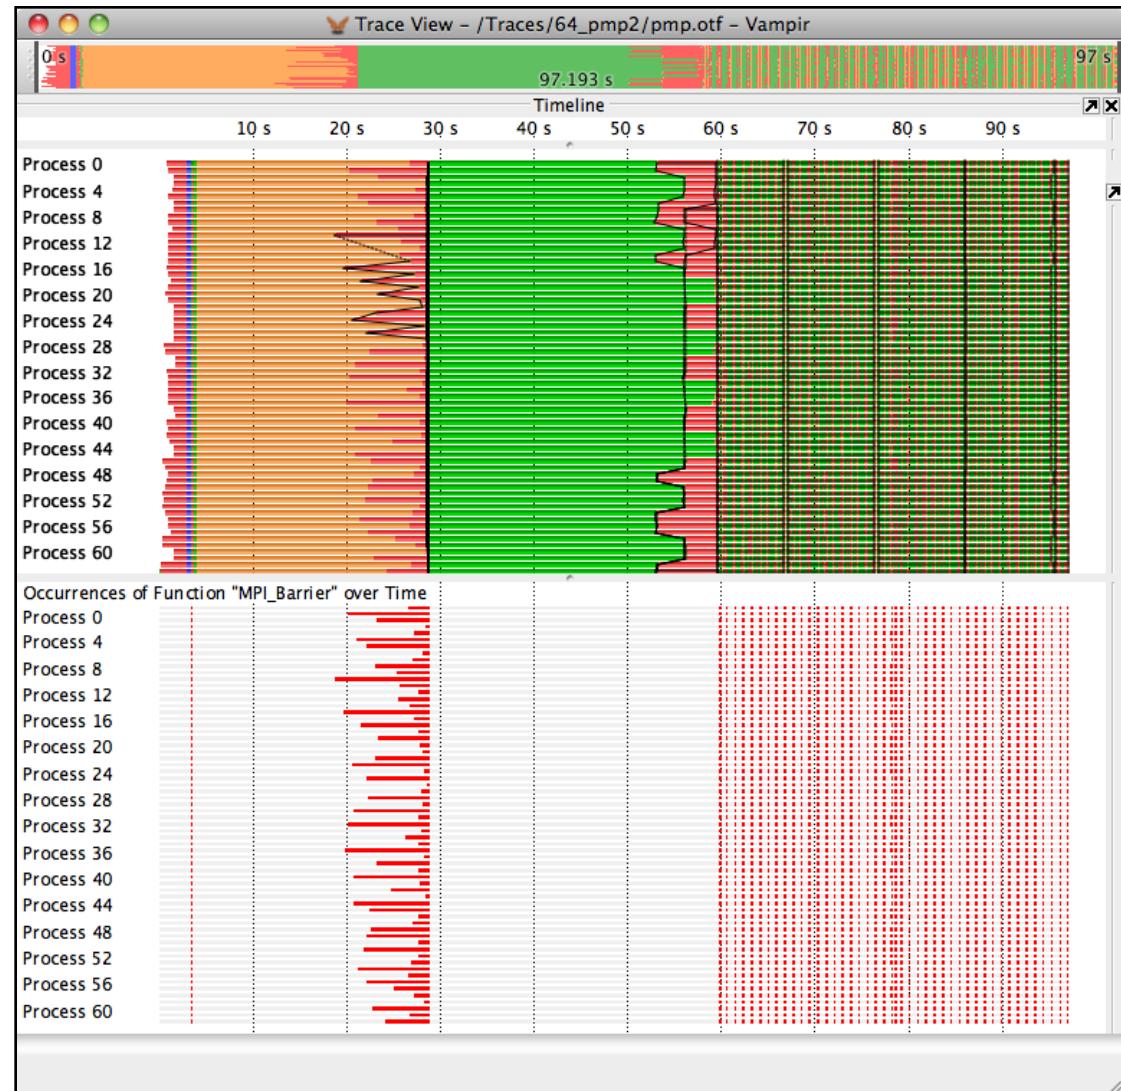
Identification of relevant spots with a heat map:

- Hardware performance counters
- Metrics derived from events, e.g. frequency of function calls, here MPI_Barrier



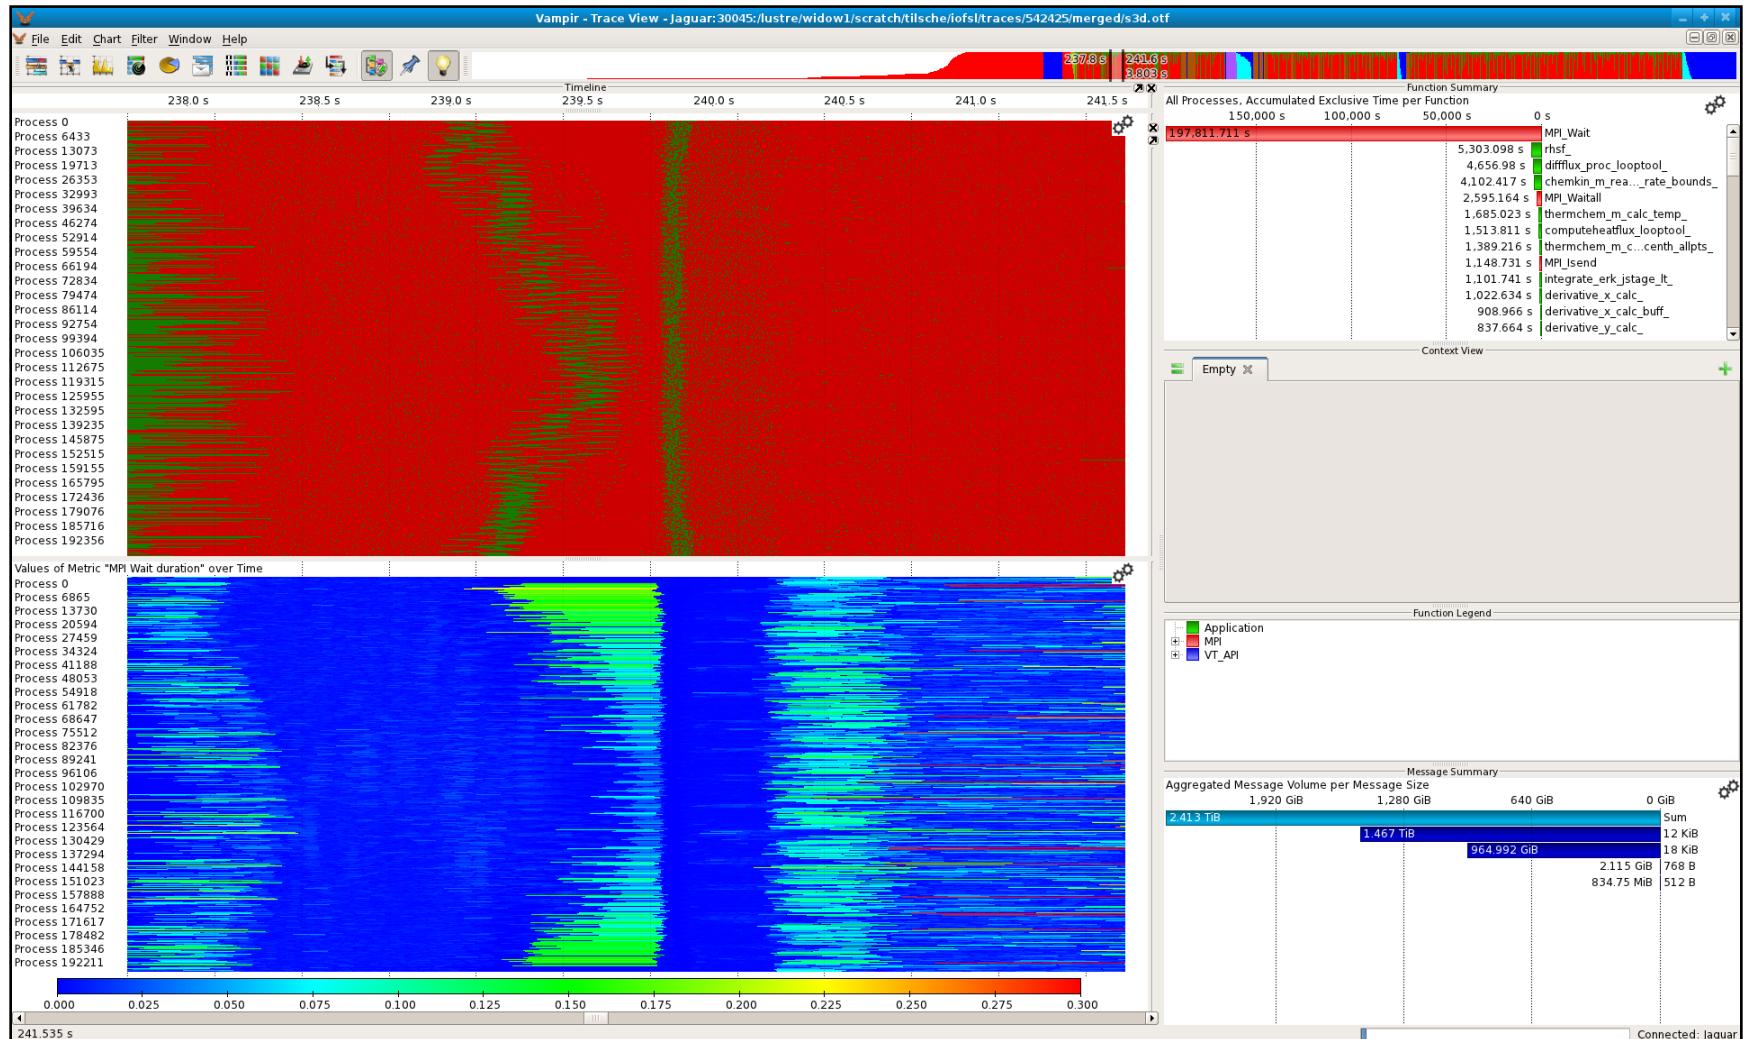
Identification of relevant spots with a heat map:

- Hardware performance counters
- Metrics derived from events, e.g. frequency of function calls, here MPI_Barrier



Scalability: 200,000+ Processes in Vampir

VI-HPS



- Introduction
- Vampir displays
- GPGPU support
- Scalability
- Performance Problems
- Conclusions & Preview on Tutorial

Types of performance problems:

- Serial vs. parallel performance
- Bottlenecks related to MPI, OpenMP, CUDA, ...
- Computation vs. communication
- Program phases: initialization/finalization, iteration, I/O, checkpointing, ...

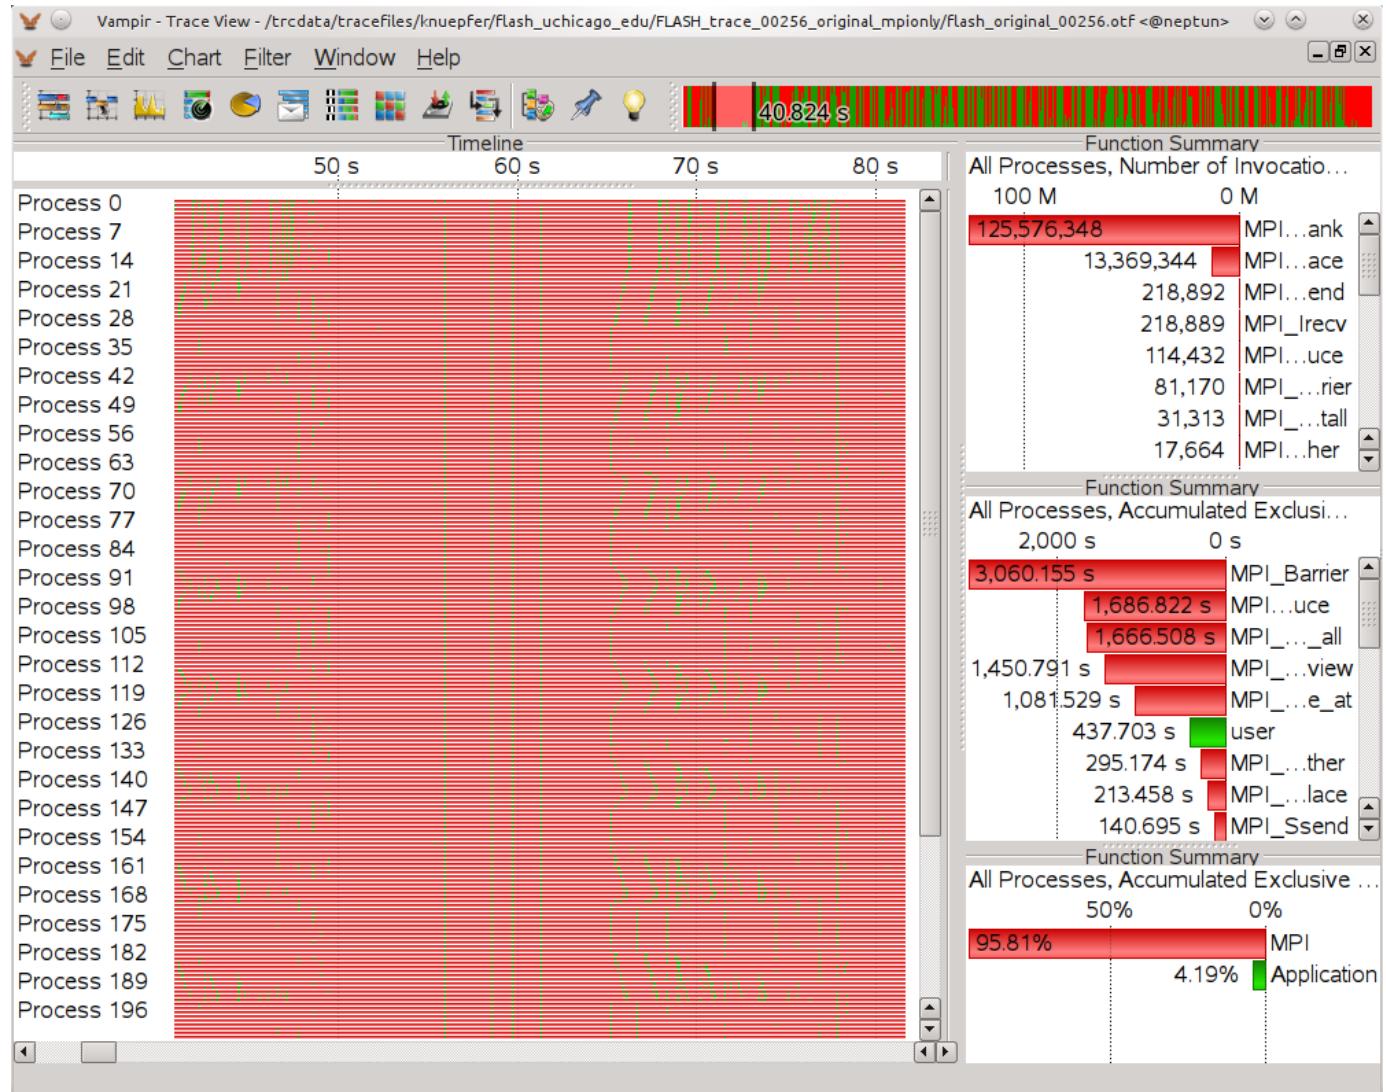
General strategy:

- Find highest consumer(s) of run-time
- Start with hypothesis about most severe performance problem, then confirm or disprove

Some examples (more in tutorial parts):

Problems with MPI

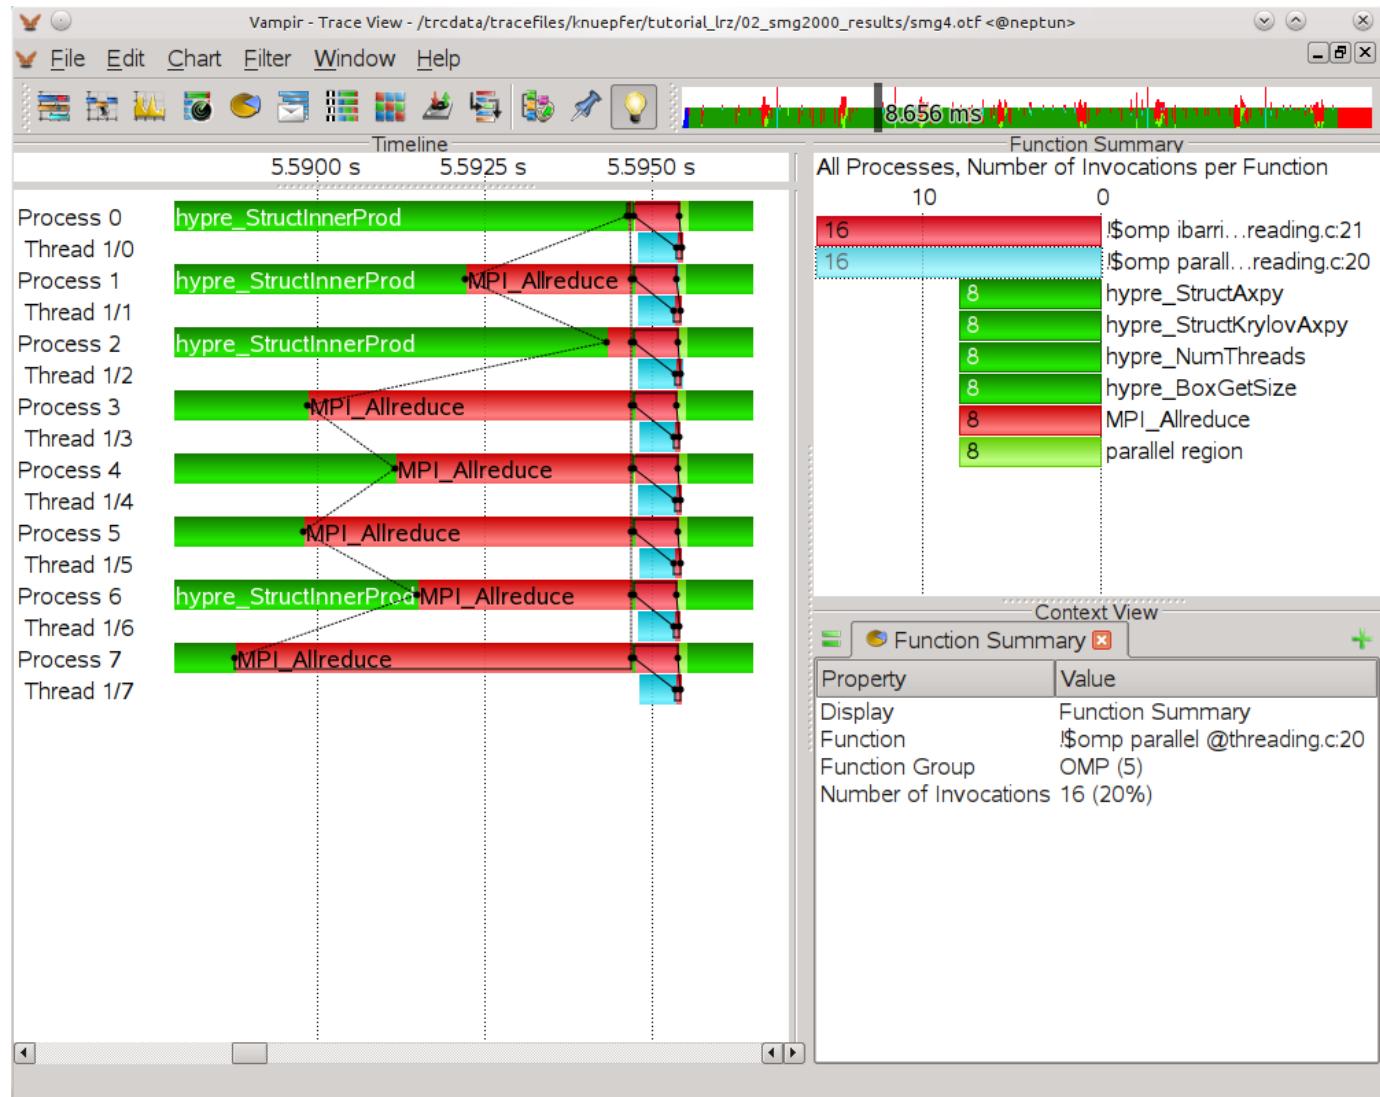
Too much runtime in MPI all-together



Problems with MPI + OpenMP

Imbalanced computation causes MPI waiting time

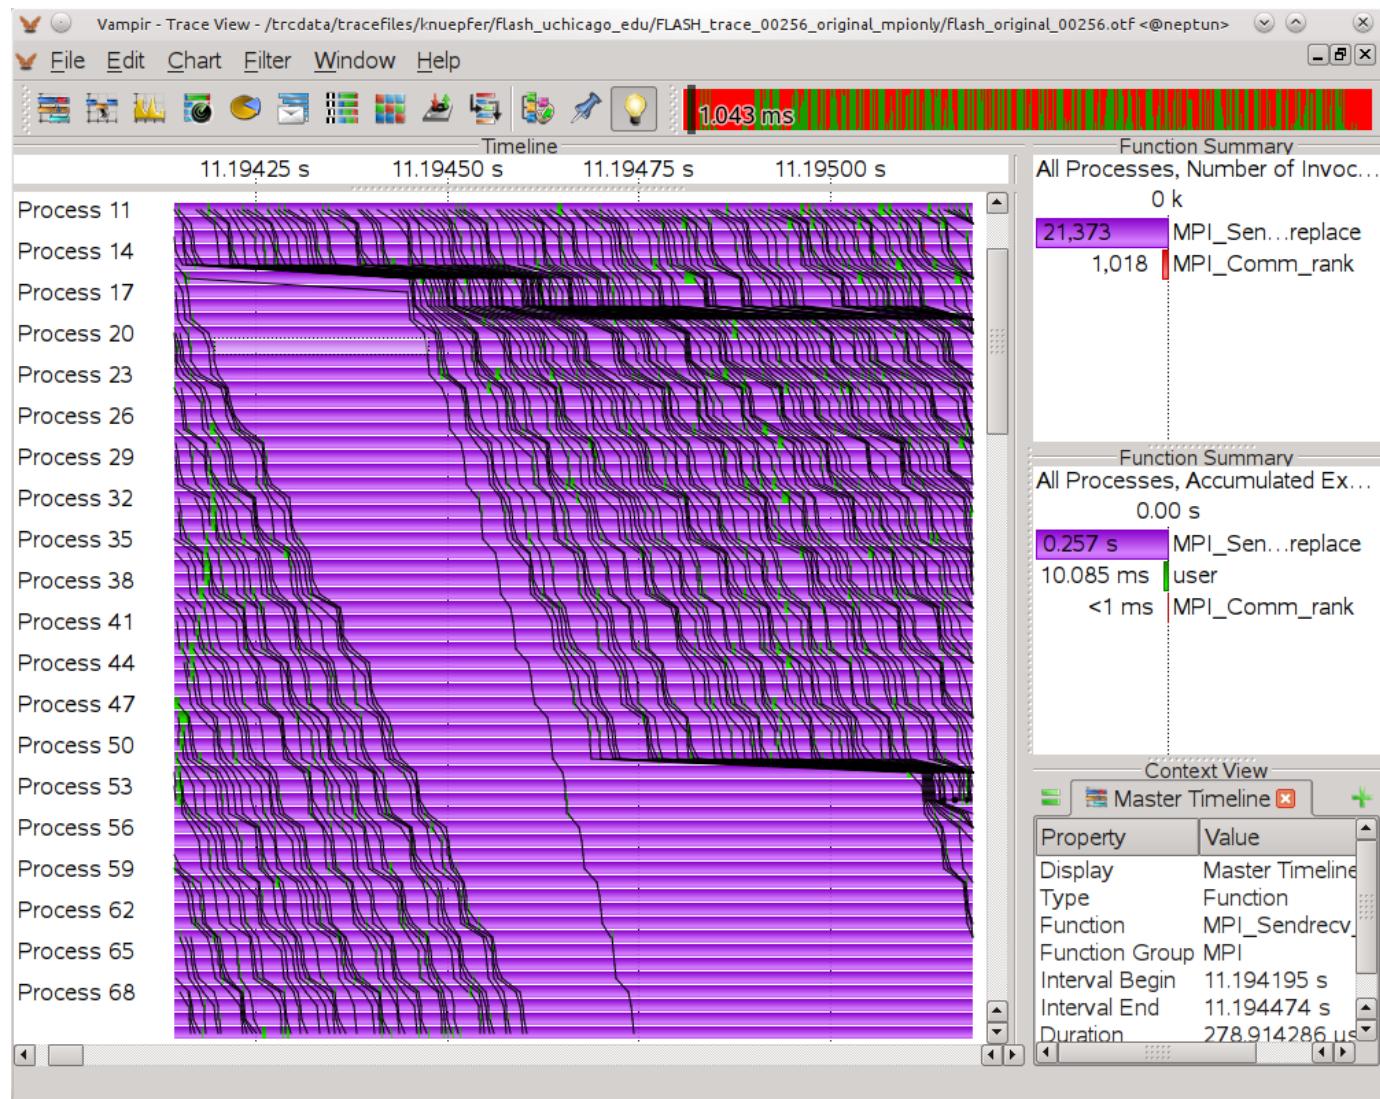
OpenMP threads mostly idle



Problems with strange MPI patterns

VI-HPS

Strange MPI
patterns
propagate
delays over
successive
MPI ranks

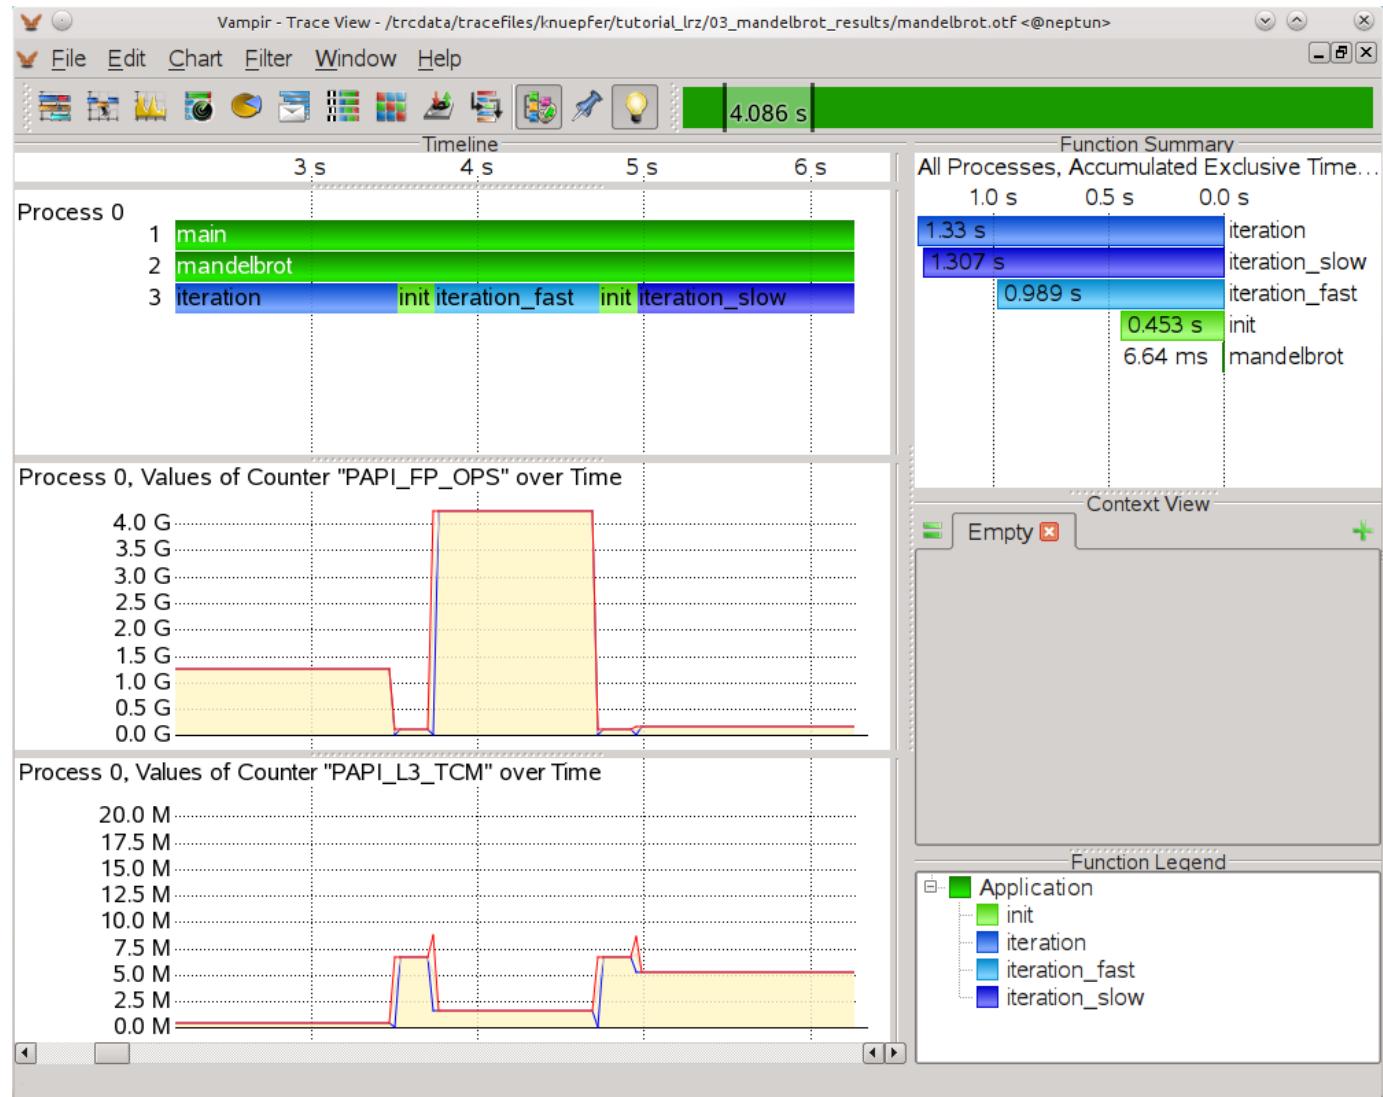


Problems for serial performance

High FP rate,
low L3 cache
miss rate

VS.

low FP rate
due to a high
L3 cache
miss rate

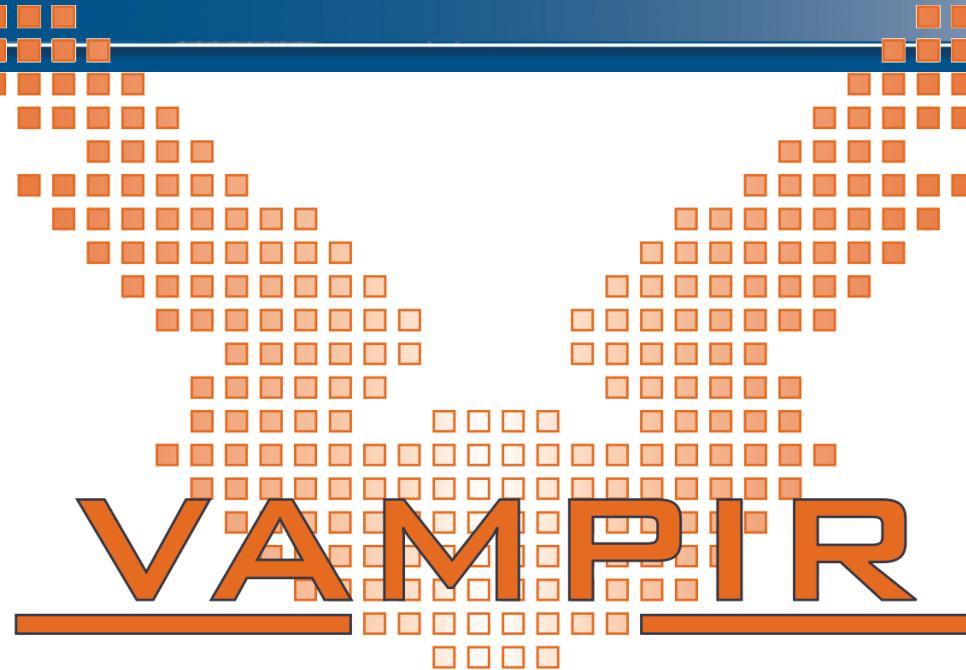


Conclusions:

- User-friendly and comprehensible visualization of complex and highly parallel run-time behavior
- Various aspects of parallel performance
- Highly scalable by itself

Tutorial on Thursday & Friday:

- Run-time measurement with Score-P + hands-on
- Analysis with Vampir, Scalasca, TAU + hands-on
- Combination of Vampir and Scalasca
- Work with your code and tool(s) of your choice
- GPGPU measurement and analysis (on request)



Matthias Jurenz, Andreas Knüpfer, Ronny Brendel, Matthias Lieber,
Jens Doleschal, Jens Domke, Holger Mickler, Daniel Hackenberg,
Michael Heyde, Thomas Ilsche, Guido Juckeland, Robert Dietrich,
Frank Winkler, Michael Kluge, Matthias Müller, Holger Brunst,
Ronald Geisler, Reinhard Neumann, Bert Wesarg, Rene Widera,
Thomas Ilsche, Matthias Weber, Thomas William,
Hartmut Mix, and Wolfgang E. Nagel

<http://www.vampir.eu>